



Aalto University
School of Science

Stream Processing and Big Data Platforms

Hong-Linh Truong

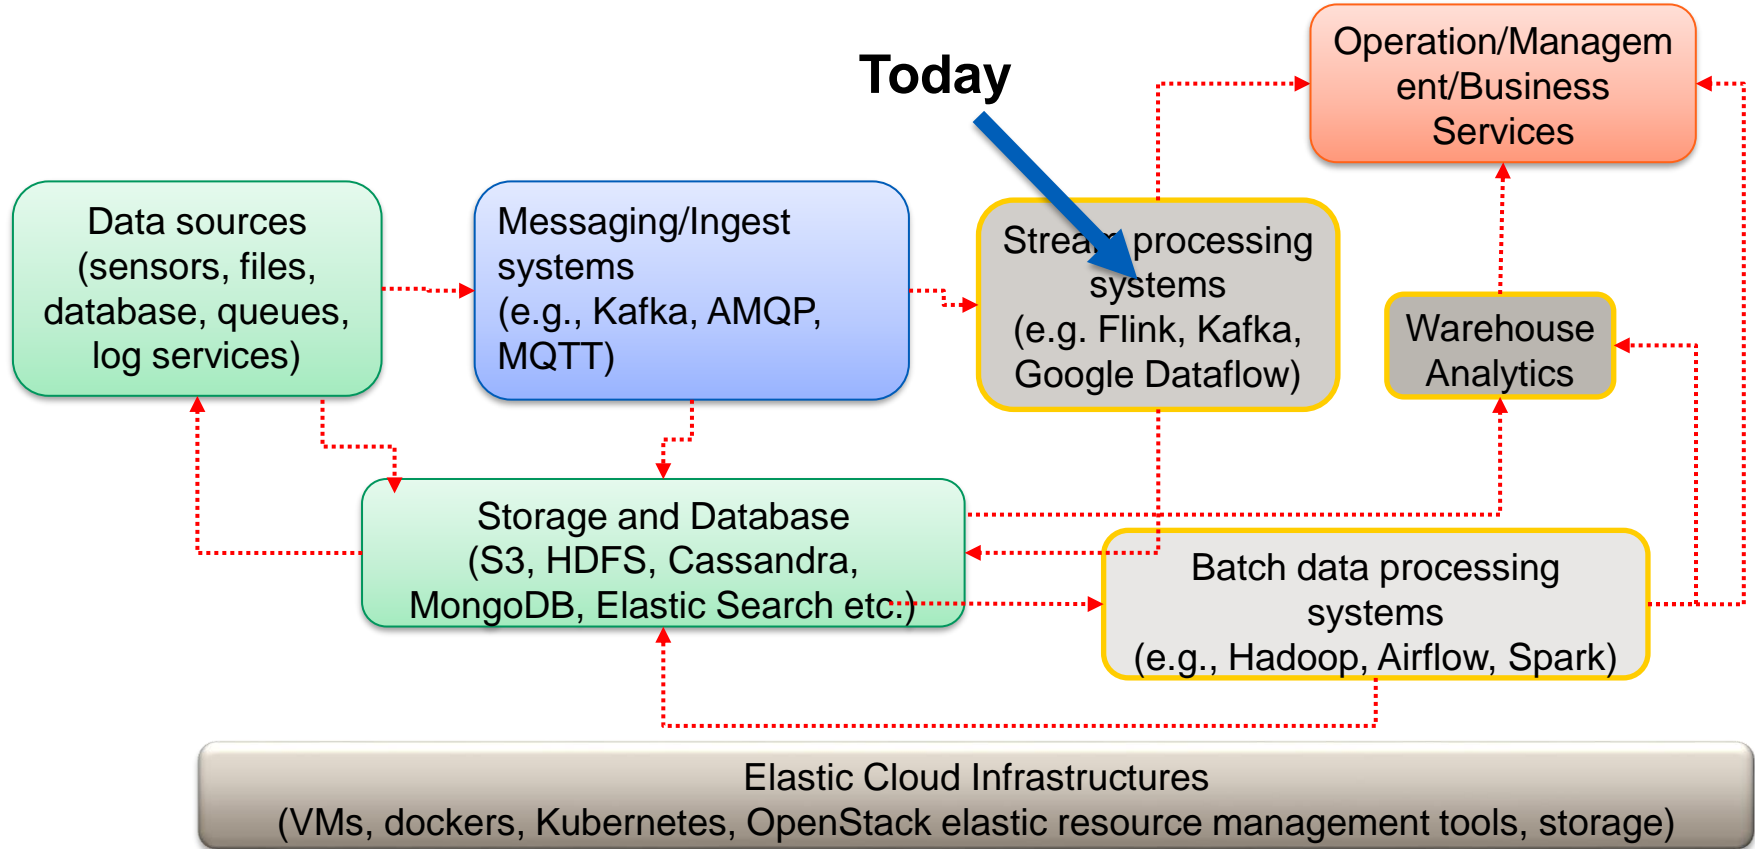
Department of Computer Science

linh.truong@aalto.fi, <https://rdsea.github.io>

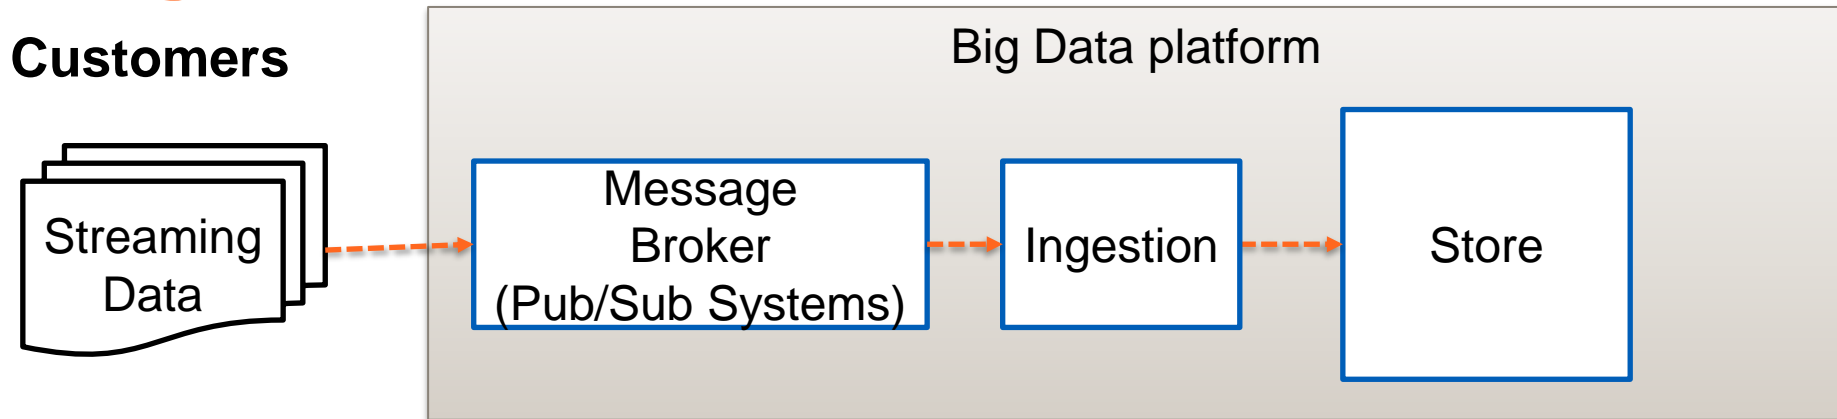
Schedule

- **Do you see the role of stream processing in big data platforms?**
- **What should we learn if you want to support stream processing in your platforms?**
 - Basic concepts and important aspects
 - Apache Flink case
- **Summary**

Big data at large-scale

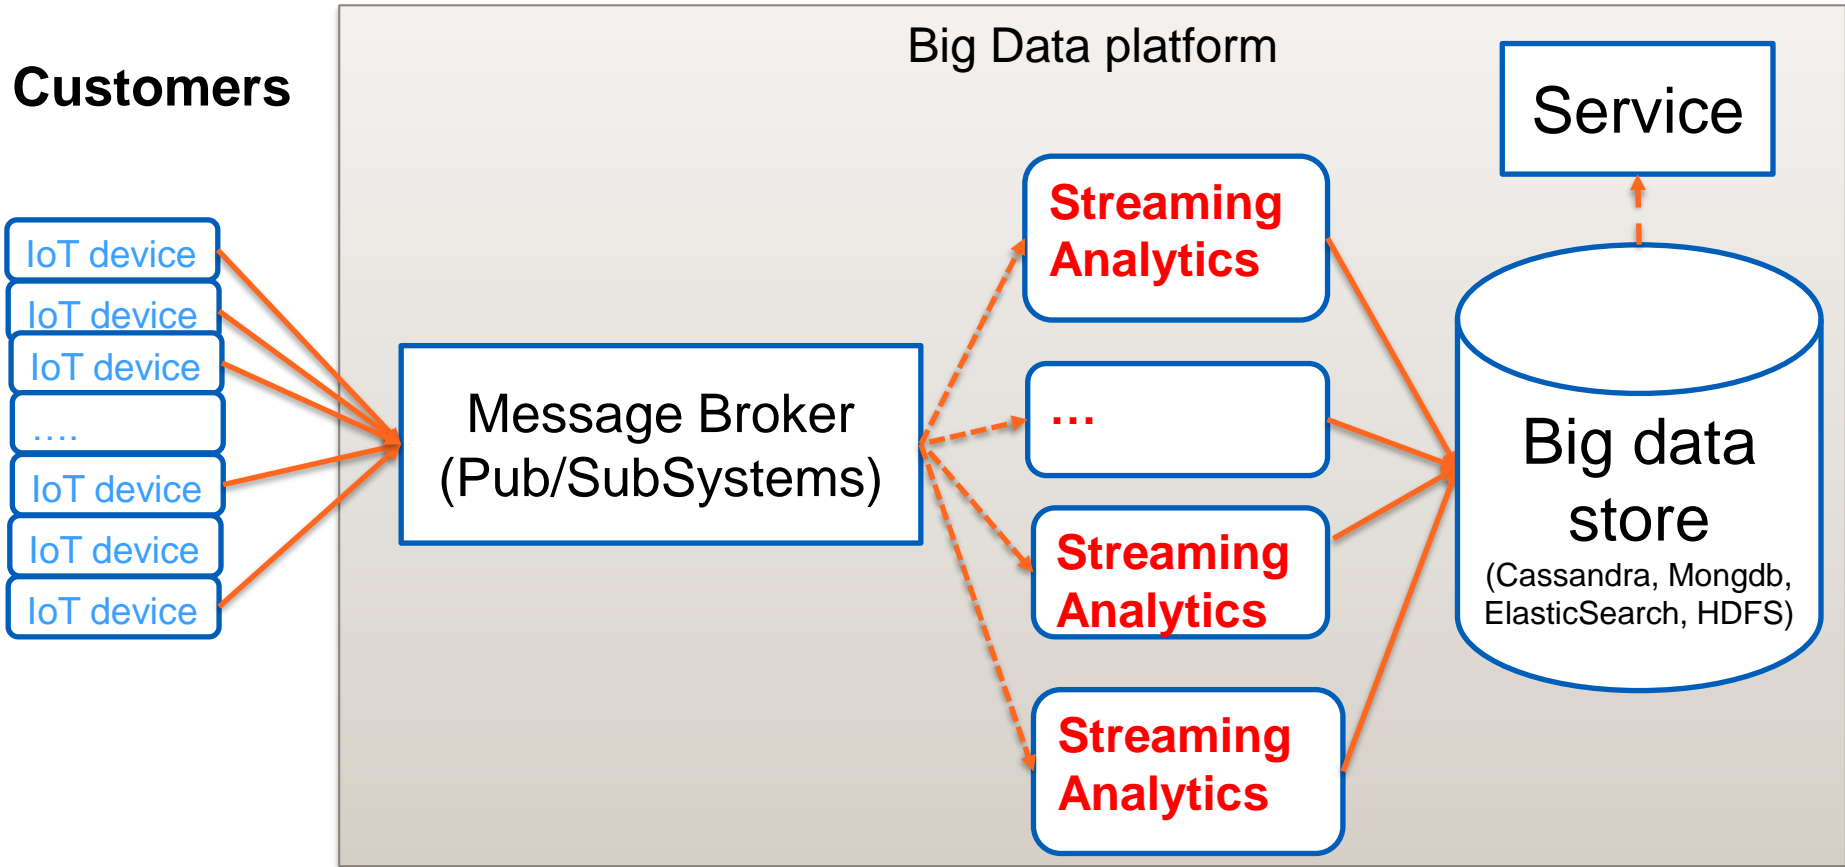


Recall: near-real time streaming data ingestion

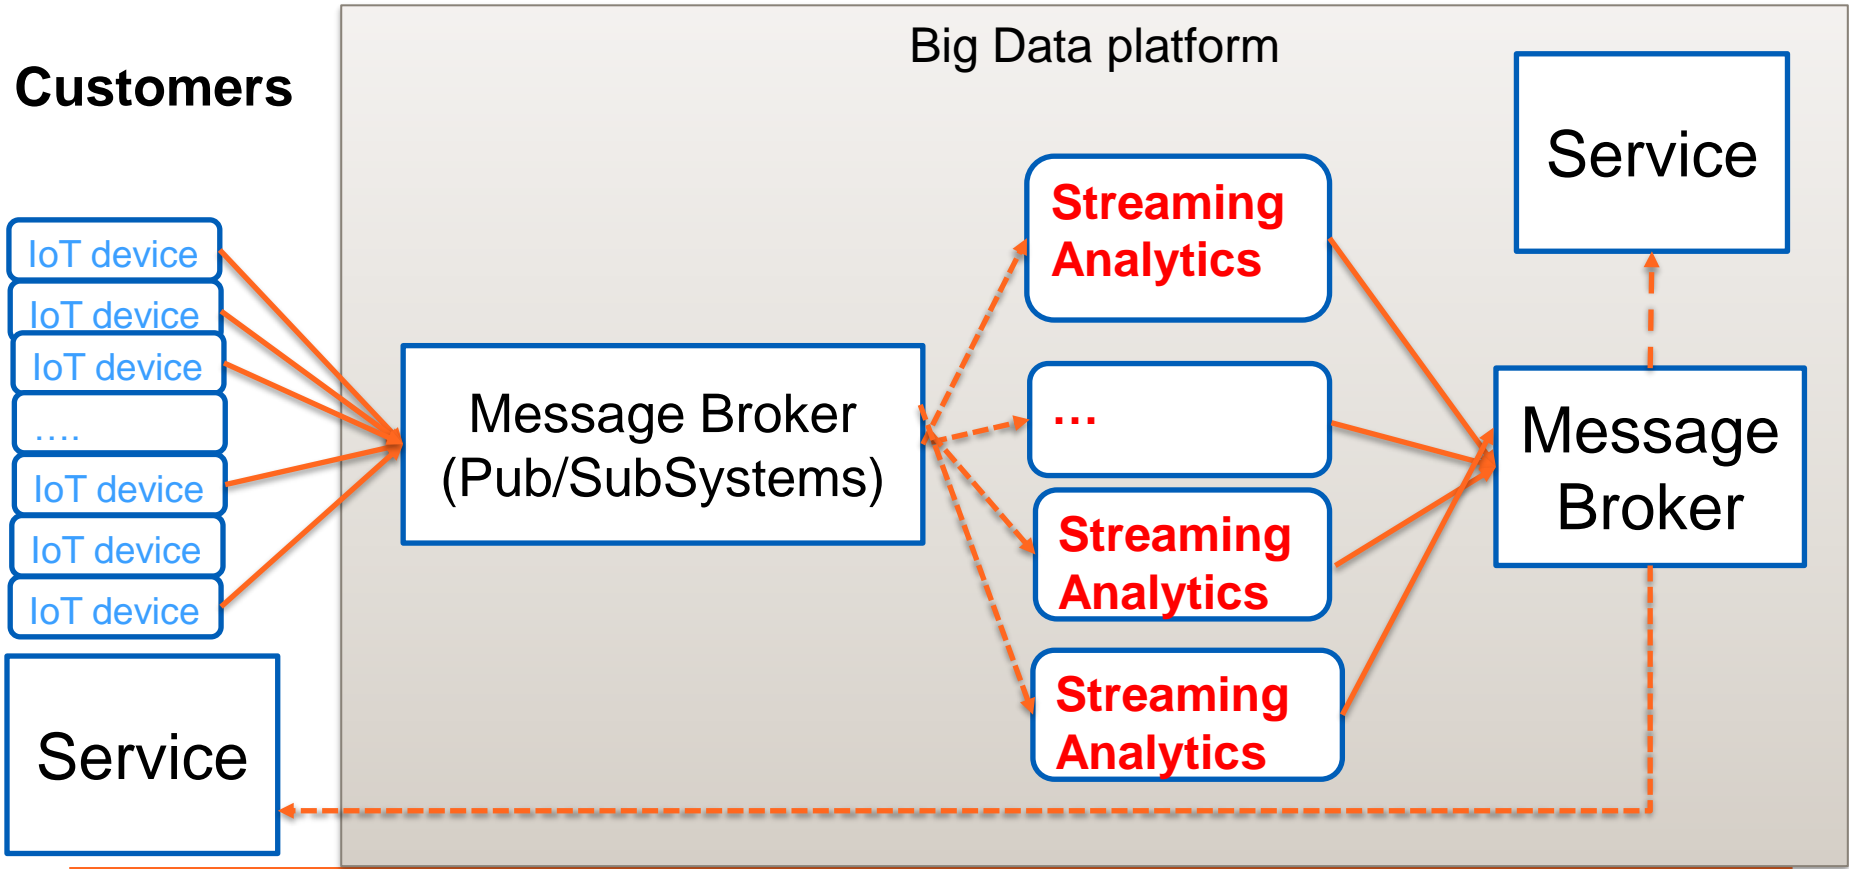


- **Mostly we ingest raw data without/little processing**
- **Data is unbounded from different places in different orders!**

Near realtime streaming data processing



Near realtime streaming data processing



Stream processing and big data platforms

- **Stream processing is part of big data platforms!**
 - A big data technology
 - Pre-processing, ingestion and high-level analytics
- **Stream processing services as big data platforms**
 - We can build a big data platform mainly based on stream processing services
 - Analytics on the fly as the first class
 - *Historical analytics results as the second class*
 - E.g., IoT analytics, e-commerce user activities, fraud detection

Example in the cloud – stream processing and big data platforms

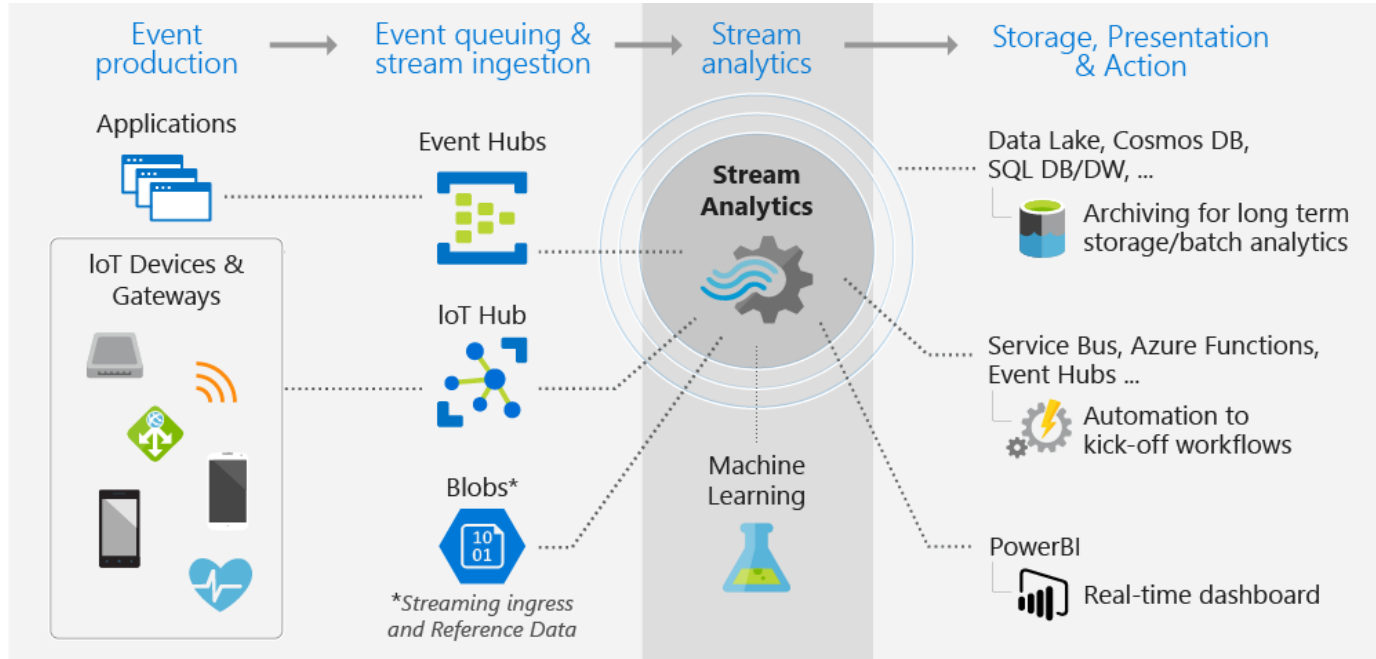


Figure source: <https://docs.microsoft.com/en-us/azure/stream-analytics/stream-analytics-introduction>

Example of frameworks for study

- **Apache Flink**
- **Apache Kafka**
- **Apache Spark**
- **Apache Storm**
- **Others:**
 - Azure Stream Analytics
 - MQTT + Node-RED

They are not toy examples! They are used in business systems and big cloud platforms

Streaming Data Processing - Concepts

Data stream programming

Data stream: a sequence/flow of data units


Data units are defined by applications: a data unit can be data described by a primitive data type or by a complex data type, a serializable object, etc.

Streaming data: produced by (near)realtime data sources as well as (big) static data sources → unbounded and bounded

- Examples of data streams
 - Continuous media (e.g., video for video analytics)
 - Discrete media (e.g., stock market events, twitter events, system monitoring events, comments, notifications)

Events/Records


In many applications: data is generated continuously and needs to be processed in near real-time



We focus on **unbounded discrete** events/records/messages

Stream processing

High level view

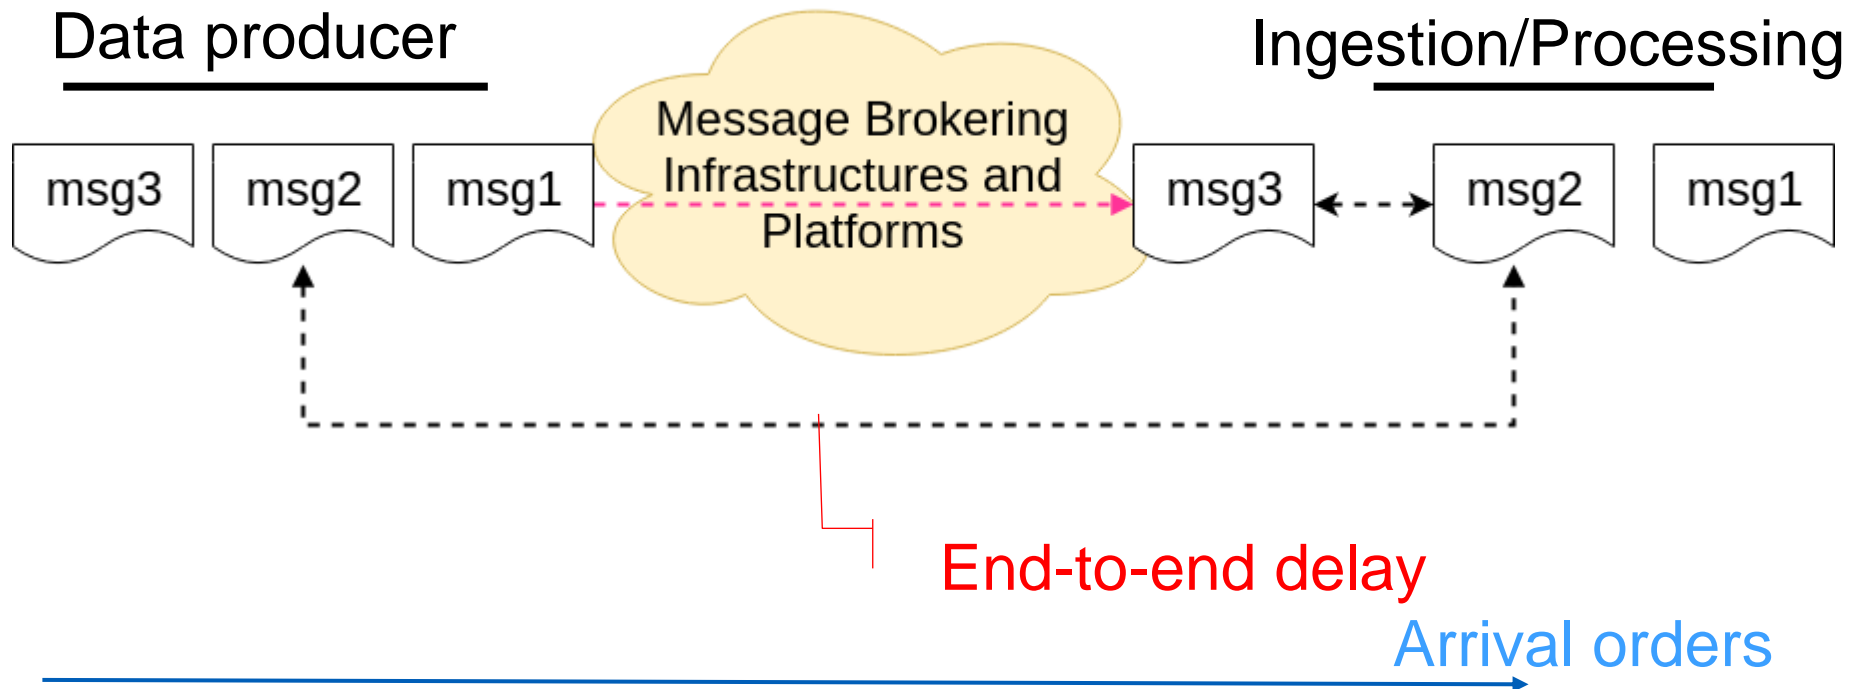


Multiple streams, a set of events

Some key issues

- **Late data, out of order data?**
- **Exactly once?**
- **Times associated with events**
- **Key-based data processing**

Key issues in streaming data: delay and out of order




Without event time, do we know the delay or out of order?

What is the consequence of delay/out of order for processing?

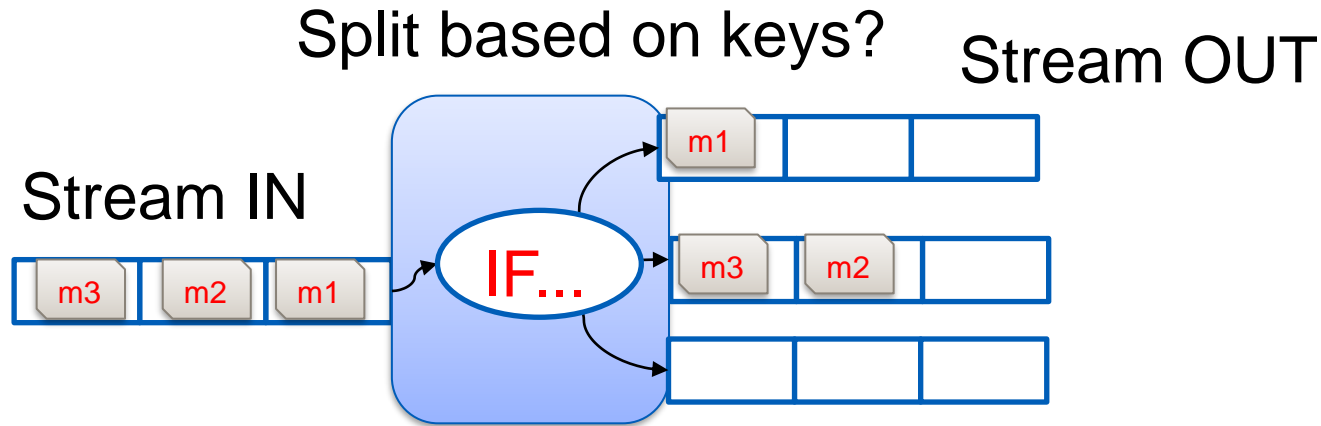
Key issues in streaming data: the notion of times

Times associated with data and processing




Which time is important for analytics (from business viewpoint)?

Partition stream data based on some keys for analytics



Windows of data



Windowing

- **Windows size: time or number of records (not popular)**
- **Tumbling window:**
 - identified by size, no gap between windows
- **Sliding window:**
 - identified by size and a sliding interval
- **Session Window:**
 - identified by “gap” between windows

Functions applied to Windows of data

If we

- specify a set of conditions for the window and events within the window

then we can

- Apply functions to events in the window that match these conditions

Example

Monitoring working hours of taxi drivers:

- Windows: 12 hours
- Partitioning data/Keyed streams: licenseID
- Function: determine working and break times and check with the law/regulation

Source: <https://www.infoworld.com/article/3293426/how-to-build-stateful-streaming-applications-with-apache-flink.html>

What if events come late into the windows?

Do we need to deal with late, out of order events?

Support lateness

- **Identify timestamp of events**
- **Identify watermark in streams**
 - A watermark is a timestamp
 - A watermark indicates that no events which are older than the watermark should be processed
 - Enable the delay of processing functions to wait for late events

Delivery guarantees

Exactly once? at least once? or at-most-once
End-to-end?

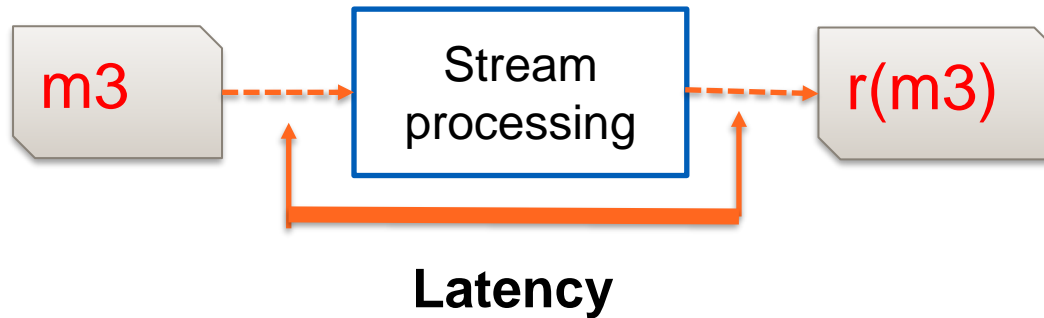



What if the stream processing fails and restarts

Message and processing guarantees

- **Message guarantees are the job of the broker**
- **Processing guarantees are the job of the stream processing frameworks**
- **They are highly connected if brokers and processing frameworks are tightly coupled (e.g. Kafka case)**

Performance metrics



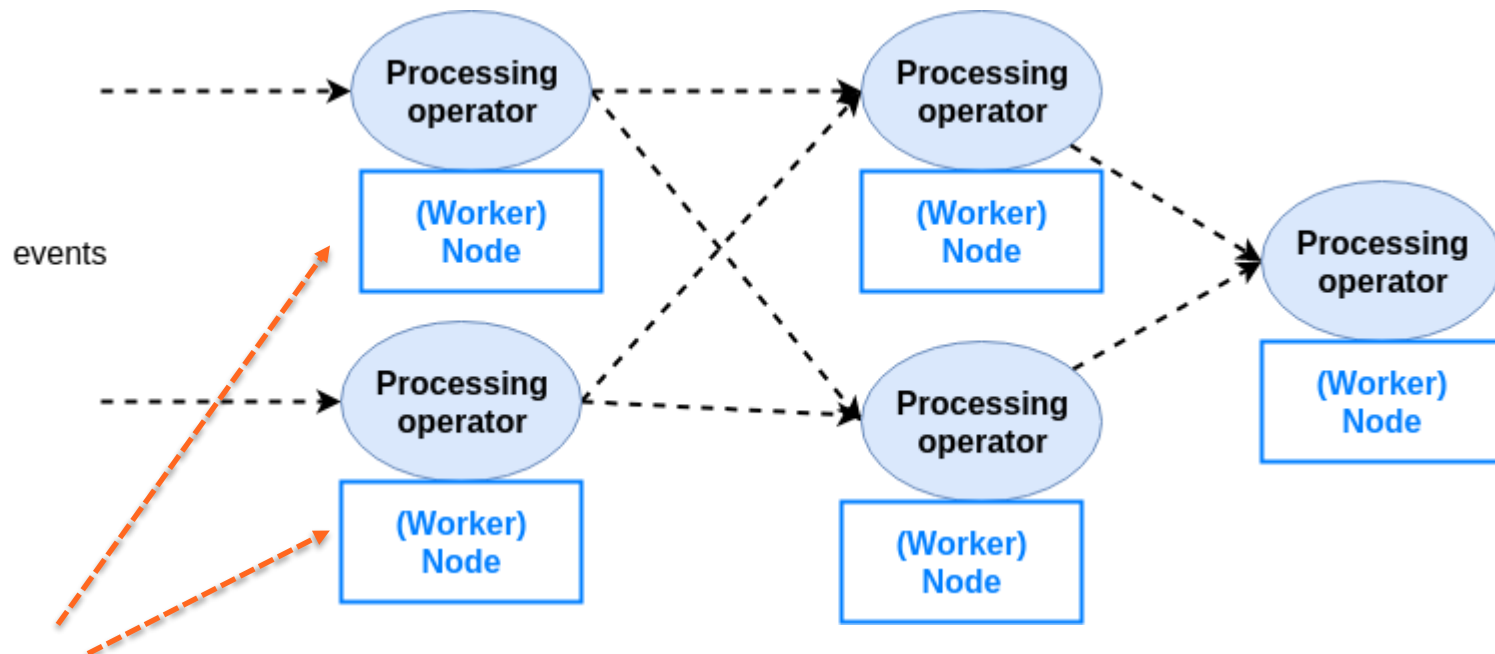
Latency and Throughput

- **Latency**
 - Subseconds! E.g., milliseconds
 - Max, min or percentile? → up to application requirements
- **Throughput**
 - How many events can be processed per second?
- **Goal: low latency and high throughput!**

Dataflow programming and stream processing

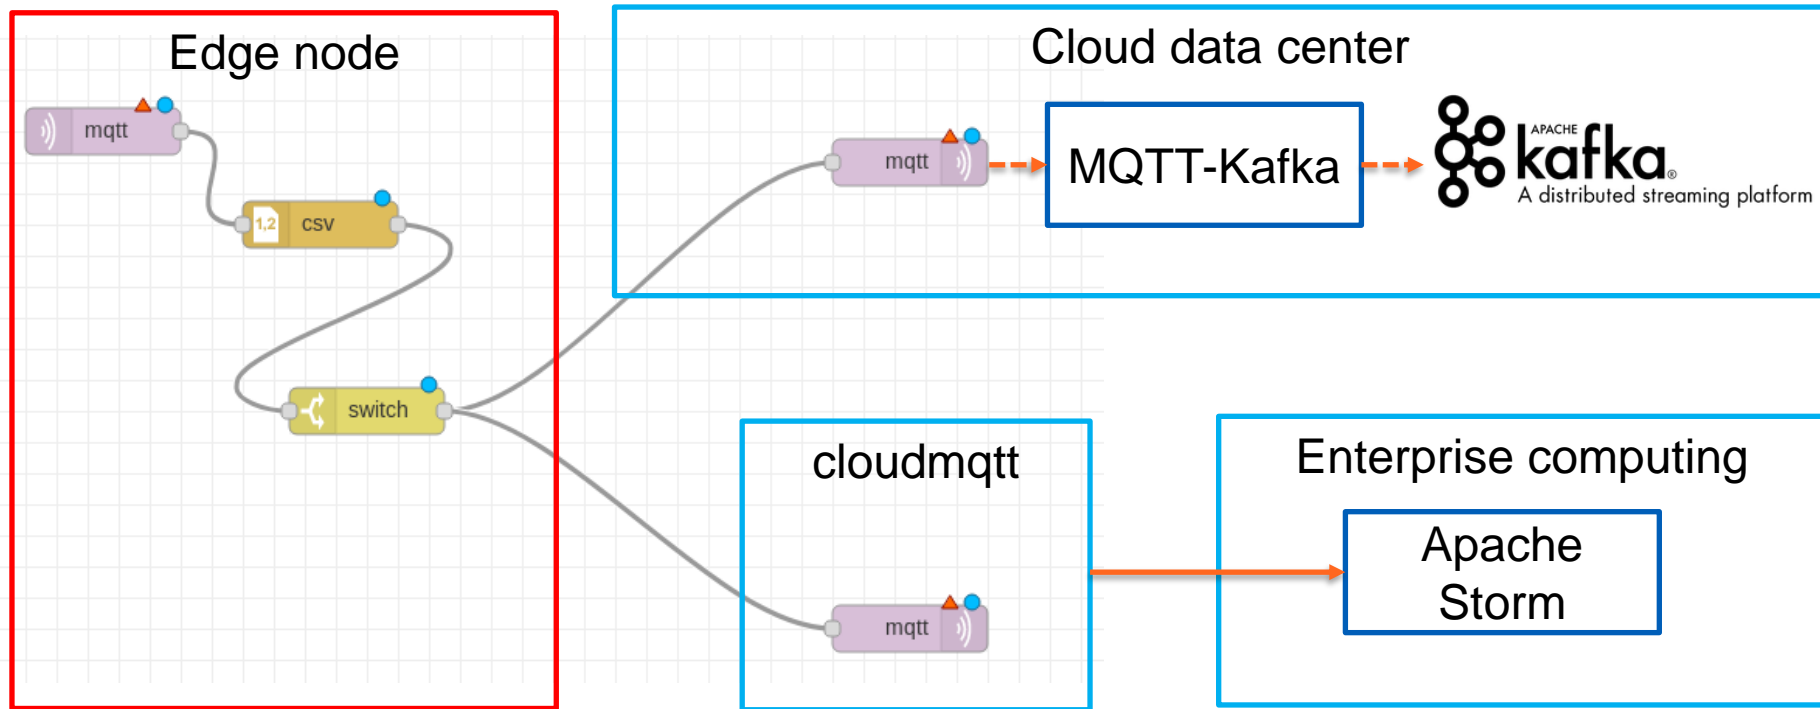
- **Data exchange between tasks**
 - Links in task graphs reflect data flows
- **Stream processing**
 - Centralized or distributed (in terms of computing resources)

Distributed processing topology in a cluster




Nodes of a cluster (VMs, containers, Kubernetes)

Distributed processing topology in cross distributed systems



Recall: syntax and semantic problems




- Like Ingestion: processing might not understand the semantics of the data
- Solution?

Event representation and streams

- **Event sources:** via message brokers, database, websocket, different IO adapters/connectors, etc.
- **Event representation & views**
 - POJO (Plain Old Java Object), CSV, Arvo format, etc.
 - SQL-alike tables
- **Event Stream**
 - Feed events from sources
- **Event Sink**
 - Feed results into sinks

Structure of streaming data processing programs



- **Data source operator:** represents a source of streams
- **Processing operators:** represents processing functions
- ***Native versus micro-batching***

Common concepts in existing frameworks

- **The way to connect data streams and obtain events**
 - Focusing very much on connector concepts and well-defined event structures (e.g., can be described JSON, POJO)
 - Assume that existing systems push the data
- **The way to specify “analytics”**
 - Statements and how they are glued together to process flows of events
 - High-level, easy to use
- **The engine to process analytics requests**
 - Centralized in the view of the user → so the user does not have to program complex distributed applications
 - Underlying it might be complex (for scalability purposes)
- **The way to push results to external components**

Common concepts in existing frameworks - programming level

- **How to streaming program?**
- **With programming languages**
 - Low level APIs
 - DSL
 - Java, Scala, Python (Spark, Flink, Kafka)
- **High-level data models**
 - KSQL
- **Flow description**
 - Node-RED

Common concepts in existing frameworks - key common concepts

- **Abstraction of streams**
- **Connector library for data sources/sinks**
 - Very important for application domains
- **Runtime elasticity**
 - Add/remove (new) operators (and underlying computing node)
- **Fault tolerance**

Why are the richness and diversity of connectors important?

Implementations

- **Apache Storm**
 - <https://storm.apache.org/>
- **Apache Spark**
 - <https://spark.apache.org/>
- **Apache Kafka Streams and KCQL**
 - strongly bounded to Kafka messaging
- **Apache Flink**
 - native, clustered, better data sources/sinks

Practical learning paths

- **Path 1: if you don't have a preference and need challenges**
 - Apache Flink Stream API (e.g., with RabbitMQ/Kafka connectors)
- **Path 2: many of you have worked with Kafka**
 - Kafka Streams DSL (everything can be done with Kafka)
- **Path 3: for those of you who are working with Spark (and Python is the main programming language)**
 - Apache Spark Streaming
- **Path 4: for those who deal with MQTT brokers**
 - Apache Storm (but also Kafka, ...): Spout and Bolt API or Stream API

Examples of Apache Flink

Apache Flink

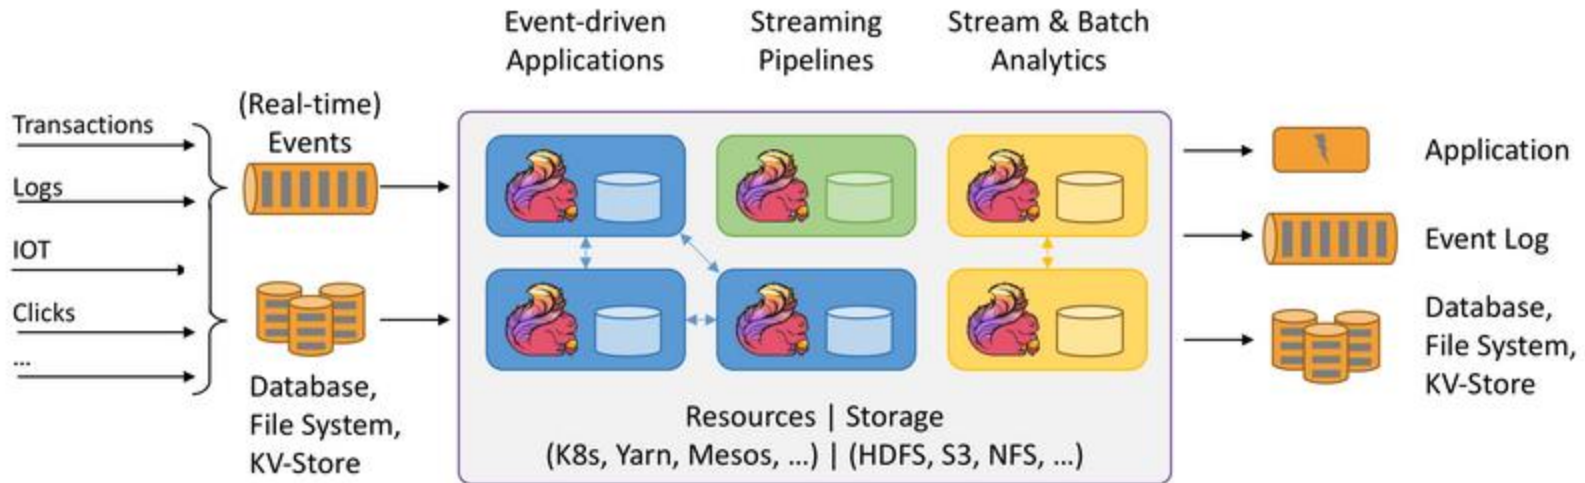



Figure source: <https://flink.apache.org/>

Apache Flink

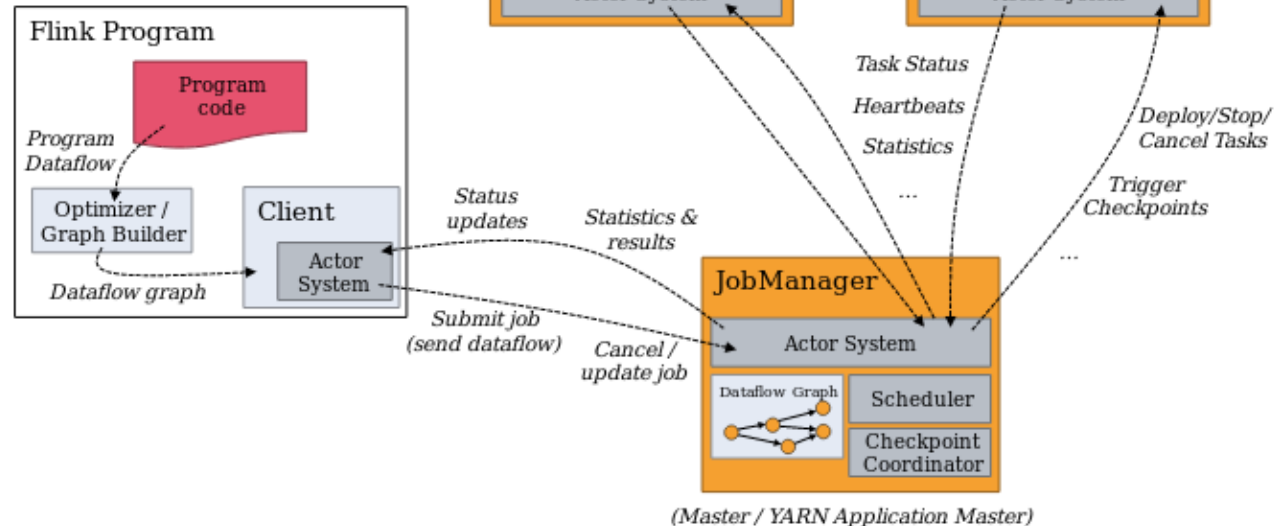


We focus only on
DataStream API for
understanding
studied concepts

Source: <https://ci.apache.org/projects/flink/flink-docs-release-1.9/internals/components.html>


Flink runtime view

- Parallelism
- Checkpointing
- Monitoring



Source: <https://ci.apache.org/projects/flink/flink-docs-release-1.9/concepts/programming-model.html>

Main elements in Flink applications



- Rich set of sources and sinks via many connectors

Connectors

- **Major systems in big data**
- **We have used many of them in our study**
 - Apache Kafka
 - Apache Cassandra
 - Elasticsearch (sink)
 - Hadoop FileSystem
 - RabbitMQ
 - Apache NiFi

Main

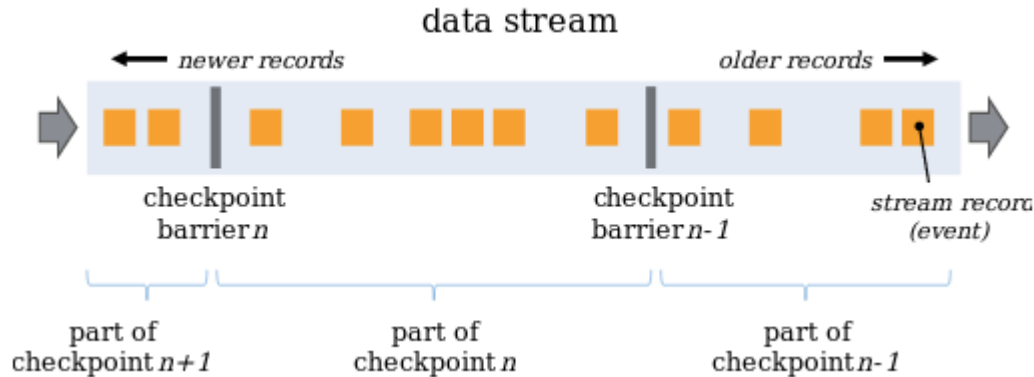
- **Setting environments**
- **Handling inputs and outputs via data streams**
- **Key functions for processing data**
- **Stream processing flows**



Bounded and unbounded streams

Fault tolerance

Can be exactly once



Principles: checkpointing, restarts operators from the latest successful checkpoints

Need support from data stream sources/sinks w.r.t. (end-to-end) exactly once message receiving and result delivery

Stream processing flows

Split streaming data into different windows with a key for analytics purposes

Keyed Windows

```
stream
  .keyBy(...)          <- keyed versus non-keyed windows
  .window(...)         <- required: "assigner"
  [.trigger(...)]      <- optional: "trigger" (else default trigger)
  [.evictor(...)]      <- optional: "evictor" (else no evictor)
  [.allowedLateness(...)] <- optional: "lateness" (else zero)
  [.sideOutputLateData(...)] <- optional: "output tag" (else no side output for late data)
  .reduce/aggregate/fold/apply() <- required: "function"
  [.getSideOutput(...)] <- optional: "output tag"
```

Source: <https://ci.apache.org/projects/flink/flink-docs-release-1.9/dev/stream/operators/windows.html>

Stream processing flows

Handling streaming data without a key for analytics purposes


Non-Keyed Windows

```
stream
  .windowAll(...)      <- required: "assigner"
  [.trigger(...)]      <- optional: "trigger" (else default trigger)
  [.evictor(...)]      <- optional: "evictor" (else no evictor)
  [.allowedLateness(...)] <- optional: "lateness" (else zero)
  [.sideOutputLateData(...)] <- optional: "output tag" (else no side output for late data)
  .reduce/aggregate/fold/apply() <- required: "function"
  [.getSideOutput(...)] <- optional: "output tag"
```

Source: <https://ci.apache.org/projects/flink/flink-docs-release-1.9/dev/stream/operators/windows.html>

Windows and Times

Windows



Times

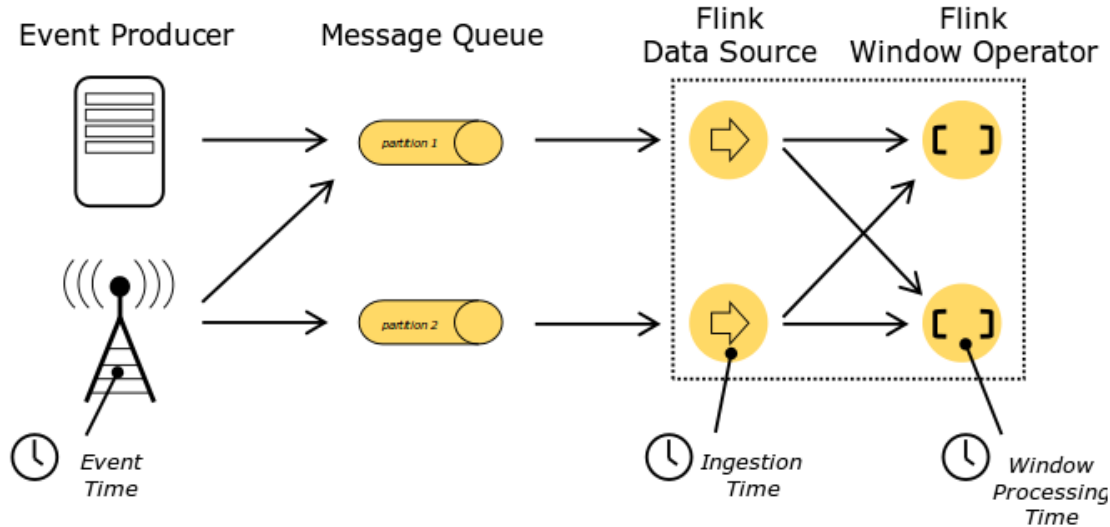



Figure source: <https://ci.apache.org/projects/flink/flink-docs-release-1.9/concepts/programming-model.html>


Batch/Tumbling Windows



Use cases:
Period computation
(e.g. stock,
temperature, IoT
data)

Source: <https://ci.apache.org/projects/flink/flink-docs-release-1.9/dev/stream/operators/windows.html>

Sliding windows




Use cases:
Moving average

Source: <https://ci.apache.org/projects/flink/flink-docs-release-1.4/dev/stream/operators/windows.html>

Session Windows

Use cases:
Web/user activities
clicks



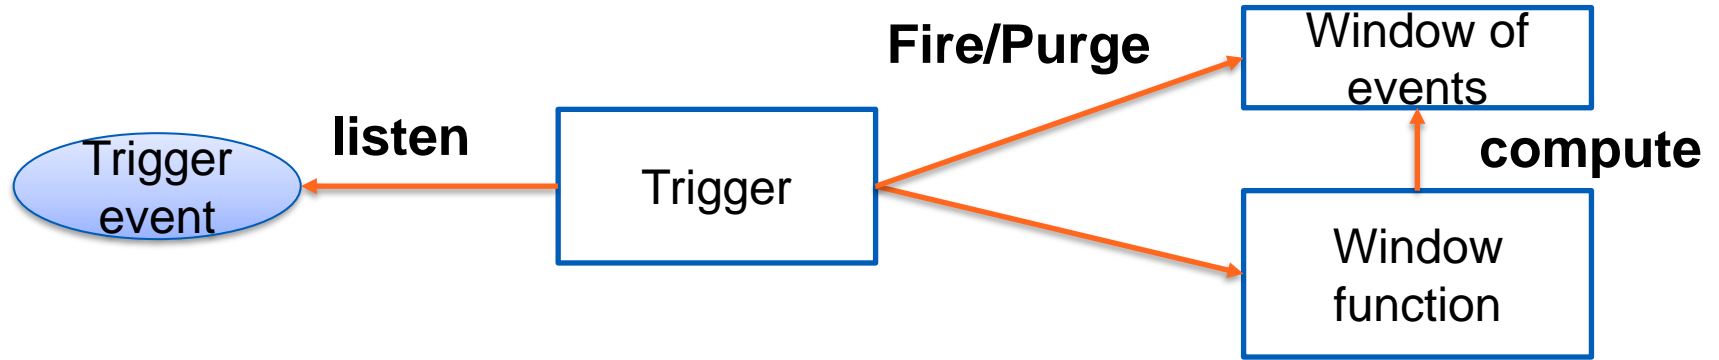
Source: <https://ci.apache.org/projects/flink/flink-docs-release-1.9/dev/stream/operators/windows.html>

Window Functions

- **Reduce Function**
 - Reduce two inputs
- **Aggregate Function**
 - Add an input into an accumulator
- **Fold Function**
 - Combine input with an output
- **ProcessWindow Function**
 - Get all elements of the windows and many other information so that you can do many tasks

Triggers & Evictor

- **Trigger:** determine if a window is ready for window functions



Evictor: actions **after** the trigger fires and **before** **and/or after** the windows function is called

Example with Base Transceiver Station

Data in our git

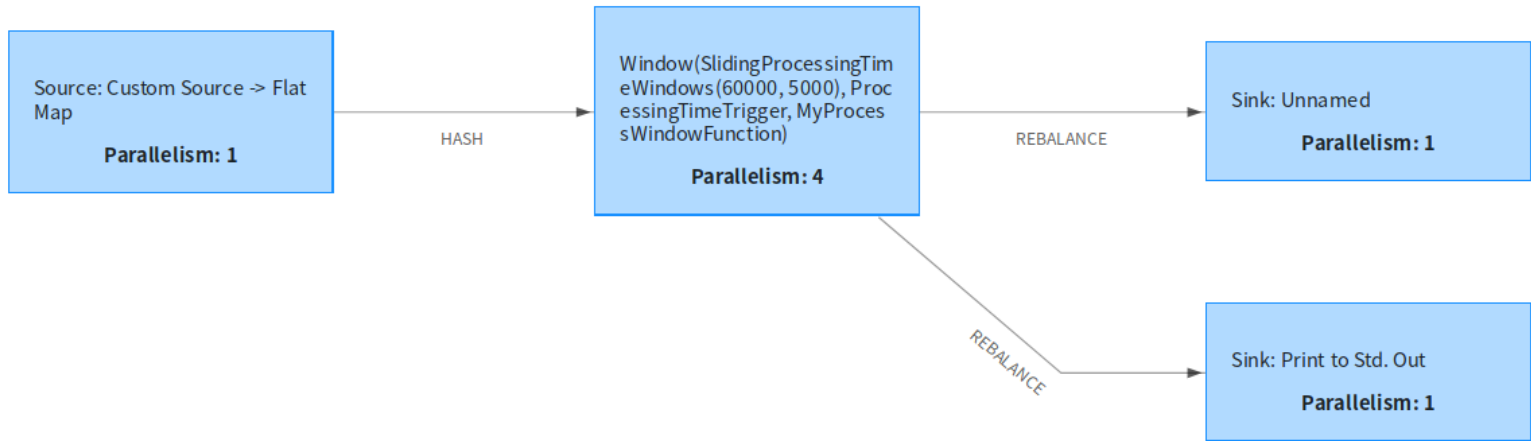
```
station_id,datapoint_id,alarm_id,event_time,value,valueThreshold,isActive,storedtime
1161115016,121,308,2017-02-18 18:28:05 UTC,240,240,false,
1161114050,143,312,2017-02-18 18:56:20 UTC,28.5,28,true,
1161115040,141,312,2017-02-18 18:22:03 UTC,56.5,56,true,
1161114008,121,308,2017-02-18 18:34:09 UTC,240,240,false,
1161115040,141,312,2017-02-18 18:20:49 UTC,56,56,false,
```

Simple example

See the code in our git:

<https://version.aalto.fi/gitlab/bigdataplatforms/cs-e4640-2019/blob/master/tutorials/streamingwithflink/simplebts/src/main/java/fi/aalto/cs/cse4640/SimpleAlarmAnalysis.java>

Simple example



Monitoring

Apache Flink Dashboard

- Overview
- Jobs
 - Running Jobs
 - Completed Jobs
- Task Managers
- Job Manager
- Submit New Job

Simple CS-E4640 BTS Flink Application

RUNNING 2

ID: 81efb959d02448b7c44328ad75a824af
Start Time: 2019-11-04 14:00:14
Duration: 55s

[Cancel Job](#)

Overview

Exceptions

TimeLine

Checkpoints

Configuration

Source: Custom Source -> Flat Map

Parallelism: 1

HASH

Window(SlidingProcessingTimeWindows(60000, 5000), ProcessingTimeTrigger, MyProcessWindowFunction) -> (Sink: Unnamed, Sink: Print to Std. Out)

Parallelism: 1

Detail

SubTasks

TaskManagers

Watermarks

Accumulators

BackPressure

Metrics

Window(SlidingProcessingTimeWindows(60000, 5000), ProcessingTimeTrigger, MyProcessWindowFunction) -> (Sink: Unnamed, Sink: Print to Std. Out)

Status: RUNNING

Task: 1

Parallelism: 1

Records Sent: 0

Start Time: 2019-11-04 14:00:14

Bytes Received: 1.64 KB

End Time: -

Records Received: 29

Duration: 55s

Bytes Sent: 0 B

Name	Status	Bytes Received	Records Received	Bytes Sent	Records Sent	Parallelism	Start Time	Duration	End Time	Tasks
Window(SlidingProcessingTimeWindows(60000, 5000), ProcessingTimeTrigger, MyProcessWindowFunction) -> (Sink: Unnamed, Sink: Print to Std. Out)	RUNNING	1.64 KB	29	0 B	0	1	2019-11-04 14:00:14	55s	-	1
Source: Custom Source -> Flat Map	RUNNING	0 B	0	1.61 KB	29	1	2019-11-04 14:00:14	55s	-	1

One of the successful project from Europe: originally from TU Berlin

Alibaba cloud:

“Flink can process over **472 million transactions per second during business peaks, which is truly astronomical”**

Source: https://www.alibabacloud.com/blog/why-did-alibaba-choose-apache-flink-anyway_595190

“Amazon Kinesis Data Analytics includes open source libraries based on Apache Flink” (From <https://aws.amazon.com/kinesis/data-analytics/>)

Summary

- **Facts:**

- Stream processing is important in big data platforms
- There are many frameworks, but they have quite common concepts
 - *You should focus on the concept and pickup a suitable one for your work*
- Core concepts can be implemented at different levels
 - *Programming languages, DSL, SQL-style*

- **Thoughts:**

- Think about combining stream with batch analytics to produce a comprehensive platform!
- There are many real world-big data needs stream processing

Summary

- **Focus:**

- Practical programming with one of the stacks:
 - *Apache Flink Stream API (with different connectors)*
 - *Apache Spark*
 - *Kafka Streams*
 - *Apache Storm*
- Check the common concepts in other tools/systems

- **Action:**

- Work on use cases where you can use stream analytics (as a user/developer) → there are many interesting analytics
- Provision services for stream processing (as a platform)

Thanks!

Hong-Linh Truong
Department of Computer Science

rdsea.github.io