



Aalto University
School of Science

Architecting Big Data Platforms

Hong-Linh Truong

Department of Computer Science

linh.truong@aalto.fi, <https://rdsea.github.io>

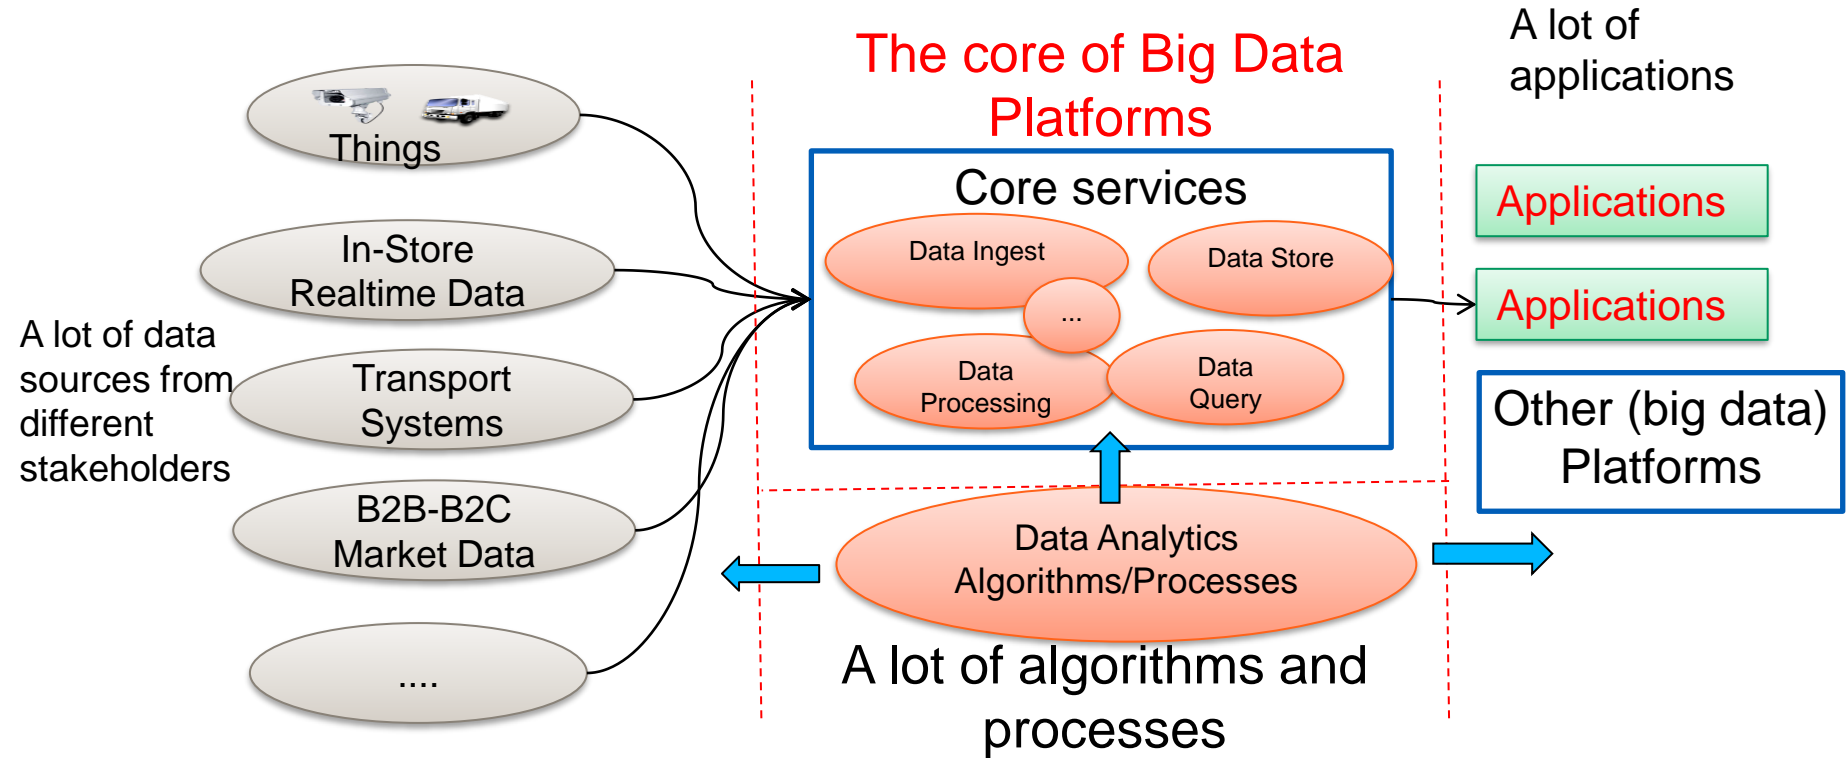
Learning objectives

- **Understand key issues in designing a big data platform**
- **Study different big data architectures**
- **Learn key architecture design issues**
 - interaction, partitioning, elasticity, API
- **Understand big data platform technologies**

Your big data platform story - an evolving scenario

“Your team has to build a big data platform for **X types of data**. Data will be generated/collected from **N sources**. We expect to have **10+ GBs/day of data to be ingested** into our platform. We will have to serve **K thousands of requests** for different types of analytics – to be determined. Our response time should be **in t milliseconds**. Our services should not be ...”

Recall - Big data platforms: system of systems view



You may have several questions?

- Do we have to support multiple types of data?
- How do data pipelines and data load look like?
- How to enable different data processing models?
- Which runtime parameters must be monitored? Which service level metrics must be guaranteed?
- To where we should distribute/deploy our components?
- Which part of the platform we must manage by ourselves and which part will be fully managed by other providers?
- How to design elastic big data infrastructures?
- Etc.

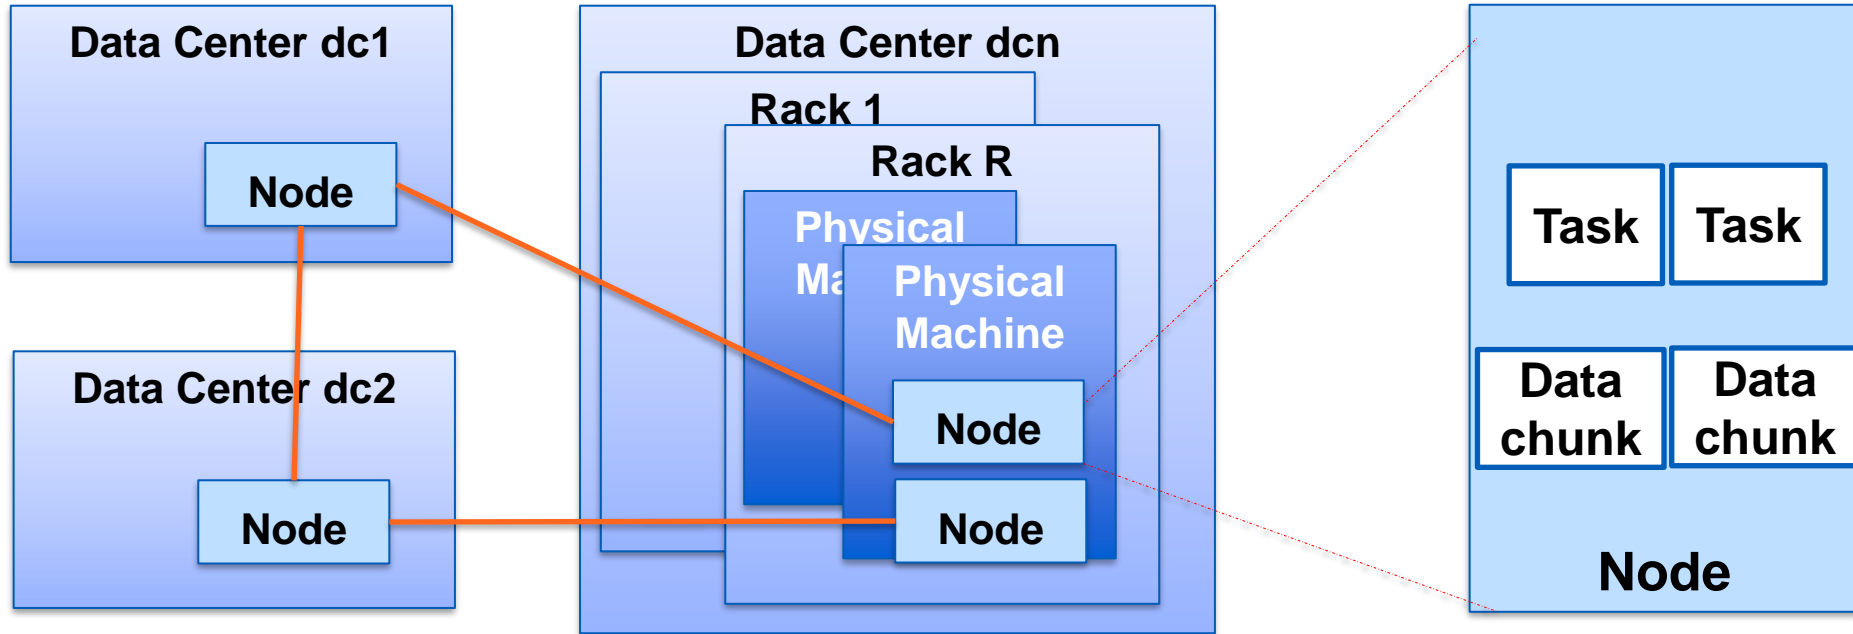
Your Big Data Platform story starts with Big Data Platform architectures!

**To architect the platform
centered around data!**

Understanding the underlying computing infrastructures

- **Computing resources and services**
 - many machines, virtual infrastructures, different types of services
- **Distributed infrastructures from different administrative domains**
 - in multiple data centers, locations and countries
 - with different security and network policies
- **Diverse service level objectives (SLO) and service level agreements (SLAs)**
 - performance, service failure, cost, privacy/security ...

Understanding the underlying computing infrastructures




Remember: large-scale distributed infrastructures!

Data-centric development & operations

- **Data ingestion and data movement**
 - from various data sources we move data into the platform
- **Data storing and management**
 - ingested data will be stored and managed using different types of storages and databases
- **Data analysis and (Machine) learning**
 - data within platforms will be processed, analyzed and learned to improve data, find insights and to create models
- **Reporting and visualization**
 - patterns/insights in data will be interpreted and presented for decision-making, reporting and creating stories

Big Data Pipelines

Multiple big data pipelines can be constructed atop a big data platform (and across distributed infrastructures)



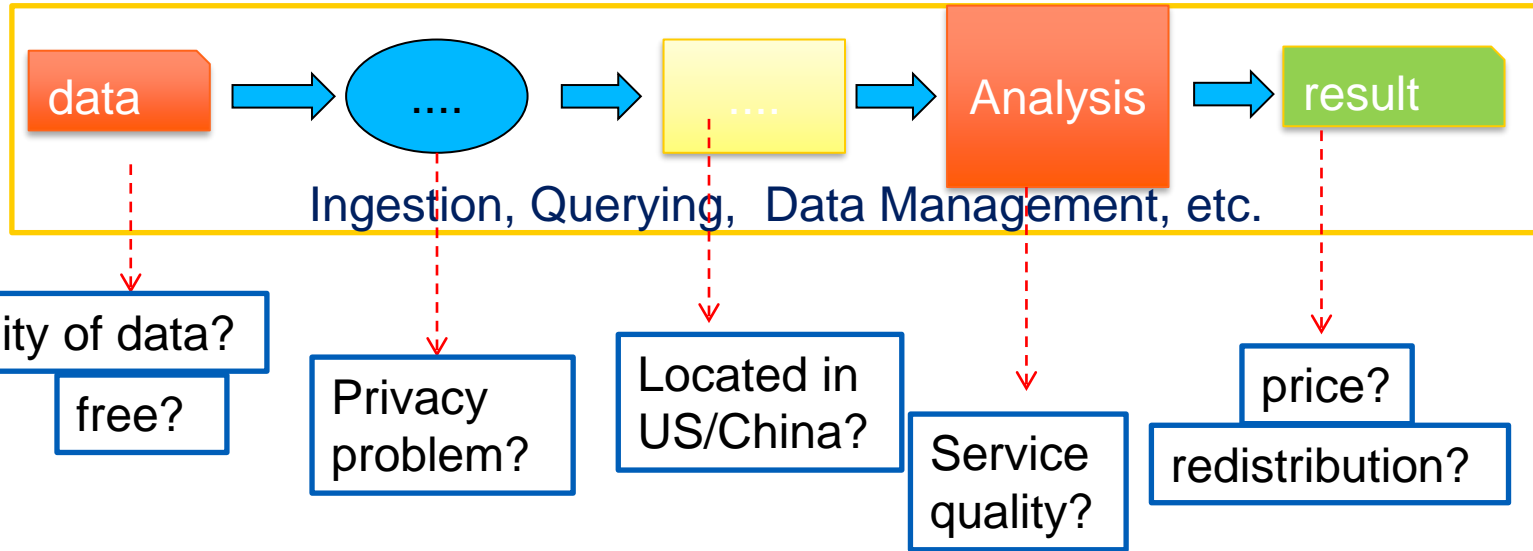
Handling multiple types of data?

- **First important aspect: you don't have to support multiple types of data**
 - but are you sure that you will not have this in the future?
- **Multiple types of data**
 - any linked models among them?
- **Any elastic solution that ensures minimum changes to support **generalization and extensibility****
 - e.g., multi-model databases, microservices of multiple of databases or data lake

Data concerns and SLAs

- **Ingesting data**
 - mapping and transforming data
 - ingestion of data under V^*
 - data validation/quality control during ingestion
- **Storing data**
 - data sharding and consistency, data backup, retention, etc.
- **SLA multitenancy versus single tenancy**
 - security, privacy, performance, reliability and maintenance?


Data concerns: data validation and quality of analytics



- Ethical consequences?
- Regulation-compliant platforms: e.g., GDPR

Fast versus slow processing

big data but not near real-time, e.g., take customer transaction files from companies and move to data centers for analytics



fast, small IoT data in near real-time flows, e.g. position of cars



But

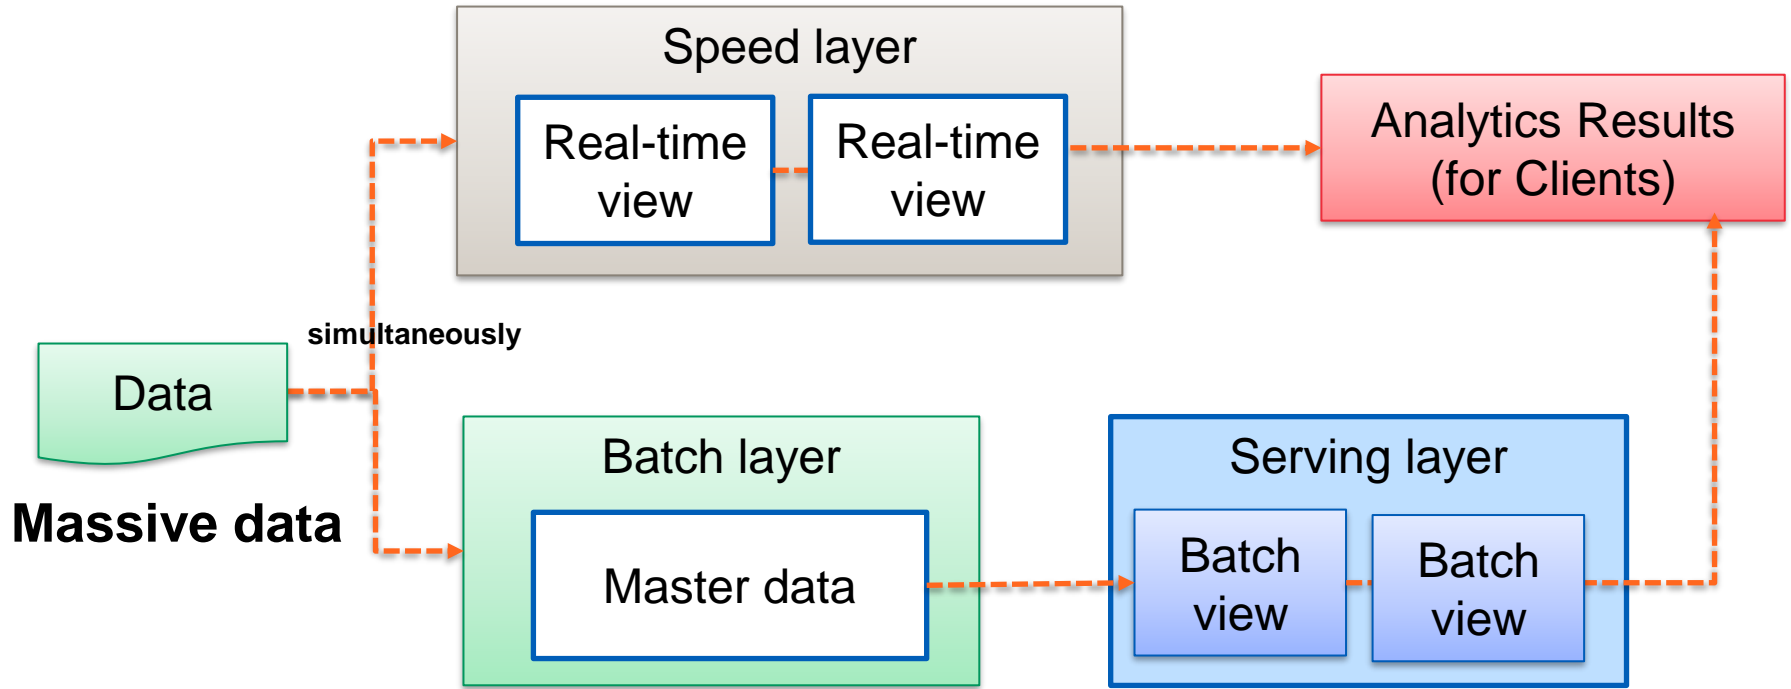
if you have mixed types of data

Or

**if you have big data you want to do analytics
with different quality of analytics (cost,
performance, quality of data)?**

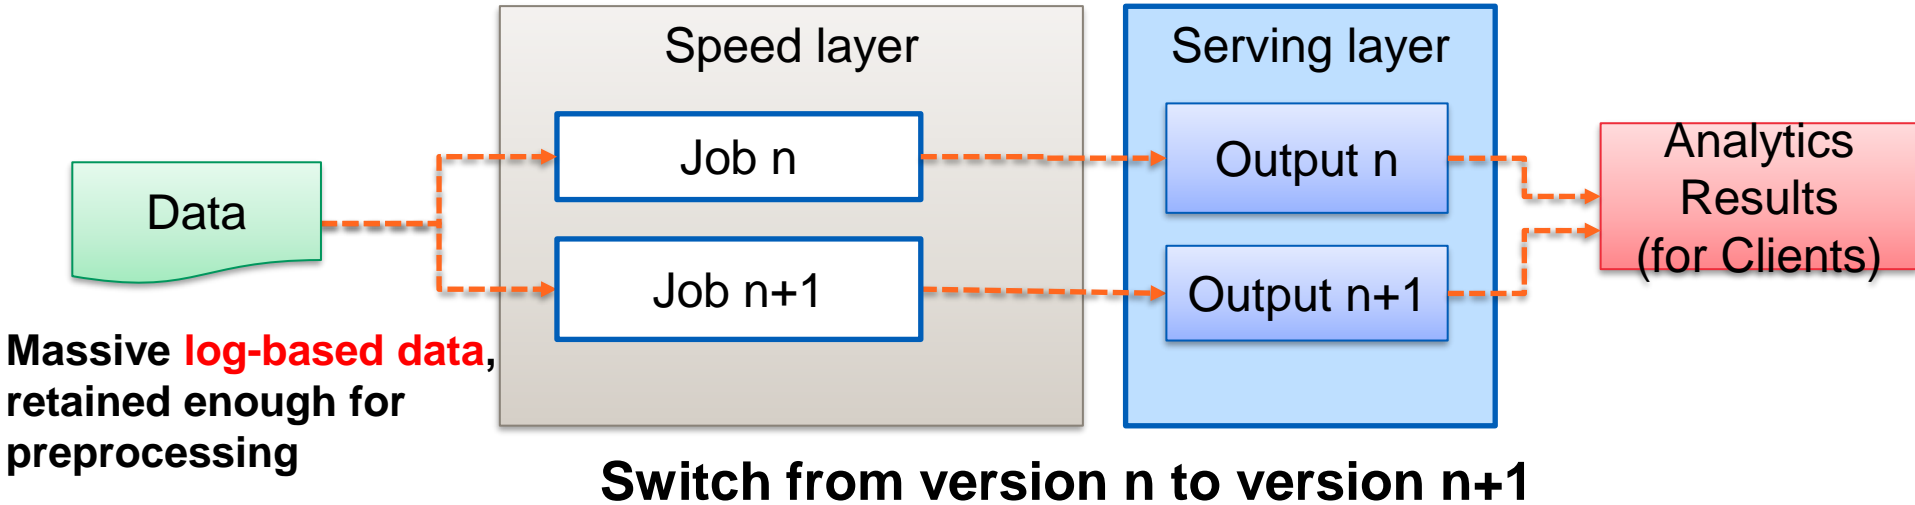
Then ?

Lambda architectural style



Check: <http://lambda-architecture.net/>

Kappa architectural style



Check: <https://milinda.pathirage.org/kappa-architecture.com/>

**The set of big data tools/frameworks
(and configurations) used is dependent
on the big data architecture**

be aware of your personal techradar!

Quick check

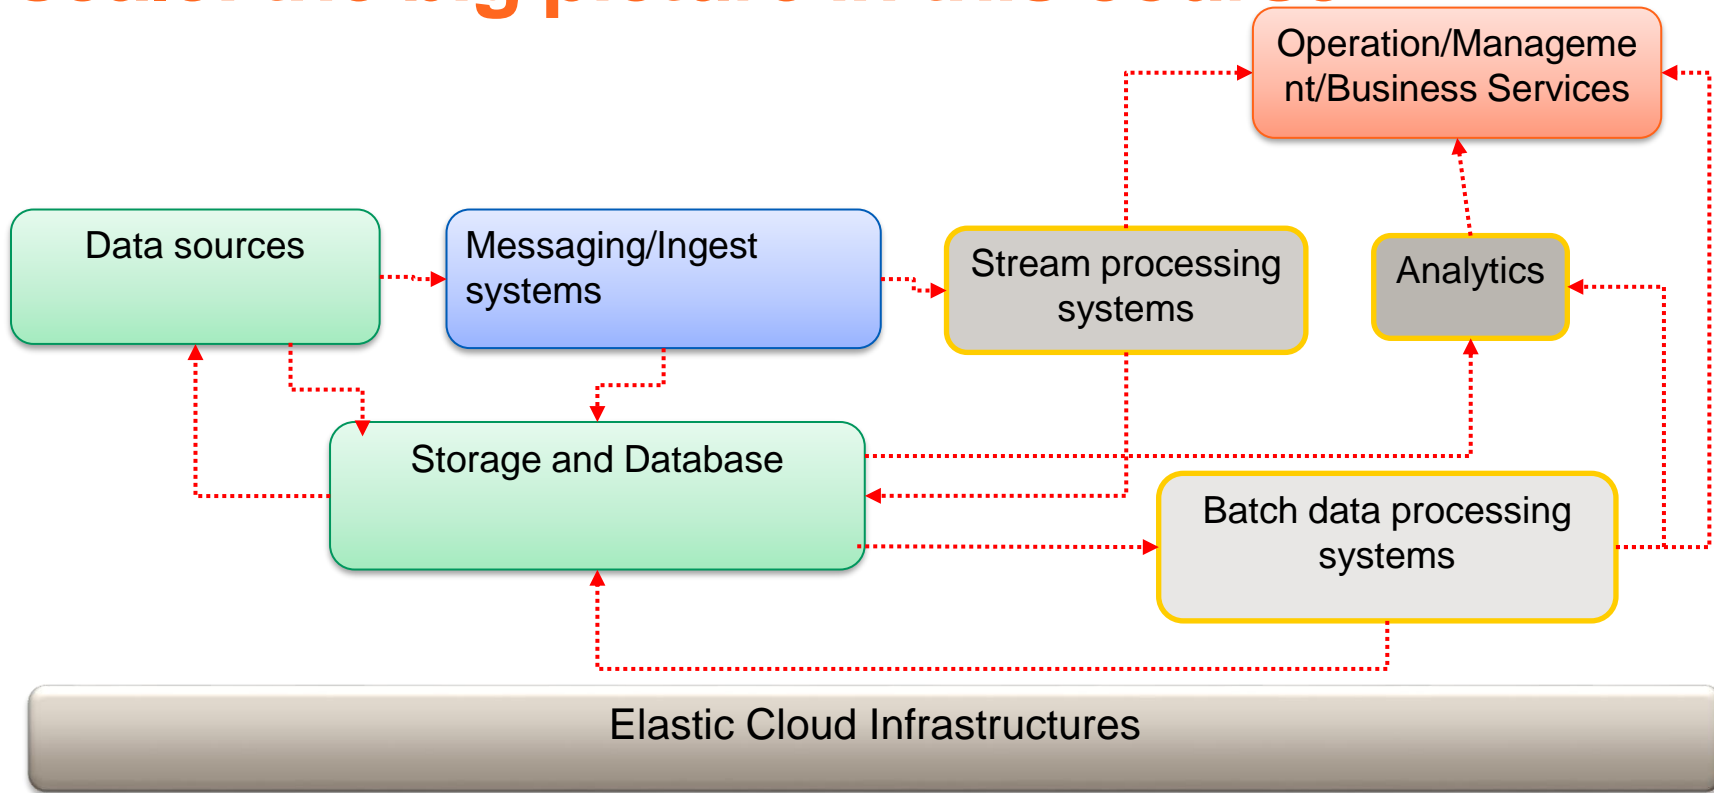
“A big data platform monitors network usage of devices from million+ customers. We have different levels: **Sensor/Customer, Node (concentrator of multiple customers), Agent (concentrator of multiple Nodes) and the whole network. In a region, the real operator can generate 1.4 billion records per day ~ 72GB per day”**

Quickcheck

**First: breakout room discussion
and then vote your choice**


<https://presemo.aalto.fi/bdp>

Basic building blocks for big data at large-scale: the big picture in this course



How to architect big data platforms and pipelines as reactive systems?

Reactive systems



Source: <https://www.reactivemanifesto.org/>

Why? For dealing with V*

- **Responsive:** quality of services
- **Resilient:** deal within failures
- **Elastic:** deal with different workload and quality of analytics
- **Message-driven:** allow loosely coupling, isolation, asynchronous among many components

Designs must address various aspects

- **Responsive:**
 - distributed computing, multi layer optimization
- **Resilient:**
 - replication, containment, isolation
- **Elastic:**
 - sharding, replication, load balancing, scale up/out
- **Message-driven:**
 - loosely coupling of services with messages, non-blocking protocols, location-independent

Partitioning: splitting functionality & data

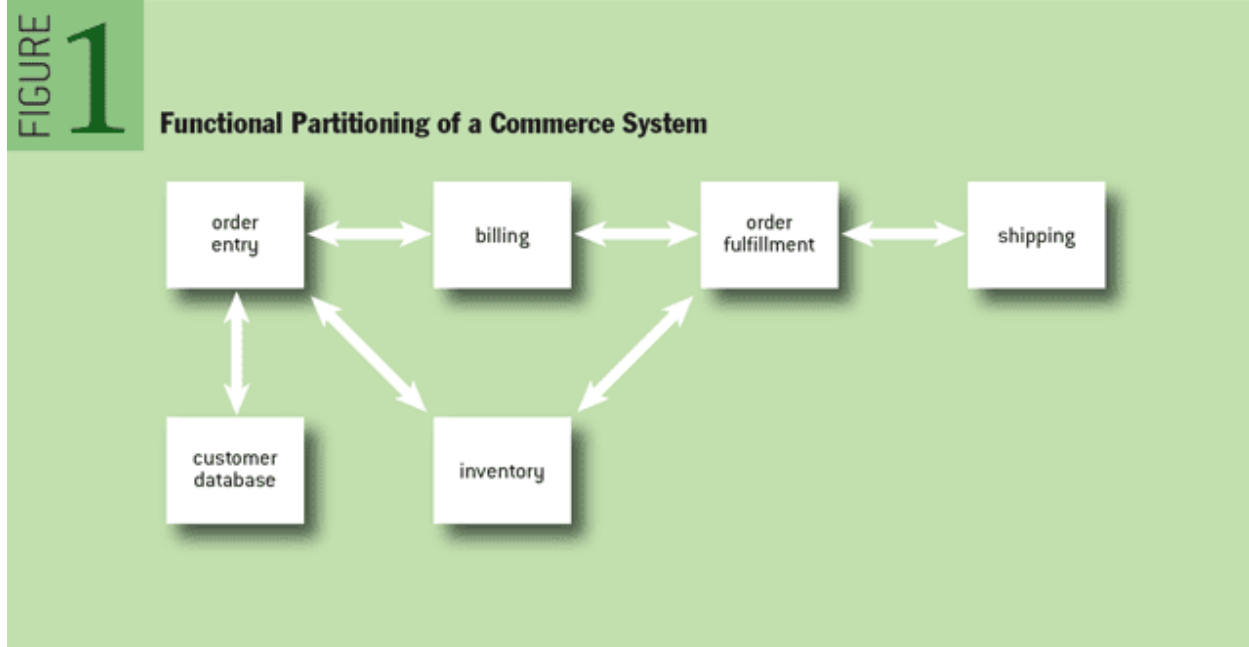
- **breakdown the complexity**
- **easy to implement, replace and compose**
- **deal with performance, scalability, security, etc.**
- **support teams in DevOps**
- **cope with technology changes**

Example of functional and data partitioning

Service-oriented
components

Microservices

Serverless functions




Figures source: <http://queue.acm.org/detail.cfm?id=1971597>

Example of functional and data partitioning

Data sharding

Multi data spaces

Multi data services



Figures source: <http://queue.acm.org/detail.cfm?id=1971597>



Aalto University
School of Science


**Distributed systems of components
are used to manage, ingest data and
process data**

Interaction: protocols & interfaces


- **Large number of communication protocols and interfaces**
- **Interaction styles, protocols and interfaces**
 - REST, gRPC, Message Passing, Stream-oriented Communication
 - your own protocols
- **Other criteria**
 - architectural styles: microservices/serverless
 - scalability, Elasticity, Performance, Monitoring, Logging, etc.

Interaction: Complex interactions

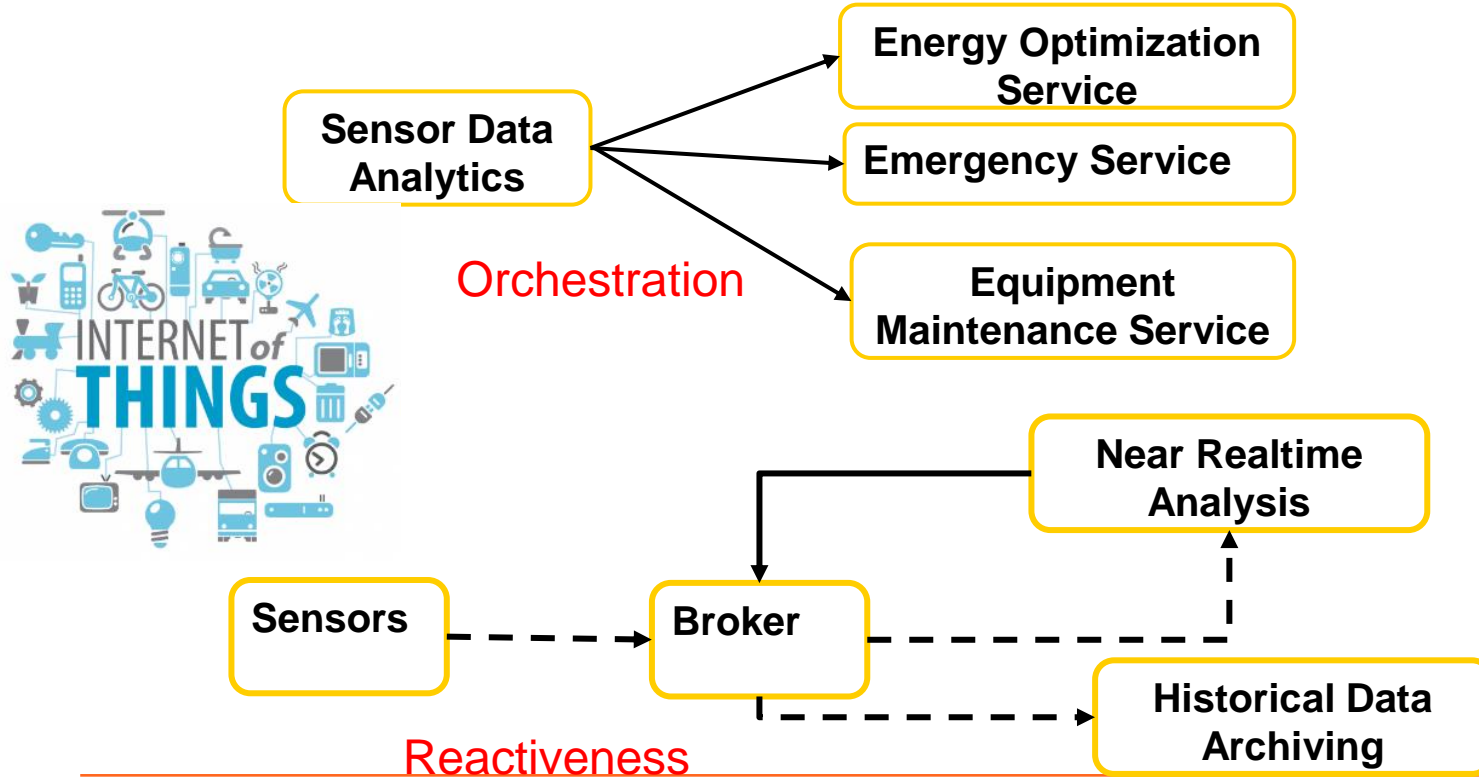
- One-to-many, many-to-one, many-to-many
- Synchronous/asynchronous calls
- Public/Subscribe, Message-oriented Middleware
- Internal data exchange versus open/external exchange



Amazon S3/MongoDB



Coordination: Orchestration and Reactiveness



Distribution: Edge or Data Centers?

**Big data & components
components can be
distributed in different
places!**

**Global deployment or
not?**

**Move analytics/work or
move data?**

Use Case 3: Video Analytics




Figure 4: Example of video analytics

Figure source: https://portal.etsi.org/portals/0/tbpages/mec/docs/mobile-edge_computing_-_introductory_technical_white_paper_v1%2018-09-14.pdf

Quick check

“A big data platform monitors network usage of devices from million+ customers. We have different levels: **Sensor/Customer, Node (concentrator of multiple customers), Agent (concentrator of multiple Nodes) and the whole network. In a region, the real operator can generate 1.4 billion records per day ~ 72GB per day”**

Quick check

<https://presemo.aalto.fi/bdp>

Scalability and elasticity: scale out

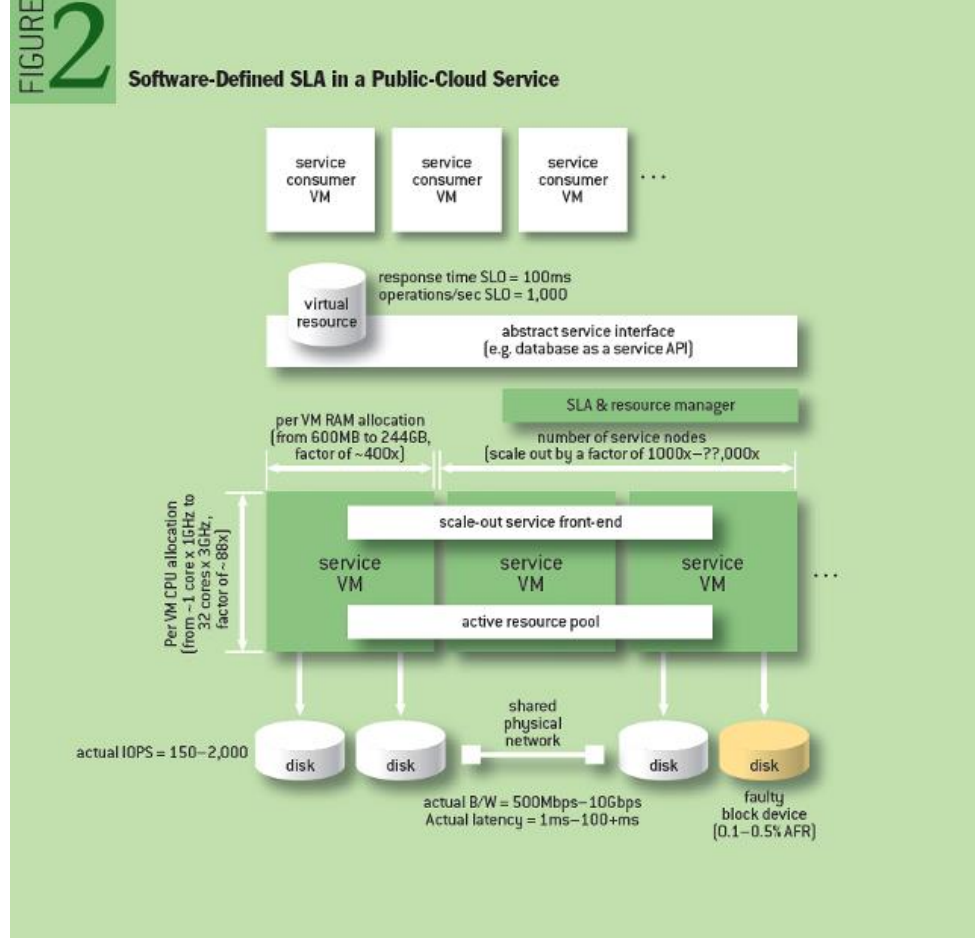


Figure source: <http://queue.acm.org/detail.cfm?id=2560948>

Scalability and elasticity: load balancing

FIGURE 3 Scalable Service Dispatch Architecture using SQL Server Service Broker

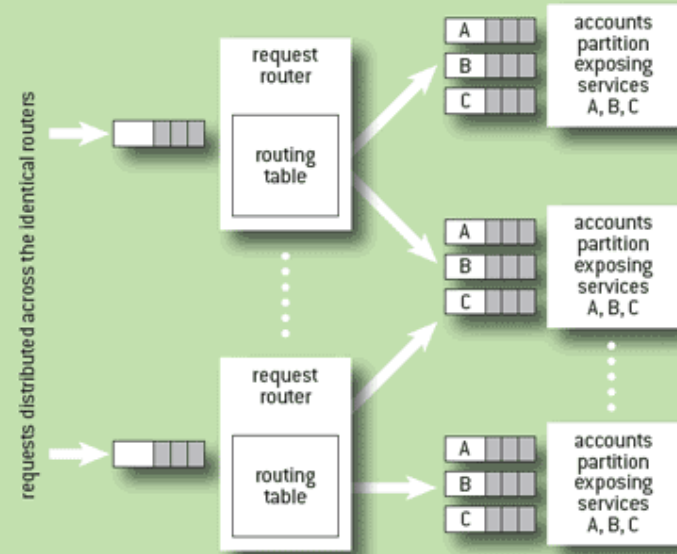



Figure source: <http://queue.acm.org/detail.cfm?id=1971597>

API for Platform as a Service

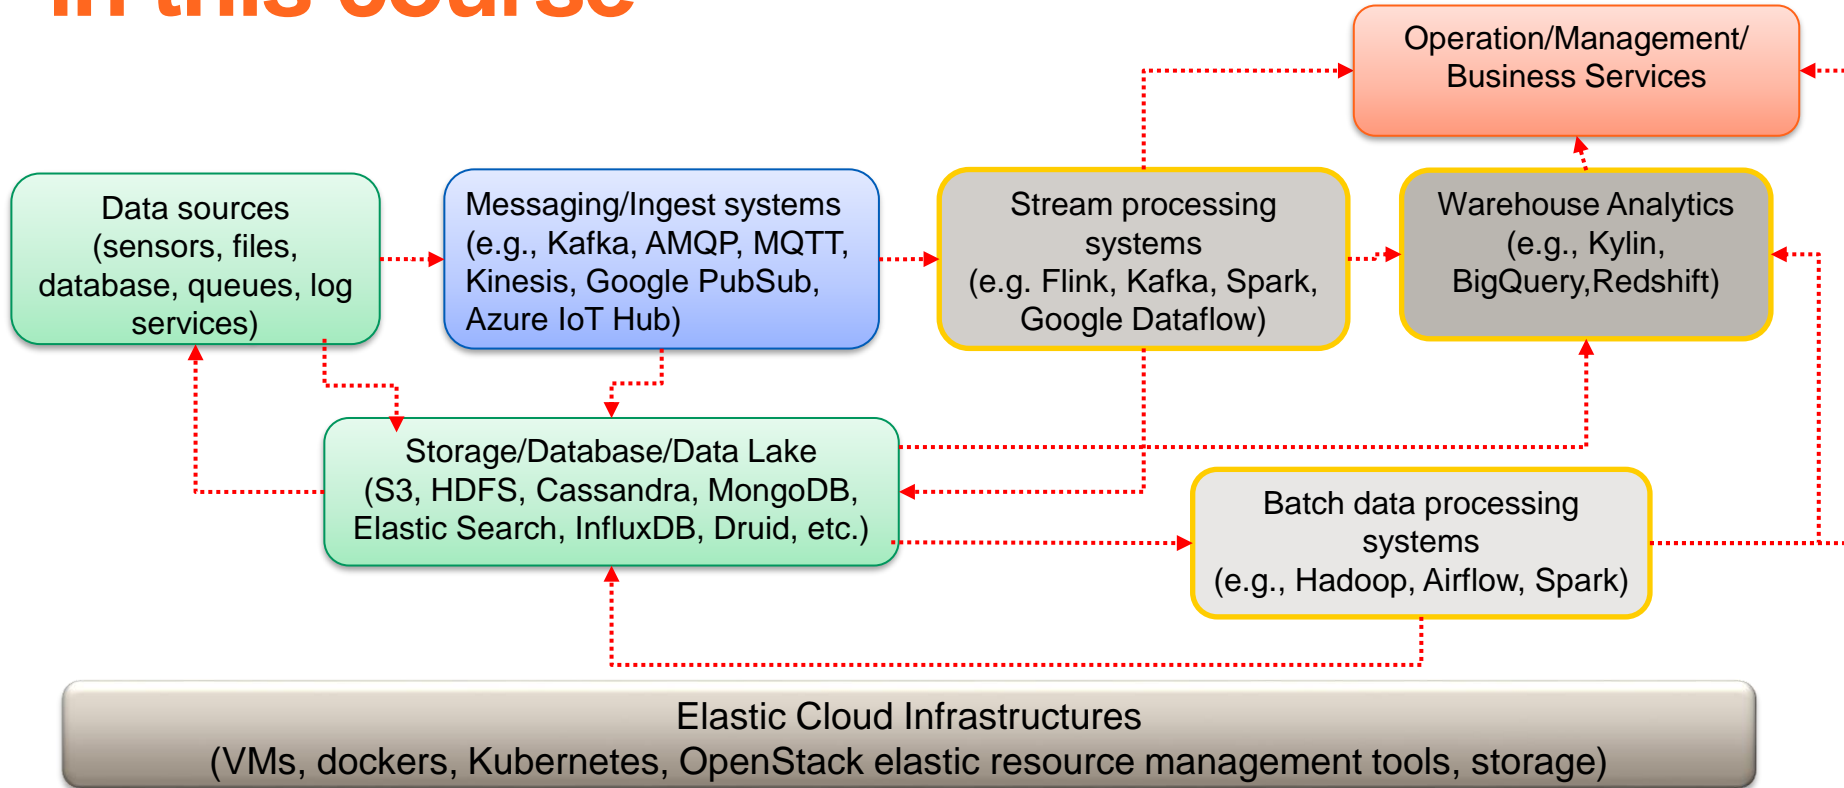
- **APIs are key! Why?**
 - **enable customers** access to data and analysis functions from your big data platforms without worrying about changes within your platforms
 - virtualization and management (hide internal, control access, throttling)



Which API would you publish? And how other concepts are related, e.g. API Gateways for Load balancing and Fault-Tolerance?

Common, high-level architecture view with popular state-of-the art technologies for our study

Big data at large-scale: the big picture in this course



Thanks!

Hong-Linh Truong
Department of Computer Science

rdsea.github.io