



Aalto University
School of Science

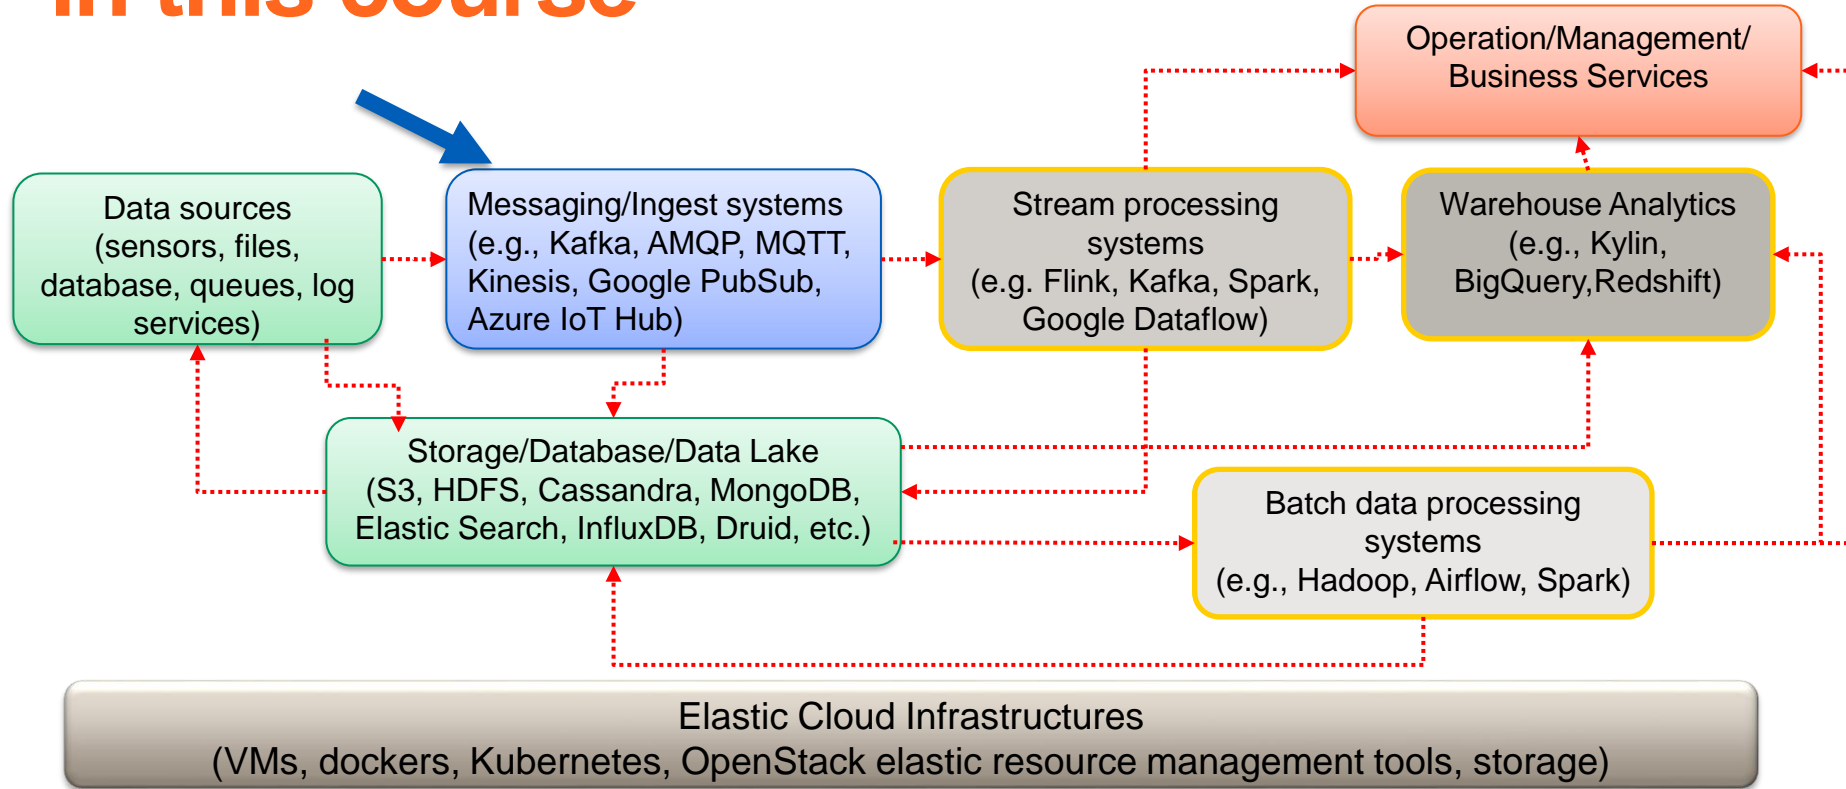
Streaming Data Ingestion with Apache Kafka

Hong-Linh Truong

Department of Computer Science

linh.truong@aalto.fi, <https://rdsea.github.io>

Big data at large-scale: the big picture in this course



Abstraction of Data Streams

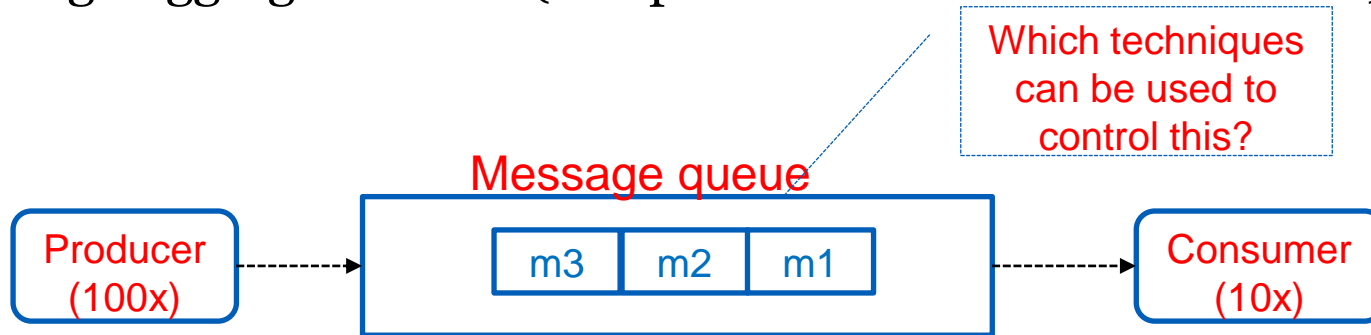
Data stream: a sequence/flow of *data units*

Data units are defined by applications: a data unit can be data described by a primitive data type or by a complex data type, a serializable object, etc.

Usually we encapsulate a data unit in a **record/message** of data

Some use cases

- Producers generate a lot of near real-time events
- Producers and consumers have different processing speeds
 - E.g. logging activities (fast producers but slow consumers)



- Rich and diverse types of events
- Dealing with cases when consumers might be on and off (fault tolerance support)

Key log-based messaging systems

- **Apache Kafka**
 - <https://kafka.apache.org/>
- **Apache Pulsar**
 - <https://pulsar.apache.org/>
- **LogDevice (Facebook)**
 - <https://logdevice.io/>

Apache Kafka


- <http://kafka.apache.org/>
 - originally from LinkedIn, not a protocol!
- **Some components are commercialized by Confluent**
 - <https://www.confluent.io/>
- **Widely used for big data use cases, including message processing in large-scale enterprise service platforms**
 - **data messages** (e.g., logs, records, historical events)
 - *It is our focus on big data platforms*
 - request/command messages (e.g., payment/database update)
 - event messages (e.g., notification of a payment due)

More than a message broker

- In **Apache Kafka**: the basic data element is **<Key,Value>** tuple
- **Messaging features**
 - for transferring messages
 - *Other frameworks in the ecosystem: RabbitMQ, Mostquitto*
- **Streaming processing**
 - streaming applications handle data from streams
 - read and write data back to Kafka messaging brokers
 - other frameworks in the ecosystem: Apache Flink and Apache Spark
- **High-level SQL-style: KSQL**
 - other possibilities: SQL-liked + Java in Apache Flink


**In the context of big data: we examine
Apache Kafka for transferring, ingesting
and processing **messages of data****

Kafka messaging design



- **Use a cluster of brokers to deliver messages**
 - usually within single data center, with high-speed networks, for a single tenant
- **Durable messages, ordered delivery via partitions**
- **Online/offline consumers**
- **Using filesystem **heavily** for message storage and caching**

Kafka design



- **A topic consists of different partitions**
- **Partitions**
 - enable parallel processing → performance
 - fault-tolerance via replication
- **Durable messages, ordered delivery** via partitions

Messages, topics and partitions

- Ordered, immutable sequence of messages
- Messages are kept in a period (regardless of consumers or not)
- Support **total order** for messages within a partition
- Partitions are distributed among server




Figure source:
<http://kafka.apache.org/documentation.html>

Consumers

- Consumer **pulls the data**
- The consumer **keeps a single pointer** indicating the position in a partition to keep track the offset of the next message being consumed
- **Why?**
 - allow customers to design their speed
 - support/optimize batching data
 - easy to implement total order over message
 - easy to implement reliable message/fault tolerance

Example of a producer

```
public SimpleProducer( String url, String inputfile, String topic ) {
    Properties props = new Properties();
    props.put("bootstrap.servers", url);
    props.put("client.id", "rdsea.io.training.demo");
    props.put("key.serializer", "org.apache.kafka.common.serialization.IntegerSerializer");
    props.put("value.serializer", "org.apache.kafka.common.serialization.StringSerializer");
    producer = new KafkaProducer<Integer,String>(props);
    this.topic = topic;
    this.inputfile =inputfile;
}

public void run() {
    int messageNo = 1;
    //read data from file:
    try {
        Reader in = new FileReader(inputfile);
        Iterable<CSVRecord> records = CSVFormat.RFC4180.withFirstRecordAsHeader().parse(in);
        for (CSVRecord record : records) {

            JsonObject event = new JsonObject();
            event.addProperty("USERPHONE", 6645);
            event.addProperty("TIME", Long.parseLong(record.get("TIME")));

            event.addProperty("lat", Float.parseFloat(record.get("LATITUDE")));
            event.addProperty("lon", Float.parseFloat(record.get("LONGITUDE")));

            event.addProperty("GSM_BIT_ERROR_RATE", Float.parseFloat(record.get("GSM_BIT_ERROR_RATE")));
            event.addProperty("GSM_SIGNAL_STRENGTH", Float.parseFloat(record.get("GSM_SIGNAL_STRENGTH")));
            //a simple way to handle missing data is to skip the record
            if (!record.get("LOC_ACCURACY").equals("")) {
                event.addProperty("LOC_ACCURACY", Float.parseFloat(record.get("LOC_ACCURACY")));
            } else {
                continue;
            }
            if (!record.get("LOC_SPEED").equals("")) {
                event.addProperty("LOC_SPEED", Float.parseFloat(record.get("LOC_SPEED")));
            } else {
                continue;
            }

            String eventString = "{\\"event\\": " + event + "}";
            try {
                producer.send(new ProducerRecord<Integer,String>(topic,messageNo,eventString)).get();
            } catch (ExecutionException e) {
                // TODO Auto-generated catch block
                e.printStackTrace();
            }
            System.out.println("Sent message : (" + messageNo + " " + eventString + ")\n");
            messageNo++;
        }
    } catch (IOException e) {
        e.printStackTrace();
    }
}
```

Example of a consumer

```
public class SimpleConsumer {
    private final KafkaConsumer<Integer, String> consumer;
    private final String topic;
    private final int pollNr;
    public SimpleConsumer(String url, String topic, int pollNr) {

        Properties props = new Properties();
        //just use standard example configuration
        props.put(ConsumerConfig.BOOTSTRAP_SERVERS_CONFIG, url);
        props.put(ConsumerConfig.GROUP_ID_CONFIG, "RDSEA Simple Consumer");
        props.put(ConsumerConfig.ENABLE_AUTO_COMMIT_CONFIG, "true");
        props.put(ConsumerConfig.AUTO_COMMIT_INTERVAL_MS_CONFIG, "1000");
        props.put(ConsumerConfig.SESSION_TIMEOUT_MS_CONFIG, "30000");
        props.put(ConsumerConfig.KEY_DESERIALIZER_CLASS_CONFIG, "org.apache.kafka.common.serialization.IntegerDeserializer");
        props.put(ConsumerConfig.VALUE_DESERIALIZER_CLASS_CONFIG, "org.apache.kafka.common.serialization.StringDeserializer");

        consumer = new KafkaConsumer<Integer, String>(props);
        this.topic = topic;
        this.pollNr = pollNr;
    }

    public void readData() {
        consumer.subscribe(Collections.singletonList(this.topic));
        ConsumerRecords<Integer, String> records = consumer.poll(pollNr);
        for (ConsumerRecord<Integer, String> record : records) {
            System.out.println("Received message: (" + record.key() + ", " + record.value() + ") at offset " + record.offset());
        }
    }

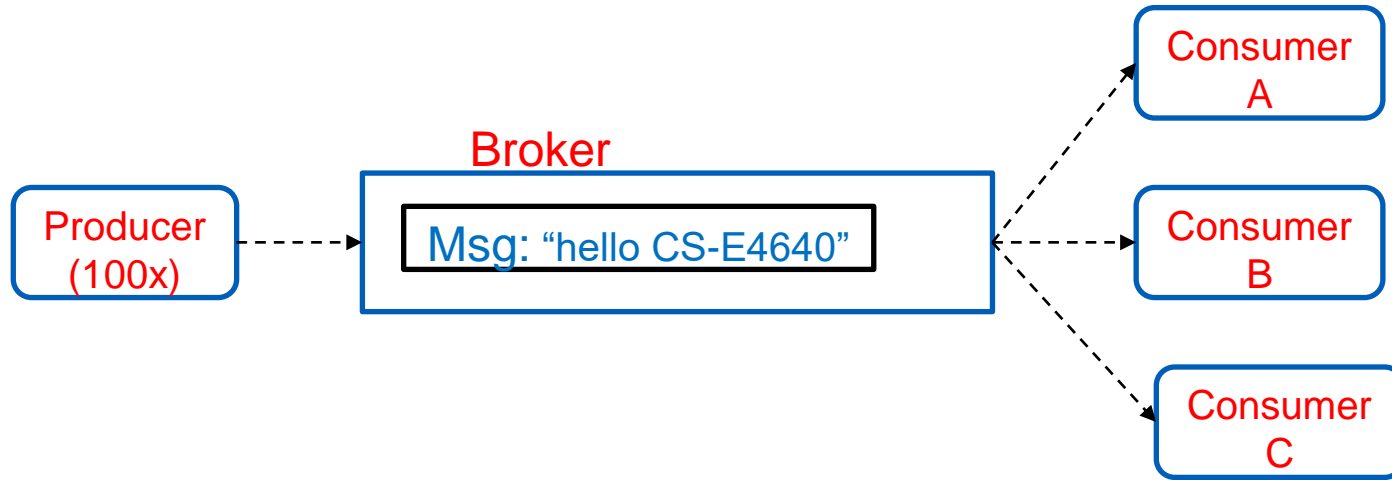
    public static void main(String[] args) {
        // TODO Auto-generated method stub
        if (args.length < 3) {
            System.out.println("Usage: SimpleProducer kafka_broker topic nr");
            System.exit(0);
        }

        int pollNr = Integer.valueOf(args[2]);
        SimpleConsumer consumer = new SimpleConsumer(args[0], args[1], pollNr);
        consumer.readData();
    }
}
```

Message delivery

- **message delivery guarantees are important for different use cases/requirements**
- **Some models**
 - At most once
 - At least once
 - Exactly once

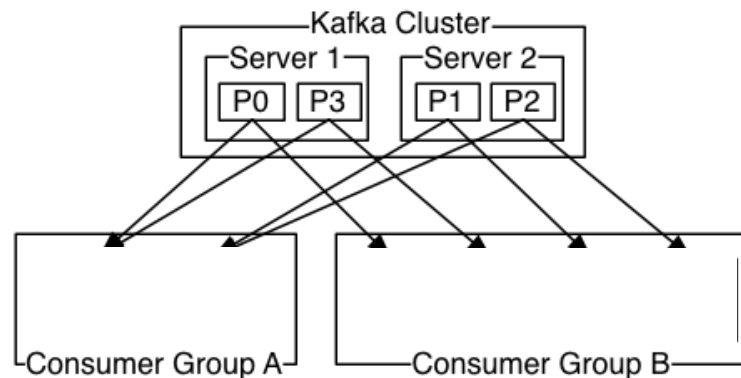
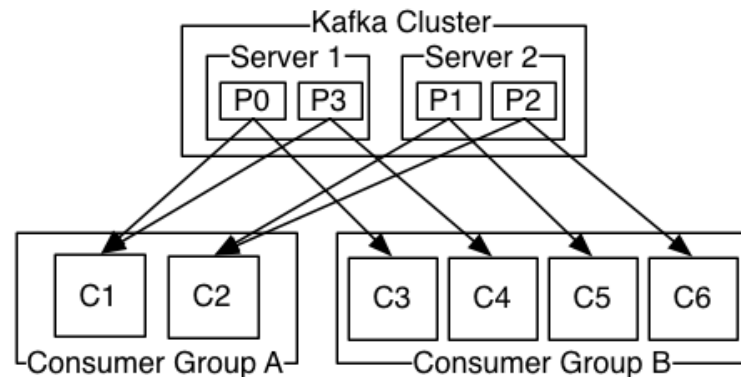
What does it mean exactly one?



- **Producer: idempotent delivery** → no duplicate entry in the log
- **Transaction-like semantics:** either message to ALL partition topics or not at all
- **Consumer behavior management**

Scalability and Fault Tolerance

- **Partitions are distributed and replicated among broker servers**
- **Consumers are organized into groups**
- **Each message is delivered to a consumer instance in a group**
- **One partition is assigned to one consumer**



Figures source: <http://kafka.apache.org/documentation.html#majordesignelements>

Partitions and partition replication

- **Why partitions?**
 - Support scalability
 - *enable arbitrary data types and sizes for a topic*
 - *enable parallelism in producing and consuming data*
- **But partitions are replicated, why?**
 - For fault tolerance

Partition Replication

Replication model: the leader-follower (primary-secondary) model!

The leader handles all read and write requests




Figure source: <http://de.slideshare.net/junrao/kafka-replication-apachecon2013>

Consumer group

- **Consumer group: a set of consumers**
 - is used to support scalability and fault tolerance
 - allows multiple consumers to read a topic
- **In one group: each partition is consumed by only consumer instance**
 - Combine „queuing“ model and „publish/subscribe“ model
- **Enable different applications receive data from the same topic.**
 - different consumers in different groups can retrieve the same data

Group rebalancing

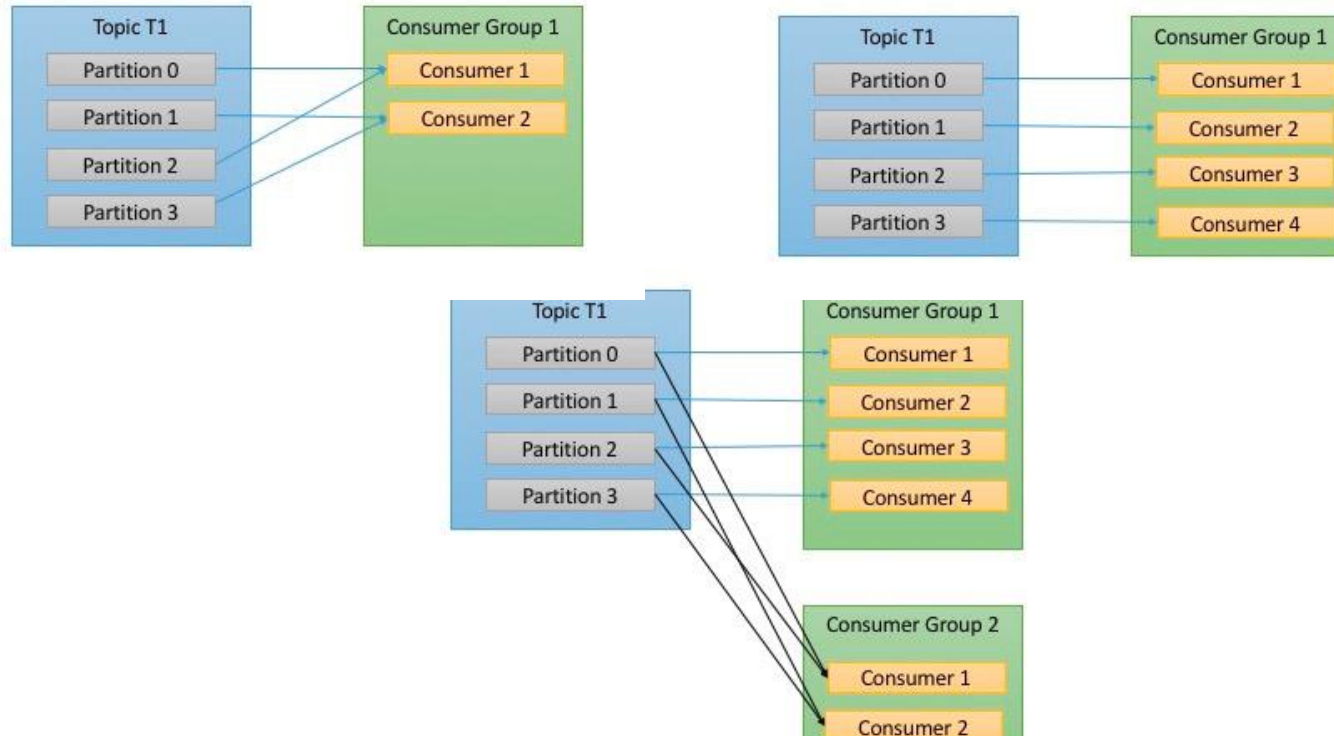


Figure source: <https://www.safaribooksonline.com/library/view/kafka-the-definitive/9781491936153/ch04.html>

Key questions/thoughts

- Why do we need partitions per topic?
 - arbitrary data handling, ordering guarantees, load balancing
- How to deal with high volume of near real-time messages for online and offline consumers?
 - partition, cluster, message storage, batch retrieval, etc.
- Queuing or publish-subscribe model?
 - check how Kafka delivers messages to consumer instances/groups

Kafka vs RabbitMQ

Figure source: Philippe Dobbelaere and Kyumars Sheykh Esmaili. 2017. Kafka versus RabbitMQ: A comparative study of two industry reference publish/subscribe implementations: Industry Paper. In Proceedings of the 11th ACM International Conference on Distributed and Event-based Systems (DEBS '17). ACM, New York, NY, USA, 227-238. DOI: <https://doi.org/10.1145/3093742.3093908>

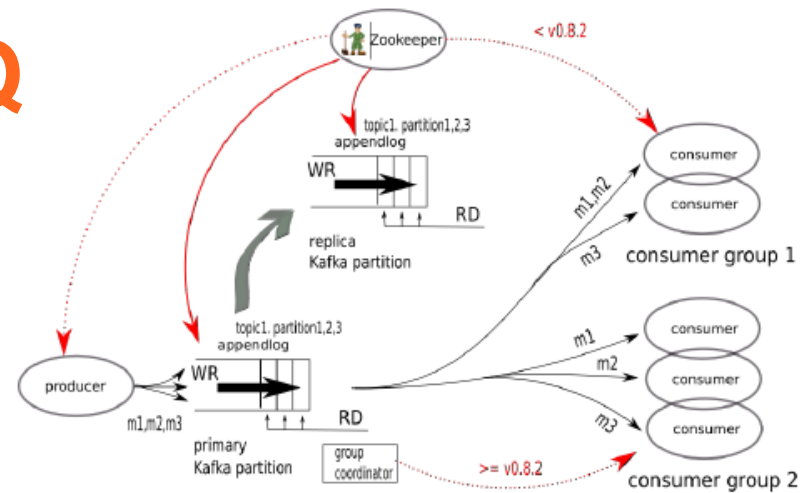


Figure 1: Kafka Architecture




Figure 2: RabbitMQ (AMQP) Architecture

Hands-on

- **Understanding the message broker systems and message delivery are key for streaming processing**
- **Check our tutorial:**
 - <https://version.aalto.fi/gitlab/bigdataplatforms/cs-e4640/-/tree/master/tutorials/basickafka>
 - <https://version.aalto.fi/gitlab/bigdataplatforms/cs-e4640/-/tree/master/tutorials/cloud-data-pipeline>