



Aalto University
School of Science

Architecting Big Data Platforms

Hong-Linh Truong

Department of Computer Science

linh.truong@aalto.fi, <https://rdsea.github.io>

Learning objectives

- **Understand key issues in designing a big data platform**
- **Learn key architecture design issues**
- **Be familiar with the course' selected big data platform technologies**

Your big data platform story - an evolving scenario

“Your team has to build a big data platform for **X types of data**. Data will be generated/collected from **N sources**. We expect to have **10+ GBs/day of data to be ingested** into our platform. We will have to serve **K thousands of requests** for different types of analytics – to be determined. Our response time should be **in t milliseconds**. Our services should not be ...”

You may have several questions?

- Do we have to support multiple types of data?
- How do data pipelines and data load look like?
- How to enable different data processing models?
- Which runtime parameters must be monitored? Which service level metrics must be guaranteed?
- To where we should distribute/deploy our components?
- Which part of the platform we must manage by ourselves and which part will be fully managed by other providers?
- How to design elastic big data infrastructures?
- Etc.

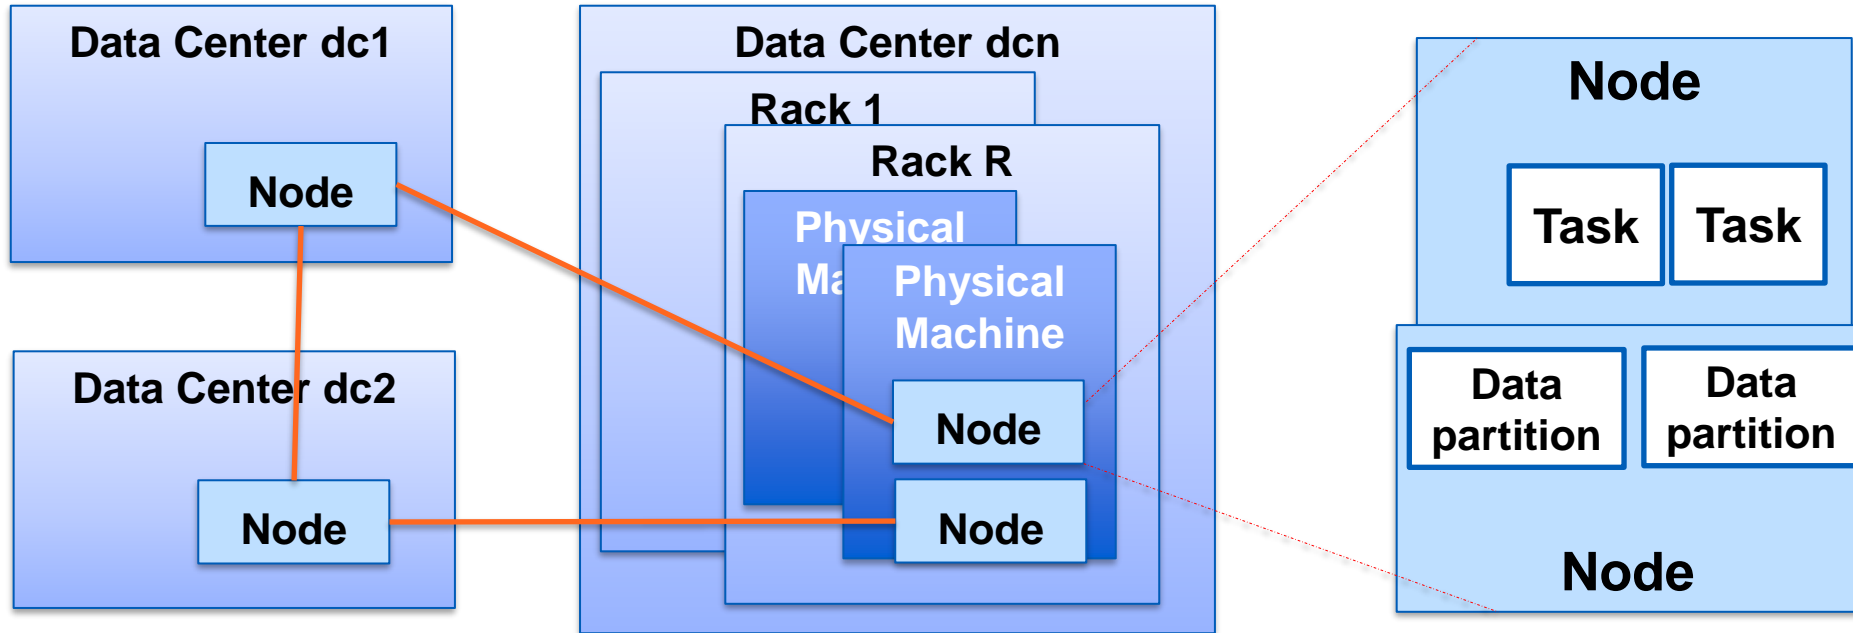
Your Big Data Platform story starts with Big Data Platform architectures!

**To architect the platform
centered around data!**

Understanding the underlying computing infrastructures

- **Computing resources and services**
 - many machines, virtual infrastructures, different types of services
- **Distributed infrastructures from different administrative domains**
 - in multiple data centers, locations and countries
 - with different security and network policies
- **Diverse service level objectives (SLO) and service level agreements (SLAs)**
 - performance, service failure, cost, privacy/security ...

Understanding the underlying infrastructures for big data platforms



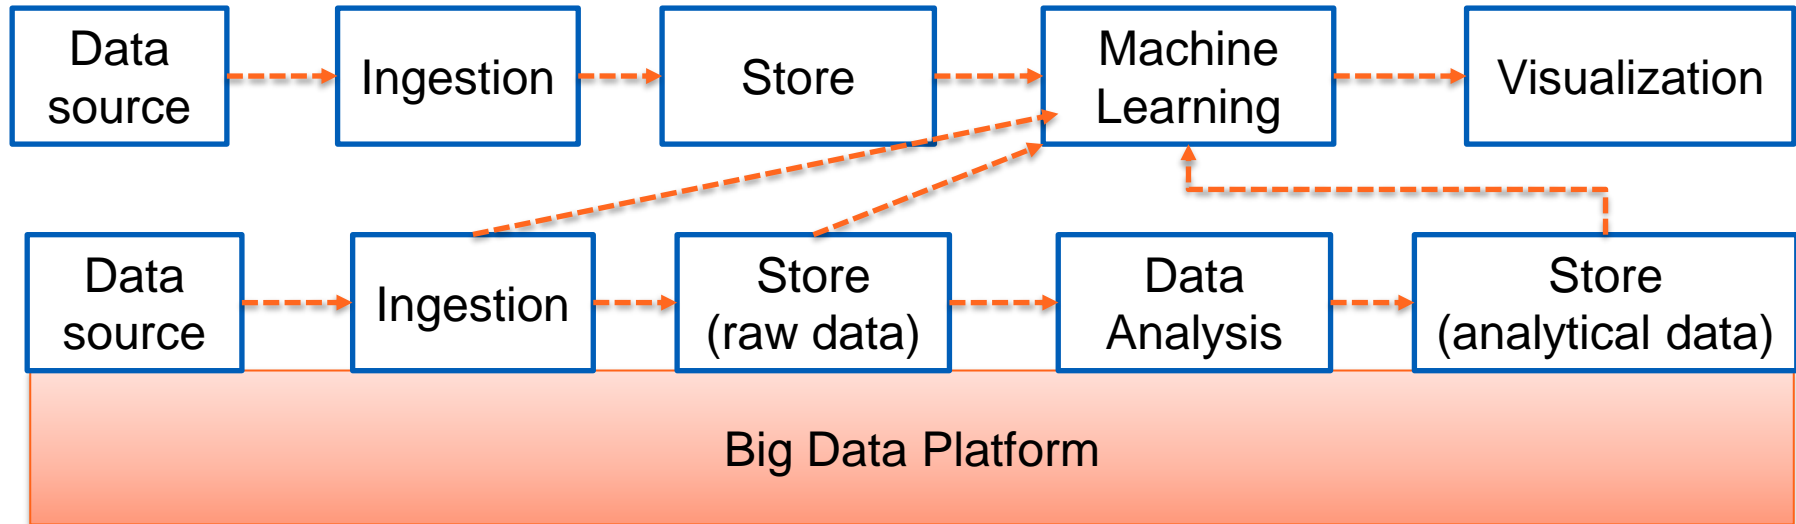
Remember: large-scale distributed infrastructures!

Data-centric development & operations

- **Data ingestion and ETL (Extract, Transform, Load)**
 - from various data sources we move data into the platform
- **Data storing and management**
 - ingested data will be stored and managed using different types of storages and databases
- **Data analysis and (Machine) learning**
 - data within platforms will be processed, analyzed and learned to improve data, find insights and to create models
 - data **at rest** vs data **in motion**
- **Reporting and visualization**
 - patterns/insights in data will be interpreted and presented for decision-making, reporting and creating stories

Big Data Pipelines

Multiple big data pipelines can be constructed atop a big data platform (and across distributed infrastructures)



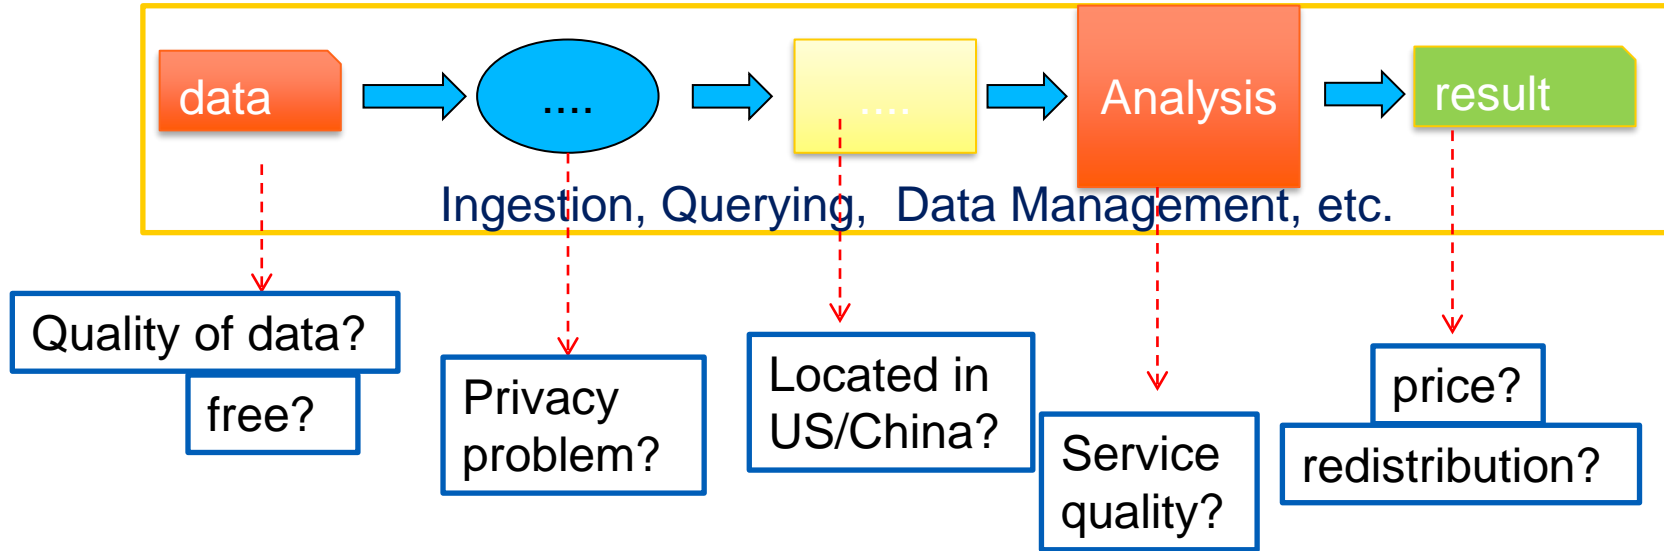
Handling multiple types of data?

- **First important aspect: you don't have to support multiple types of data**
 - but are you sure that you will not have this in the future?
- **Multiple types of data**
 - any linked models among them?
- **Any elastic solution that ensures minimum changes to support **generalization and extensibility****
 - e.g., multi-model databases, microservices of multiple of databases or data lake

Data concerns and SLAs

- **Ingesting data**
 - mapping and transforming data
 - ingestion of data under V^*
 - data validation/quality control during ingestion
- **Storing data**
 - data sharding and consistency, data backup, retention, etc.
 - the impact of the rights to remove data
- **SLA multitenancy versus single tenancy**
 - security, privacy, performance, reliability and maintenance?

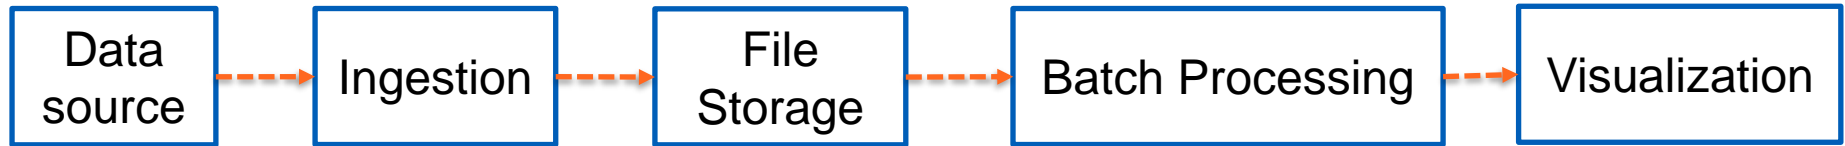
Data concerns: example



- Ethical consequences?
- Regulation-compliant platforms: e.g., GDPR

Fast/slow reliable processing

big data but not near real-time, e.g., take customer transaction files from companies and move to data centers for analytics



fast, small IoT data in near real-time flows, e.g. position of cars



Design goals

For dealing with V^*

- **Responsive: guarantee quality of services**
- **Resilient: deal within failures**
- **Elastic: deal with different workload and quality of analytics**
- **Loosely coupling: support reusability, composition, and extensibility**

Designs must address various aspects

- **Responsive:**

- distributed computing, multi layer optimization

- **Resilient:**

- replication, containment, isolation

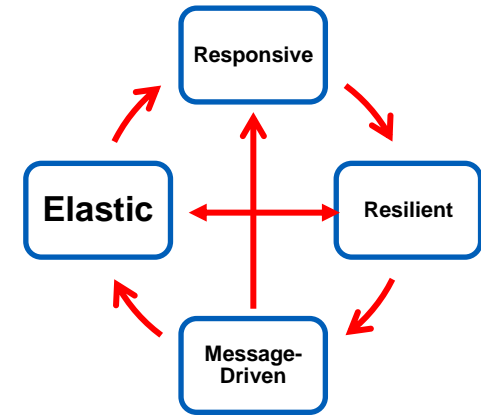
- **Elastic:**

- sharding, replication, load balancing, scale up/out

- **Message-driven:**

- loosely coupling with messages, non-blocking protocols, location-independent

Reactive systems



Source: <https://www.reactivemanifesto.org/>

Efficiency and sustainability

- Highly efficient might not be good?
- Sustainability within big data platforms
 - energy consumption, reusability, extensibility

Design and implementation



Partitioning: splitting functionality & data

- **breakdown the complexity**
- **easy to implement, replace and compose**
- **deal with performance, scalability, security, etc.**
- **support teams in DevOps**
- **cope with technology changes**
- **Many things are related to the current trend of microservices**

Example of functional and data partitioning

Service-oriented components

Microservices and domain-oriented microservices

Serverless functions/function as-a service

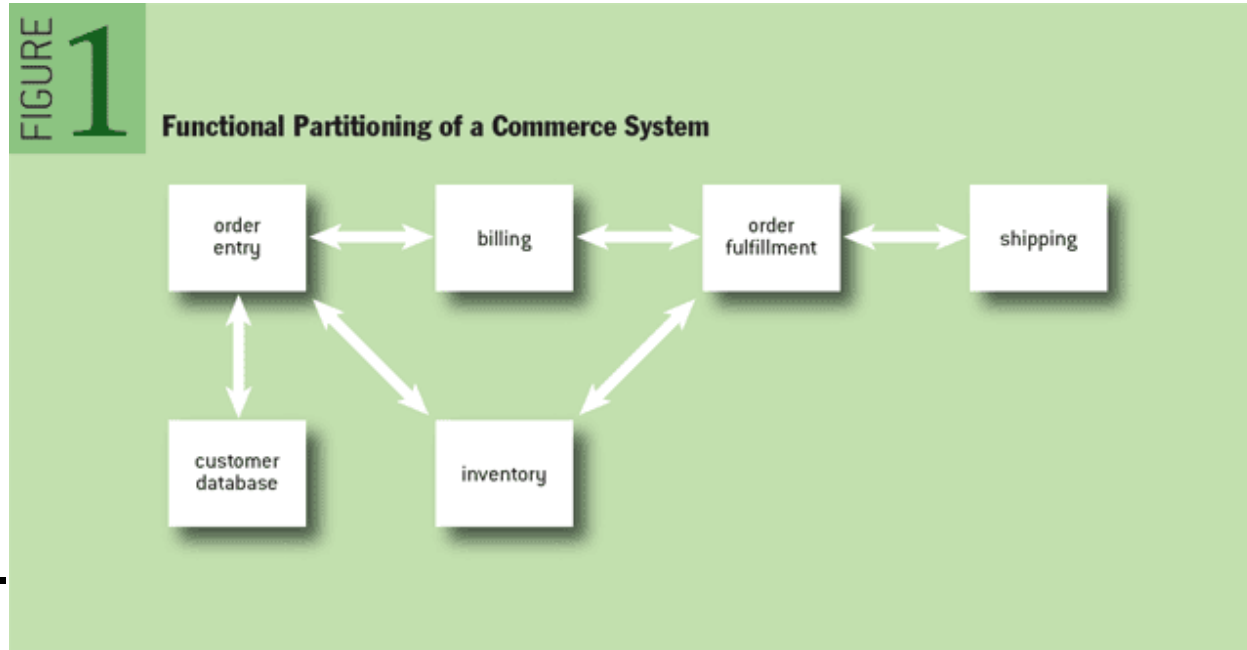


Figure source: <http://queue.acm.org/detail.cfm?id=1971597>

Example of functional and data partitioning

Data sharding

Multi data spaces

Multi data services

Multiple data infrastructures

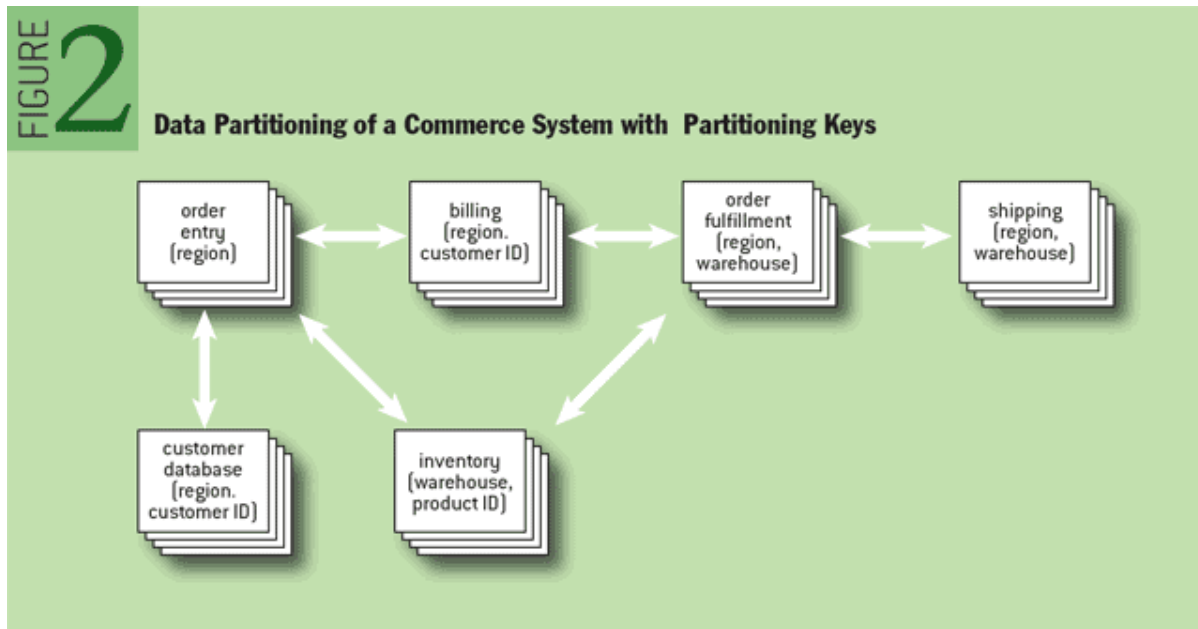


Figure source: <http://queue.acm.org/detail.cfm?id=1971597>



Aalto University
School of Science

**Distributed systems of components
are used to manage, ingest data and
process data**

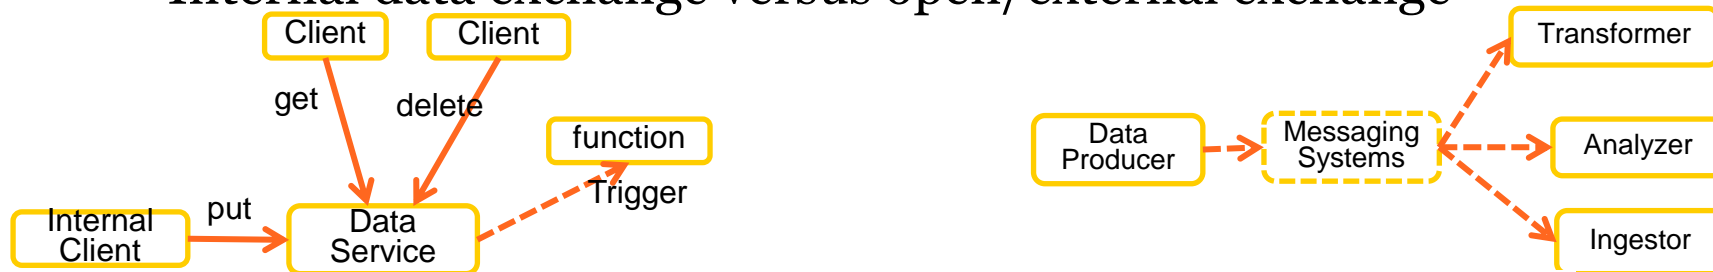
Interaction: Complex interactions

- **Protocols:**

- REST, gRPC, Message Passing, Stream-oriented Communication

- **Models**

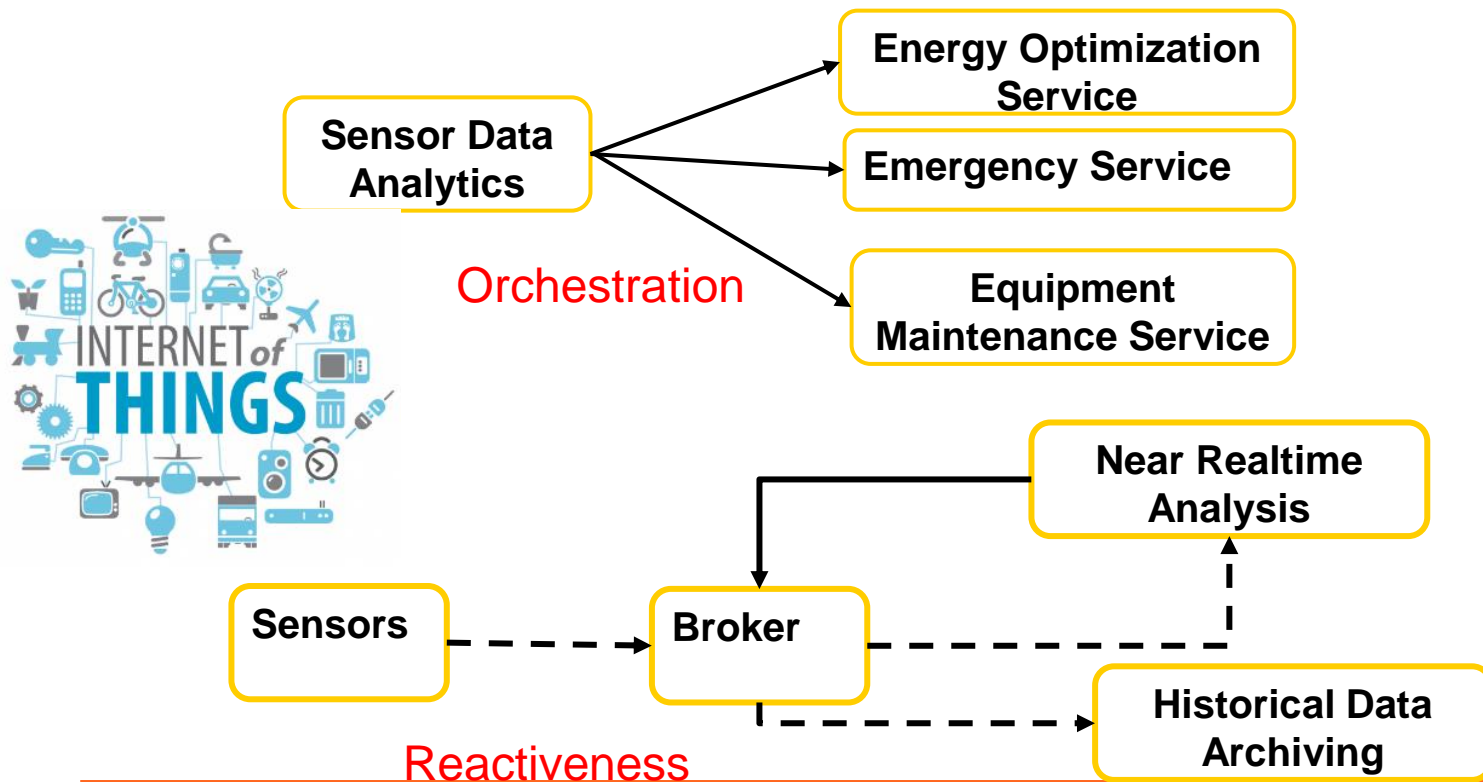
- One-to-many, many-to-one, many-to-many
- Synchronous/asynchronous calls
- Internal data exchange versus open/external exchange



Amazon S3/MongoDB

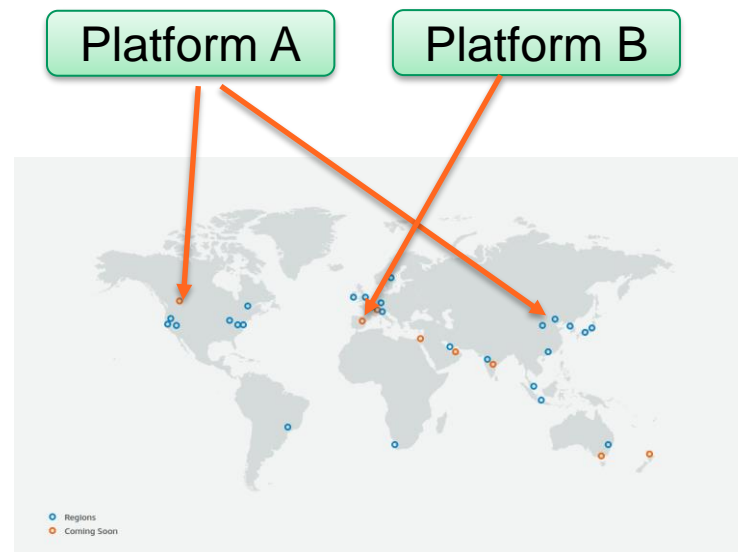


Coordination: Orchestration and Reactiveness



Distribution: Edge or Data Centers?

- Big data & components components can be distributed in different places!
- Global deployment or not?
- Move analytics/work or move data?



Map of AWS infrastructure (08.01.2022)

Source: <https://aws.amazon.com/about-aws/global-infrastructure/>

An outage can lead to a huge problem. Example:

<https://www.thousandeyes.com/blog/aws-outage-analysis-dec-7-2021>

Scalability & elasticity

Presto: <https://prestodb.io/>

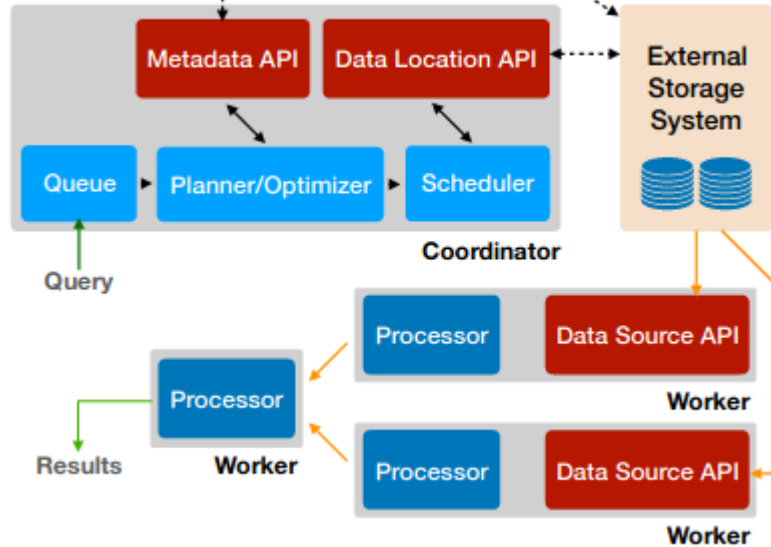


Figure source: Presto: SQL on Everything
<https://ieeexplore.ieee.org/stamp/stamp.jsp?tp=&arnumber=8731547&tag=1>

Lyft Presto Gateway

“As of today we have 60 PB of query-able event data stored in an S3 based data lake and about 10 PB of raw data is being scanned every day using Presto”

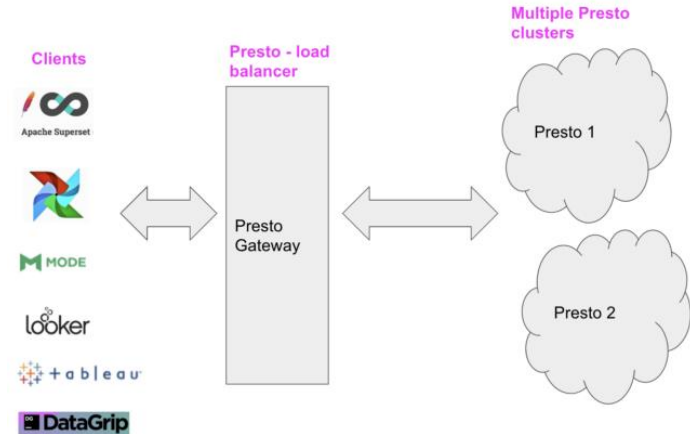
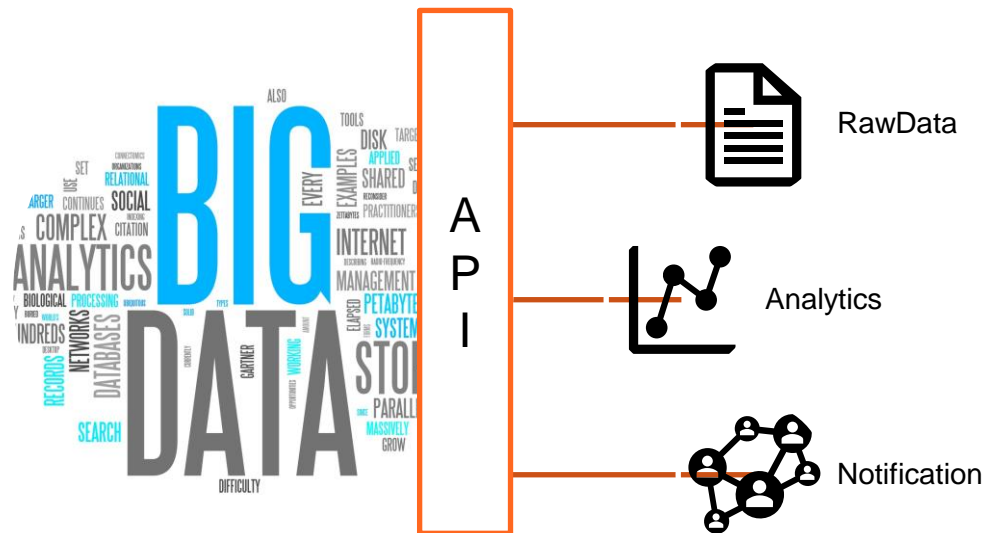


Figure and text source: <https://eng.lyft.com/presto-infrastructure-at-lyft-b10adb9db01>

API for Platform as a Service

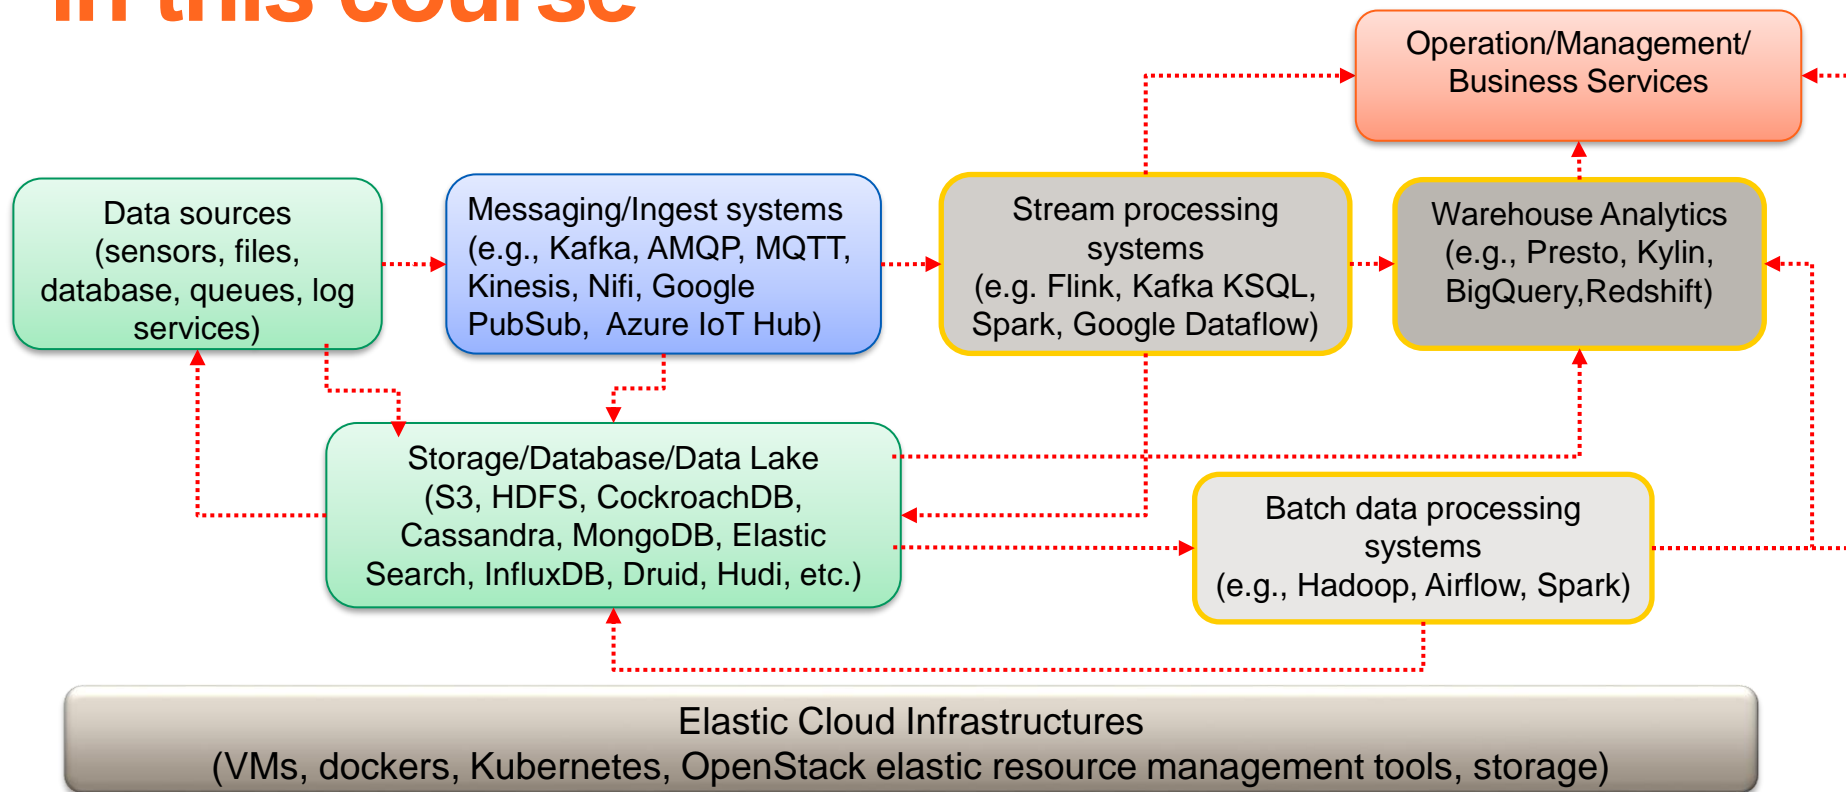
- **APIs are key! Why?**
 - **enable customers** access to data and analysis functions from your big data platforms without worrying about changes within your platforms
 - virtualization and management (hide internal, control access, throttling)



Which API would you publish? And how other concepts are related, e.g. API Gateways for Load balancing and Fault-Tolerance?

Common, high-level architecture view with popular state-of-the art technologies for our study

Big data at large-scale: the big picture in this course



Thanks!

Hong-Linh Truong
Department of Computer Science

rdsea.github.io