



Aalto University  
School of Science

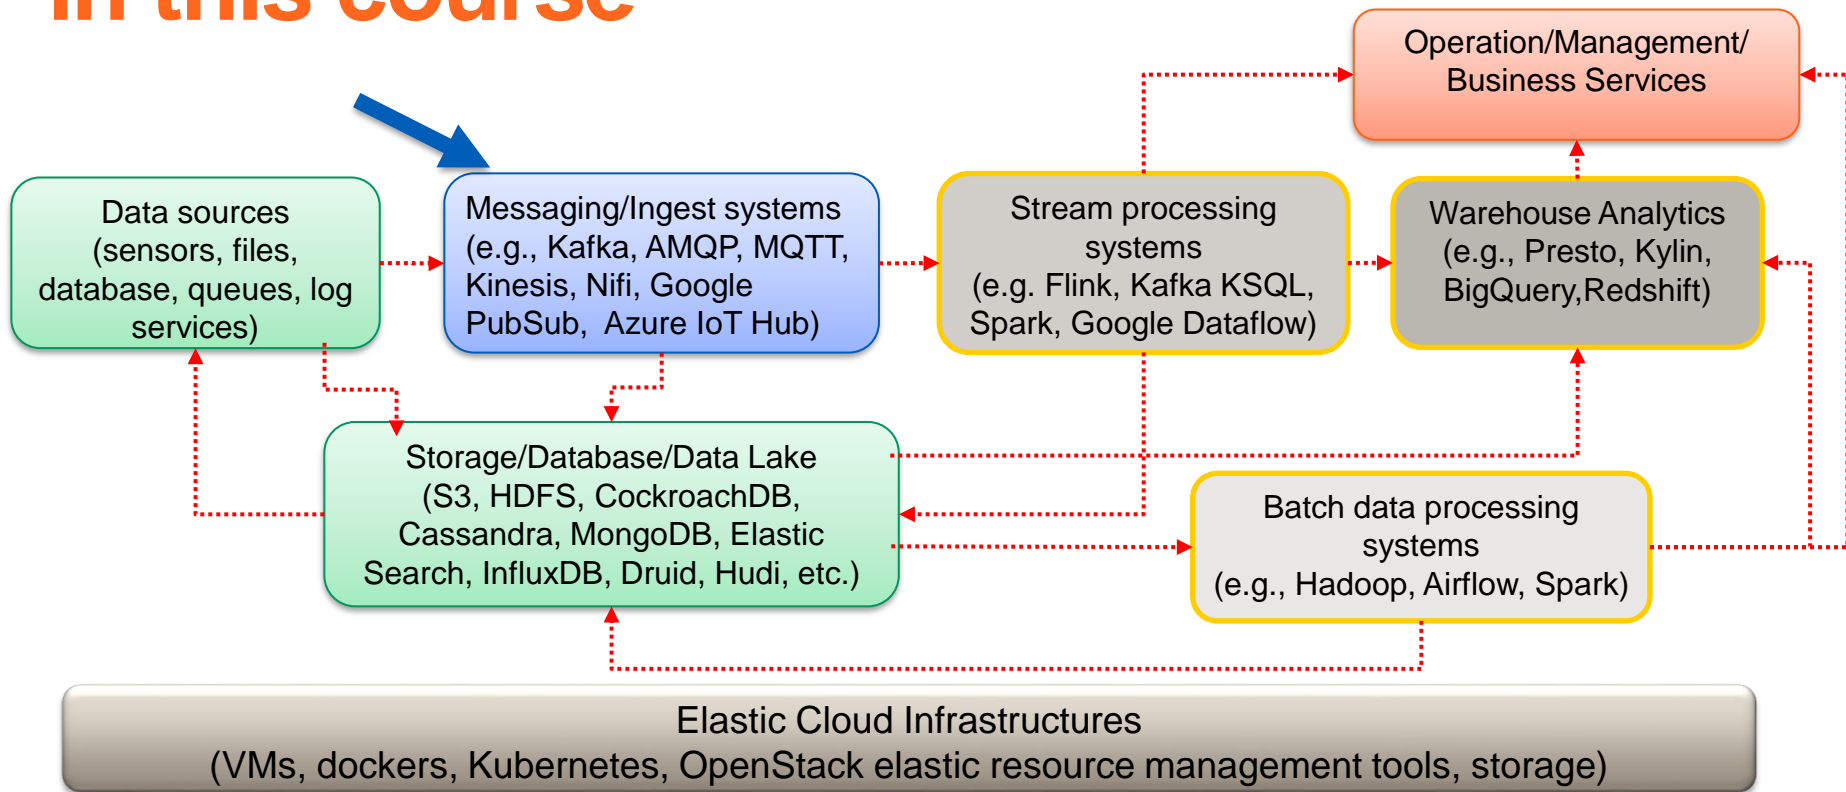
# Streaming Data Ingestion with Apache Kafka

*Hong-Linh Truong*

*Department of Computer Science*

*[linh.truong@aalto.fi](mailto:linh.truong@aalto.fi), <https://rdsea.github.io>*

# Big data at large-scale: the big picture in this course



# Abstraction of Data Streams

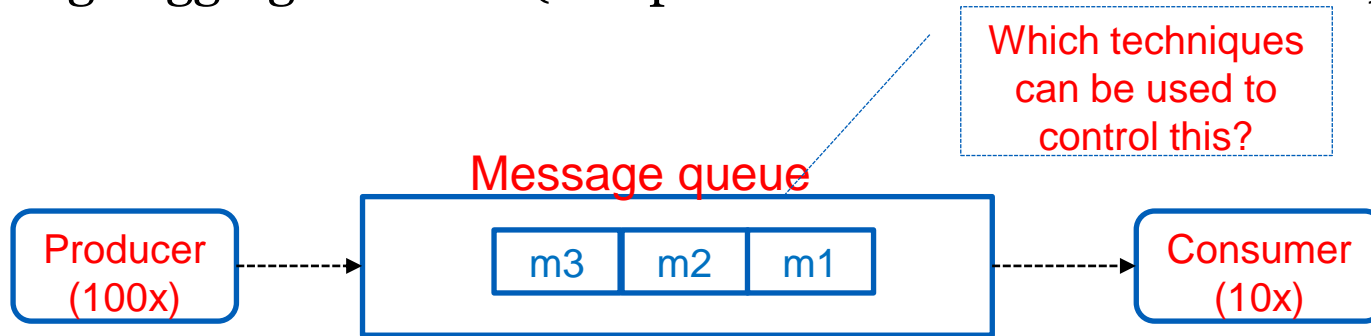
**Data stream:** a sequence/flow of *data units*

*Data units* are defined by applications: a data unit can be data described by a primitive data type or by a complex data type, a serializable object, etc.

Usually we encapsulate a data unit in a **record/message** of data

# Some use cases

- Producers generate a lot of near real-time events
- Producers and consumers have different processing speeds
  - E.g. logging activities (fast producers but slow consumers)



- Rich and diverse types of events
- Dealing with cases when consumers might be on and off (fault tolerance support)

# Key log-based messaging systems

- **Apache Kafka**
  - <https://kafka.apache.org/>
- **Apache Pulsar**
  - <https://pulsar.apache.org/>
- **LogDevice (Facebook)**
  - <https://logdevice.io/>

# Apache Kafka

- <http://kafka.apache.org/>
  - originally from LinkedIn, not a protocol!
- **Some components are commercialized by Confluent**
  - <https://www.confluent.io/>
- **Widely used for big data use cases, including message processing in large-scale enterprise service platforms**
  - **data messages** (e.g., logs, records, historical events)
    - *It is our focus on big data platforms*
  - request/command messages (e.g., payment/database update)
  - event messages (e.g., notification of a payment due)


# More than a message broker

- In **Apache Kafka**: the basic data element is <Key,Value> tuple
- **Messaging features**
  - for transferring messages
    - *Other frameworks in the ecosystem: RabbitMQ, Mostquitto*
- **Streaming processing**
  - streaming applications handle data from streams
  - read and write data back to Kafka messaging brokers
  - other frameworks in the ecosystem: Apache Flink and Apache Spark
- **High-level SQL-style: KSQL**
  - other possibilities: SQL-liked + Java in Apache Flink

**In the context of big data: we examine  
Apache Kafka for transferring, ingesting  
and processing **messages of data****




# Kafka messaging design



- **Use a cluster of brokers to deliver messages**
  - usually within single data center, with high-speed networks, for a single tenant
- **Durable messages, ordered delivery via partitions**
- **Online/offline consumers**
- **Using filesystem **heavily** for message storage and caching**

# Kafka design



- A topic consists of different partitions
- Partitions
  - enable parallel processing → performance
  - fault-tolerance via replication
- **Durable messages, ordered delivery** via partitions

# Messages, topics and partitions

- Ordered, immutable sequence of messages
- Messages are kept in a period (regardless of consumers or not)
- Support **total order** for messages within a partition
- Partitions are distributed among server




Figure source:  
<http://kafka.apache.org/documentation.html>

# Consumers

- Consumer **pulls the data**
- The consumer **keeps a single pointer** indicating the position in a partition to keep track the offset of the next message being consumed
- **Why?**
  - allow customers to design their speed
  - support/optimize batching data
  - easy to implement total order over message
  - easy to implement reliable message/fault tolerance

# Example of a producer

[https://version.aalto.fi/gitlab/bigdataplatfroms/cs-e4640/-/blob/master/tutorials/basic kafka/code/simple\\_kafka\\_producer.py](https://version.aalto.fi/gitlab/bigdataplatfroms/cs-e4640/-/blob/master/tutorials/basic kafka/code/simple_kafka_producer.py)

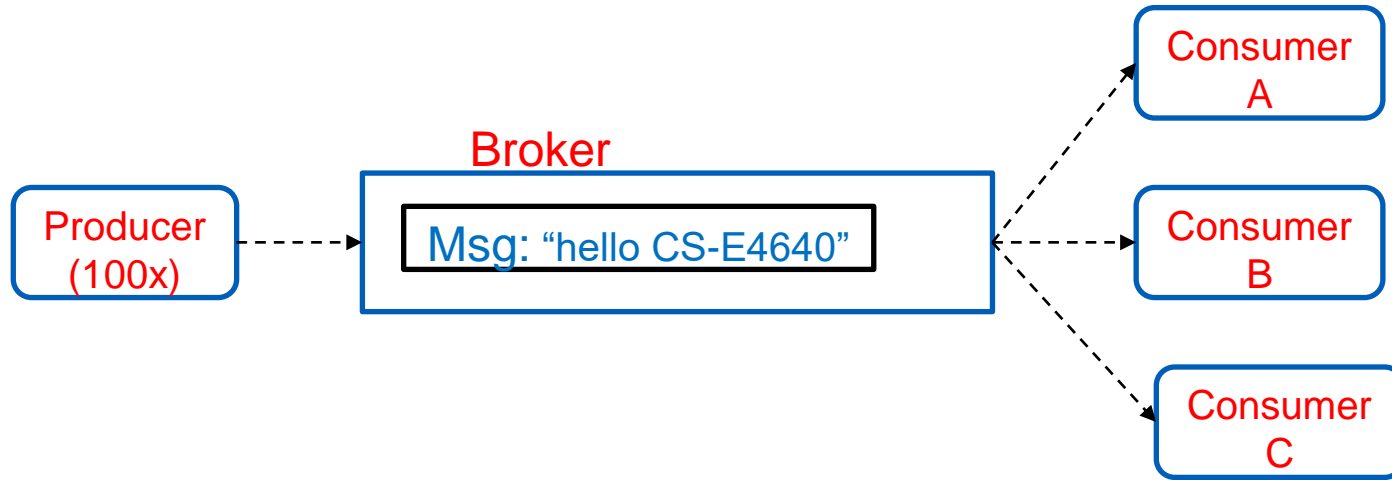
# Example of a consumer

[https://version.aalto.fi/gitlab/bigdataplatforms/cs-e4640/-/blob/master/tutorials/basickafka/code/simple\\_kafka\\_consumer.py](https://version.aalto.fi/gitlab/bigdataplatforms/cs-e4640/-/blob/master/tutorials/basickafka/code/simple_kafka_consumer.py)

# Message delivery

- **message delivery guarantees are important for different use cases/requirements**
- **Some models**
  - At most once
  - At least once
  - Exactly once

# What does it mean exactly one?

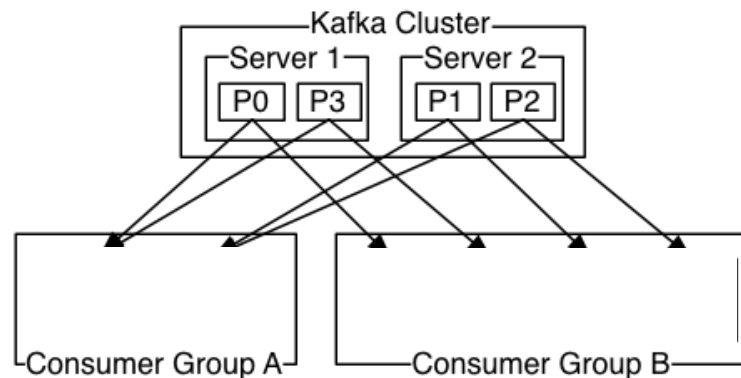
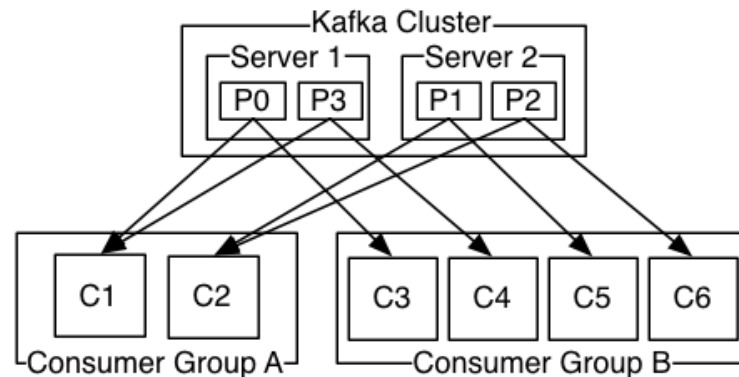


- **Producer: idempotent delivery → no duplicate entry in the log**
- **Transaction-like semantics: either message to ALL partition topics or not at all**
- **Consumer behavior management**



# Scalability and Fault Tolerance

- **Partitions are distributed and replicated among broker servers**
- **Consumers are organized into groups**
- **Each message is delivered to a consumer instance in a group**
- **One partition is assigned to one consumer**



Figures source: <http://kafka.apache.org/documentation.html#majordesignelements>

# Partitions and partition replication

- **Why partitions?**
  - Support scalability
    - *enable arbitrary data types and sizes for a topic*
    - *enable parallelism in producing and consuming data*
- **But partitions are replicated, why?**
  - For fault tolerance

# Partition Replication

**Replication model: the leader-follower (primary-secondary) model!**

**The leader handles all read and write requests**




Figure source: <http://de.slideshare.net/junrao/kafka-replication-apachecon2013>

# Consumer group

- **Consumer group: a set of consumers**
  - is used to support scalability and fault tolerance
  - allows multiple consumers to read a topic
- **In one group: each partition is consumed by only consumer instance**
  - Combine „queuing“ model and „publish/subscribe“ model
- **Enable different applications receive data from the same topic.**
  - different consumers in different groups can retrieve the same data

# Group rebalancing

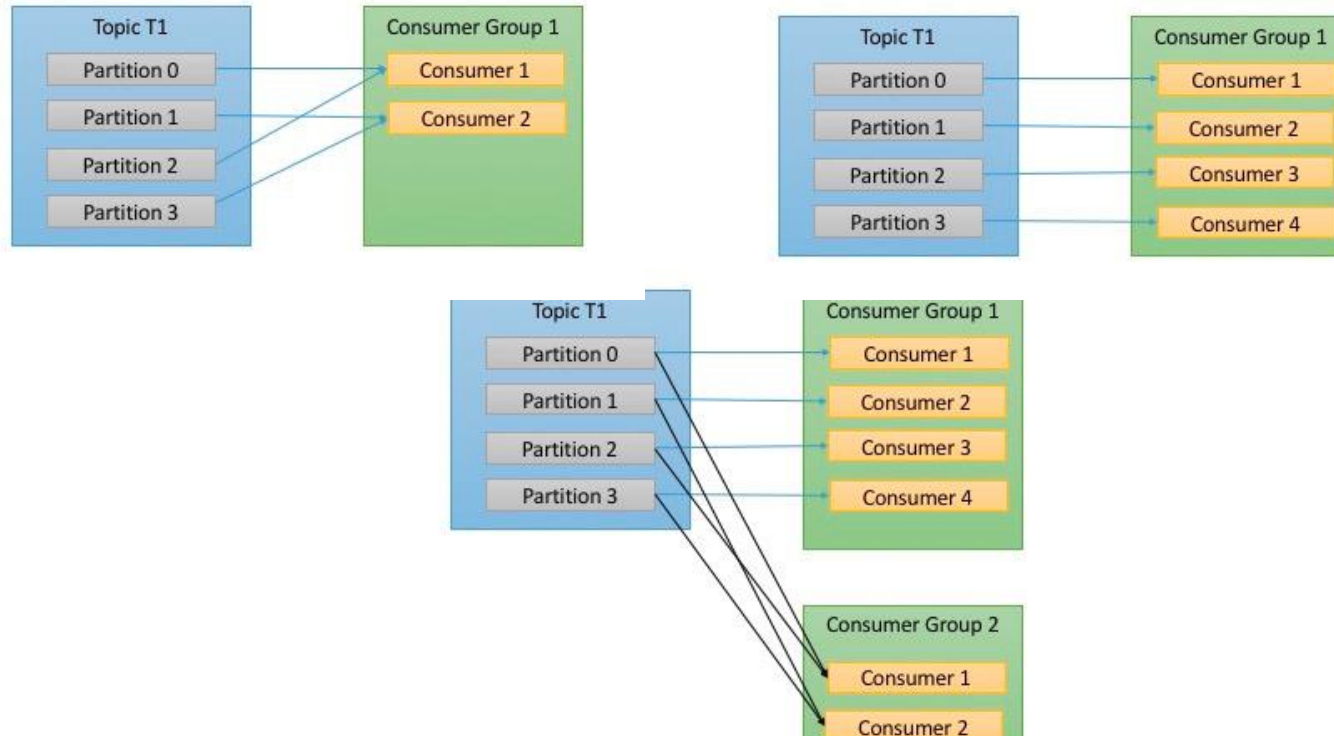


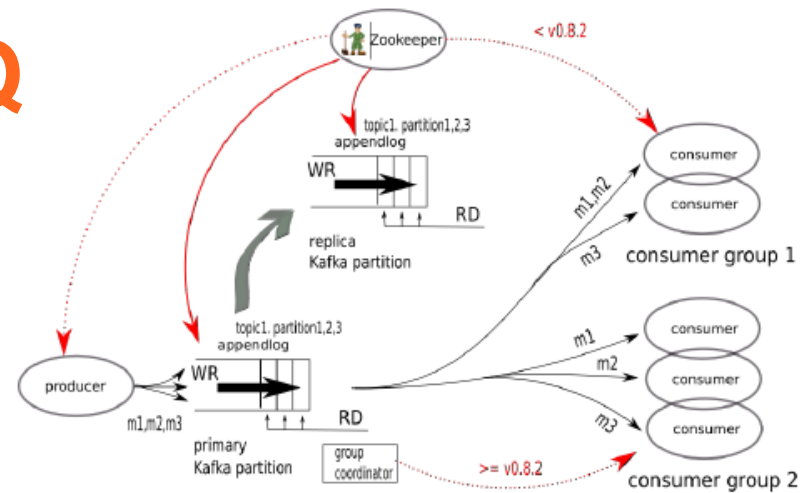
Figure source: <https://www.safaribooksonline.com/library/view/kafka-the-definitive/9781491936153/ch04.html>

# Key questions/thoughts


- Why do we need partitions per topic?
  - arbitrary data handling, ordering guarantees, load balancing
- How to deal with high volume of near real-time messages for online and offline consumers?
  - partition, cluster, message storage, batch retrieval, etc.
- Queuing or publish-subscribe model?
  - check how Kafka delivers messages to consumer instances/groups

# Kafka vs RabbitMQ

Figure source: Philippe Dobbelaere and Kyumars Sheykh Esmaili. 2017. Kafka versus RabbitMQ: A comparative study of two industry reference publish/subscribe implementations: Industry Paper. In Proceedings of the 11th ACM International Conference on Distributed and Event-based Systems (DEBS '17). ACM, New York, NY, USA, 227-238. DOI: <https://doi.org/10.1145/3093742.3093908>



**Figure 1: Kafka Architecture**



**Figure 2: RabbitMQ (AMQP) Architecture**

# Hands-on

- **Understanding the message broker systems and message delivery are key for streaming processing**
- **Check our tutorial:**
  - <https://version.aalto.fi/gitlab/bigdataplatforms/cs-e4640/-/tree/master/tutorials/basickafka>
  - <https://version.aalto.fi/gitlab/bigdataplatforms/cs-e4640/-/tree/master/tutorials/cloud-data-pipeline>