



Aalto University
School of Science

Stream Processing and Big Data Platforms

Hong-Linh Truong

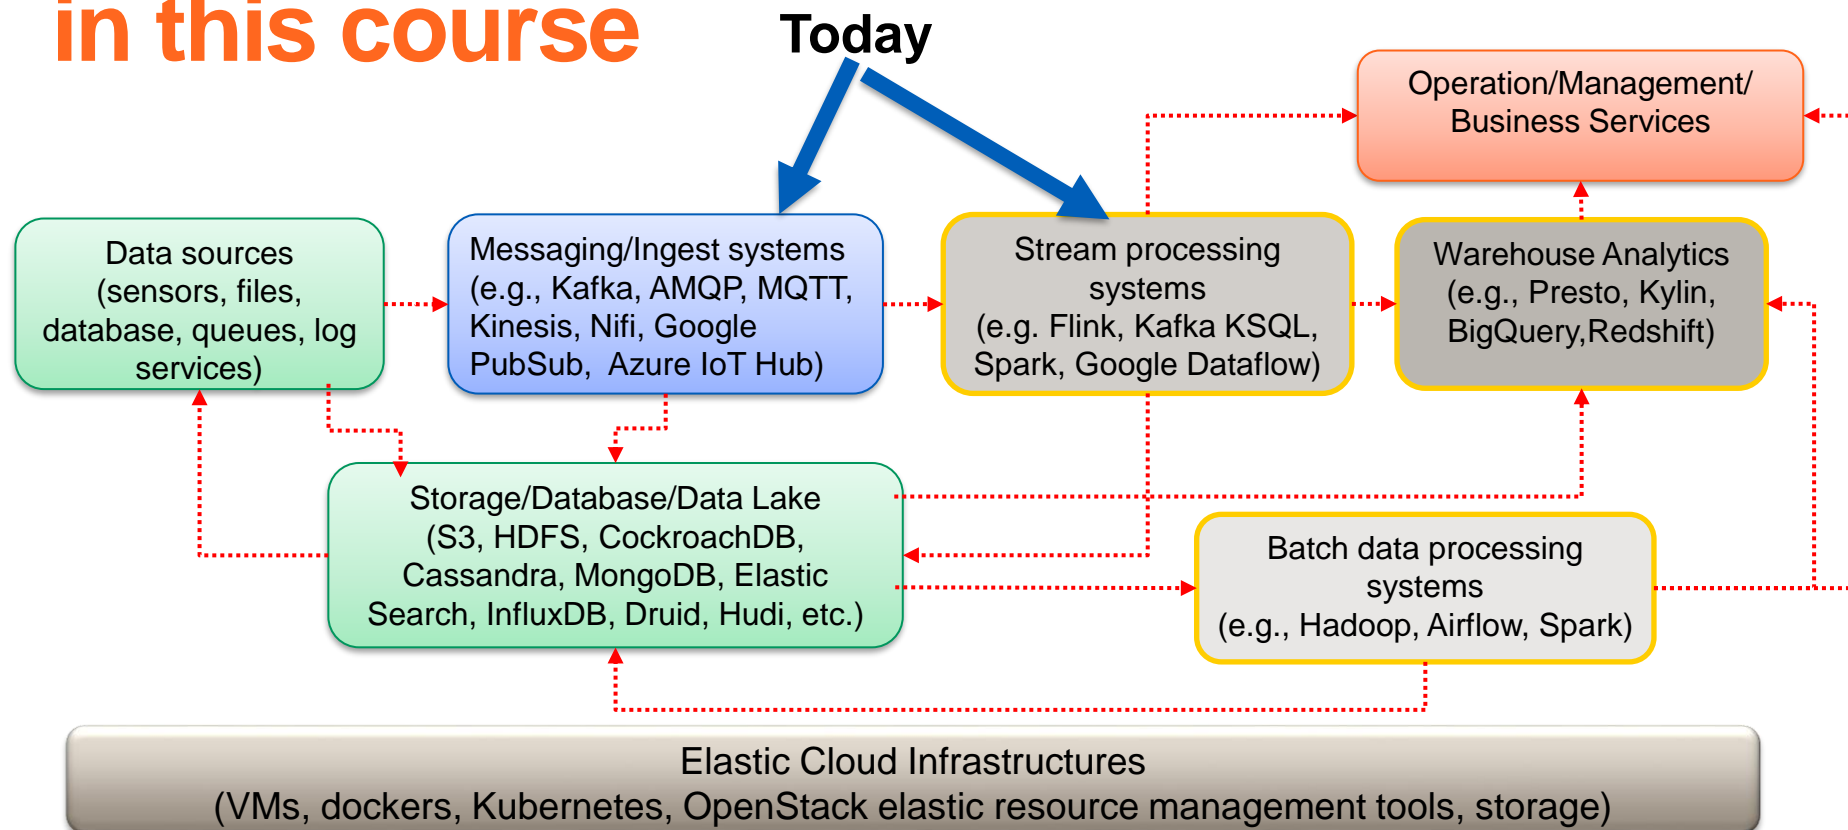
Department of Computer Science

linh.truong@aalto.fi, <https://rdsea.github.io>

Learning objectives

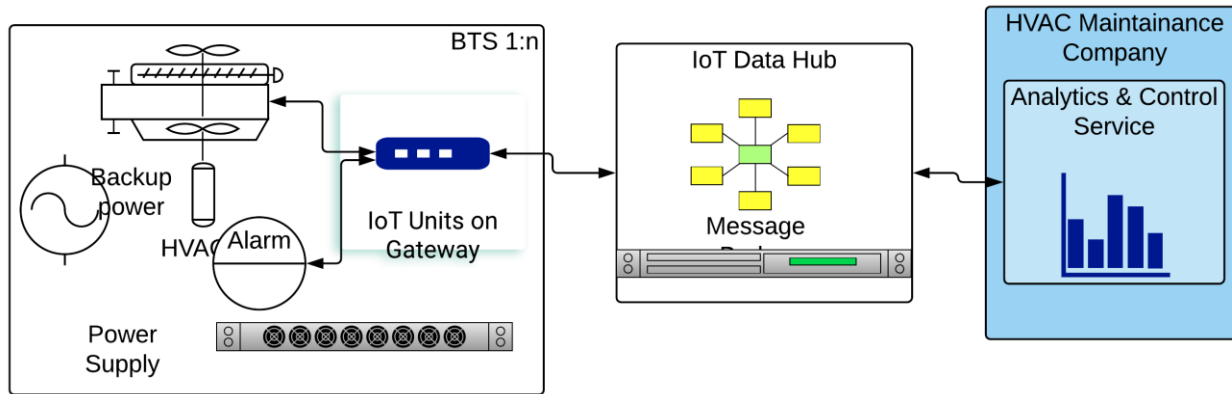
- **Understand fundamental concepts and techniques in stream processing in big data**
- **Able to design stream processing analytics in big data platforms and applications**
- **Able to select and use common stream processing frameworks**

Big data at large-scale: the big picture in this course



Motivating examples

Near real-time monitoring and anomaly detection for equipment and sites: what if you have 200K of Base Stations (BTS)



We need to analyze streaming data in a near-real time manner

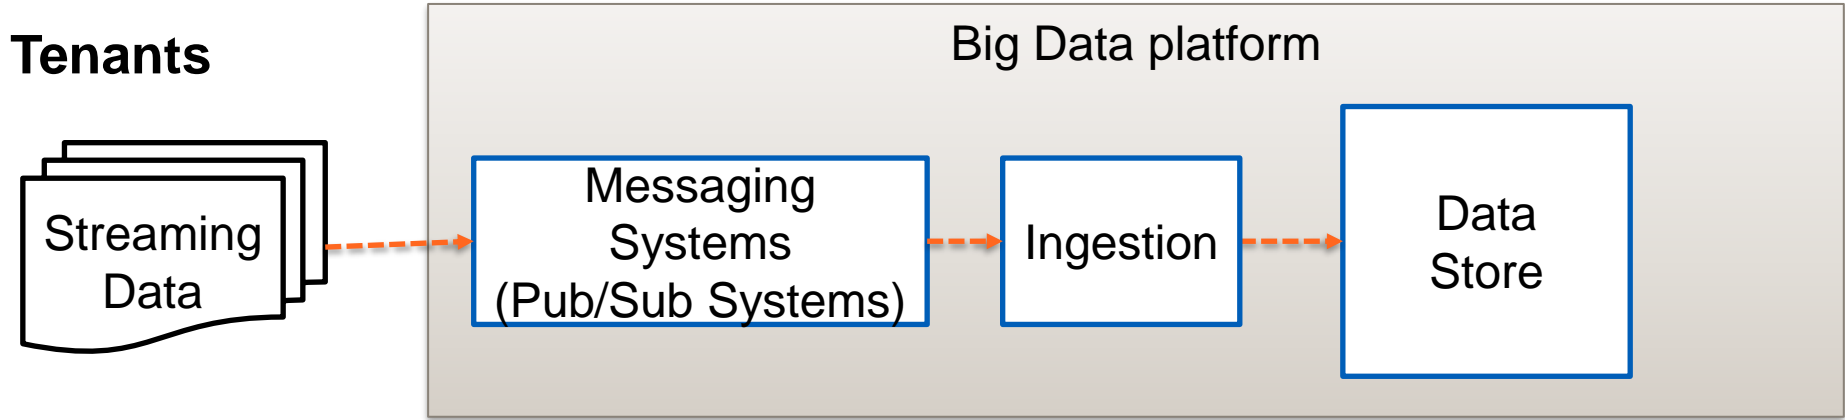
Many other scenarios: fraud detection in online payment, stock market monitoring, traffic detection, etc.



Aalto University
School of Science

Basic stream analytics for data in motion

Recall: near-real time streaming data ingestion



- **Ingestion**
 - Mostly we ingest raw data without/little processing, e.g., IoT data
 - Data is **unbounded** from different places in different orders!


Stream processing and big data platforms

- **Stream processing is a component of big data platforms**
 - A big data technology for pre-processing, ingestion and high-level analytics
- **Stream processing services as big data platforms**
 - a big data platform offers mainly stream processing services for streaming analytics
 - Analytics on the fly as the first class
 - *Historical analytics results as the second class*
 - E.g., IoT analytics, e-commerce user activities, fraud detection

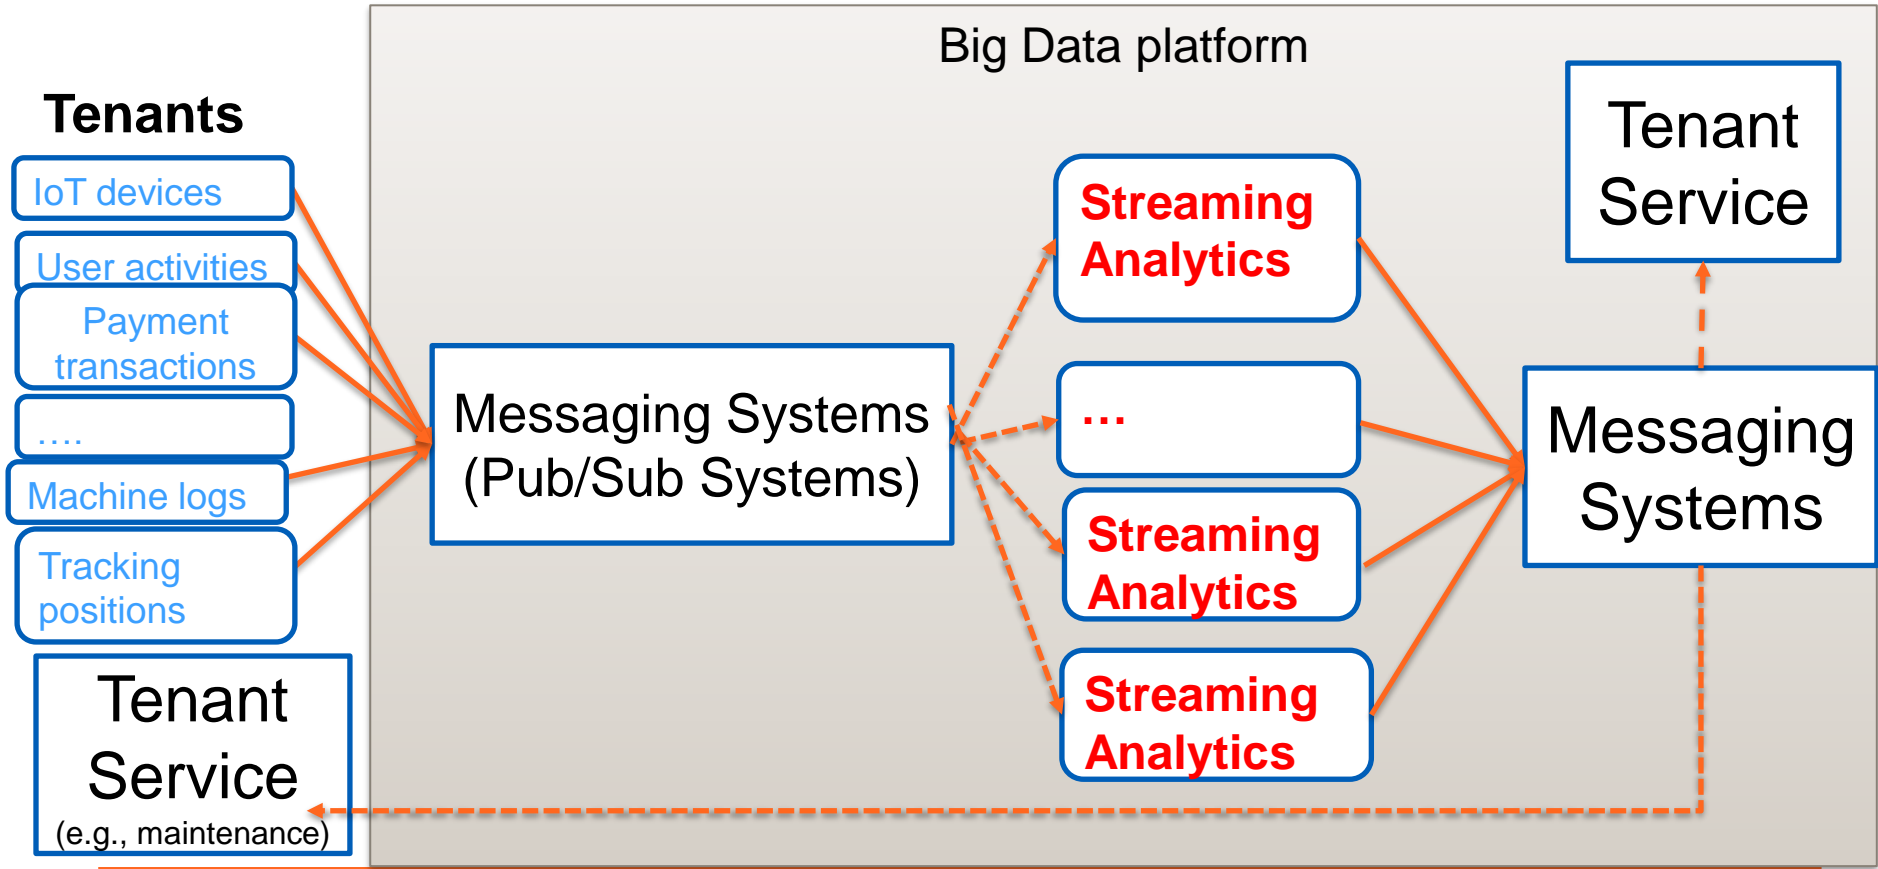
Stream processing in big data

- **Processing big data coming from streams at near real time**
 - The data element is “small” but voluminous and delivered in a near-real time manner
 - High and volatile throughput, low service latency
- **Require large-scale computing infrastructures and many other platform services**
 - Task parallelism: multiple tasks processing data
 - Data parallelism: data is partitioned into concurrent/parallel data streams → not like in data at rest

Near realtime streaming data processing



Near realtime streaming data processing



Example in the cloud – stream processing and big data platforms

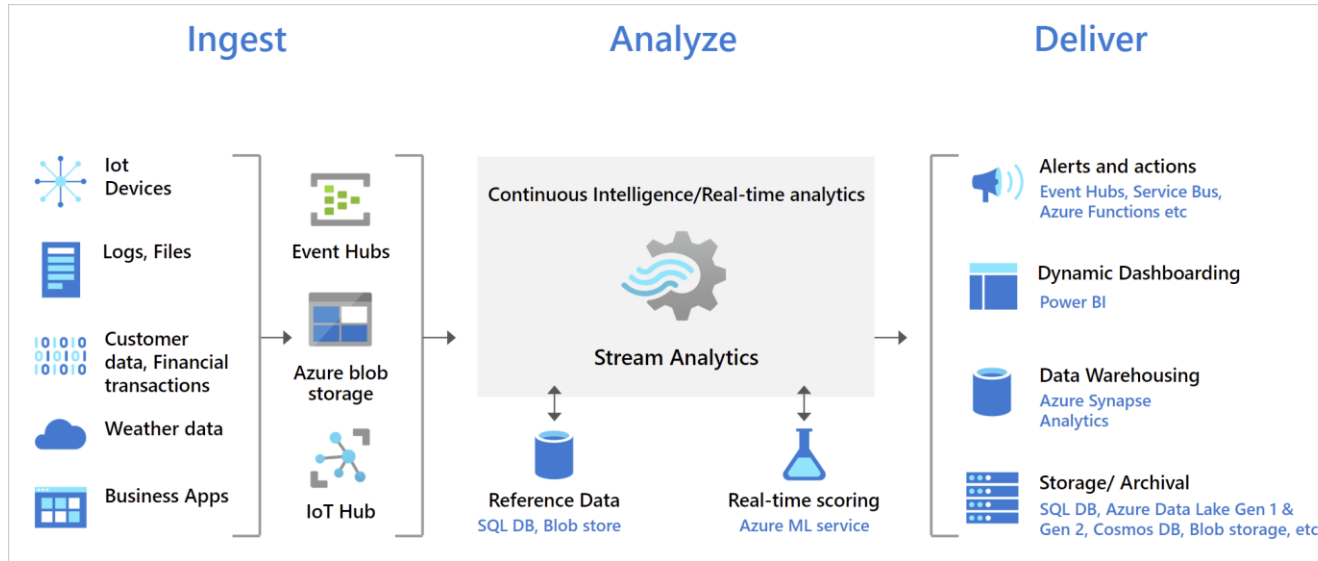


Figure source:

<https://docs.microsoft.com/en-us/azure/stream-analytics/stream-analytics-introduction>

Long history, e.g., complex event processing (CEP) from enterprise computing



Esper CEP



Apache Apex™



Our practices focus on Apache Flink, Apache Kafka and Apache Spark which are used intensively in business systems and big cloud platforms

Stream Processing – Key concepts

Common concepts

- **The way to connect data streams and obtain events/messages/records**
 - Focusing very much on *connector concepts* and well-defined event structures
 - The data can be pulled/pushed via connectors
- **The way to specify/programming the “analytics” logic**
 - Analytics functions, statements and how they are glued together to process flows of events
 - High-level, easy to use
- **The engine to process analytics tasks/operators**
 - Centralized in the view of the user → so the user does not have to program complex distributed applications
- **The way to push results to external components**

Data stream programming

Data stream: a sequence/flow of data units


Data units are defined by applications: a data unit can be data described by a primitive data type or by a complex data type, a serializable object, etc.

Streaming data: produced by (near)realtime data sources as well as (big) static data sources → unbounded and bounded

- Examples of data streams
 - Continuous media (e.g., video for video analytics)
 - Discrete media (e.g., stock market events, twitter events, system monitoring events, comments, notifications)

Events/Records

In many applications: data is generated continuously and needs to be processed in near real-time




We focus on **unbounded discrete** events/records/messages

Event/record representation and streams

- **Data Sources:**
 - via message brokers, database, websocket, different IO adapters/connectors, etc.
- **Data Sink**
 - Message systems, databases, files, etc.
- **Data Representation & views**
 - POJO (Plain Old Java Object), CSV, JSON, Arvo format, etc.
 - SQL-alike tables

Stream processing

High level view

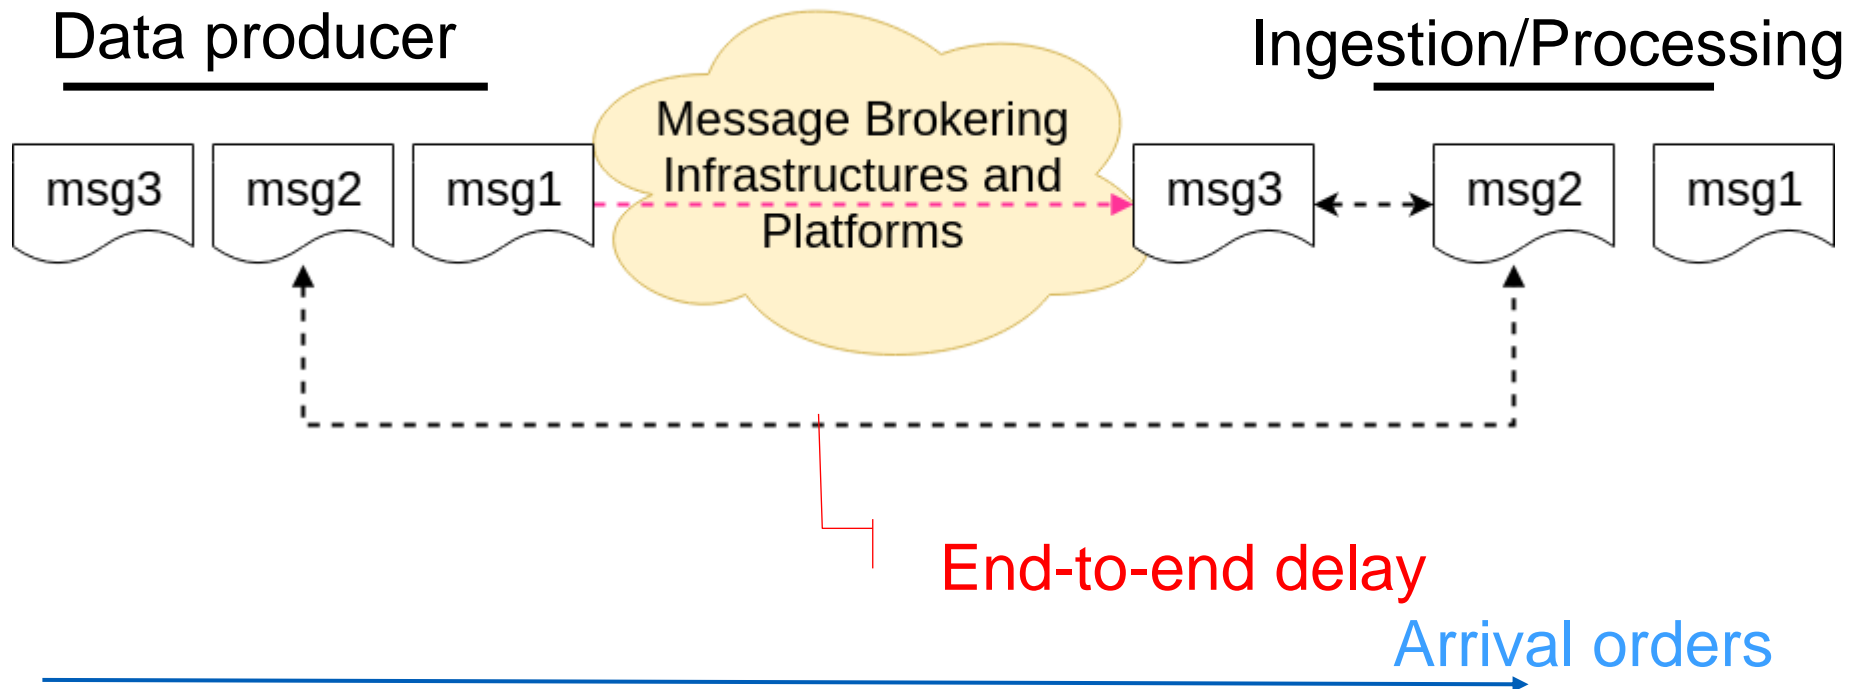


Multiple streams, a set of events

Some key issues

- **Data order & delivery**
 - Late data, out of order data
- **Times associated with events and processing**
- **Data parallelism**
 - Key-based data processing
- **Task parallelism**
 - Stateful vs stateless processing

Key issues in streaming data: delay and out of order

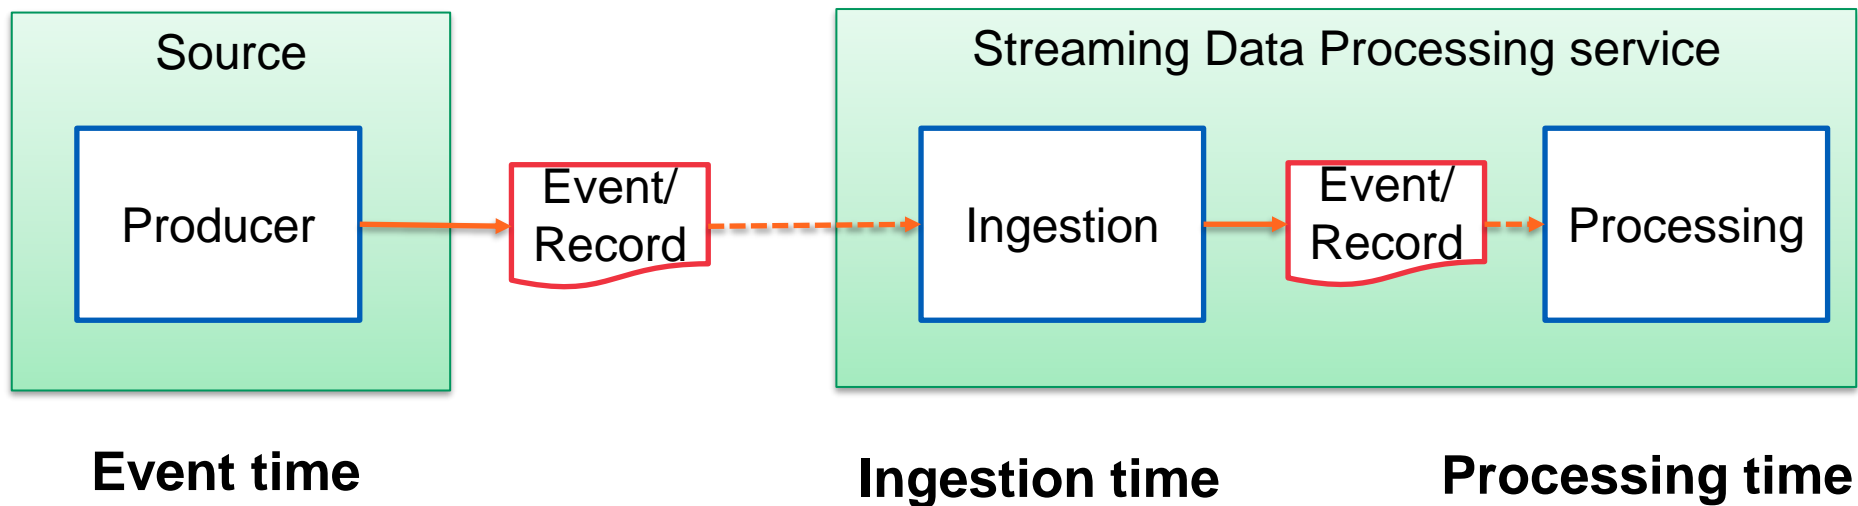


Without event/record time, do we know the delay or out of order?

What is the consequence of delay/out of order for processing?

Key issues in streaming data: the notion of times

Times associated with data and processing




Which time is important for analytics (from business viewpoint)?

Data parallelism: partition stream data based on some keys for analytics

Split based on keys?

One-record vs batch of records




With **keyed data**: enable parallel processing based on the keys

Windows of data

Window is used to group data for processing:

Which constraints are used to determine a window?



a stream of events

Sliding/Tumble window size: period of time or number of events/records

Arrival order

Windowing

- **Windows size: time or number of records (not popular)**
- **Tumbling window:**
 - identified by size, no gap between windows
- **Sliding window:**
 - identified by size and a sliding interval
- **Session Window:**
 - identified by “gap” between windows (e.g., the gap of events is used to mark “sessions”)

Functions applied to Windows of data

If we

- **specify a set of conditions** for the window and events within the window

then we can

- **Apply functions to events in** the window that match these conditions

Task parallelism: we can have a lot of such functions executed in parallel in multiple compute nodes

Example

Monitoring working hours of (taxi/truck) drivers (assume events about pickup/drop captured at near realtime):

- **Windows: 12 hours**
- **Partitioning data/Keyed streams: licenseID**
- **Function: determine working and break times and check with the law/regulation**

Source: <https://www.infoworld.com/article/3293426/how-to-build-stateful-streaming-applications-with-apache-flink.html>


What if events/records come late into the windows?

Do we need to deal with late, out of order events/records?

correctness and completeness issues


Support lateness

- Identify timestamp of events/records
- Identify watermark in streams
 - A watermark is a timestamp
 - A watermark indicates that no events which are older than the watermark should be processed
 - Enable the delay of processing functions to wait for late events
- Using watermark to ignore late data → computing under “incompleteness assumption”



Delivery guarantees

Exactly once? at least once? or at-most-once
End-to-end?



What if the stream processing fails and restarts

Message and processing guarantees

- **Message guarantees are the job of the broker/messaging system**
- **Processing guarantees are the job of the stream processing frameworks**
- **They are highly connected if messaging systems and processing frameworks are tightly coupled (e.g., Kafka case)**

End-to-end exactly once

- **Exactly once for processing is not enough**
- **Messaging systems must support**
 - redeliver messages/data, recoverable data
- **Sink and output must support exactly one**
 - idempotent results, roll back
- **Coordination among various components**



Further reading:

<https://flink.apache.org/features/2018/03/01/end-to-end-exactly-once-apache-flink.html>

<https://www.confluent.io/blog/simplified-robust-exactly-one-semantics-in-kafka-2-5/>

<https://docs.microsoft.com/en-us/azure/hdinsight/spark/apache-spark-streaming-exactly-once>

Performance metrics



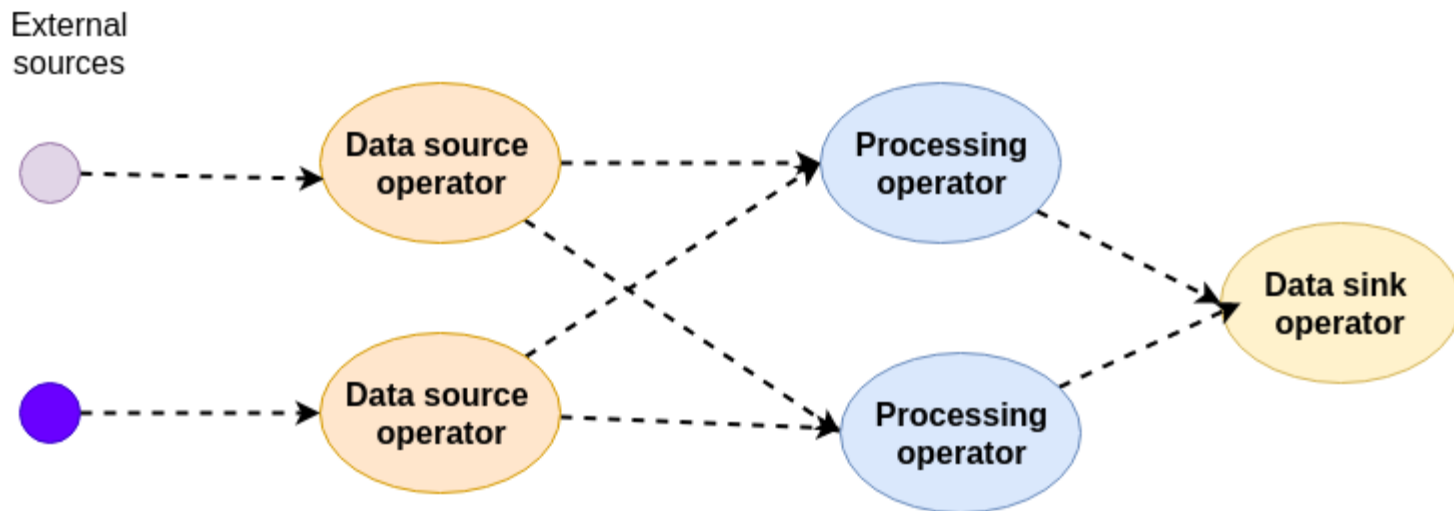
Latency and Throughput

- **Service latency**
 - Subseconds! E.g., milliseconds
 - Max, min or percentile? → up to application requirements
- **Throughput**
 - How many events can be processed per second?
- **Goal: low latency and high throughput!**

Structure of streaming data processing programs (1)

- **We have multiple streams of data, different functions for processing data, multiple computing nodes**
- **Data exchange between tasks**
 - Links in task graphs reflect data flows
- **Stream processing**
 - Centralized or distributed (in terms of computing resources)

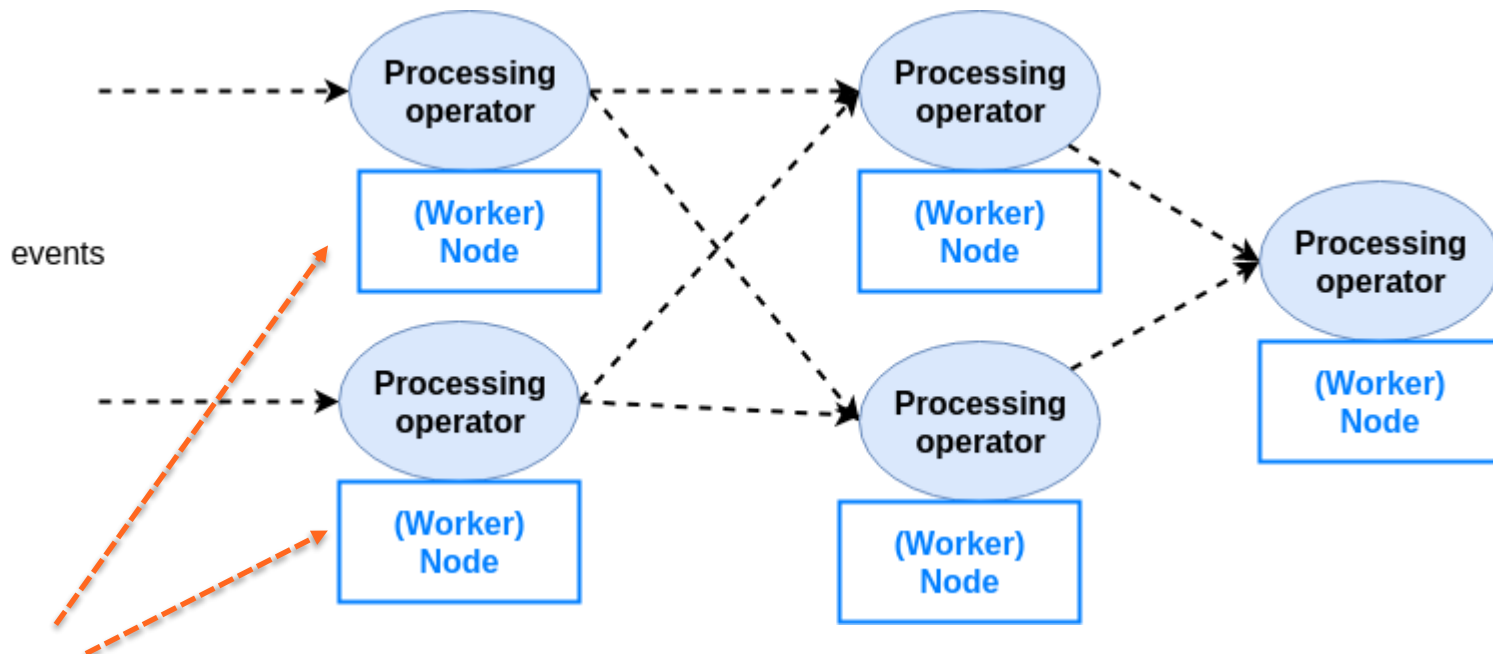
Structure of streaming data processing programs (2)



- **Data source operator:** represents a source of streams
- **Processing operators:** represents processing functions
- ***Native versus micro-batching***

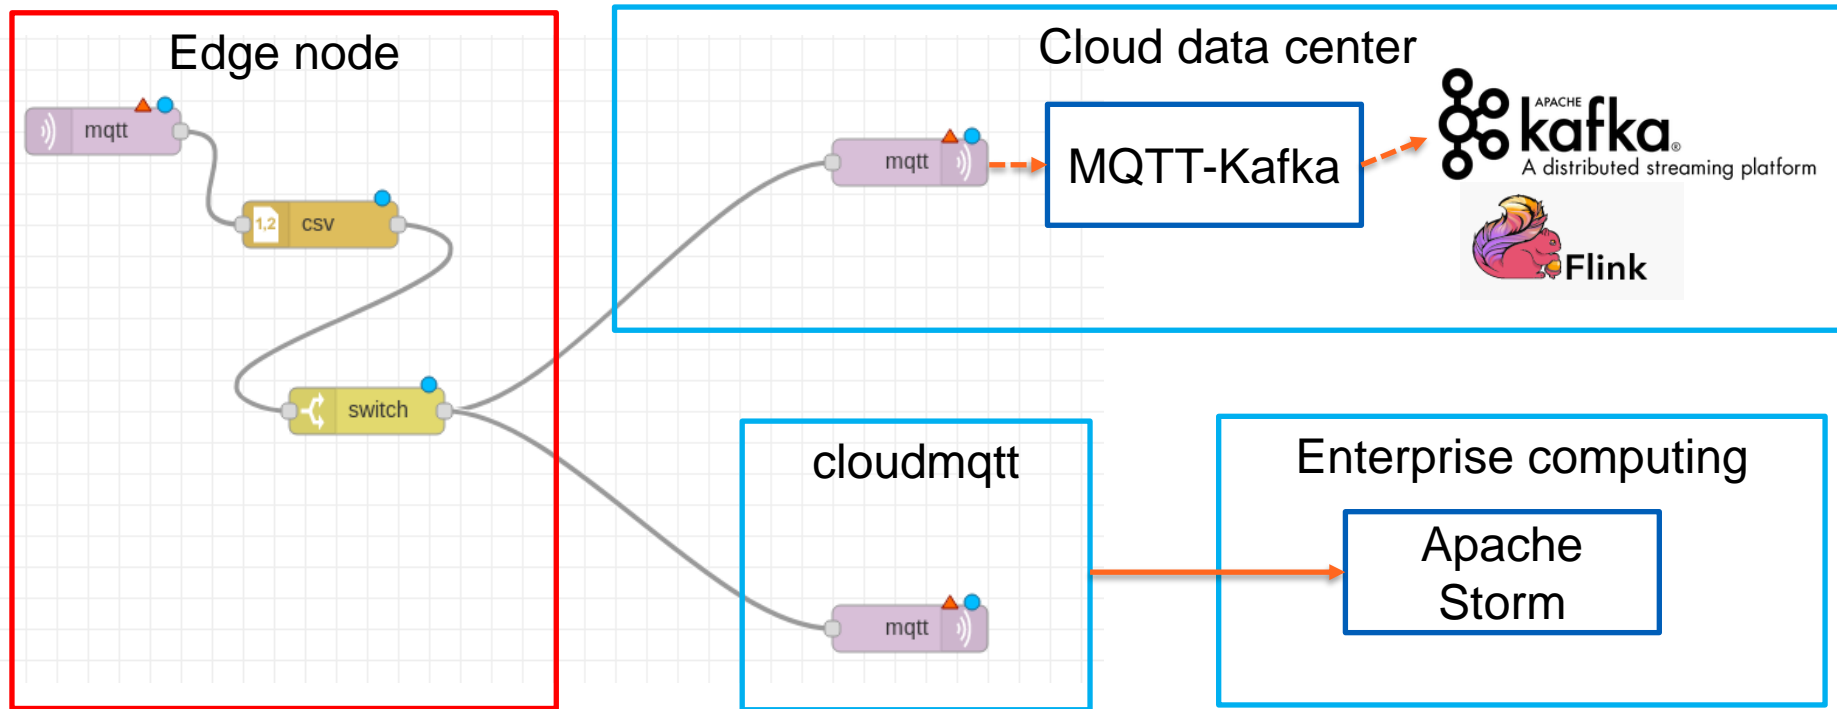
Distributed processing topology in a cluster

A graph of tasks (running operators); all tasks are running



Nodes of a cluster (VMs, containers, Kubernetes)

Distributed processing topology in cross distributed sites



Common concepts in existing frameworks - programming level

- **How to write streaming program?**
- **With programming languages**
 - Low level APIs
 - DSL
 - Java, Scala, Python (Spark, Flink, Kafka)
- **High-level data models**
 - KSQL
- **Flow/pipeline description**
 - Node-RED/GUI-based flow editors

Common concepts in existing frameworks - key common concepts

- **Abstraction of streams**
- **Connector library for data sources/sinks**
 - Very important for application domains
- **Runtime elasticity**
 - Add/remove (new) operators
 - Add/remove underlying computing nodes
- **Fault tolerance**

Why are the richness and the diversity of connectors important?

Where do you find most of concepts that we have discussed

- **Apache Storm**
 - *<https://storm.apache.org/>*
- **Apache Spark (Structured Streaming)**
 - *<https://spark.apache.org/>*
- **Apache Kafka Streams and KSQL**
 - strongly bounded to Kafka messaging
- **Apache Flink (Stream Analytics)**
 - native, clustered, better data sources/sinks
- **Apache Beam (<https://beam.apache.org/>)**
 - Unifying programming models for batch and stream processing

Practical learning paths

- **Path 1: if you don't have a preference and need challenges**
 - Apache Flink Stream API (e.g., with RabbitMQ/Kafka connectors)
- **Path 2: many of you have worked with Kafka**
 - Kafka Streams DSL (everything can be done with Kafka)
- **Path 3: for those of you who are working with Spark (and Python is the main programming language)**
 - Apache Spark Structured Streaming
- **Path 4: for those who deal with MQTT brokers**
 - Apache Storm (but also Kafka, ...): Spout and Bolt API or Stream API

Examples of Apache Flink

Apache Flink


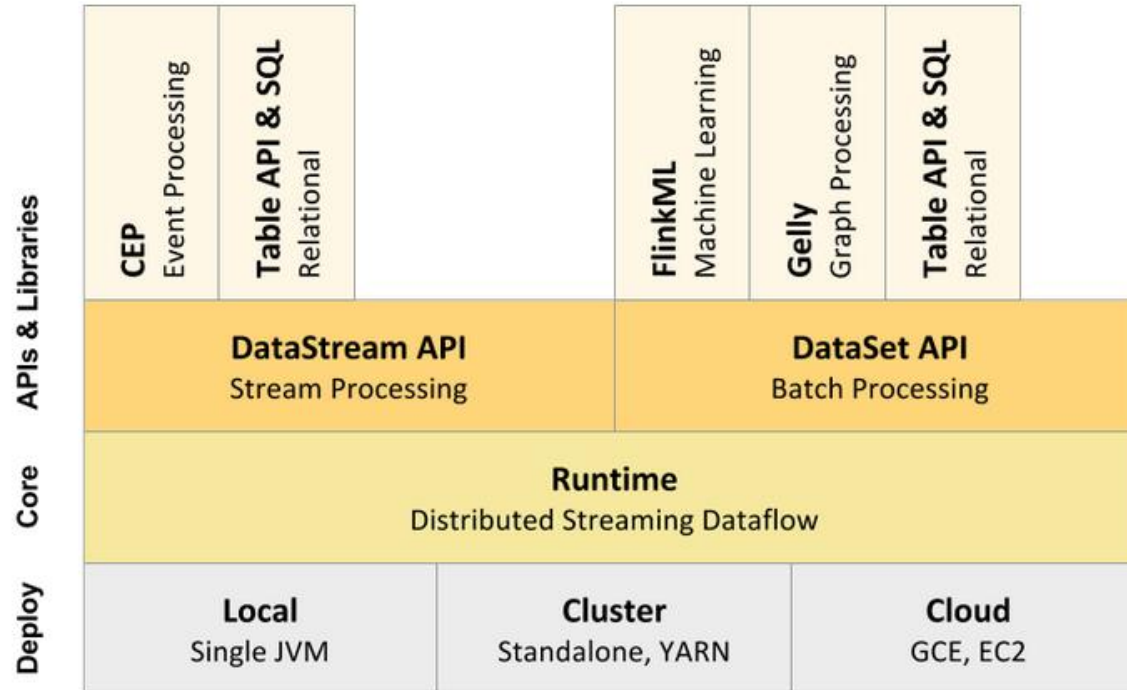


Figure source: <https://flink.apache.org/>

Apache Flink

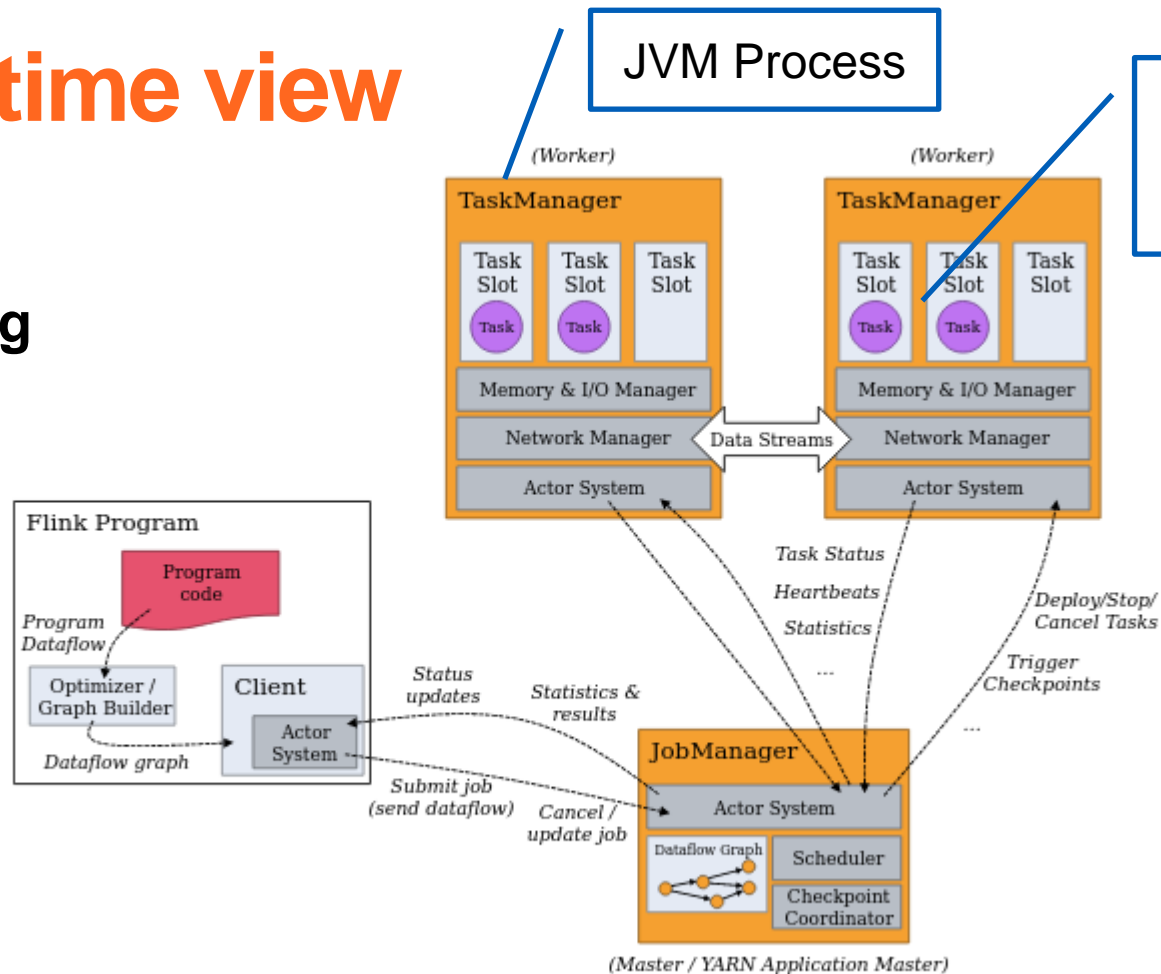


We focus only on
DataStream API for
understanding
studied concepts

Source: <https://ci.apache.org/projects/flink/flink-docs-release-1.9/internals/components.html>


Flink runtime view

- Parallelism
- Checkpointing
- Monitoring



Source <https://nightlies.apache.org/flink/flink-docs-release-1.14/docs/concepts/flink-architecture/>

Main elements in Flink applications



- Rich set of sources and sinks via many connectors

Connectors

- **Major systems in big data**
- **We have used many of them in our study**
 - Apache Kafka
 - Apache Cassandra
 - Elasticsearch (sink)
 - Hadoop FileSystem
 - RabbitMQ
 - Apache NiFi
 - Google PubSub

Main

- **Setting environments**
- **Handling inputs and outputs via data streams**
- **Key functions for processing data**
- **Stream processing flows**



Bounded and unbounded streams

Stream processing flows

Split streaming data into different windows with a key for analytics purposes

Keyed Windows

```
stream
  .keyBy(...)          <- keyed versus non-keyed windows
  .window(...)         <- required: "assigner"
  [.trigger(...)]      <- optional: "trigger" (else default trigger)
  [.evictor(...)]      <- optional: "evictor" (else no evictor)
  [.allowedLateness(...)] <- optional: "lateness" (else zero)
  [.sideOutputLateData(...)] <- optional: "output tag" (else no side output for late data)
  .reduce/aggregate/fold/apply() <- required: "function"
  [.getSideOutput(...)] <- optional: "output tag"
```

Source: <https://ci.apache.org/projects/flink/flink-docs-release-1.12/dev/stream/operators/windows.html>

Stream processing flows

Handling streaming data without a key for analytics purposes


Non-Keyed Windows

```
stream
  .windowAll(...)      <- required: "assigner"
  [.trigger(...)]      <- optional: "trigger" (else default trigger)
  [.evictor(...)]      <- optional: "evictor" (else no evictor)
  [.allowedLateness(...)] <- optional: "lateness" (else zero)
  [.sideOutputLateData(...)] <- optional: "output tag" (else no side output for late data)
  .reduce/aggregate/fold/apply() <- required: "function"
  [.getSideOutput(...)] <- optional: "output tag"
```

Source: <https://ci.apache.org/projects/flink/flink-docs-release-1.9/dev/stream/operators/windows.html>

Windows and Times

Windows



Times

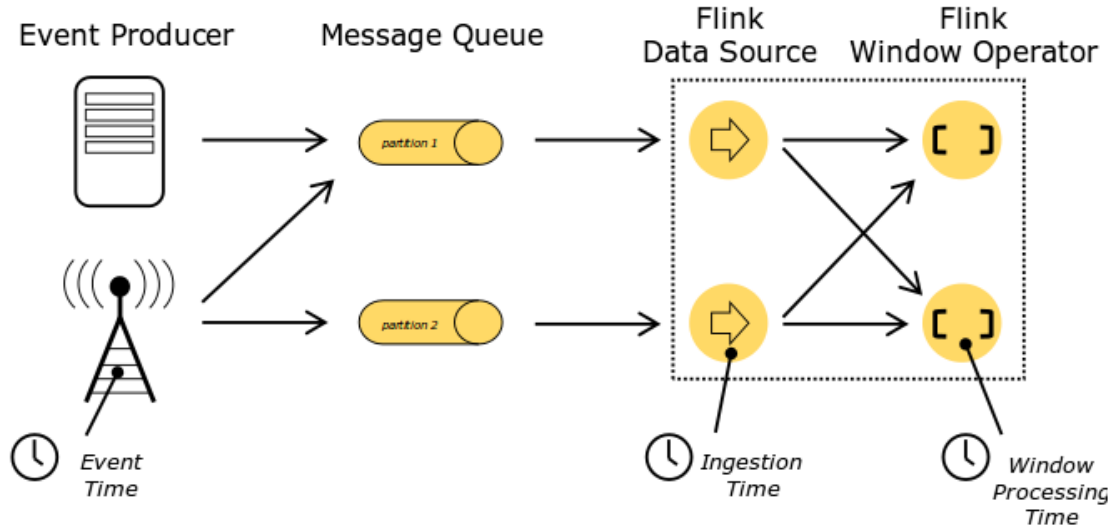



Figure source: <https://nightlies.apache.org/flink/flink-docs-release-1.14/docs/concepts/time/>


Batch/Tumbling Windows



Use cases:
Period computation
(e.g. stock,
temperature, IoT
data)

Figure source: <https://nightlies.apache.org/flink/flink-docs-release-1.14/docs/dev/datastream/operators/windows/>

Sliding windows



Use cases:
Moving average

Figure source: <https://nightlies.apache.org/flink/flink-docs-release-1.14/docs/dev/datastream/operators/windows/>

Session Windows

Use cases:
Web/user activities
clicks




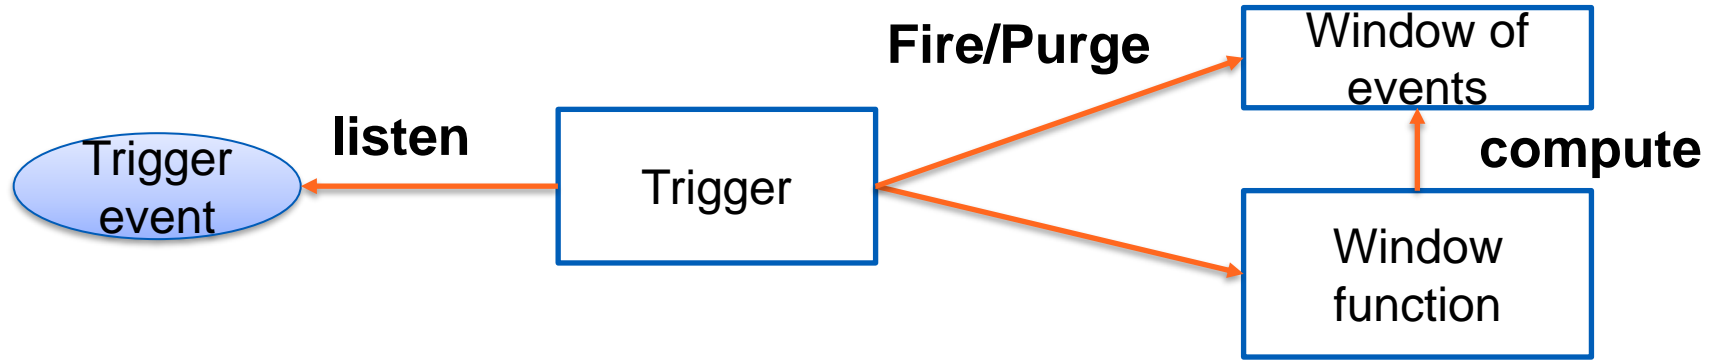
Figure source: <https://nightlies.apache.org/flink/flink-docs-release-1.14/docs/dev/datastream/operators/windows/>

Window Functions

- **Reduce Function**
 - Reduce two inputs
- **Aggregate Function**
 - Add an input into an accumulator
- **Fold Function**
 - Combine input with an output
- **ProcessWindow Function**
 - Get all elements of the windows and many other information so that you can do many tasks

Triggers & Evictor

- **Trigger:** determine if a window is ready for window functions



Evictor: actions **after** the trigger fires and **before** **and/or after** the windows function is called

Fault tolerance

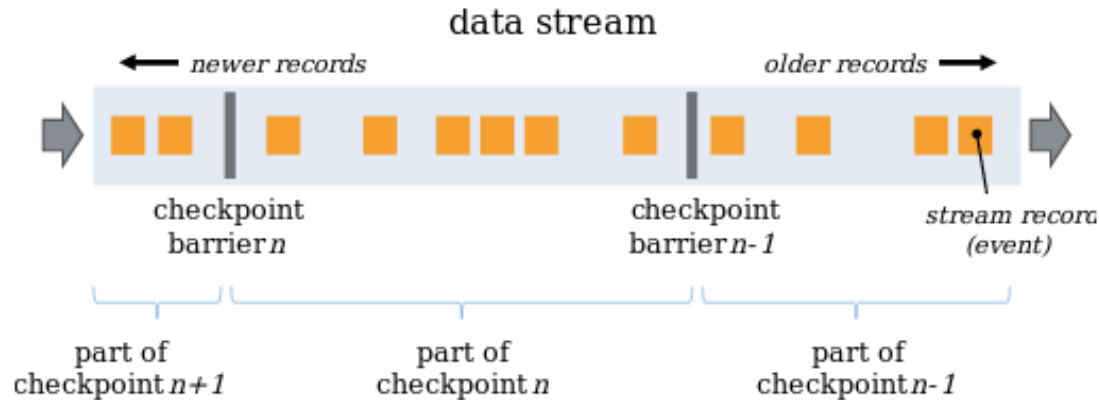


Figure source: https://nightlies.apache.org/flink/flink-docs-master/docs/learn-flink/fault_tolerance/

- **Principles: checkpointing, restarts operators from the latest successful checkpoints**
- **Need support from data stream sources/sinks w.r.t. (end-to-end) exactly once message receiving and result delivery**

Example with Base Transceiver Station

Data in our git


```
station_id,datapoint_id,alarm_id,event_time,value,valueThreshold,isActive,storedtime
1161115016,121,308,2017-02-18 18:28:05 UTC,240,240,false,
1161114050,143,312,2017-02-18 18:56:20 UTC,28.5,28,true,
1161115040,141,312,2017-02-18 18:22:03 UTC,56.5,56,true,
1161114008,121,308,2017-02-18 18:34:09 UTC,240,240,false,
1161115040,141,312,2017-02-18 18:20:49 UTC,56,56,false,
```

Simple example

See the code in our git:

<https://version.aalto.fi/gitlab/bigdataplatforms/cs-e4640/-/blob/master/tutorials/streamingwithflink/>

Simple example



Monitoring

Apache Flink Dashboard

Overview

Jobs

Running Jobs

Completed Jobs

Task Managers

Job Manager

Submit New Job

Simple CS-E4640 BTS Flink Application RUNNING 2

ID: 81efb959d02448b7c44328ad75a824af | Start Time: 2019-11-04 14:00:14 | Duration: 55s

[Overview](#) | [Exceptions](#) | [TimeLine](#) | [Checkpoints](#) | [Configuration](#)

Source: Custom Source -> Flat Map
Parallelism: 1

HASH

Window(SlidingProcessingTimeWindows(60000, 5000), ProcessingTimeTrigger, MyProcessWindowFunction) -> (Sink: Unnamed, Sink: Print to Std. Out)
Parallelism: 1

Detail

SubTasks

TaskManagers

Watermarks

Accumulators

BackPressure

Metrics

Window(SlidingProcessingTimeWindows(60000, 5000), ProcessingTimeTrigger, MyProcessWindowFunction) -> (Sink: Unnamed, Sink: Print to Std. Out)

Status: RUNNING

Task: 1

Parallelism: 1

Records Sent: 0

Start Time: 2019-11-04 14:00:14

Bytes Received: 1.64 KB

End Time: -

Records Received: 29

Duration: 55s

Bytes Sent: 0 B

Name	Status	Bytes Received	Records Received	Bytes Sent	Records Sent	Parallelism	Start Time	Duration	End Time	Tasks
Window(SlidingProcessingTimeWindows(60000, 5000), ProcessingTimeTrigger, MyProcessWindowFunction) -> (Sink: Unnamed, Sink: Print to Std. Out)	RUNNING	1.64 KB	29	0 B	0	1	2019-11-04 14:00:14	55s	-	1
Source: Custom Source -> Flat Map	RUNNING	0 B	0	1.61 KB	29	1	2019-11-04 14:00:14	55s	-	1

Cancel Job

One of the successful projects from Europe: originally from TU Berlin

Alibaba cloud:

“Flink can process over **472 million transactions per second during business peaks, which is truly astronomical”**

Source: https://www.alibabacloud.com/blog/why-did-alibaba-choose-apache-flink-anyway_595190

Four Billion Records per Second! What is Behind Alibaba Double 11 - Flink
Stream-Batch Unification Practice during Double 11 for the Very First Time

Alibaba Clouder December 2, 2020 👁 2,607 💬 0

Source: <https://bit.ly/3sUTwqT>

“Amazon Kinesis Data Analytics is the easiest way to transform and analyze streaming data in real time with Apache Flink” (From <https://aws.amazon.com/kinesis/data-analytics/>)

Interested in a use case?

Check:

https://www.alibabacloud.com/blog/application-of-real-time-compute-for-apache-flink-in-weibo_597955

Summary

- **Focus:**
 - Practical programming with one of the stacks:
 - *Apache Flink Stream API (with different connectors)*
 - *Apache Spark*
 - *Kafka Streams*
 - Check the common concepts in other tools/systems
- **Action:**
 - Work on use cases where you can use stream analytics (as a user/developer) → there are many interesting analytics
 - Provision services for stream processing (as a platform)

Thanks!

Hong-Linh Truong
Department of Computer Science

rdsea.github.io