



Aalto University
School of Science

Hadoop and its Big Data Ecosystems

Hong-Linh Truong

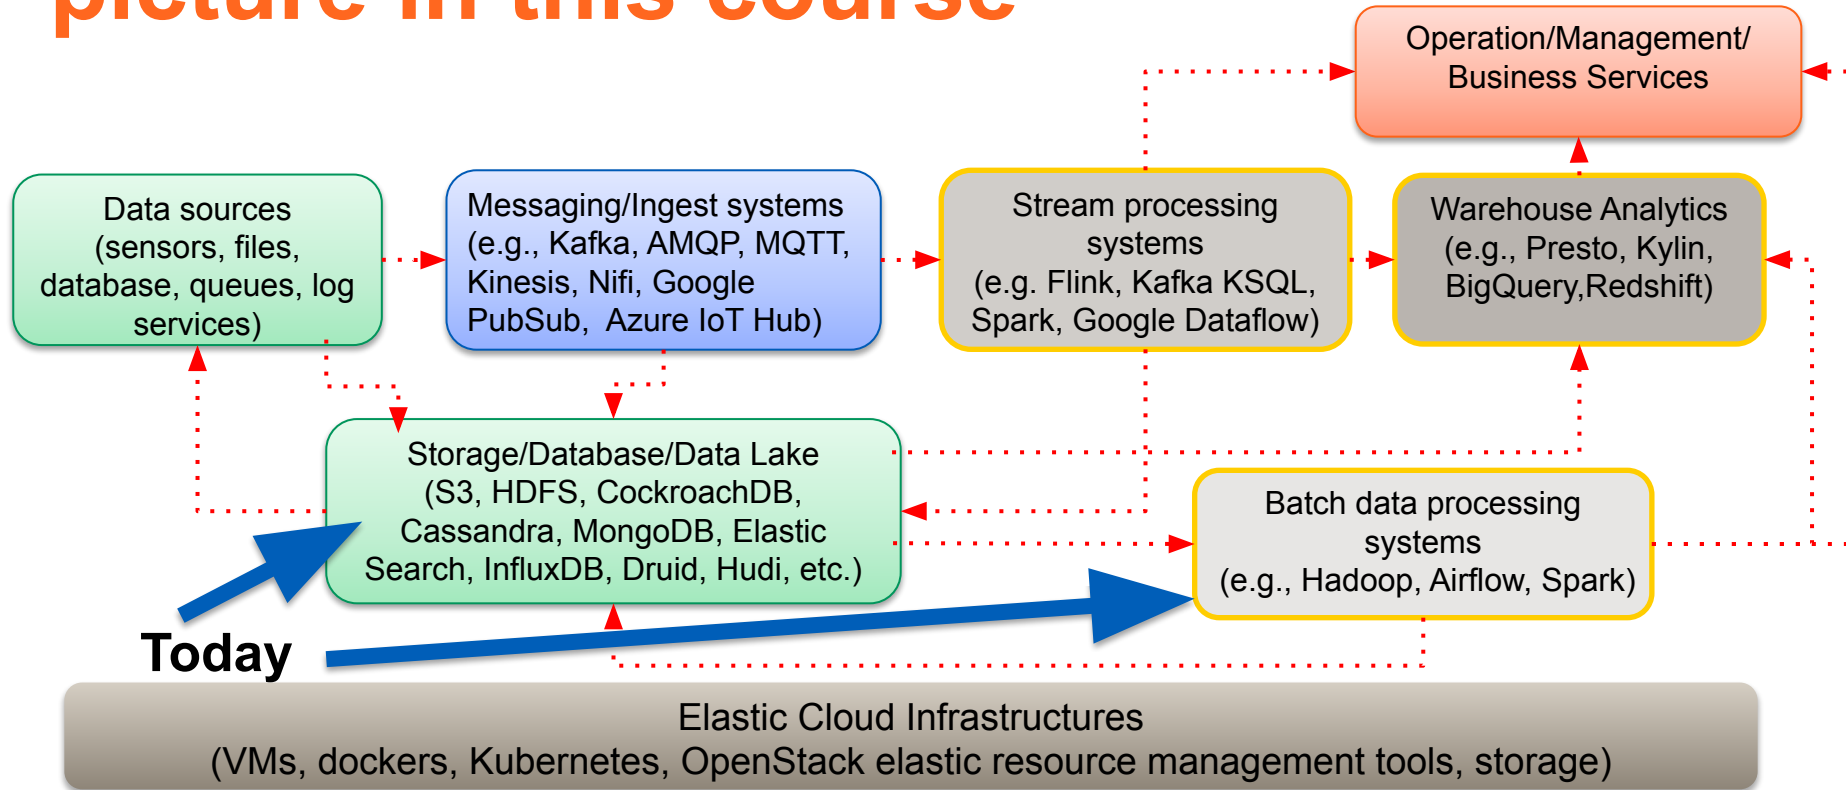
Department of Computer Science

linh.truong@aalto.fi, *<https://rdsea.github.io>*

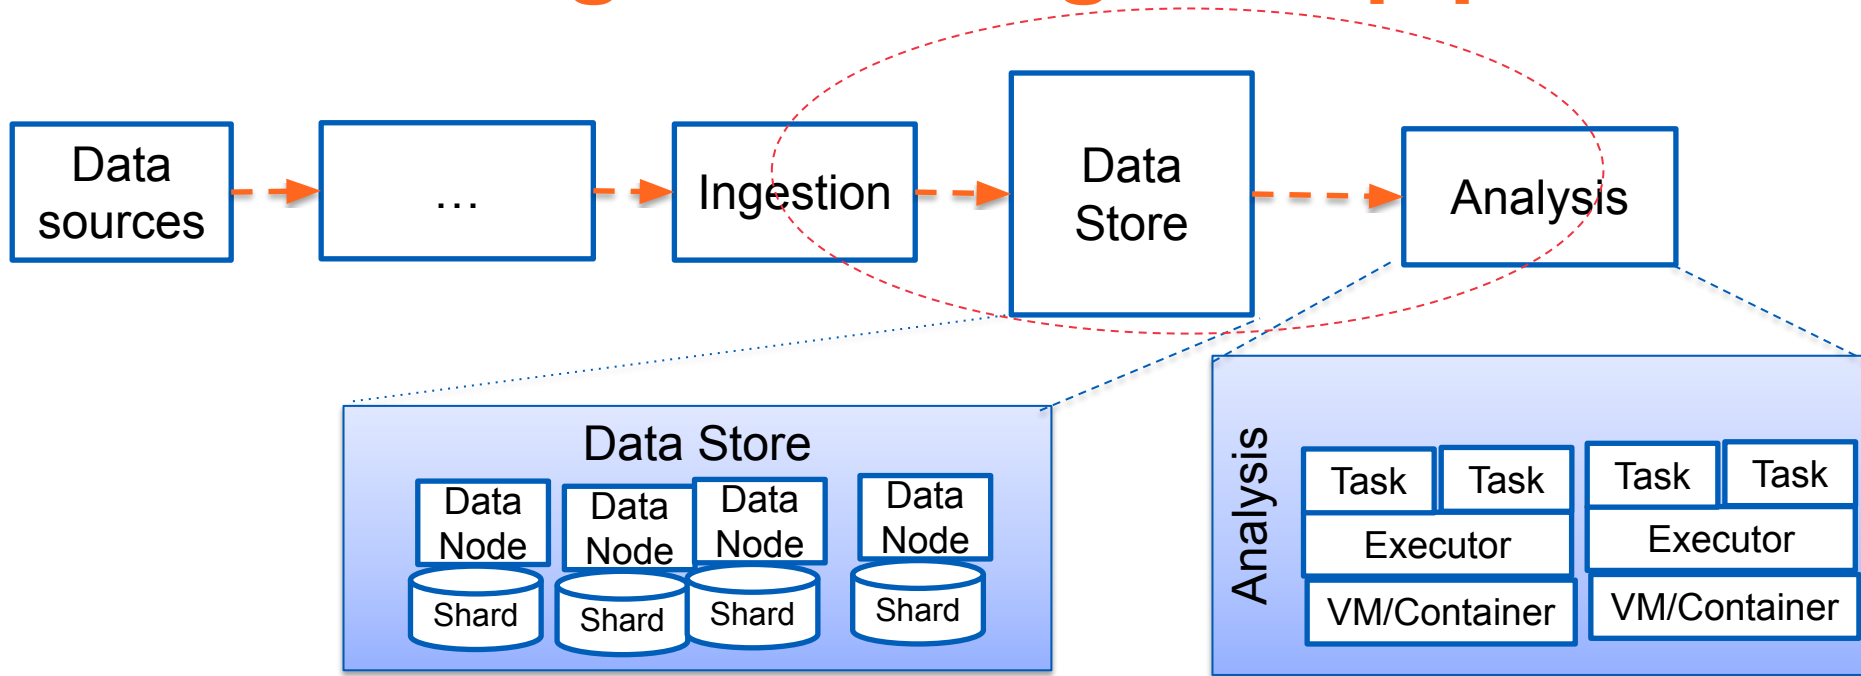
Learning objectives

- Understand massive scale data management and processing with Hadoop
- Understand and able to use Hadoop components for big data platform designs
- Able to integrate Hadoop with other frameworks for data ingestion and analytics systems

Big data at large-scale: the big picture in this course

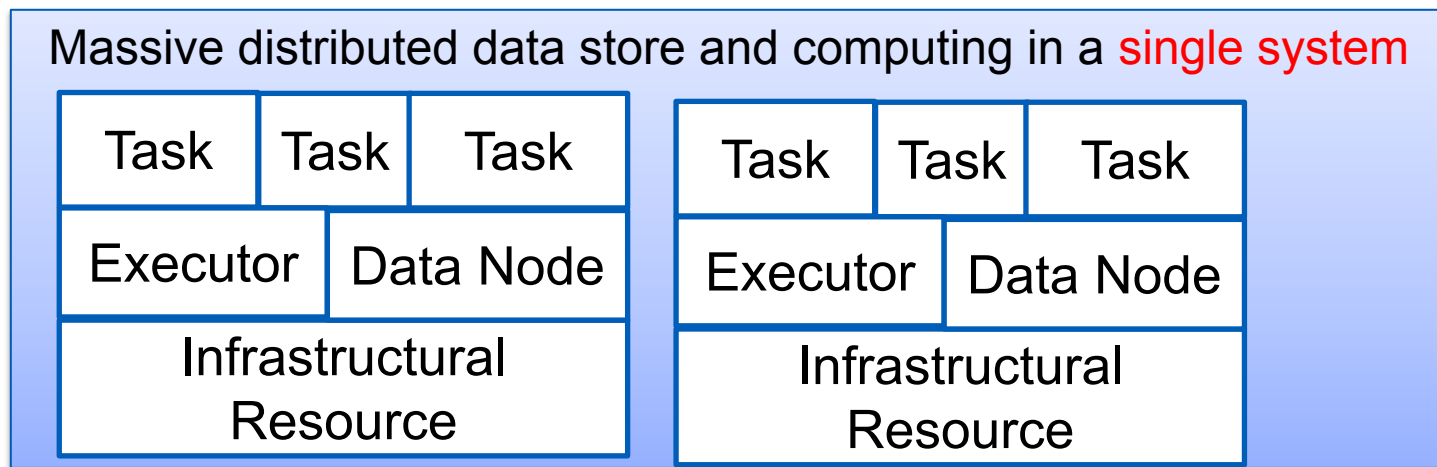


Consider again our big data pipelines



Combining data storage and analysis

Data locality, massive parallel and distributed data processing with large-scale data store




Single system: **many machines connected into a type of “cluster”**

Shared nothing architectures

- **Machine as a node:**
 - completely independent from other nodes
 - no shared physical resources (CPUS, memory, hard disk, ...)
- **A distributed system of machines with shared nothing architectures**
 - allow to scale and extend physical resources easily
 - enable high fault tolerance designs
 - quite common and easy-to-be-setup for enterprises
 - *think about commodity clusters*

Distributed Big Data in Shared Nothing Architectures



Benefit

- Reduce cost and increase performance when moving data is much more expensive than moving computation
- Consolidate and integrate various types of data
- Save shared infrastructural resources
- Provide suitable solutions for different types of customers/clients and different infrastructures

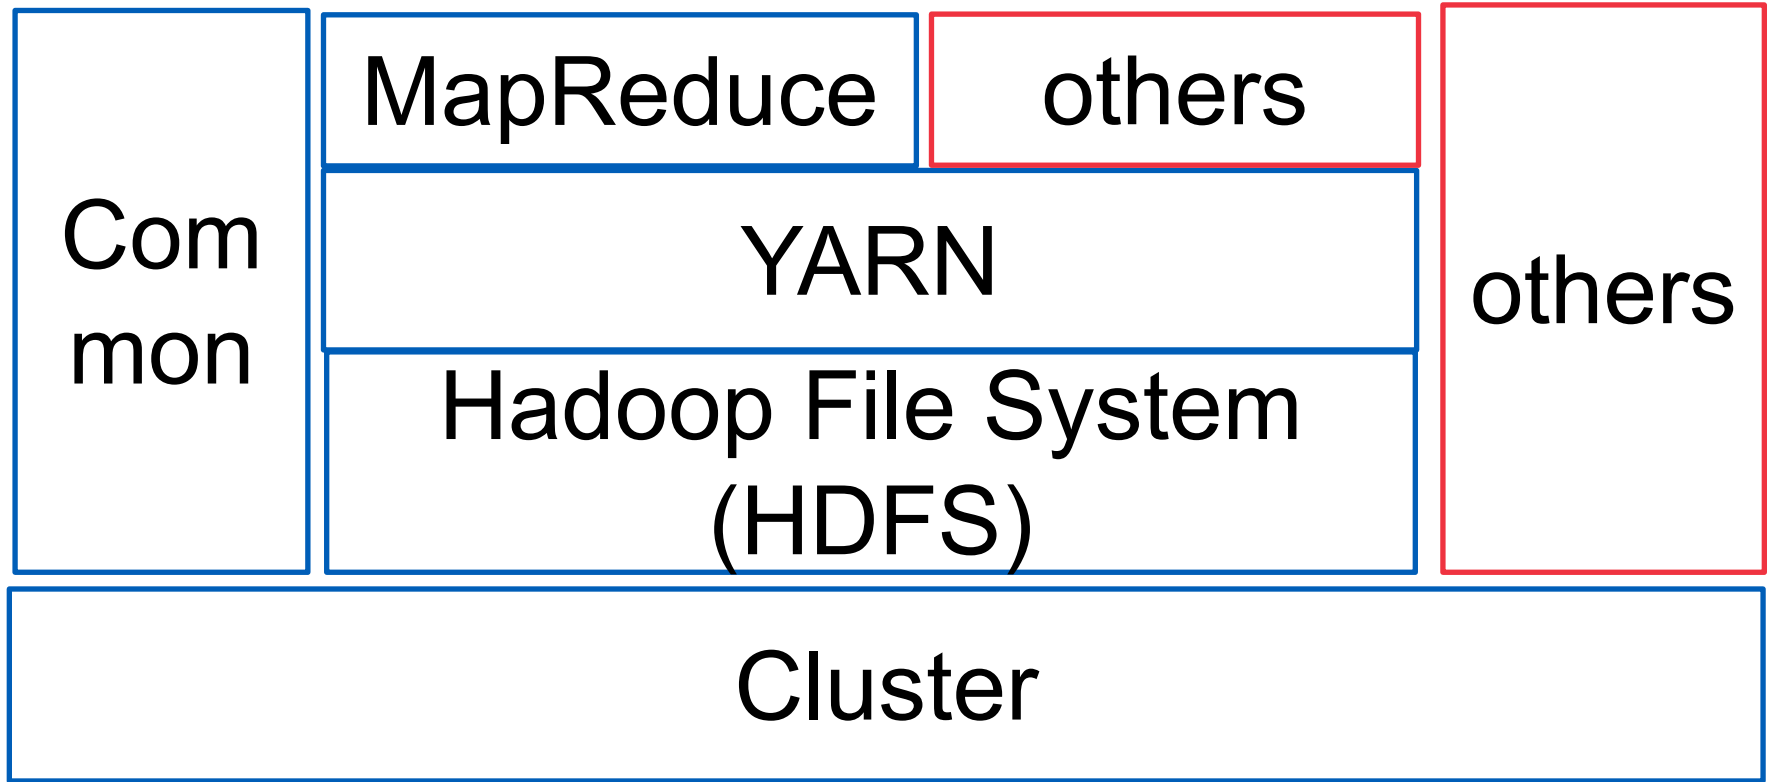
Benefit - Variety

- **Multiple data formats**
- **In big data we have a lot of files**
 - should we always transform them into a single format?
 - must always be in a database form?
- **Processing files directly is a benefit!**
- **The datalake concept:**
 - multiple types of big data analytics with high concurrent/parallel data writes/reads
 - dealing with different data access frequencies: **hot, warm and cold data**

Hadoop

- <http://hadoop.apache.org/>, original from Yahoo
- The goal is to combine storage and processing in the same cluster system
- Designed for *massive scale of data and computing with shared nothing architecture*
 - commodity hardware, highly scalability, fault tolerance, easy to extend
- Suitable for both on-premise and clouds
- There are very rich software ecosystems centered around Hadoop

Hadoop: layers



Hadoop key components

- **HDFS as a distributed file system**
 - for managing data
- **YARN as a resource management system**
 - for executing and managing tasks
- **MapReduce as one programming model**
 - for MapReduce applications
- **Coordination (ZooKeeper)**
 - for fault tolerance and metadata

Hadoop File System (HDFS)

- **For handling very big data files**
 - GBs of data within a single file
- **Assume model of data**
 - write-once-read-many
 - not suitable for random-access update
- **Deal with hardware failures, support data locality, reliability**

Example

e.g., 112M rows

NYC

OpenData

Home

Data

About ▾

Learn ▾

Alerts

Contact Us

Blog

Q

Sign In

2018 Yellow Taxi Trip Data

Transportation

View Data

Visualize ▾

Export

API

...

The yellow and green taxi trip records include fields capturing pick-up and drop-off dates/times, pick-up and drop-off locations, trip distances, itemized fares, rate types, payment types, and driver-reported passenger counts. The data used in the attached datasets were collected and provided to the NYC Taxi and Limousine Commission (TLC) by

More

Updated

February 8, 2020

Data Provided by

Taxi and Limousine Commission (TLC)

About this Dataset

Mute Dataset

Updated

February 8, 2020

Data Last Updated

April 5, 2019

Metadata Last Updated

February 8, 2020

Date Created

September 24, 2018

Views

33.1K

Downloads

3,703

Data Provided by

Taxi and Limousine Commission (TLC)

Dataset Owner

NYC OpenData

Update

Update Frequency

Historical Data

Automation

No

Date Made Public

10/19/2018

Dataset Information

Agency

Taxi and Limousine Commission (TLC)

Attachments

data_dictionary_trip_records_yellow.pdf

Topics

Category

Transportation

Tags

This dataset does not have any tags

Snapshot from: <https://data.cityofnewyork.us/Transportation/2018-Yellow-Taxi-Trip-Data/t29m-gskq>

HDFS - data blocks

- **Files are stored in many nodes**
 - but we access them just like “typical file systems”
- **A file includes many blocks**
- **File blocks are replicated and distributed across nodes**
- **Conventional way of access data**
 - naming resolving: **hdfs://**
 - common operations: list, put, get, ...

Crucial components



A **Aalto University**
School of Science

File blocks, metadata and data replication

- **Block size is 128MB (default)**
 - can be configurable but should not be small size
 - all blocks of the same file are the same, except the last one
- **Data is replicated across the cluster**
 - usually, replication factor is 3
- **NameNode manages file system metadata**

HDFS fault tolerance

- **Data blocks**
 - file blocks are replicated and distributed across nodes
- **Rack awareness**
 - avoid communication problems between nodes in different racks
- **Monitoring**
 - DataNode reports to NameNode
- **Read and write**
 - using NameNode for metadata and for information of DataNodes
 - NameNode has replication (master-worker)

Compatible file systems with HDFS

- **For integration and analysis purpose: many file systems are compatible with HDFS**
- **Amazon S3**
- **Azure Blob Storage**
- **Azure Data Lake Storage**
- **OpenStack Swift**

Is that Hadoop dead/unattractive?

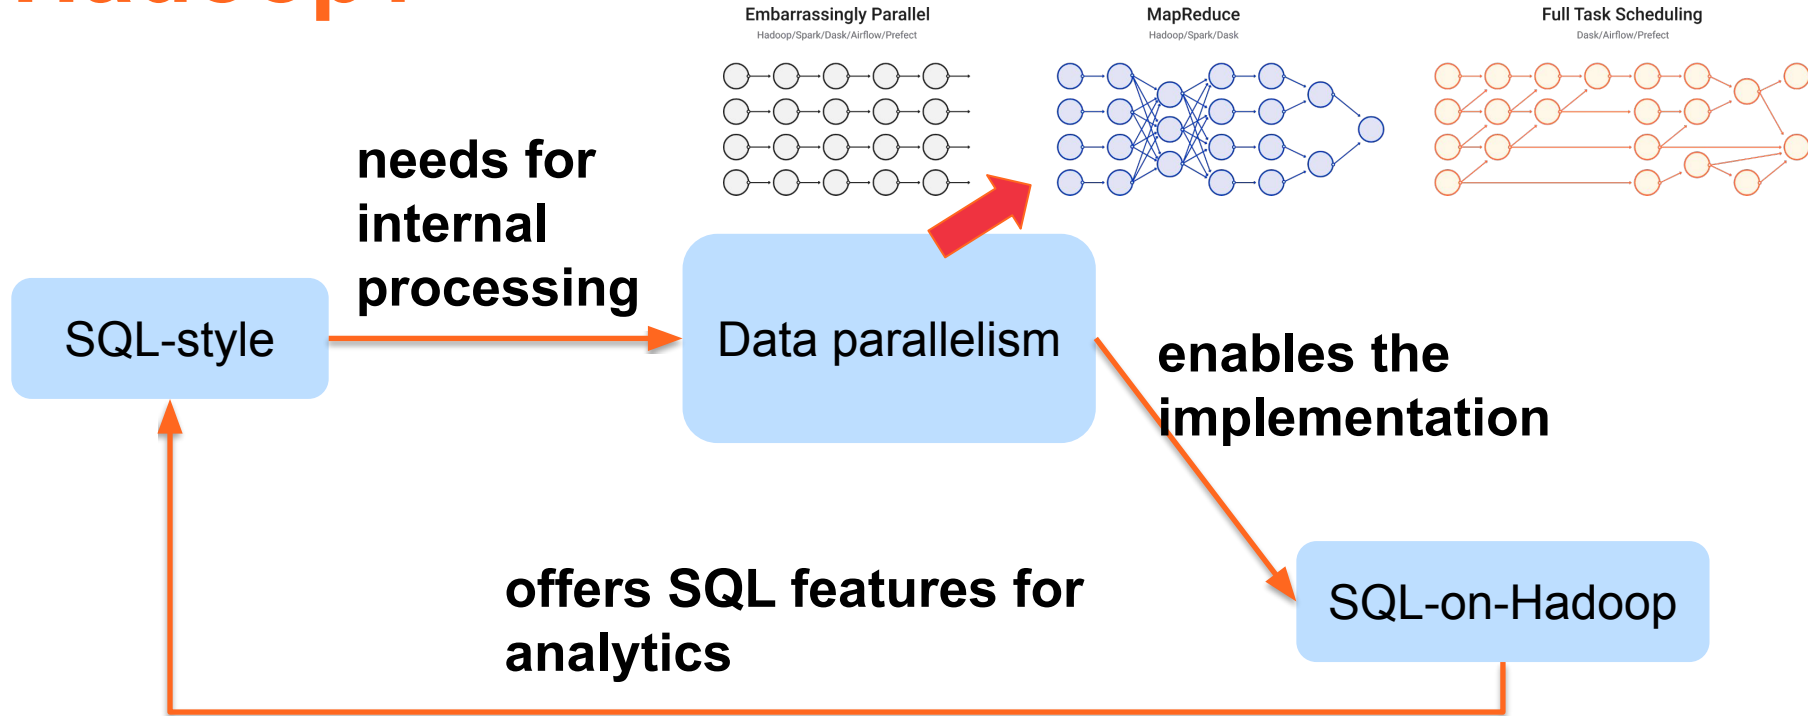
- **Hadoop is for shared nothing architecture**
- **Shared nothing architectures are not “morden”?**
 - new, advanced developments w.r.t. memory, networks, ...
 - in-memory processing, move data to computing nodes
- **Many tools seem not powerful for data science/ML?**
 - Java vs Python development for data science and ML
 - services in ecosystem
 - near realtime analytics use cases
- **So why do we still study it?**
 - foundation issues & still many Hadoop services are important (data warehouse, deep storage)
 - well-supported by many cloud providers for big data workloads

Why do we still need to learn Hadoop?

- **Hadoop file systems as a storage for many big data services**
 - databases: Accumulo, Apache Druid
 - data warehouses/Datalakes: Hive, Hudi
- **Support different models of access/analytics**
 - via SQL styles atop big data platforms: You can do extract/transform/load (ETL), reporting, and data analysis using SQL styles
 - large-scale parallel programming (e.g., Spark)
- **Still many needs given shared nothing architectures**

Why do we still need to learn Hadoop?

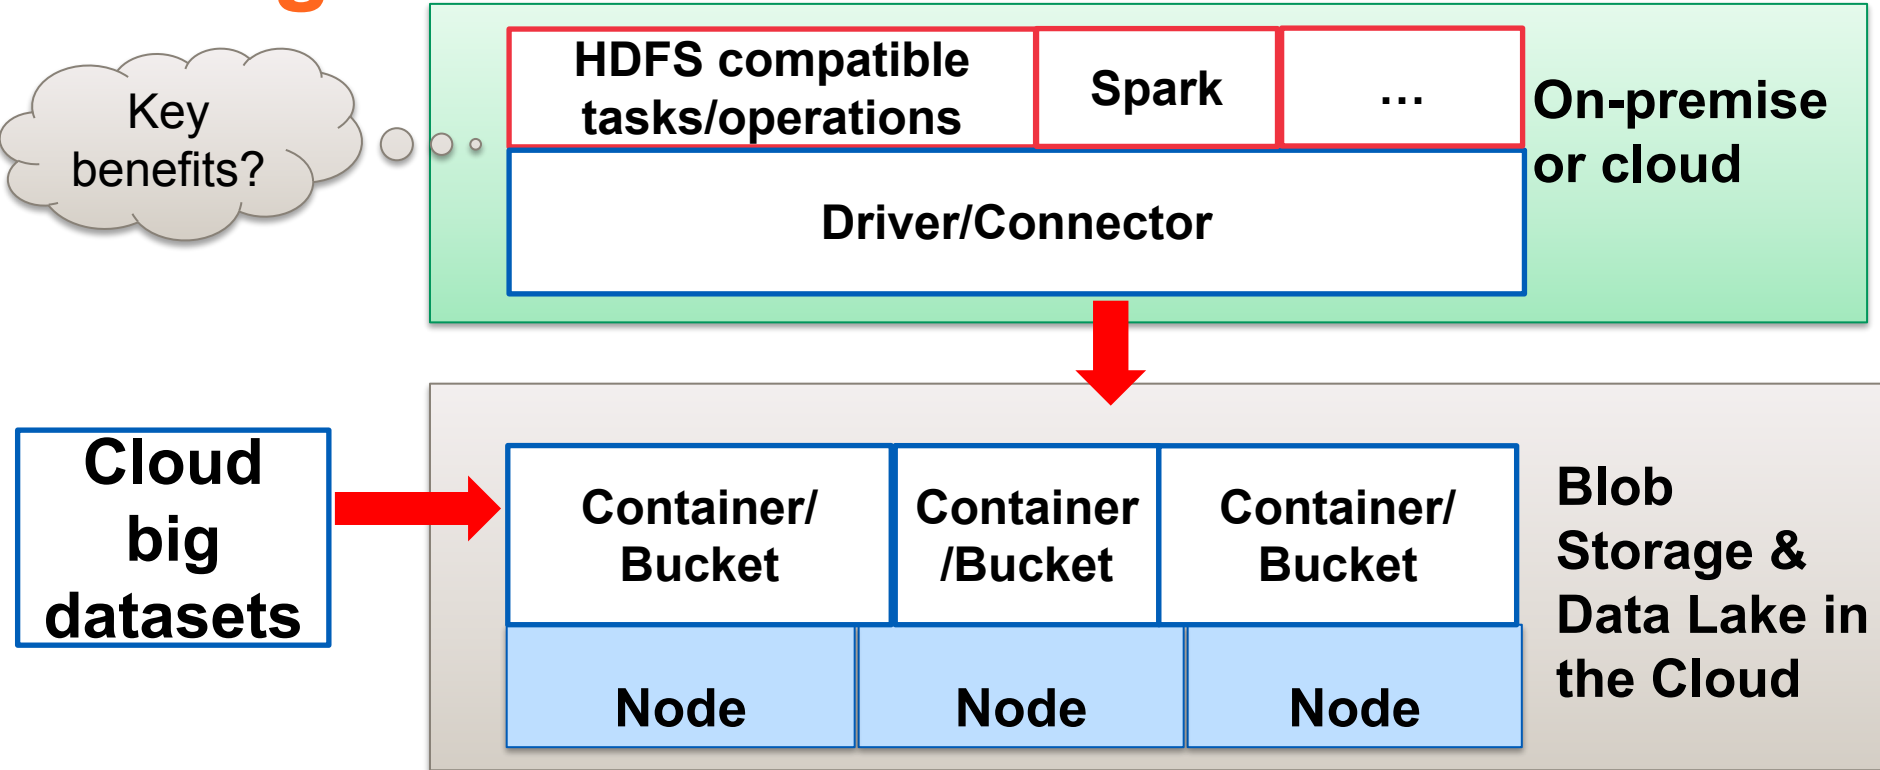
Figure source: <https://docs.dask.org/en/stable/graphs.html>



Enabling analysis big data using SQL style

- Provide command line tools & JDBC and server for integration
- “SQL-on-Hadoop”
- Examples
 - Apache Hive, <http://hive.apache.org/>, on top of Hadoop
 - *Data warehouse, access data in HDFS or Hbase*
 - Apache Druid (using HDFS as a deep storage)
 - Spark SQL for various types of files

Integration models with other cloud storage services



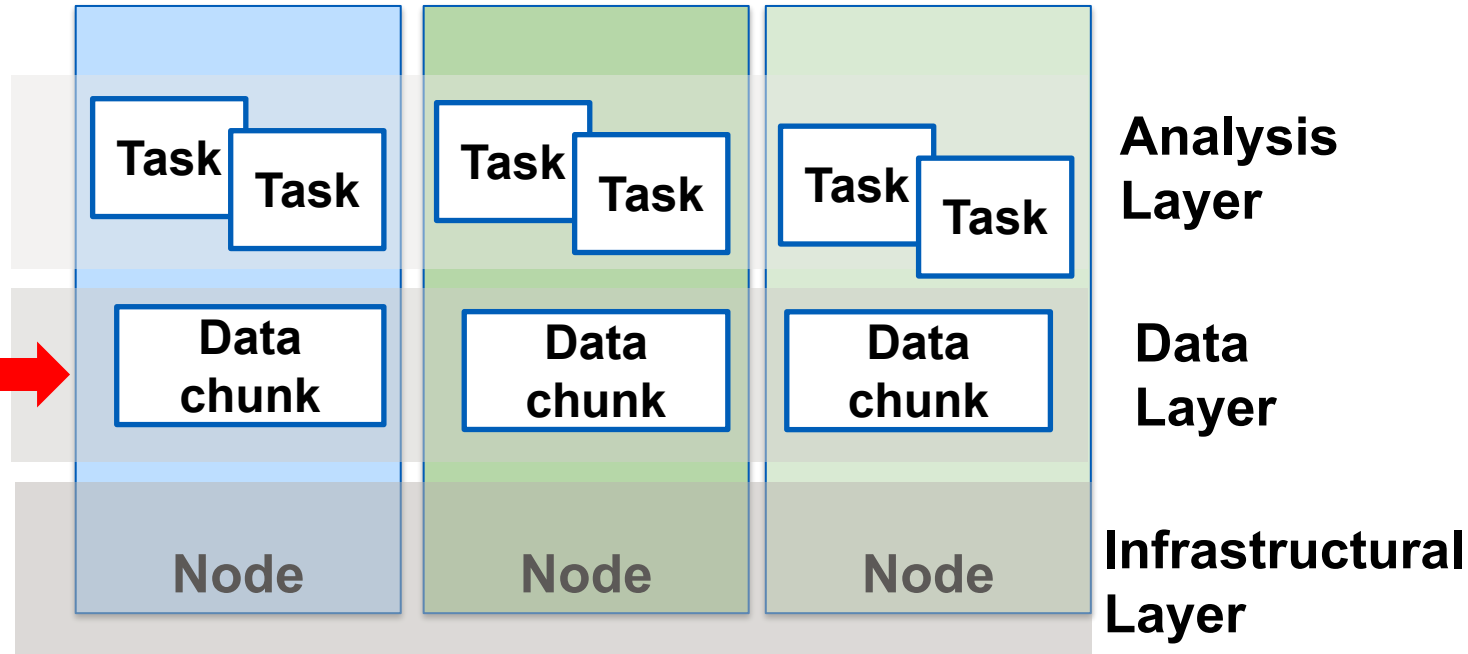
YARN (Yet Another Resource Negotiator)

If HDFS can be used to store different files, what would be the good way to enable “data processing” atop HDFS?

Key requirements:

scalability/elasticity, high utilization of infrastructural resources, serviceability for multiple users/application types

Take a look again



Big data set



How to leverage the same infrastructure and data management layers to perform data analysis?

How to enable different (distributed) programming models for data analytics?

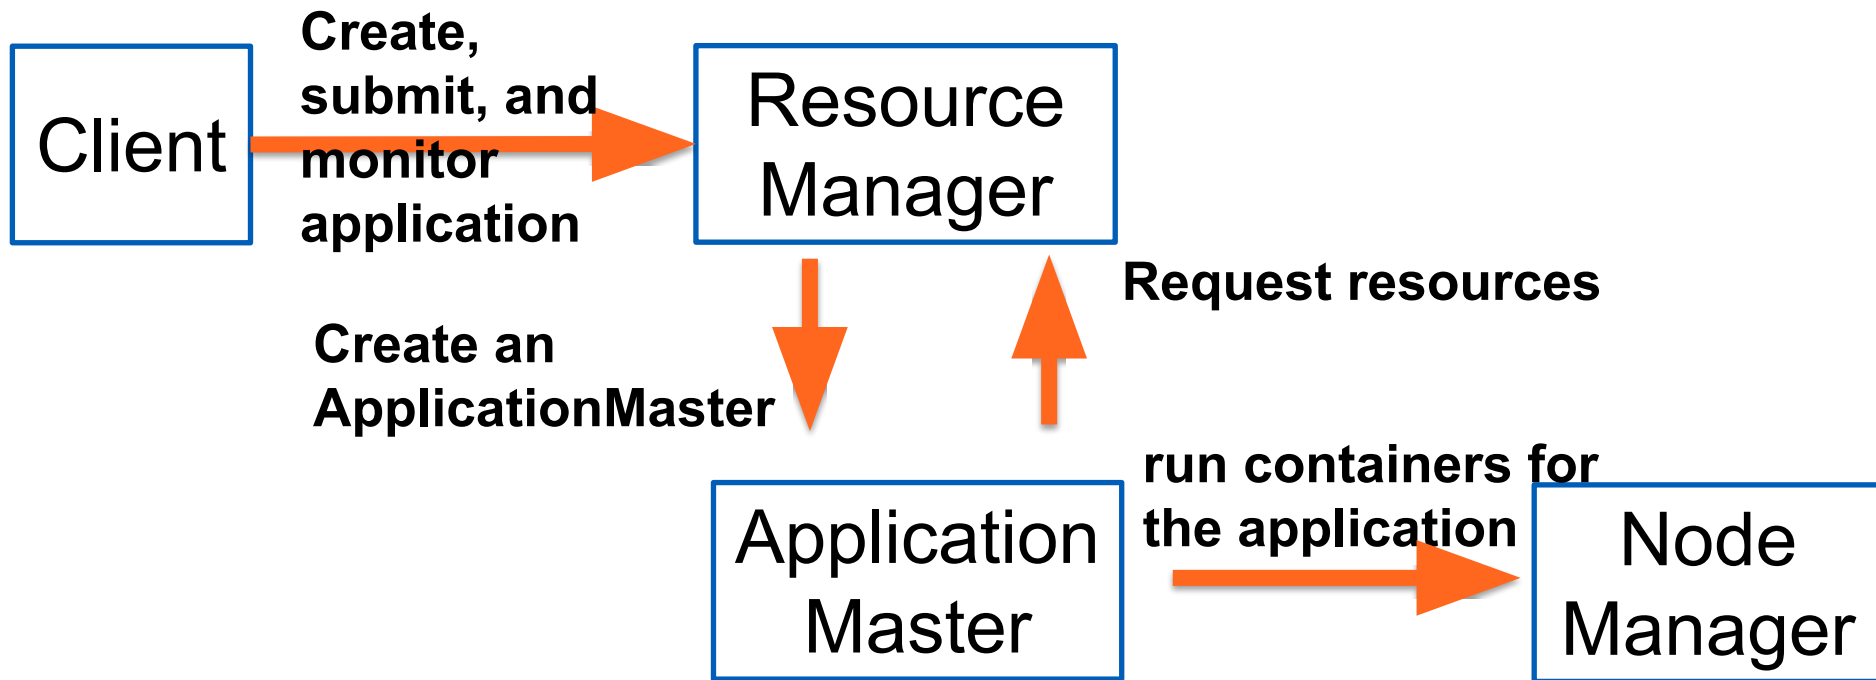
YARN (Yet Another Resource Negotiator)

- **Manage resources for processing tasks**
 - each node in the cluster provides resources for executing tasks
- **Resource types:**
 - CPU, Memory and Disks
 - also support GPU and FPGA Node
- **Resources are abstracted into “Containers”**
 - don't be confused: it is not a (Docker) container!
- **Multi-tenancy support**

YARN Components

- **Resource Manager**
 - scheduler: how to schedule tasks atop resources
- **NodeManager**
 - for managing resources of execution tasks in a node
- **ApplicationMaster**
 - application-specific manager for each application
 - work with Resource Manager and NodeManager provisioning resources and manage application execution
 - handle application-specific tasks

YARN basic model



YARN Architecture

User applications running within “Containers”:
This model **decouples** from specific types of applications

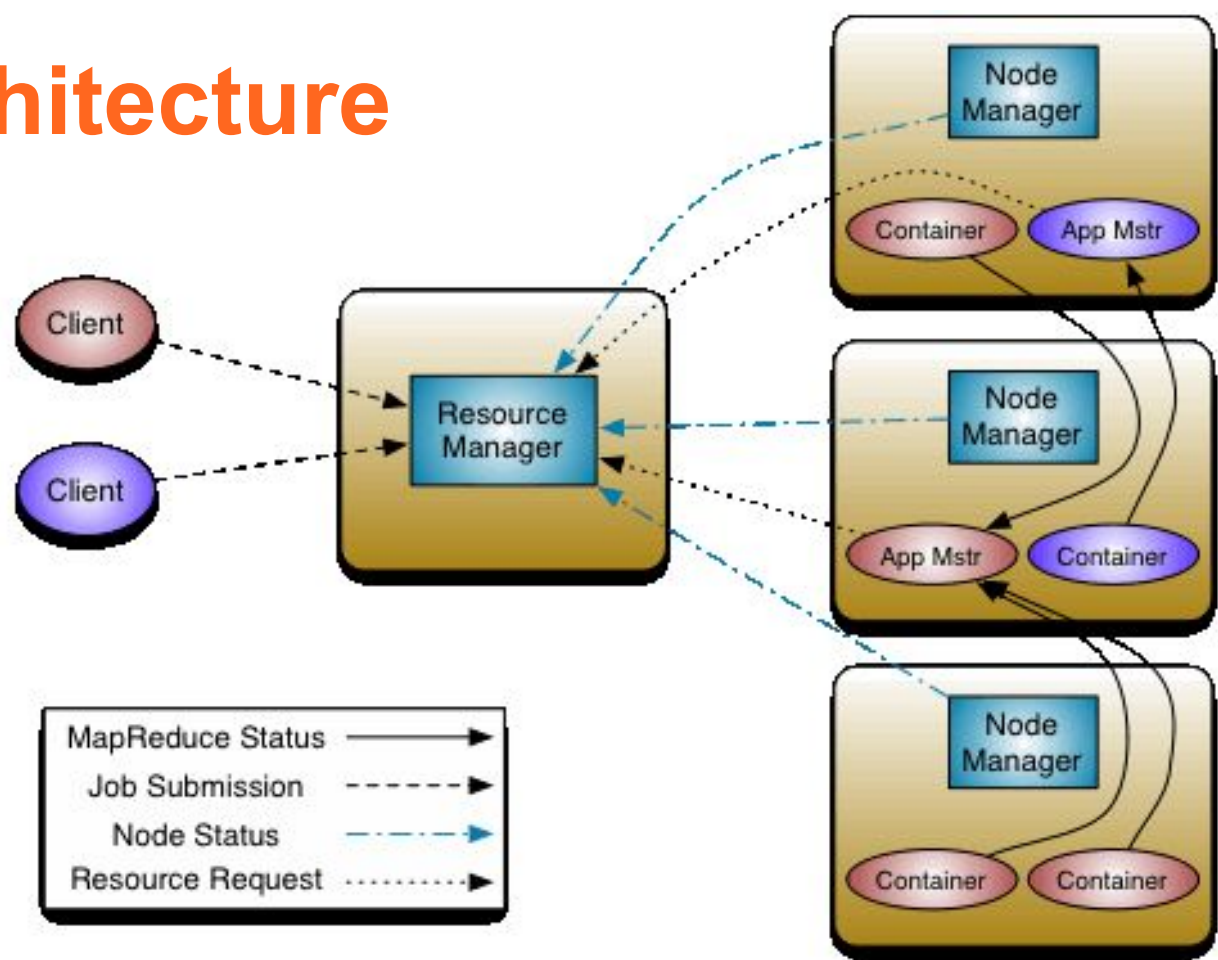


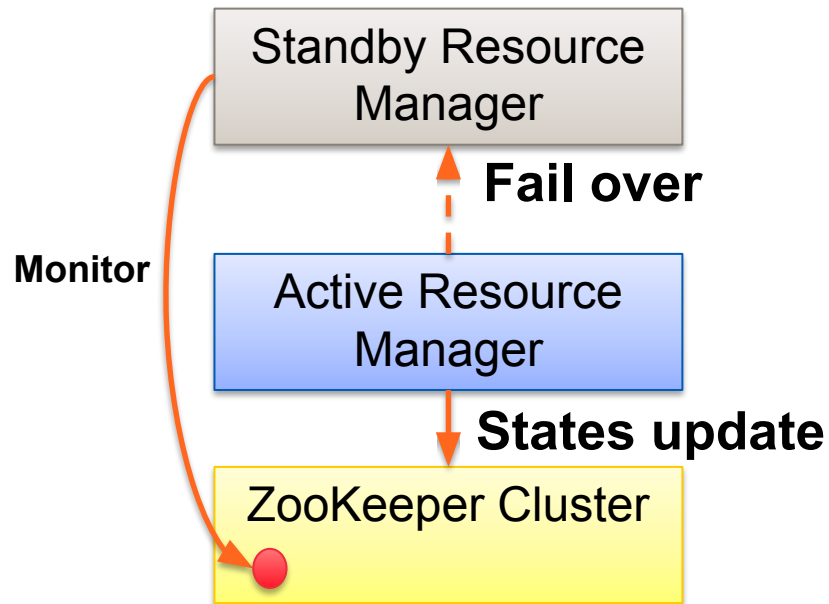
Figure source: <https://hadoop.apache.org/docs/stable/hadoop-yarn/hadoop-yarn-site/YARN.html>

Scheduling of tasks

- **Scheduling**
 - FIFO
 - CapacityScheduler
 - *Use multiple queues, each with a limit of resources*
 - FairScheduler
 - *All apps will get an average share of resources over time*
- **You can research and add new types of scheduling algorithms**

Fault tolerance

- **Resource Manager is a critical component**
 - active-passive Resource Manager
 - Zookeeper quorum failover
- **ApplicatonMaster**
 - ApplicationMaster is application-specific
 - Resource Manager restarts AM
- **Node**
 - remove out of the cluster



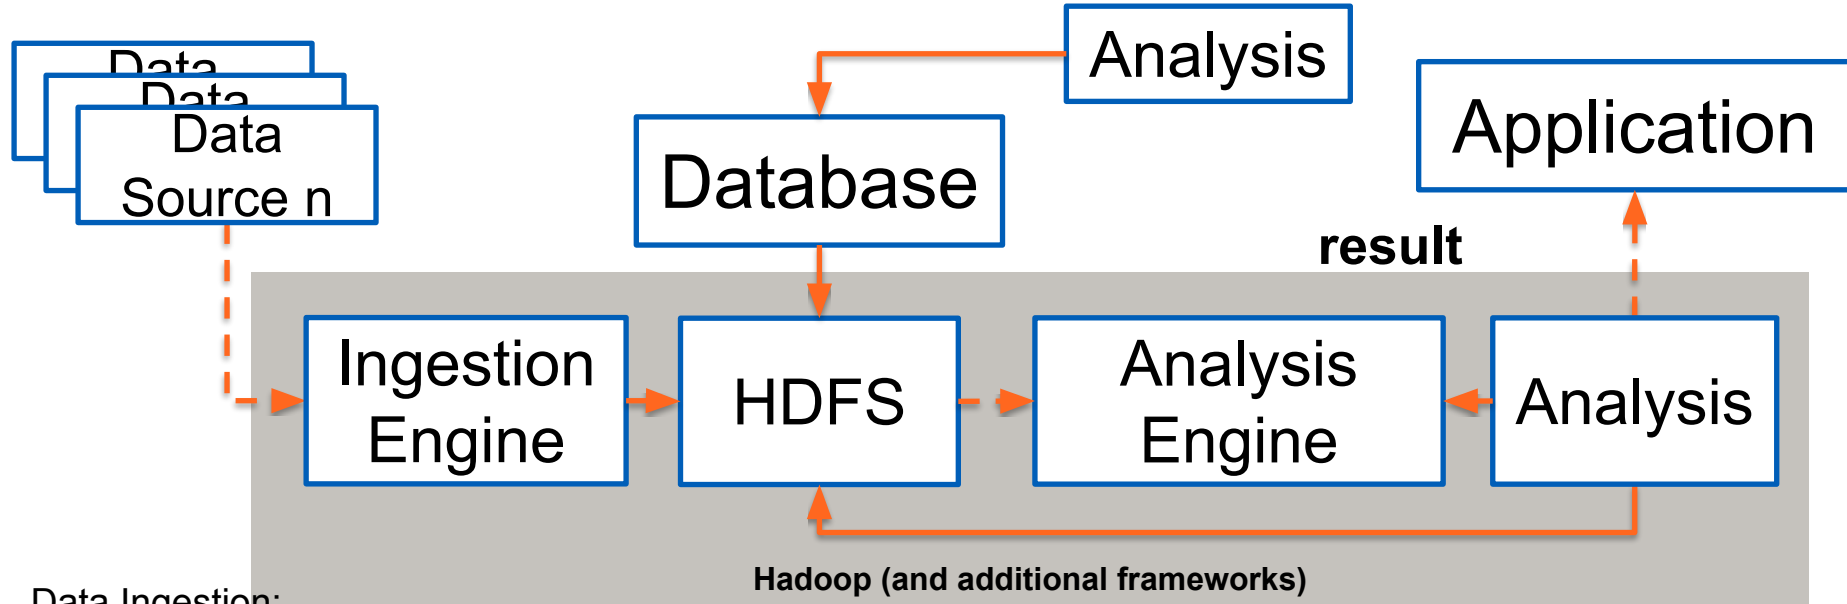
Programming models

- **YARN allows different programming models for applications**
 - MapReduce
 - Apache Spark
 - Workflows
 - *E.g., Apache Tez*

Integration models

- **Using Hadoop for developing large-scale data analysis**
 - Apache Spark, HBase, Hive, Apache Tez
- **Using Hadoop HDFS as components in a big data system**
 - Hadoop HDFS can become data store
 - emerging datalake models, combined batch and stream ingestions for incremental data processing
 - *e.g., Apache Hudi*

ETL and Analytics with Hadoop/HDFS



Data Ingestion:

- HDFS Client/Hadoop Streaming
- Spark Streaming
- Kafka Connect
- Apache Nifi

HDFS as storage for databases

- Accumulo, Druid, etc.

Computing/Data Processing Framework

- Apache Spark
- Hadoop MapReduce
- Apache Tez

Hadoop-native big database/data warehouse systems

HBase

- **NoSQL database atop Hadoop**
 - use HDFS for storing data
 - use YARN for running jobs on data
- **Follow a master-based architecture**

Reading – Why HBase?

<https://engineering.fb.com/2010/11/15/core-data/the-underlying-technology-of-messages/>

<https://engineering.fb.com/2014/06/05/core-data/hydrabase-the-evolution-of-hbase-facebook/>

<https://engineering.fb.com/2018/06/26/core-data/migrating-messenger-storage-to-optimize-performance/>

Recall: Big data: column-family data model

Many situations we aggregate and scan few columns of million rows of data \Rightarrow store big data in columns enable fast scan/retrieval/aggregation

Column Family = (Column, Column, ...): for similar type of data

Column Key = Family: qualifier

Data = (Key, Value) where Key =(Row Key, Column Key, Timestamp)

Example of a data model in HBase

Row key

Column

Column family (e.g., birdinfo)

Cell: versioning

1	species	birdinfo:country	birdinfo:english_name	duration
	Aberti	US	...	3

2	species	birdinfo:country	birdinfo:english_name	duration	name	url	latitude	longitude	Text

Row

Example of a data model

- Example with families: birdinfo, songinfo, location

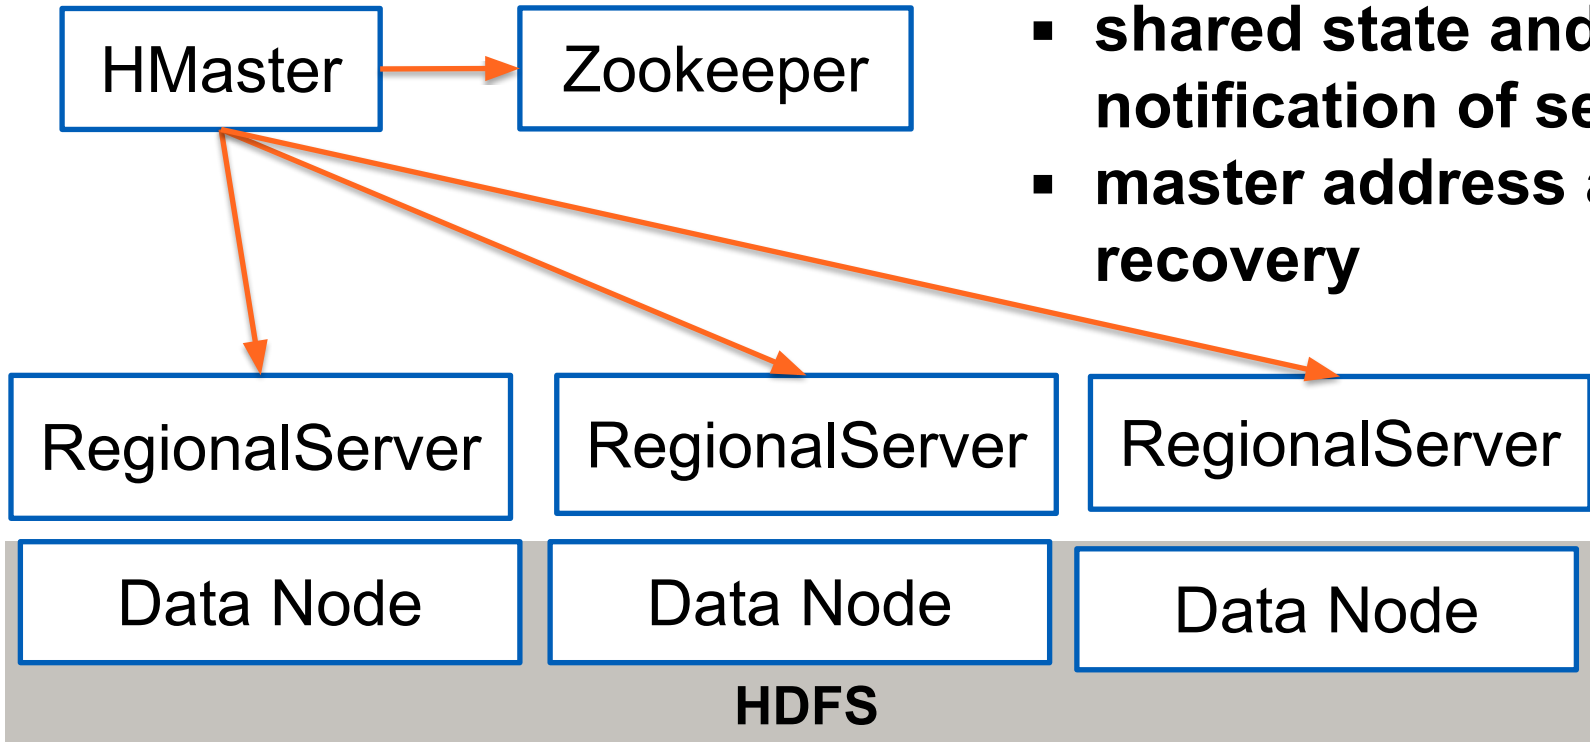
```
hbase:038:0> scan 'hbird0'
ROW                                COLUMN+CELL
17804                             column=birdinfo:english_cname, timestamp=2022-02-05T11:56:09.114, value=Aberts Towhee
17804                             column=songinfo:duration, timestamp=2022-02-05T11:56:09.154, value=3
17804                             column=songinfo:file_id, timestamp=2022-02-05T11:56:09.178, value=17804
17804                             column=songinfo:file_name, timestamp=2022-02-05T11:56:09.205, value=XC17804.mp3
71852                             column=birdinfo:country, timestamp=2022-02-05T11:56:09.415, value=Mexico
71852                             column=birdinfo:english_cname, timestamp=2022-02-05T11:56:09.445, value=Ash-throated Flycatcher
71852                             column=birdinfo:species, timestamp=2022-02-05T11:56:09.488, value=cinerascens
71852                             column=location:latitude, timestamp=2022-02-05T11:56:09.504, value=32.156
71852                             column=location:longitude, timestamp=2022-02-05T11:56:12.253, value=-115.79299999999999
71852                             column=songinfo:duration, timestamp=2022-02-05T11:56:09.430, value=28
71852                             column=songinfo:file_id, timestamp=2022-02-05T11:56:09.459, value=71852
2 row(s)
Took 0.0206 seconds
```

Enable analytics based on column families (as well as data management)

Data model – sharding and storage

- **Table** includes multiple **Regions**
 - a **Region** keeps related row data of a **Table** (partitioning)
- **Auto-sharding**
 - **Regions** are spitted based on policies
- **Region** has multiple **column families**
 - Different **column families** will be stored in different files
 - **HFiles** are used to store real data (also include index data)

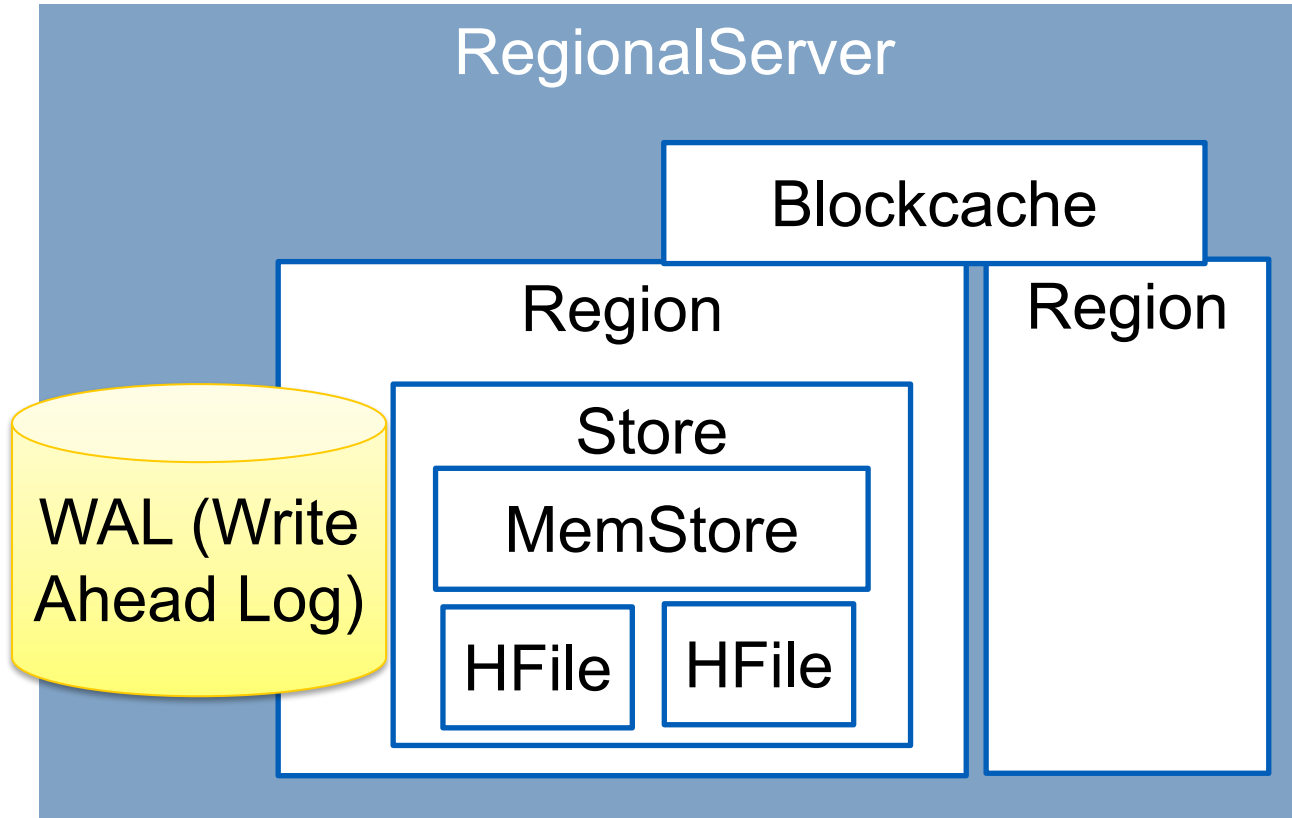
HBase architecture



Zookeeper:

- **shared state and failure notification of servers**
- **master address and recovery**

HBase architecture



MemStore: write cache for data in memory before written into files

BlockCache: for read cache

WAL is for durability

ACID

- **Atomic within a row**
- **Consistency**
 - can be programmed: e.g.,
 - read: STRONG (read performed by the primary region) and TIMELINE (primary region first, if not, then the secondary region)
- **Durability**
 - can be programmed
 - WAL (write ahead log)

Example of using HBase

Open Scalable Open Time Series Database (OpenTSDB)




Figure source: <http://opentsdb.net/overview.html>

Salesforce

- **Salesforce Shield:** for supporting data security/compliance
- **Field Audit Trail:** for track changes
- **Event Monitoring:** log/event
- **Argus:** time series data and alerts

Source:

<https://engineering.salesforce.com/investing-in-big-data-apache-hbase-420edfba2d30/>

Apache Hive

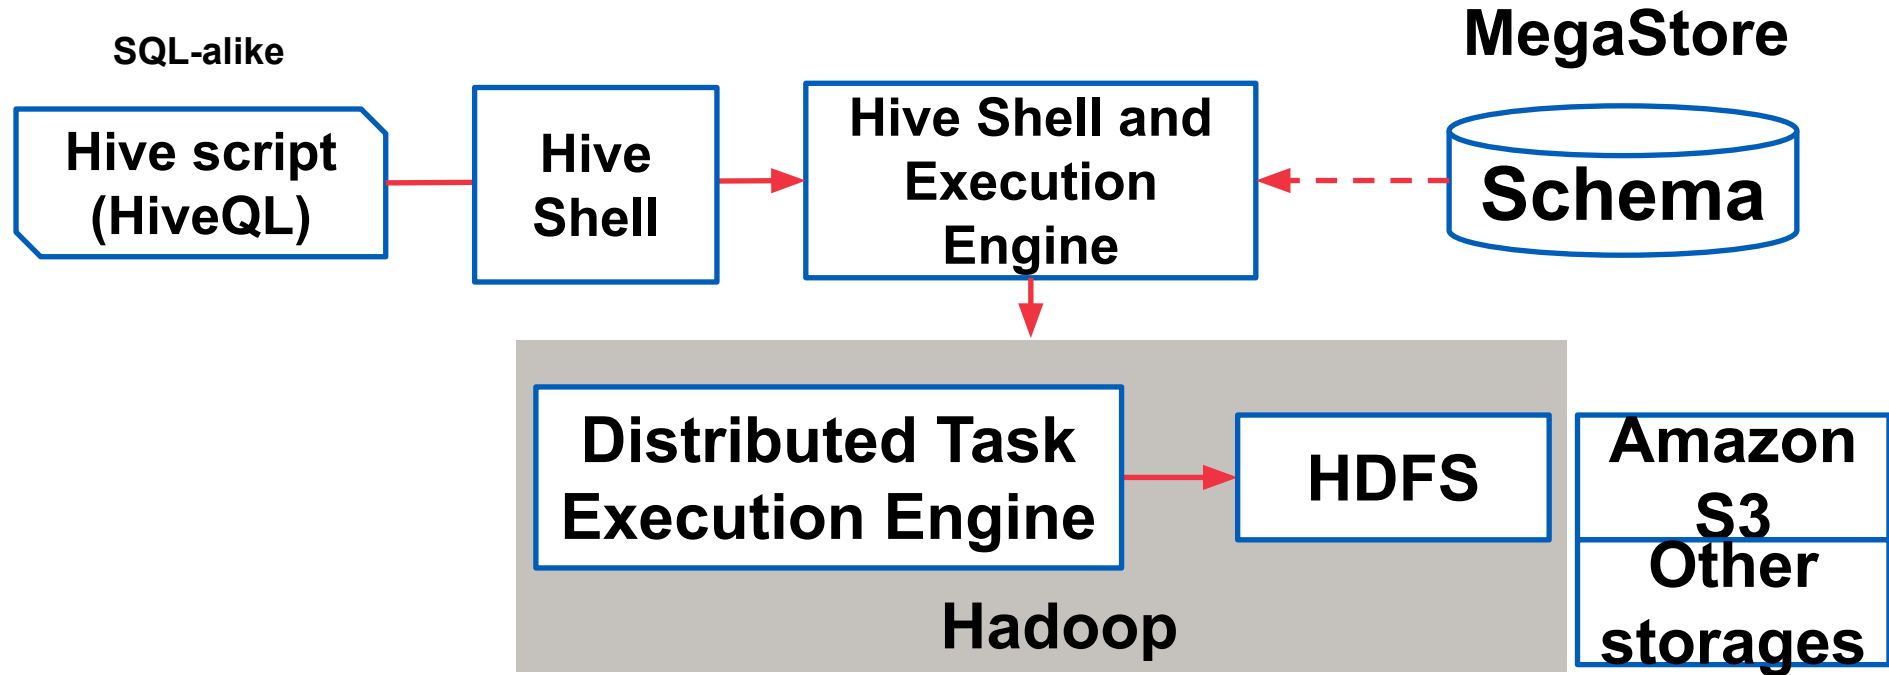
Apache Hive

- **<http://hive.apache.org/>, on top of Hadoop**
 - data warehouse at a very large-scale
 - access data in large-scale storage like HDFS or S3
- **Support access to data via SQL styles**
 - extract/transform/load (ETL), reporting, and data analysis using SQL styles
- **Provide command line tools & JDBC and server for integration**

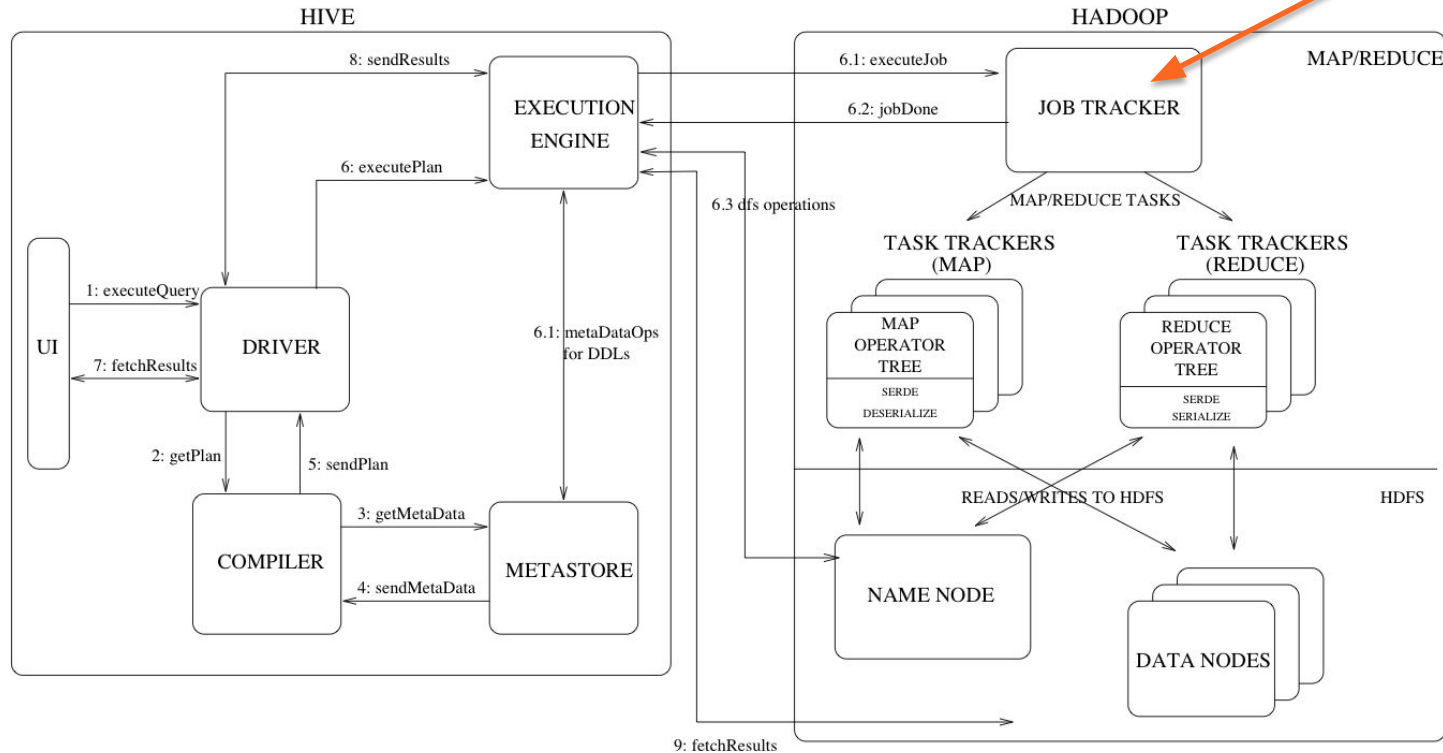
Reading – Hive in Facebook

<https://engineering.fb.com/2009/06/10/web/hive-a-petabyte-scale-data-warehouse-using-hadoop/>

High-level data flow language & programs



Hive building blocks



Distributed tasks with MapReduce, Tez (Workflow) or Spark

Figure source: <https://cwiki.apache.org/confluence/display/Hive/Design>

Hive data organization

- **Databases**
- **Table**
 - **Managed** table versus **External** tables
 - External table: data is referenced so delete only table metadata but not the data)
 - Table is mapped to a directory in HDFS
- **Low-level files: CSV/Parquet, ORC, ...**

Example

Large tables lead
to performance
issues

```
0: jdbc:hive2://localhost:10000> describe taxiinfo;
```

col_name	data_type	comment
vendorid	int	
tpep_pickup_datetime	timestamp	
tpep_dropoff_datetime	timestamp	
passenger_count	int	
trip_distance	float	
ratecodeid	int	
store_and_fwd_flag	int	
pu_locationid	string	
do_locationid	int	
payment_type	int	
fare_amount	float	
extra	float	
mta_tax	float	
tip_amount	float	
tolls_amount	float	
improvement_surcharge	float	
total_amount	float	

Hive data organization for performance optimization

- **Partitioning: using value of a column as a partitioning key**
 - partition keys determine how data in Table will be divided
 - *E.g. date or countries*
 - each partition is stored as a subdirectory
 - *Avoid many sub directories!*
- **Buckets: using a hash function of a column for grouping records into the same bucket**
 - avoid large number of small partitions
 - each bucket is stored in a file
 - bucket columns should be optimized for join/filter operations

Example of partitions

```
CREATE TABLE taxiinfo1 ( ... )  
PARTITIONED BY (year int, month int)  
...;
```

Indicate partition info

Define partition names

```
LOAD DATA LOCAL INPATH .... INTO TABLE taxiinfo1  
PARTITION (year=2019, month=11);
```

```
truong@aaltosea:/opt/hadoop$ bin/hdfs dfs -ls /user/hive/warehouse/taxiinfo1
Found 4 items
drwxr-xr-x - truong supergroup 0 2021-03-02 22:37 /user/hive/warehouse/taxiinfo1/year=2017
drwxr-xr-x - truong supergroup 0 2021-03-02 22:37 /user/hive/warehouse/taxiinfo1/year=2018
drwxr-xr-x - truong supergroup 0 2021-03-02 22:36 /user/hive/warehouse/taxiinfo1/year=2019
drwxr-xr-x - truong supergroup 0 2021-03-02 22:33 /user/hive/warehouse/taxiinfo1/year=__HIVE_DEFAULT_PA
truong@aaltosea:/opt/hadoop$ bin/hdfs dfs -ls /user/hive/warehouse/taxiinfo1/year=2019
Found 2 items
drwxr-xr-x - truong supergroup 0 2021-03-02 22:36 /user/hive/warehouse/taxiinfo1/year=2019/month=11
drwxr-xr-x - truong supergroup 0 2021-03-02 22:36 /user/hive/warehouse/taxiinfo1/year=2019/month=12
```


Example of buckets

```
CREATE TABLE taxiinfo2 (VendorID int, ....)
  CLUSTERED BY (VendorID) INTO 2 BUCKETS
.....;
```

Identify bucket column



```
truong@aaltosea:/opt/hadoop$ bin/hdfs dfs -ls /user/hive/warehouse/taxiinfo2
Found 2 items
-rw-r--r--  1 truong supergroup    66784 2021-03-02 22:54 /user/hive/warehouse/taxiinfo2/000000_0
-rw-r--r--  1 truong supergroup    31339 2021-03-02 22:54 /user/hive/warehouse/taxiinfo2/000001_0
```

Combine partitions with buckets



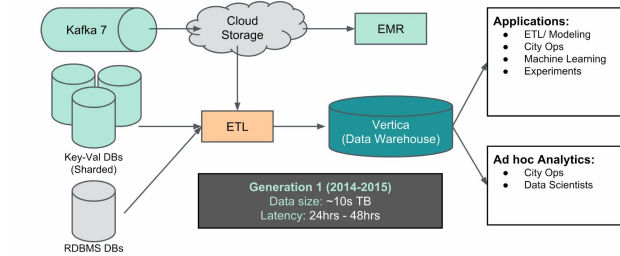
```
truong@aaltosea:/opt/hadoop$ bin/hdfs dfs -ls /user/hive/warehouse/taxiinfo3/year=2019/month=11
Found 2 items
-rw-r--r--  1 truong supergroup    66784 2021-03-02 23:04 /user/hive/warehouse/taxiinfo3/year=2019/month=11/000000_0
-rw-r--r--  1 truong supergroup    31339 2021-03-02 23:04 /user/hive/warehouse/taxiinfo3/year=2019/month=11/000001_0
```


ACID

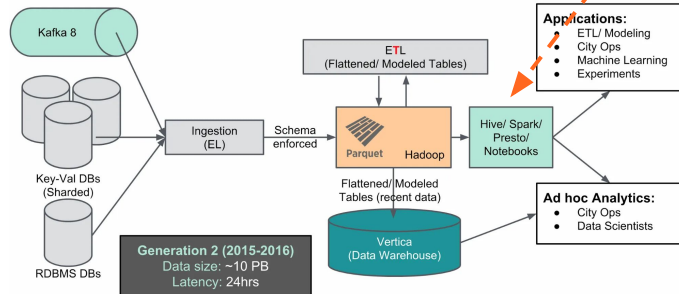
- **Full ACID support**
 - row-level transaction
- **Locks are used for data isolation**
 - shared lock: for concurrent read of tables/partitions
 - exclusive lock: for modifying table/partition

Uber example

Generation 1 (2014-2015) - The beginning of Big Data at Uber

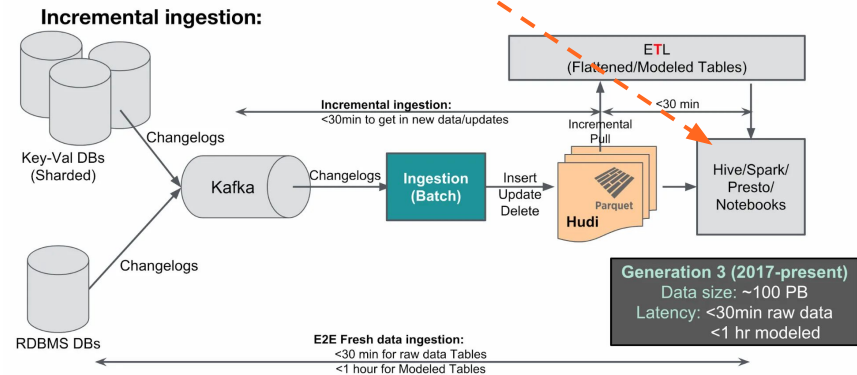


Generation 2 (2015-2016) - The arrival of Hadoop



Hive for analytics: “extremely large queries”, “20,000 Hive queries per day” (see the source)

Generation 3 (2017-present) - Let's rebuild for long term



Figures source:

<https://www.uber.com/en-FI/blog/uber-big-data-platform/>

Summary

- **Hadoop software ecosystem is very powerful**
 - many applications and use cases have been developed
 - Hadoop File System (HDFS) is a crucial subject
 - **Managed Hadoop ecosystem services by cloud providers**
 - try to look at Azure HDInsight, Google Dataproc, and Amazon EMR
 - **High-level distributed query engineering using Hadoop components (HDFS, Hive)**
 - **Understand the combination of data management with data processing techniques in the same system with Hadoop that simplify your big data tasks**
-

Thanks!

Hong-Linh Truong
Department of Computer Science

rdsea.github.io