



Aalto University  
School of Science

# Big Data Storage and Database Services – common systems & integration problems

*Hong-Linh Truong*

*Department of Computer Science*

*[linh.truong@aalto.fi](mailto:linh.truong@aalto.fi)*, *<https://rdsea.github.io>*



Aalto University  
School of Science

# Common data models and data storage/database systems

# Common data models

- **File**
- **Relational data model**
- **Key-Value data model**
- **Document-oriented model**
- **Column family model**
- **Graph model**
- **Vectorization model**

# Special file formats in big data

- **For analytical big data, remember the patterns about data**
  - write once, read many
  - analytics queries often access data based on “columns” (e.g., sum of all “trip payments”)
- **File formats**
  - compression, columnar representation for column-based queries/accesses, encryption
  - suitable with big data analytics (e.g., Spark, Hadoop)
- **Examples**
  - Apache ORC (<https://orc.apache.org/>)
  - Apache Parquet (<https://parquet.apache.org/>)
- **They are the file formats under many big data systems**

# Blob data

## Big files:


- Pictures, documents, big log files, images, video, backup data

## Storage

- File systems or blob storage

## Implementations

- File systems: NFS, GPFS, Lustre (<http://lustre.org/>), Hadoop File systems
- Storage: Amazon S3, Azure Blob storage, OpenStack Swift, Minio
- Simple API for direct access



# Example - Amazon S3

## Store blob files and their metadata

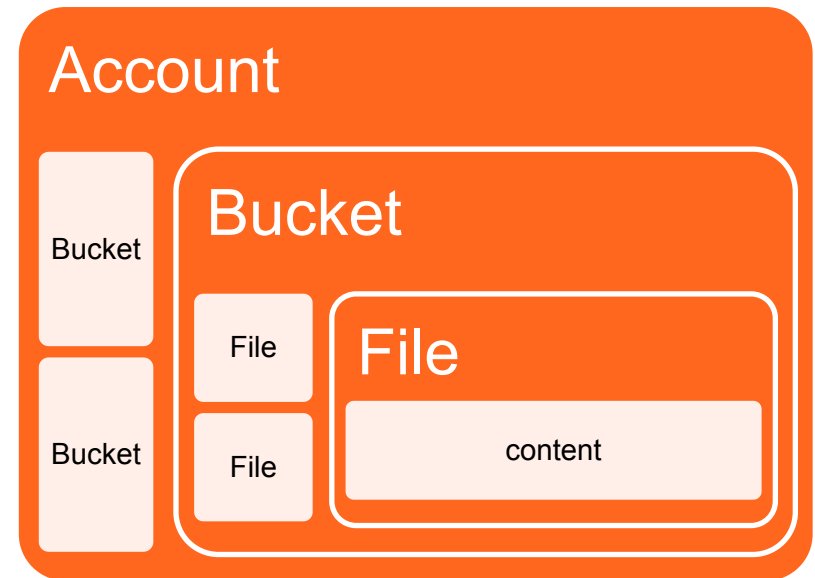
- Max 5TB per file
- A file is identified by a key

## Structure

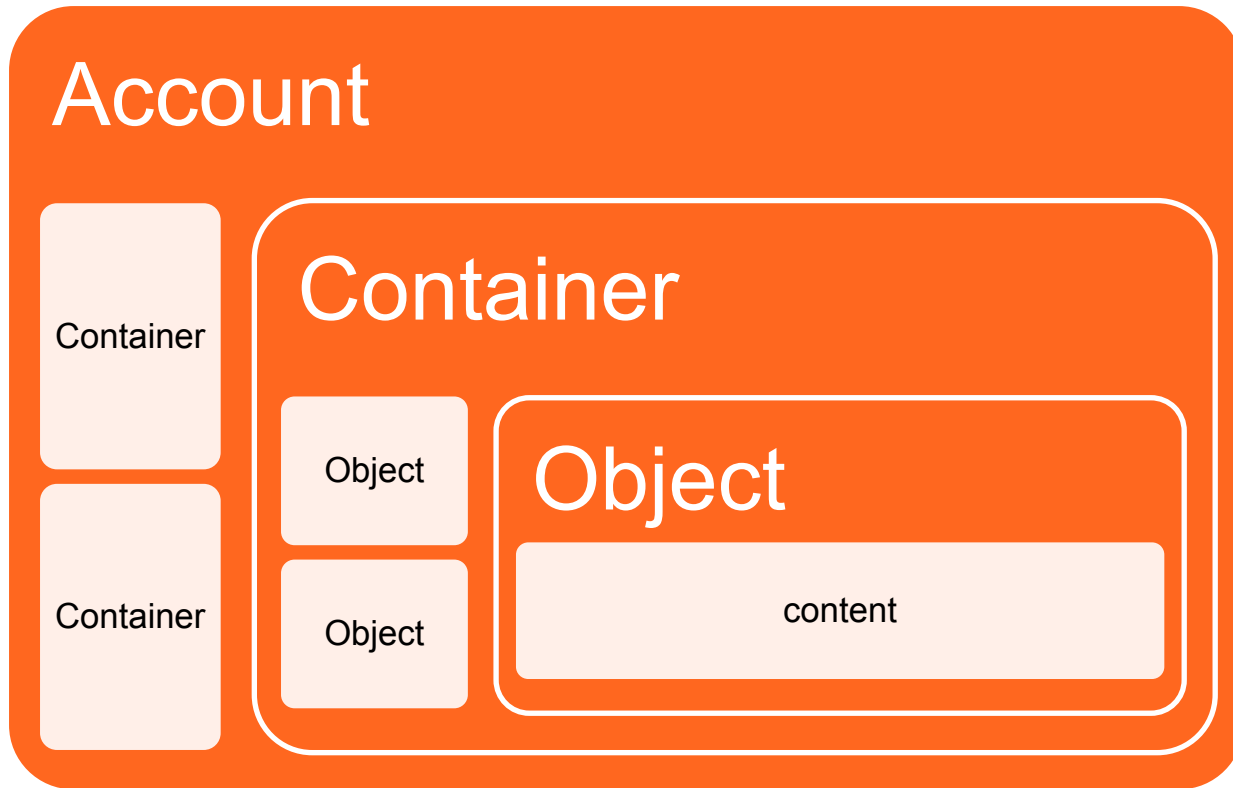
- File = Object
- Object: name and metadata
- Objects are organized into Buckets

## Simple APIs

- REST



# OpenStack Swift



<http://docs.openstack.org/developer/swift/>

# Minio

- Check <https://min.io/>
- For different deployment models: Kubernetes, VMs, edge-cloud
- S3 compatibility



# Relational Model

- **Well-known, long history**
- **Tables with rows and columns**
  - Strict schema requirements
- **Powerful querying & strong consistency support**
  - E.g.: Oracle Database, MySQL Server, PostgreSQL, MariaDB

# Example: Alarm in BigQuery

stationdescription

Schema Details Preview

| Field name  | Type    | Mode     |
|-------------|---------|----------|
| station_id  | INTEGER | NULLABLE |
| code        | STRING  | NULLABLE |
| name        | STRING  | NULLABLE |
| address     | STRING  | NULLABLE |
| description | STRING  | NULLABLE |
| latitude    | STRING  | NULLABLE |
| longitude   | STRING  | NULLABLE |

stationparameters

Schema Details Preview

| Field name   | Type      | Mode     |
|--------------|-----------|----------|
| reading_time | TIMESTAMP | NULLABLE |
| value        | FLOAT     | NULLABLE |
| station_id   | INTEGER   | NULLABLE |
| parameter_id | INTEGER   | NULLABLE |

stationalarms

Schema Details Preview

| Field name   | Type      | Mode     |
|--------------|-----------|----------|
| station_id   | INTEGER   | NULLABLE |
| alarm_id     | INTEGER   | NULLABLE |
| parameter_id | INTEGER   | NULLABLE |
| start_time   | TIMESTAMP | NULLABLE |
| end_time     | TIMESTAMP | NULLABLE |
| value        | FLOAT     | NULLABLE |
| threshold    | INTEGER   | NULLABLE |

# Relational Databases for big data scenarios

- **Relational database at very large-scale**
  - Amazon Aurora, CockroachDB, Microsoft Azure SQL Data Warehouse
- **We said ACID is hard with big data**
  - relational big database must address replication, distribution, and scalability issues
- **Examples of Amazon Aurora (reading list)**
  - based on MySQL/InnoDB but change the architecture, separate storage from engine, support cloud scale and replication, etc.

# Key-Value Model

- **Tuple = (key, value)**
  - Values can be based on different structures
- **Scalable and performance**
- **Primary use case: caching (pages, sessions, frequently access data, distributed lock)**
  - Simple, very efficient but limited querying capabilities
- **Implementation:**
  - Memcached, Riak, Redis, Apache Accumulo


# Example: Redis

- <http://redis.io/>
- **In-memory cache service**
  - *Store (key,value) tuples in memory but persistent back to database*
- Simple APIs
  - Well support with many programming languages
  - Widely used in big data ecosystems
- **Learning**
  - *<https://app.redislabs.com/#/login> provides a free account*

# Example: Redis

<http://redis.io/topics/benchmarks>

# Document-oriented model – simple analogy



# Document-oriented Model

- **Documents**

- flexible schema (schemaless) with flexible content
- data fields can be complex for sub documents
- use collections, each collection is a set of documents

- **Primary use cases**

- large amounts of semi-structured data
- collection of data with different structures



# Examples: MongoDB.Atlas

<https://www.mongodb.com/atlas/database>

# Graph-oriented model

- **Data is represented as a graph**
  - nodes or vertices represent objects
  - an edge describes a relationship between nodes
  - properties associated with nodes and edge provide other information
- **Use cases**
  - when searching data is mainly based on relations (social networks, asset relationship, knowledge graph)

# Working with graph databases

- **Graph databases**

- Azure CosmosDB, ArgangoDB, Titan, TypeDB, Neo4J, OrientDB

- **Query languages:**

- Gremlin, SPARQL, Cypher

- **Graph computing frameworks (analysis)**

- Apache TinkerPop, Apache Spark GraphX

# Example

**<https://github.com/vaticle/typedb>**

# Column-family data model

**Motivation: scalable, distributed storage for multi-dimensional sparse sorted map data**

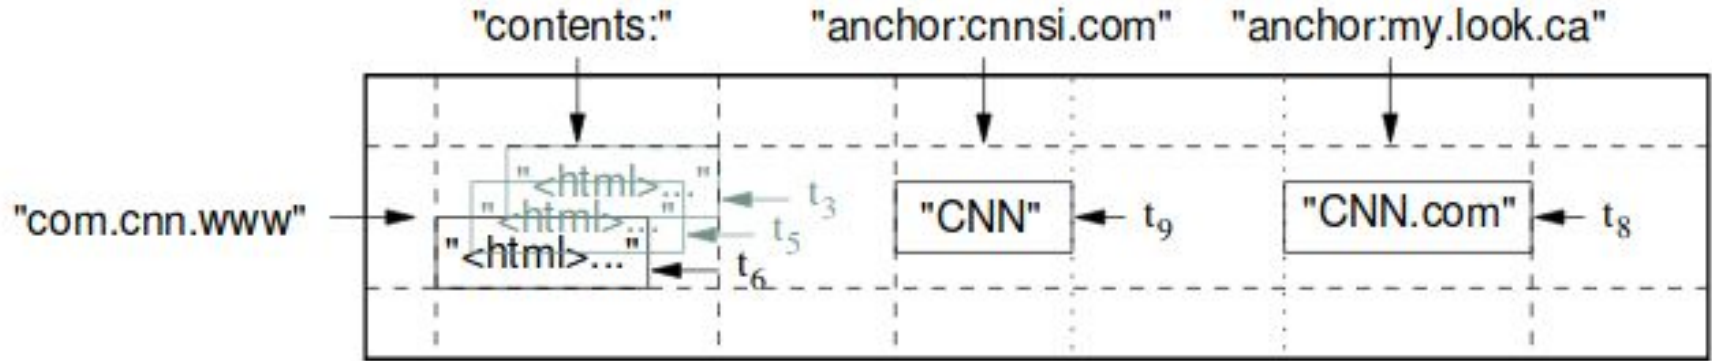


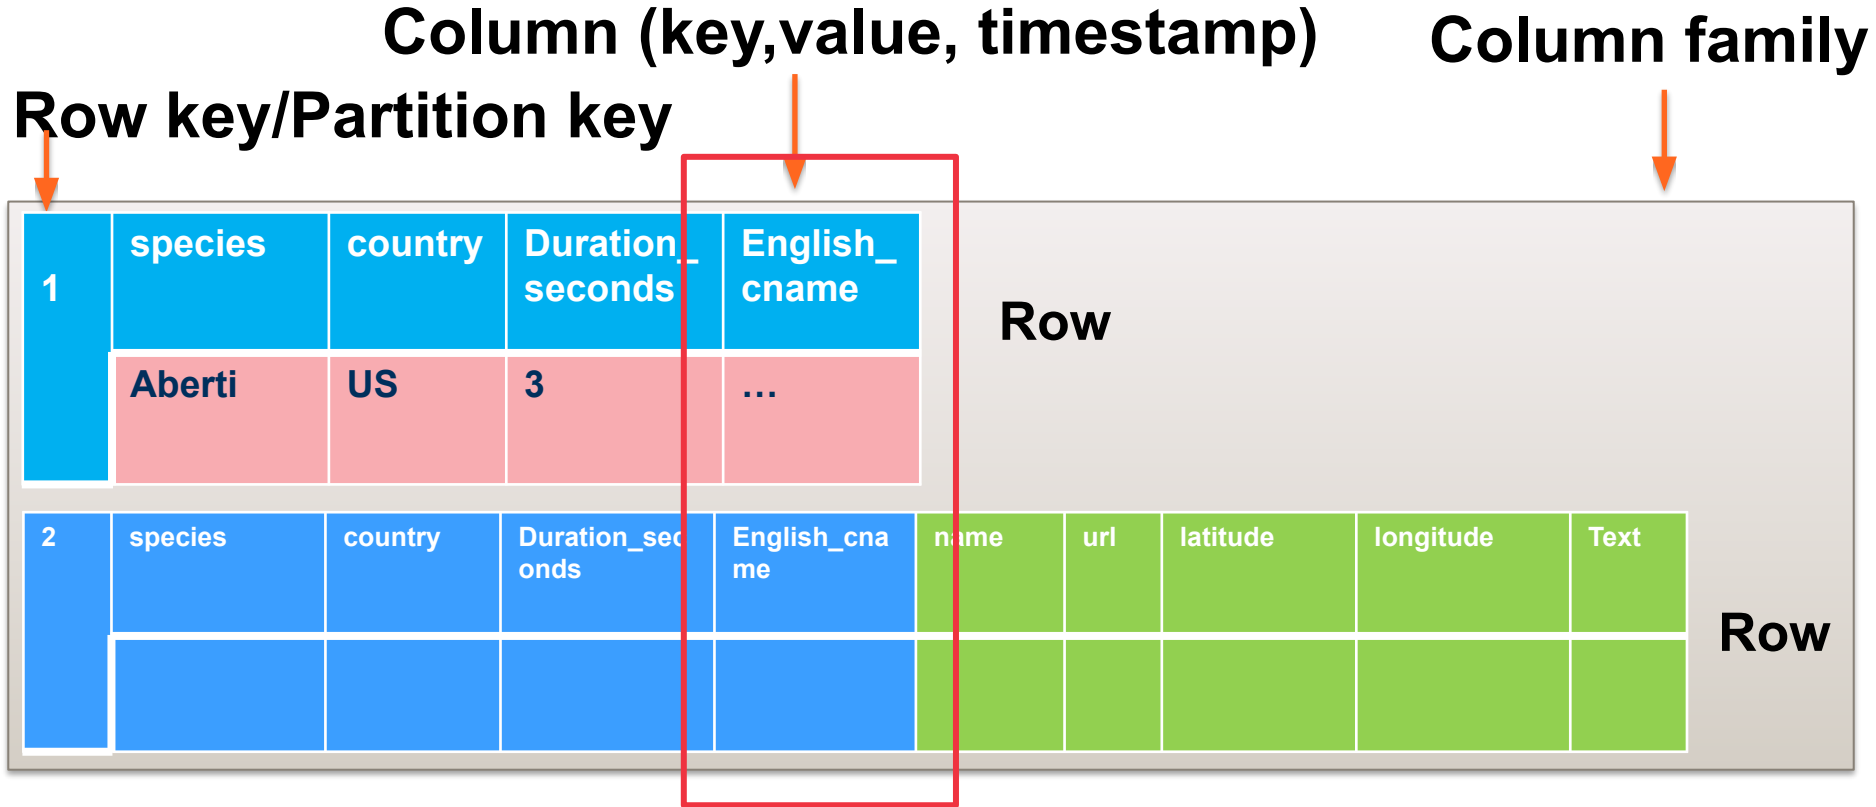
Figure source: Fay Chang, Jeffrey Dean, Sanjay Ghemawat, Wilson C. Hsieh, Deborah A. Wallach, Mike Burrows, Tushar Chandra, Andrew Fikes, and Robert E. Gruber. 2006. Bigtable: a distributed storage system for structured data. In Proceedings of the 7th symposium on Operating systems design and implementation (OSDI '06). USENIX Association, Berkeley, CA, USA, 205-218.

# Column-family data model

- **Data Model**

- Table consists of rows
- Row consists of a key and one or more columns
- Columns (column name, value, timestamp)
- Columns are grouped into column families
- Columns can be different in rows
  - flexible, wide columns → save spaces

# Example of a data model in Cassandra



# Examples

## Examples of rows

Column (name, value, timestamp)

| english_cname            | writetime(english_cname) |
|--------------------------|--------------------------|
| Black-tailed Gnatcatcher | 1569966171073228         |

(1 rows)

```
cassandra@cqlsh> select * from tutorial12345.bird2;
```

@ Row 1

| species          | melanura                 |
|------------------|--------------------------|
| country          | Mexico                   |
| duration_seconds | 29                       |
| english_cname    | Black-tailed Gnatcatcher |
| file_id          | 71907                    |
| latitude         | 32.156                   |
| longitude        | -115.793                 |

@ Row 2

| species          | melanura                 |
|------------------|--------------------------|
| country          | United States            |
| duration_seconds | 29                       |
| english_cname    | Black-tailed Gnatcatcher |
| file_id          | 358907                   |
| latitude         | 33.7329                  |
| longitude        | -115.8023                |

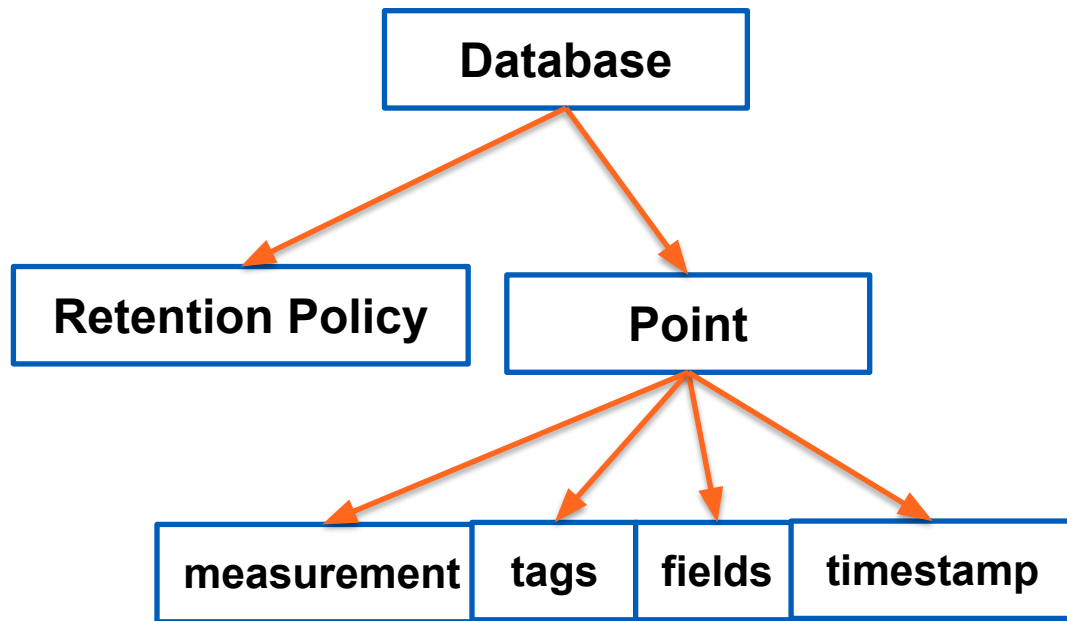


# Time Series Database

- **So many types of data in big data are time series**
  - IoT measurements, session data, log, etc.
- **Of course you can also use other databases**
  - e.g., Cassandra, ElasticSearch, BigTable
- **Time Series Databases specially designed for time series data**
  - *examples: Riak TS (Time Series), InfluxDB, Apache Druid*

# Example: InfluxDB

- <https://www.influxdata.com/>
- High-level query, SQL-alike Language
- Retention policy for data storage, sharding and replication



# An example of InfluxDB

```
> show measurements
```

```
name: measurements
```

```
name
```

```
----
```

```
stationalarm
```

```
stationparameter
```

```
> select * from stationalarm;
```

```
name: stationalarm
```

| time             | alarm_id | datapoint_id | station_id | value | valueThreshold |
|------------------|----------|--------------|------------|-------|----------------|
| ----             | -----    | -----        | -----      | ----  | -----          |
| 1487444343000000 | 308      | 121          | 1161115016 | 240   | 240            |

# In-memory databases

- **Databases use machine memory for storage**
  - Persist data on disks
  - Require very powerful machines
- **In principle it is not just about data models but also data management, data processing, software and hardware optimization, e.g.,**
  - SAP HANA, VoltDB: in memory relational databases
- **Why are in-memory databases important?**



Aalto University  
School of Science

# Interfaces between a data storage/databases system and its external analysis systems



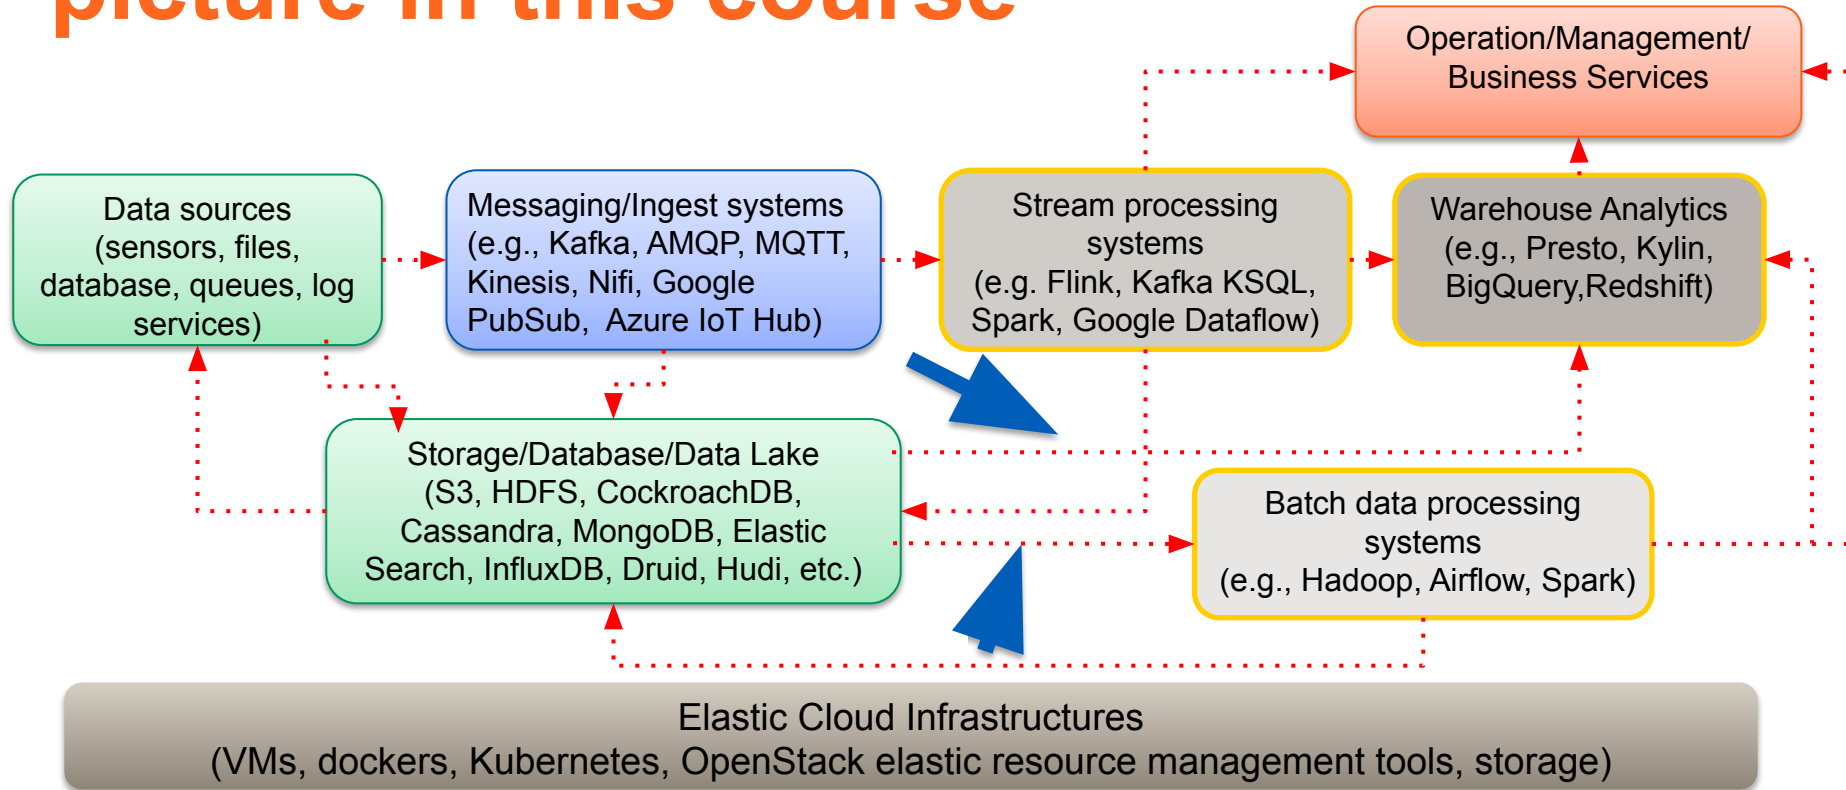
Aalto University  
School of Science

**In many cases: the data in data storage/database service must be made available for large-scale analysis:**

**large-scale analytics and data are managed by different systems**

**an important consideration in big data platforms design!**

# Big data at large-scale: the big picture in this course



# Making data available to the analytics

- **Data layer must map/provide data to processing layer**
  - maximize the analytics possibilities
- **Key issues**
  - avoid data movement as much as possible
  - avoid contention between the data management and the data analytics system
- **Techniques**
  - “mount”, specific connectors/drivers, copy-process-remove activities




# Mount/"Fuse"

- **Mapping a remote storage as a local file system**
  - Blobfuse (Microsoft Azure), gcsfuse (Google Storage)
  - the network performance is important

# Connectors

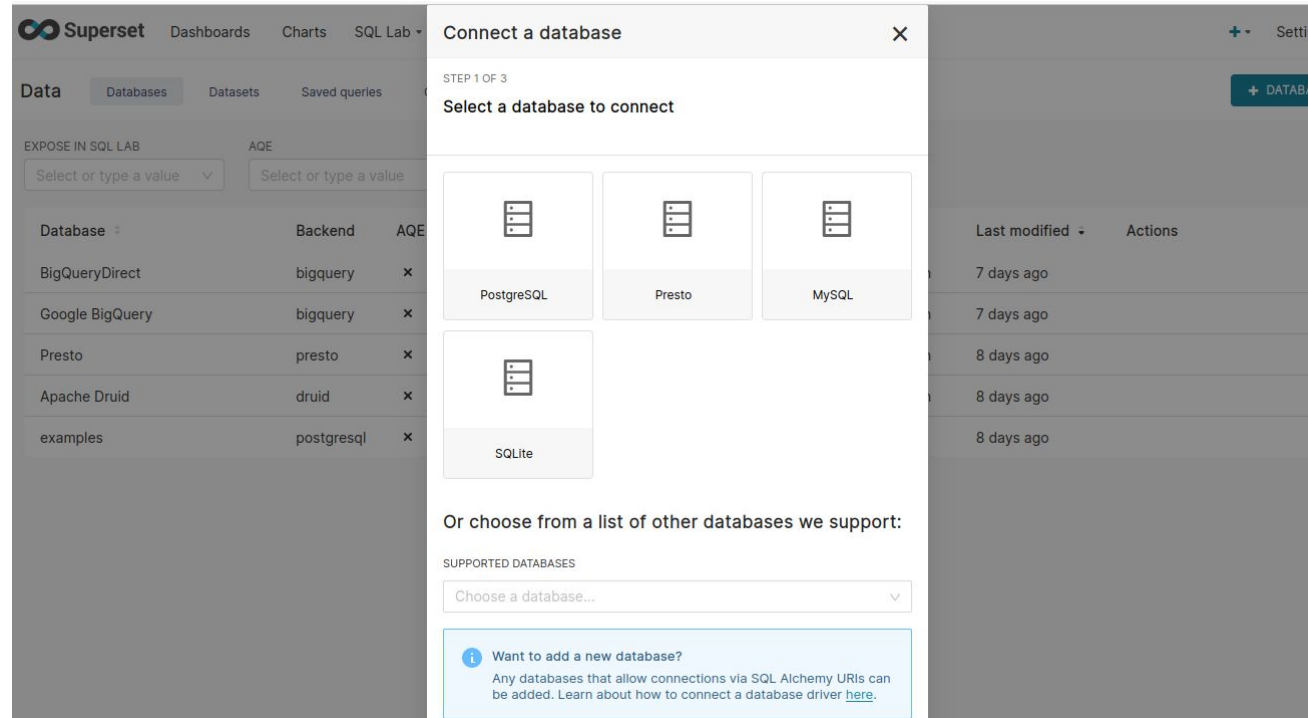
**ODBC or other specific protocol connectors to enable data access and ingestion!**



# Example: Superset connectors

<https://superset.apache.org/>


Connectors to different types of databases/datasets to retrieve and analyze data



# Example: Druid

<https://druid.apache.org/>

Different types of connectors (e.g., Kafka, Files, S3, etc.) to allow data ingestion into the database



# Analytics and Cloud Storage

- Various connectors for making data in cloud storages available for analytics
- Apache Hadoop/Spark (data analysis) can work with Amazon S3, OpenStack Swift, Google Cloud Storage
- Examples:
  - <https://github.com/GoogleCloudDataproc/hadoop-connectors>
  - <https://spark.apache.org/docs/latest/cloud-integration.html>

# “Copy and Process”

Client libraries are used to move data from storages and databases to processing places

Examples:

```
import pandas as pd
from sklearn.tree import DecisionTreeClassifier
from cassandra.cluster import Cluster

cluster = Cluster(contact_points=hosts, port=9042, auth_provider=auth_provider)
session = cluster.connect("tutorial12345")
sql_query = "SELECT * FROM tutorial12345.bird1234;"
df = pd.DataFrame()
rows = session.execute(sql_query)
df = rows._current_rows
print(df)
```

# Thanks!

**Hong-Linh Truong**  
**Department of Computer Science**

**rdsea.github.io**