



Aalto University  
School of Science

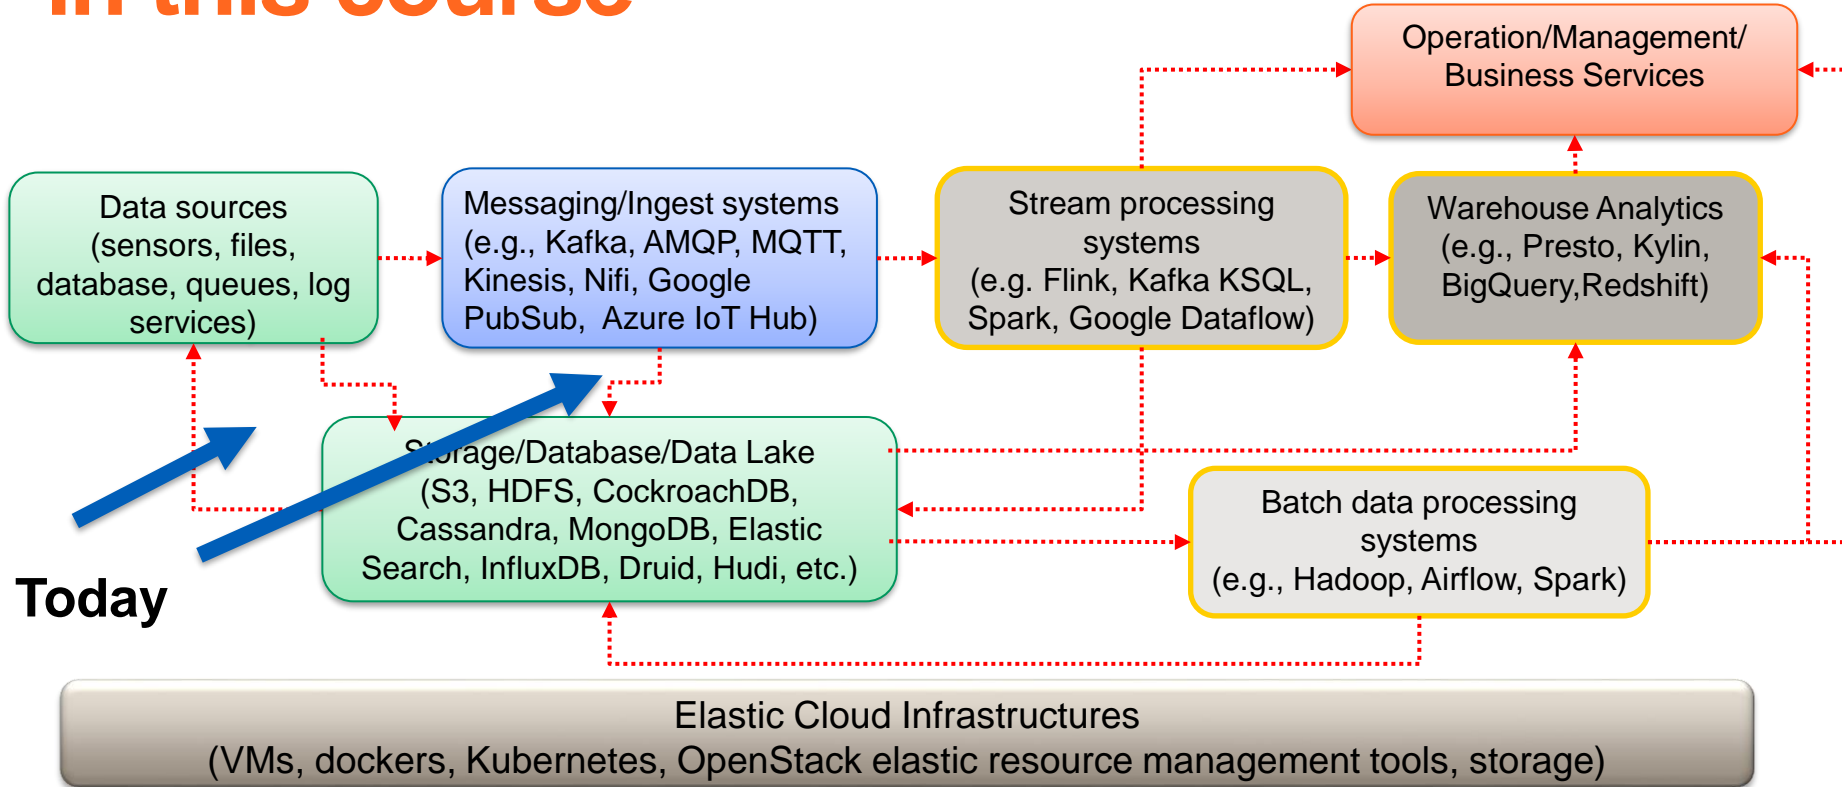
# Big Data Ingestion

*Hong-Linh Truong*  
*Department of Computer Science*  
*[linh.truong@aalto.fi](mailto:linh.truong@aalto.fi), <https://rdsea.github.io>*

# Learning objectives

- **Understand the overall design of data ingestion**
- **Study common tasks in data ingestion**
- **Understand and design efficient, robust data ingestion pipelines/processes**
- **Learn existing technologies/frameworks for your own design**

# Big data at large-scale: the big picture in this course



# Ingest big data into platforms

Data ingestion: Move data from different sources into the big data platform



Big data platform

e.g.

- logs of machines
- sell receipt  
transaction records
- IoT measurements

**Two important aspects:**

- requirements and tasks
- architectures, pipelines and service models

**Reusability and extensibility are very important!**

# Big Data Ingestion

- **Relation with ETL (Extract, Transform, Load)**
  - During ingestion, some transformation tasks might be needed
  - ETL has many operations to deal with the semantics/syntax of data and the business of data
- **Transformation within ingestion**
- **Transformation done within the (target) platform**

**Correctness and quality assurance are hard!**

# Fundamental ingestion models

- **Batch ingestion**

- Data is in files
- Ingestion can be done in batches of files or batches of parts of files

- **Files**

- CSV, Text, JSON, ARVO
- Other typical formats (video, images, etc.)

- **(Near) real-time ingestion**

- Data is encapsulated into messages
- Ingest data as soon as the data is available
- Messaging systems are needed

- **Messages**

- Text/CSV/JSON, ARVO
- Application-specific designs

# Data source and sinks

## Data sources

Input files

REST  
Services

Messaging  
Systems  
(MQTT,  
KafKa, etc)

Databases

Ingestion  
Tasks

## Data sinks

Storage/File  
Systems

Output  
files

Big Database

Big data store systems

## Big data platform examples:

Hadoop File systems  
Google Storage  
Amazon Storage

Druid  
Google BigQuery  
Hive  
MongoDB  
ElasticSearch  
Cassandra  
InfluxDB  
Hudi

# Requirements from V\* of big data

- **Requirements from access API and protocols**
  - REST API, ODBC, SFTP, specific client libs
  - MQTT, AMQP, CoAP, HTTP, ...
- **Requirements from data**
  - structured, unstructured and semi-structured
  - speed, volume, accuracy, confidentiality, data regulation
- **How deep a platform can support?**
  - able to go into inside of data elements (understanding the syntax and semantics of data)?



# Ingestion tasks: common tasks and requirements

# Main tasks in ingestion

- **Key categories of tasks**
  - Data access and extraction
  - Data routing
  - Data wrangling
  - Data storing
  - Quality assurance/governance (quality check, anonymizing data)
- **Customer/user tasks vs platform tasks**
- **Other supports: compression, end-to-end security**

# Data access and extraction tasks

- **Access**
  - Obtaining/copy data from sources, change data capture (CDC)
  - Often built based on common protocols and APIs
  - Reusability is important!
- **Encryption, masking/anonymization**
  - Might need to be done when accessing and extracting data
  - Also during transfers of data
  - data security requirements, personally identifiable information

# Change data capture becomes important for big data ingestion

- **The principles:**
  - Capture and ingest only new data by listening data changes
  - Leverage many features of databases (update, query, insert operations), data stream offsets and status notification (e.g., the availability of new files)
- **You see implementation in different tools like Redhat Debezium, Hudi DeltaStreamer, Kafka connect**

# Dealing with data structures

- Remember that the data sender and the receiver are **diverse**
  - In many cases, they are not in the same organization
    - *You need to guarantee the message syntax and semantics*
- **Solutions**
  - Agreed in advance → in the implementation or with a standard
  - Know and use tools to deal with **syntax differences**
- **Understanding the syntax allows some automatic transformations/quality check**
- **But semantics are domain/application-specific**

# Example: Arvo

Syntax specification

<https://avro.apache.org/>

```
{  
  "namespace": "bdp.courses.aalto.fi",  
  "type": "record",  
  "name": "event",  
  "fields": [  
    {"name": "station_id", "type": "string"},  
    {"name": "datapoint_id", "type": "int"},  
    {"name": "alarm_id", "type": "int"},  
    {"name": "event_time", "type": "int"},  
    {"name": "value", "type": "float"},  
    {"name": "valueThreshold", "type": "float"},  
    {"name": "isActive", "type": "boolean"}  
  ]  
}
```

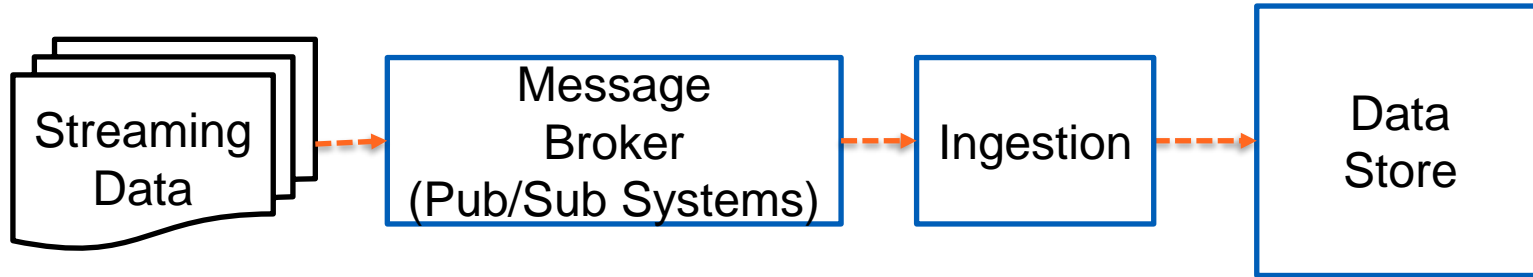


# Target file formats in big data storage

- **Parquet, <https://parquet.apache.org/>**
  - Columnar storage (optimizing for reading columns), big files, compression features
  - In Hadoop ecosystem/Spark (thus also available in Druid, Hudi), Azure, S3, etc.
- **ORC, <https://orc.apache.org/>**
  - Large-scale files, self-describing data and metadata, available in Hive, support ACID, multiple-level of indexes and complex types
- **Many big databases/storage and datalakes use them as the storage level**
  - Still allow SQL-style or other types of analytics

Reading: an Uber blog bout file formats and performance: <https://eng.uber.com/cost-efficiency-big-data/>

# How do we move streaming data into big data databases/storage?



**Protocol**

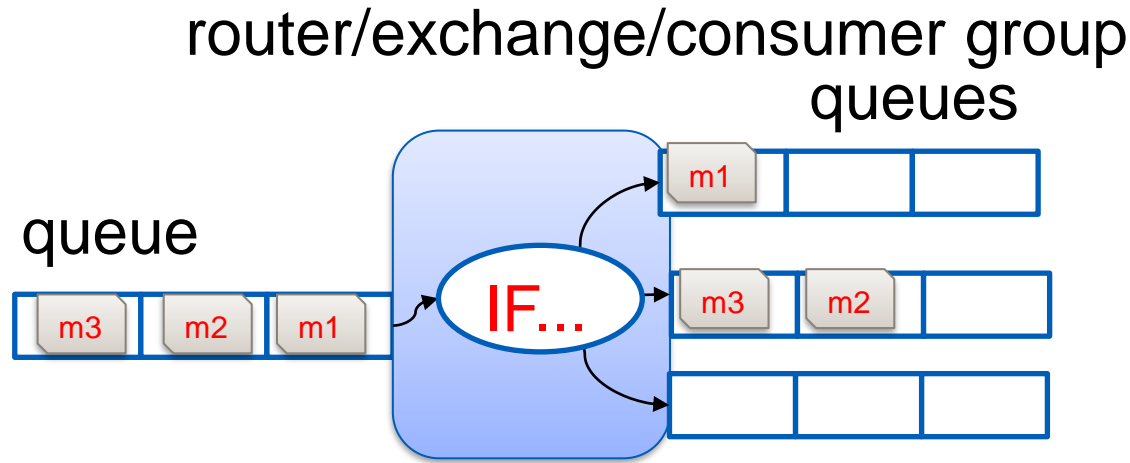
**Data format**

**Message structure**

**Basic streaming data processing techniques**

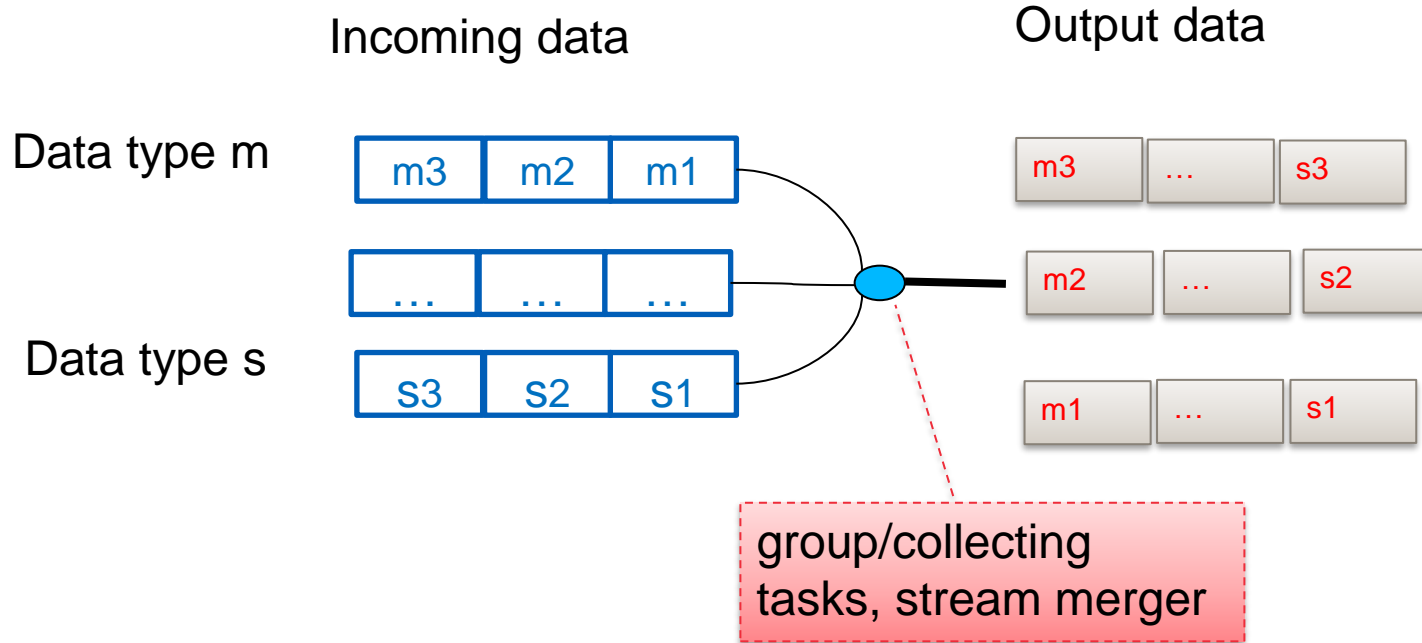


# Use split tasks/distributor patterns to separate data for data parallelism processing



**Read the famous book: “Enterprise Integration Patterns”**  
**<https://www.enterpriseintegrationpatterns.com/patterns/messaging/>**

# Data routing: grouping data/collector pattern using stream merging/common topics



# Data wrangling

- **Convert data from one form to another**
  - Cleaning, filtering, merging and reshaping data
- **Require access to the data!**
- **Key design choices:**
  - do you support it during the ingestion or after the ingestion?
  - as a platform provider: are you able to do this?

# Data wrangling

- **In the context of big data platforms**
  - Automatic data wrangling: write pipelines/programs which do the wrangling
- **Wrangling programs provided by customers**
  - Needs the platform to support debugging, monitoring and exception handling
  - Runtime management for wrangling
- **Wrangling programs provided by platforms**
  - Constraints in dealing with customer data

# Quality control/data regulation assurance

## Data sources

Responsible data: profiling, sampling, measuring quality and inspecting data  
→ implications on data products

Log file

Transaction records

User-provided data

Collect data

Parse & profile data

Patterns/rules/AI

Databases  
Data monitoring

Hot issues: misinformation, GDPR, data quality, inappropriate content

# Examples: Logstash Grok – a kind of domain specific language?

**Grok is for parsing unstructured log data text patterns into something that matches your logs.**

**Grok pattern syntax: `%{SYNTAX:SEMANTIC}`**

**Regular and custom patterns**

**A lot of exiting patterns:**

- <https://github.com/logstash-plugins/logstash-patterns-core/tree/master/patterns>

**Debug Tools: <http://grokdebug.herokuapp.com/>**

# Example with NETACT Log

```
29869;10/01/2017 00:57:56;;Major;PLMN-PLMN/BSC-xxxxxx/BCF-xxx/BTS-  
xxx;XYZ01N;ABC08;DEF081;BTS OPERATION DEGRADED;00 00 00 83 11  
11;Processing
```

## Simple Grok

```
1 input {  
2   file {  
3     path => "/tmp/alarmtest2.txt"  
4     start_position => "beginning"  
5   }  
6 }  
7 filter {  
8   grok {  
9     match => {"message" => "%{NUMBER:AlarmID};%{DATESTAMP:Start};%{DATESTAMP:End};%{WORD:Severity};%{NOTSPACE:NetworkType};%{NOTSPACE:BSCName};%{NOTSPACE:Sta  
10  }  
11 }  
12 output {  
13   stdout {}  
14   csv {  
15     fields => ['AlarmID', 'Start', 'Stop', 'Severity', 'NetworkType', 'BSCName', 'StationName', 'CellName', 'AlarmInfo', 'Extra', 'AlarmStatus']  
16     path => "/tmp/test-%{+YYYY-MM-dd}.txt"  
17   }  
18 }
```

# Examples

Write your  
own code with  
Pandas and  
Data frame?



Automatically  
generate code  
for wrangling?

```
Alarms={}
with open(sys.argv[1], 'rb') as csvfile:
    reader = csv.DictReader(csvfile)
    for row in reader:
        try:
            #print row['Started']
            alarm_time = datetime.strptime(row['Started'], '%d.%m.%Y %H:%M:%S')
            #diff = start_time - alarm_time
            #print "different time is ",diff
            if alarm_time >=start_time:
                #print(row['RNW Object Name'], row['Severity'])
                typeOfAlarm = 0
                cleanSeverity = re.sub('\W+', '',row['Severity'])
                if (cleanSeverity in mobifone.AlarmSeverity.keys()):
                    typeOfAlarm = mobifone.AlarmSeverity[cleanSeverity]
                #print ("Type of Alarm: ",typeOfAlarm)

                if row['RNW Object Name'] in Alarms:
                    #print "Again"
                    severies =Alarms[row['RNW Object Name']];
                    severies[typeOfAlarm]=severies[typeOfAlarm]+1
                else:
                    severies =[row['RNW Object Name'],0,0,0,0,0,0]
                    severies[typeOfAlarm]=severies[typeOfAlarm]+1
                    Alarms[row['RNW Object Name']]=severies;

        except:
            print "Entry has some problem"
            print row
            #timestamp =long(row['TIME'])
            #times.append(datetime.datetime.fromtimestamp(timestamp/1000))
            #times.append(long(row['TIME']))
            #signals.append(float(row['GSM_SIGNAL_STRENGTH']))
dataframe =pd.DataFrame(Alarms,index=mobifone.AlarmSeverityIndex).transpose()
alarmdata =dataframe.as_matrix();
#TODO print Alarms to file
#only for debugging
print dataframe
dataframe.to_csv(outputFile, index=False)
```



# Ingestion tasks implemented as extensible, composable connectors

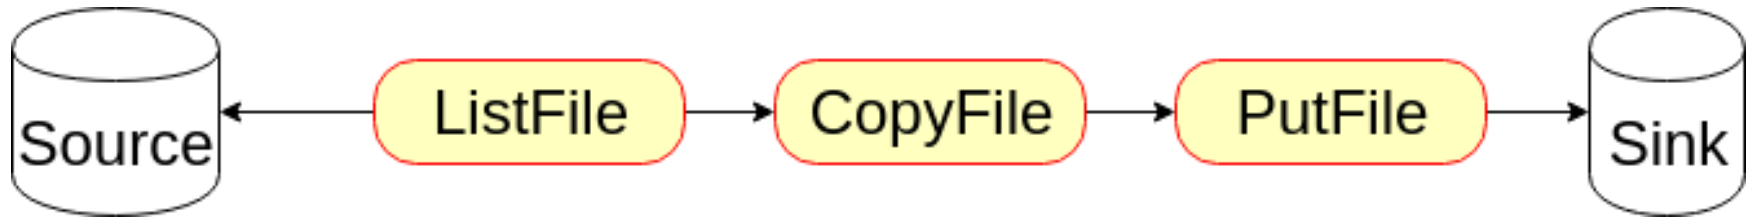
- **Basic tasks for big data ingestion can be used in different cases**
- **Support end-user tasks**
  - Platform enables the user to do many tasks through configurations
- **Enable pluggable approaches is important**
  - Input data plugin/component → filter/extract/convert → output data plugin/component
- **Data compression and security must be considered**

# Ingestion is not a single task!

## Ingestion pipelines/processes: architectures and tools

# Complex deployment and composition models

- Understanding strong dependencies between protocols/APIs, **security, performance and management**

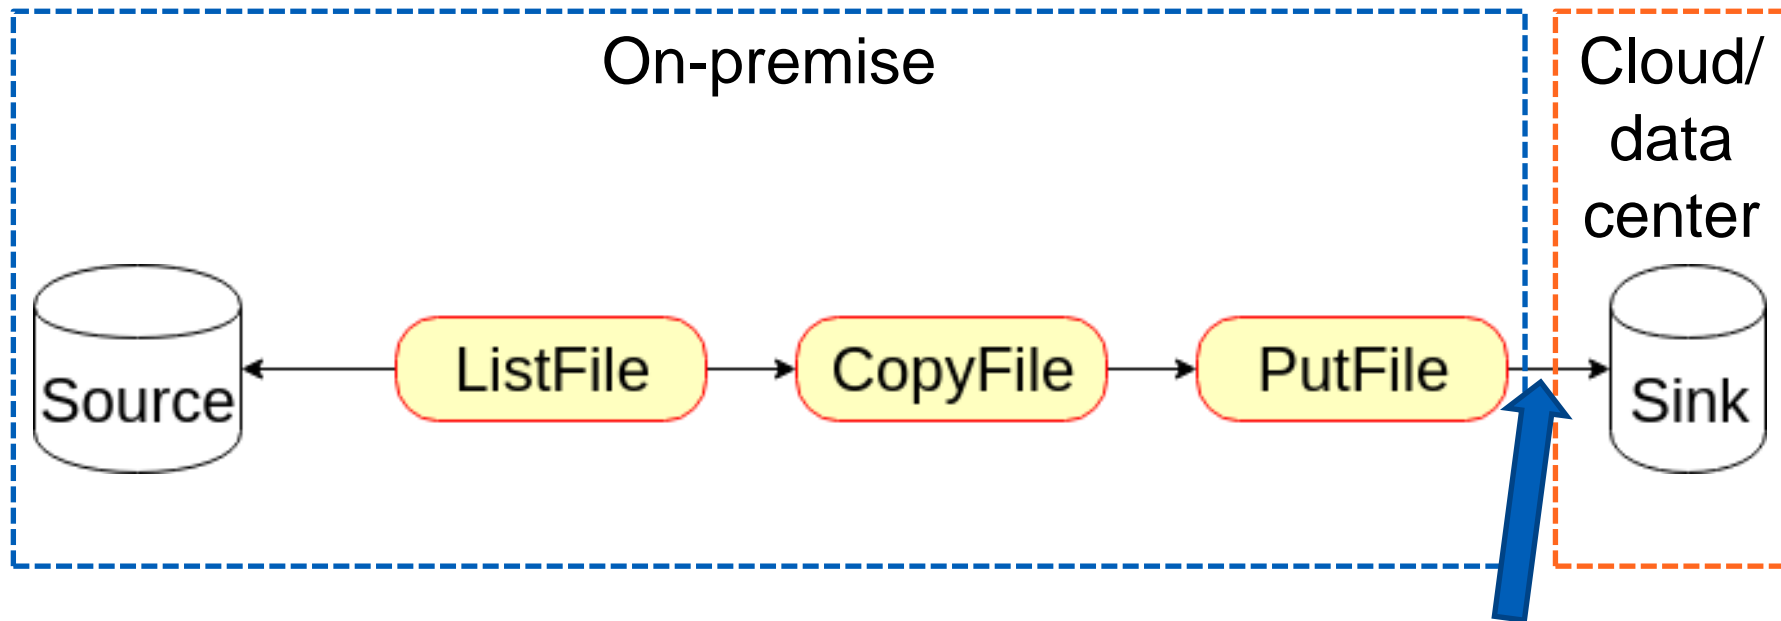


**Customer**

**Ingestion pipeline developer  
( for whom?)**

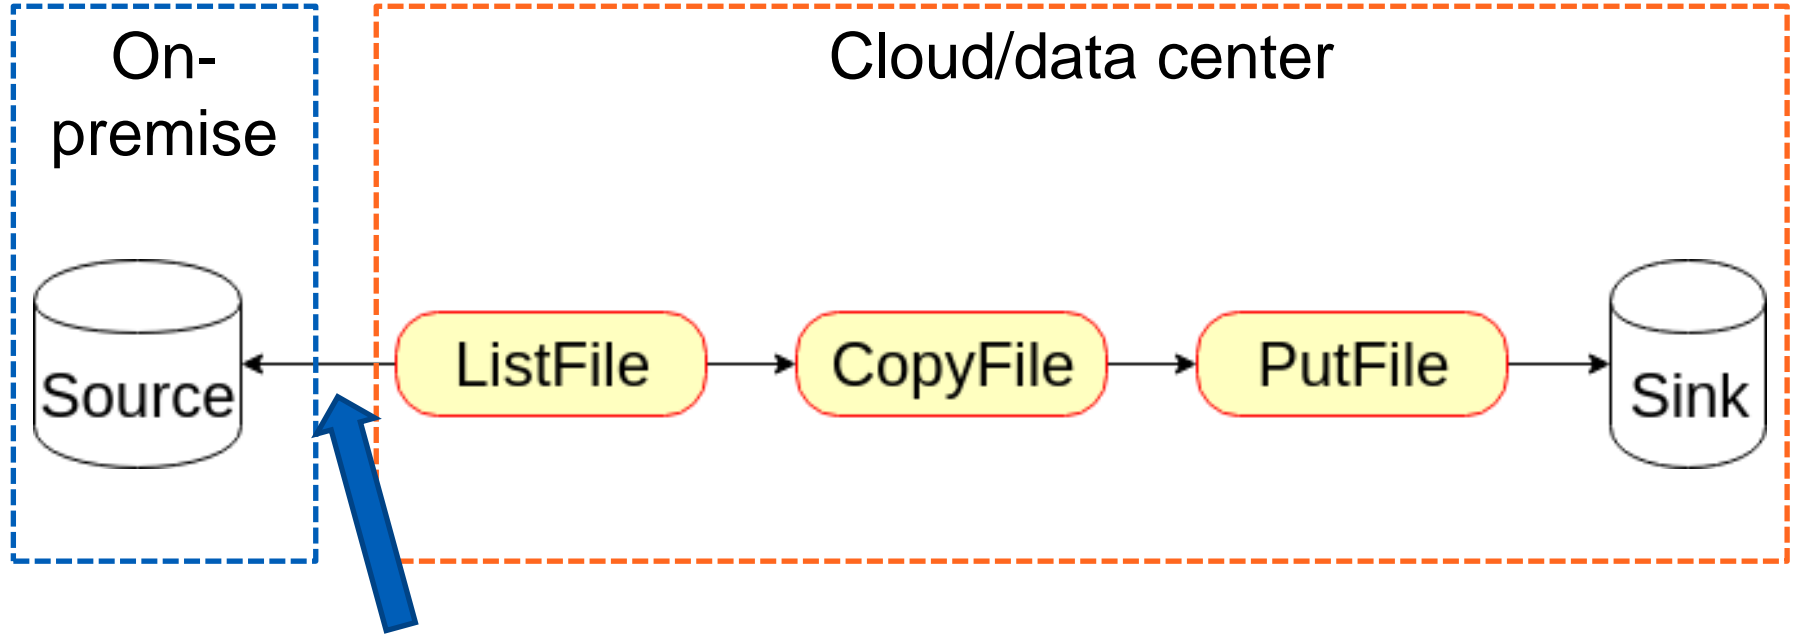
**Data  
store/platform  
provider**

# Complex deployment and composition models



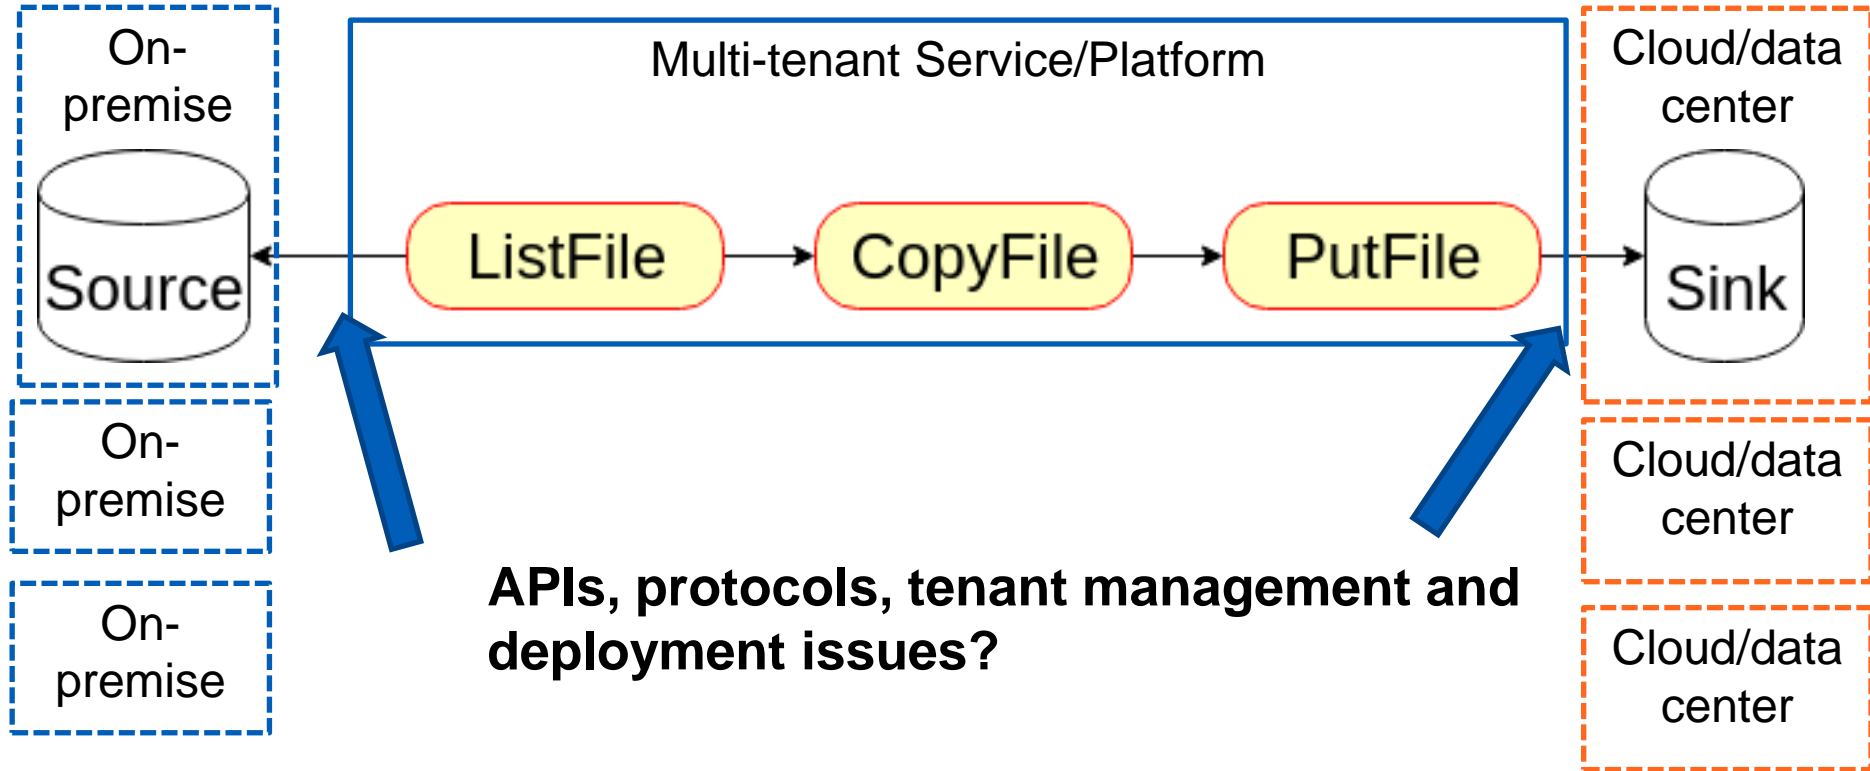
**APIs, protocols and deployment issues?**

# Complex deployment and composition models



**APIs, protocols and deployment issues?**

# Complex deployment and composition models



# Pipeline designs and execution models

# Architecture requirements

- **Data source integration**
  - The richness and extensibility of data sources and data sinks
- **Batch ingestion and near real-time ingestion requirements**
- **Integration between different ingestion processes across distributed places**
- **The architecture addresses “big data” properties**




# Batch ingestion processes

- **Data to be ingested is bounded**
  - files or messages are finite
- **Ingestion architectural styles**
  - (1) Direct APIs, (2) reactive pipelines, (3) workflows
- **Incremental ingestion**
  - Dealing with the same data source but the data in the source has been changed over the time (related to change data capture)
- **Parallel and distributed execution**
  - Use workflows and distributed processing

# Simple, direct APIs for ingestion

**Pull model:** register webhook/API

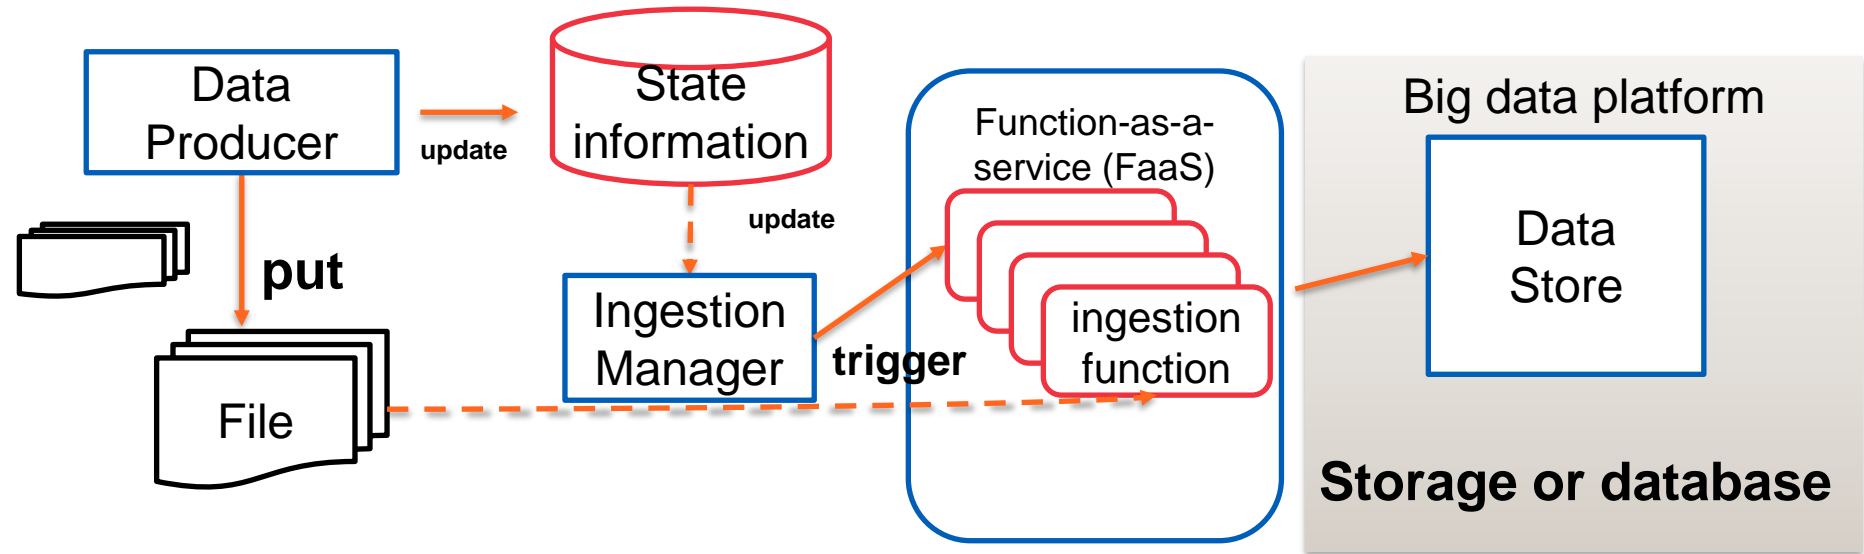


**Push model**  
When?



## Try to analyze pros and cons for your platform?

# Reactive with function-as-a-service



## Who develops which components?

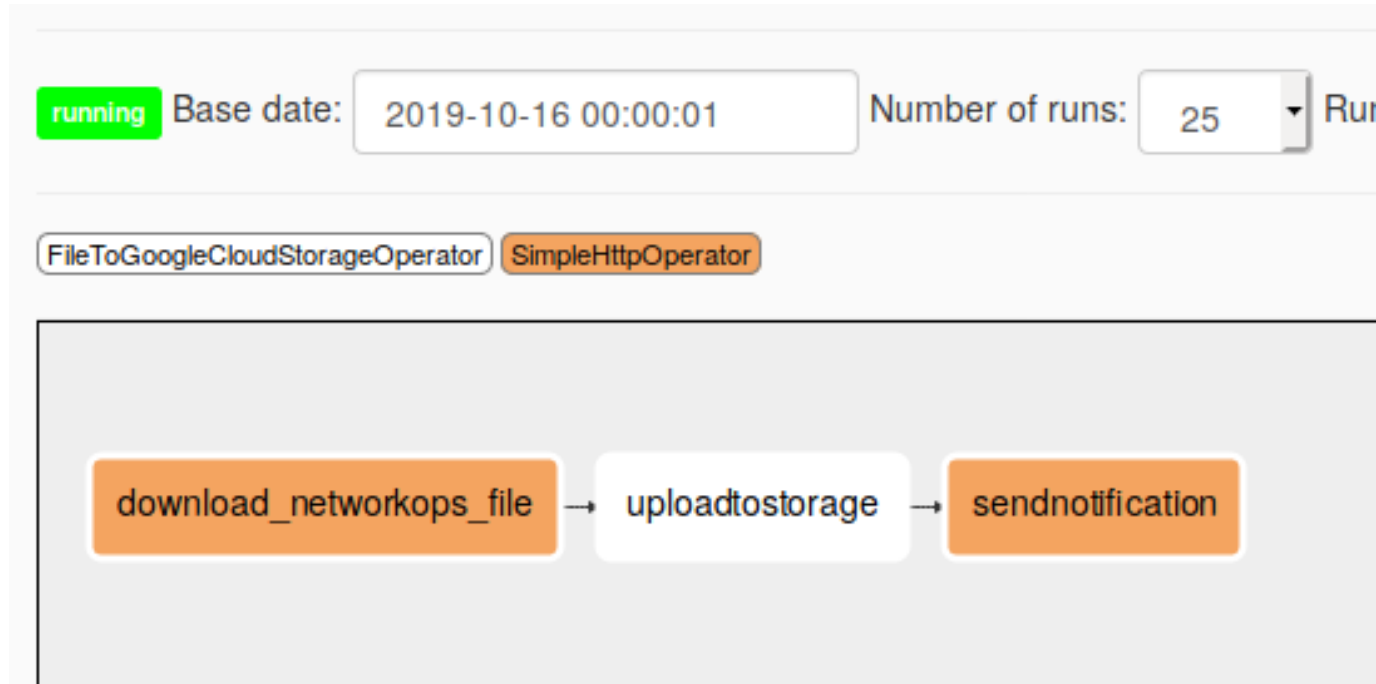
Remember?:

<https://version.aalto.fi/gitlab/bigdataplatforms/cs-e4640/-/tree/master/tutorials/queuebaseddataingestion>

# Orchestrating ingestion workflow

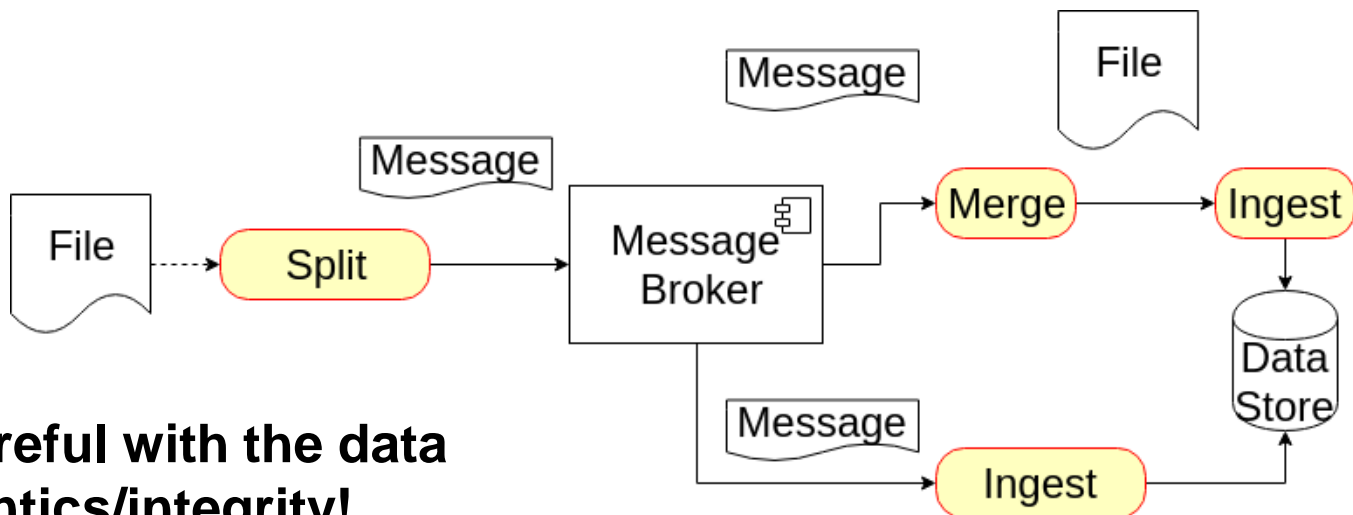
- **Different tasks for**
  - Access and copy, extract, covert, quality check, and write data
  - Tasks can be connected based on data or control flows
- **Workflows**
  - A set of connected tasks is executed by an engine
  - Tasks can be scheduled and executed in different places
- **Bulk ingestion can be done using workflows**

# E.g., workflow based on scheduled time, with Apache Airflow



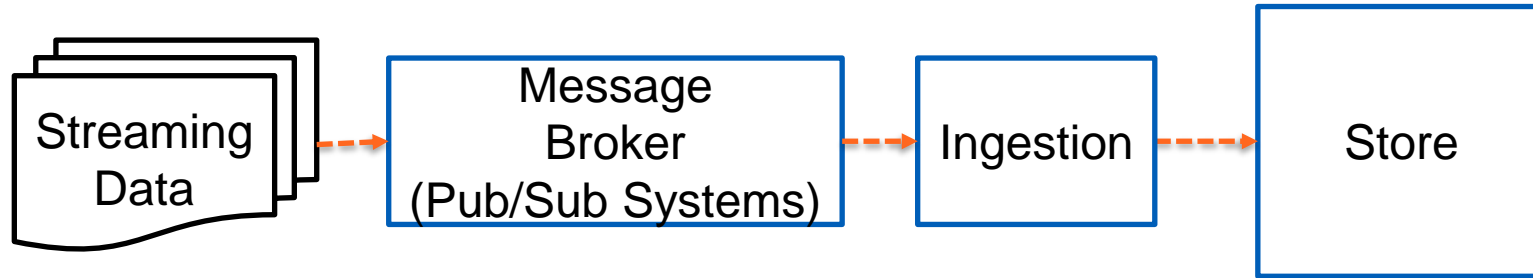
# Microbatching for ingestion

- **Data is split into different chunks ingested using a batch**
  - Using “streaming” to send chunks
  - Chunks are ingested into the system, or merged and then ingested



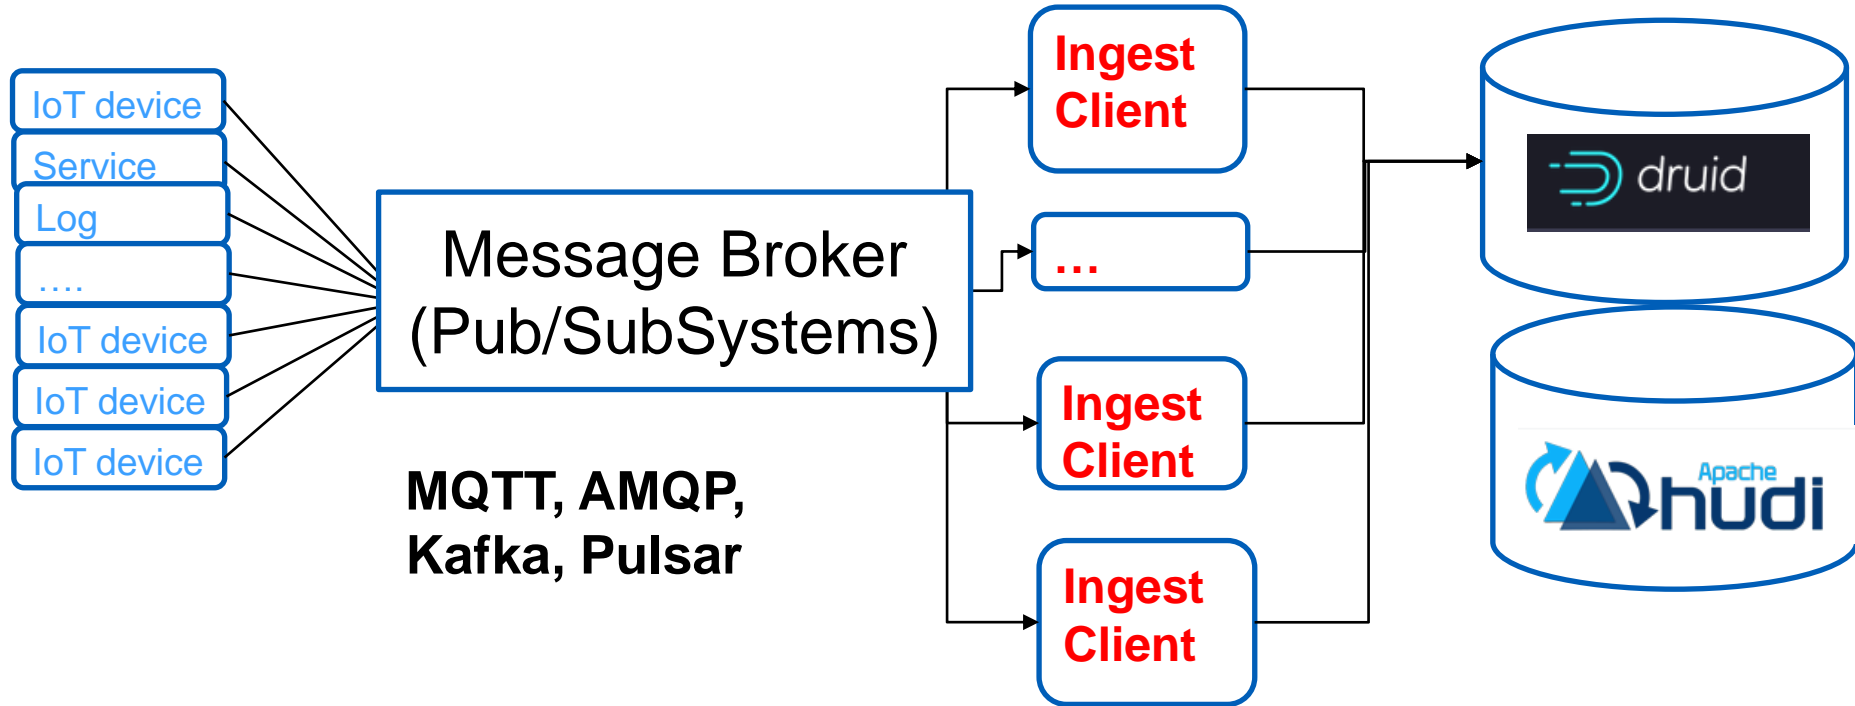
**Be careful with the data semantics/integrity!**

# Near-real time ingestion processes



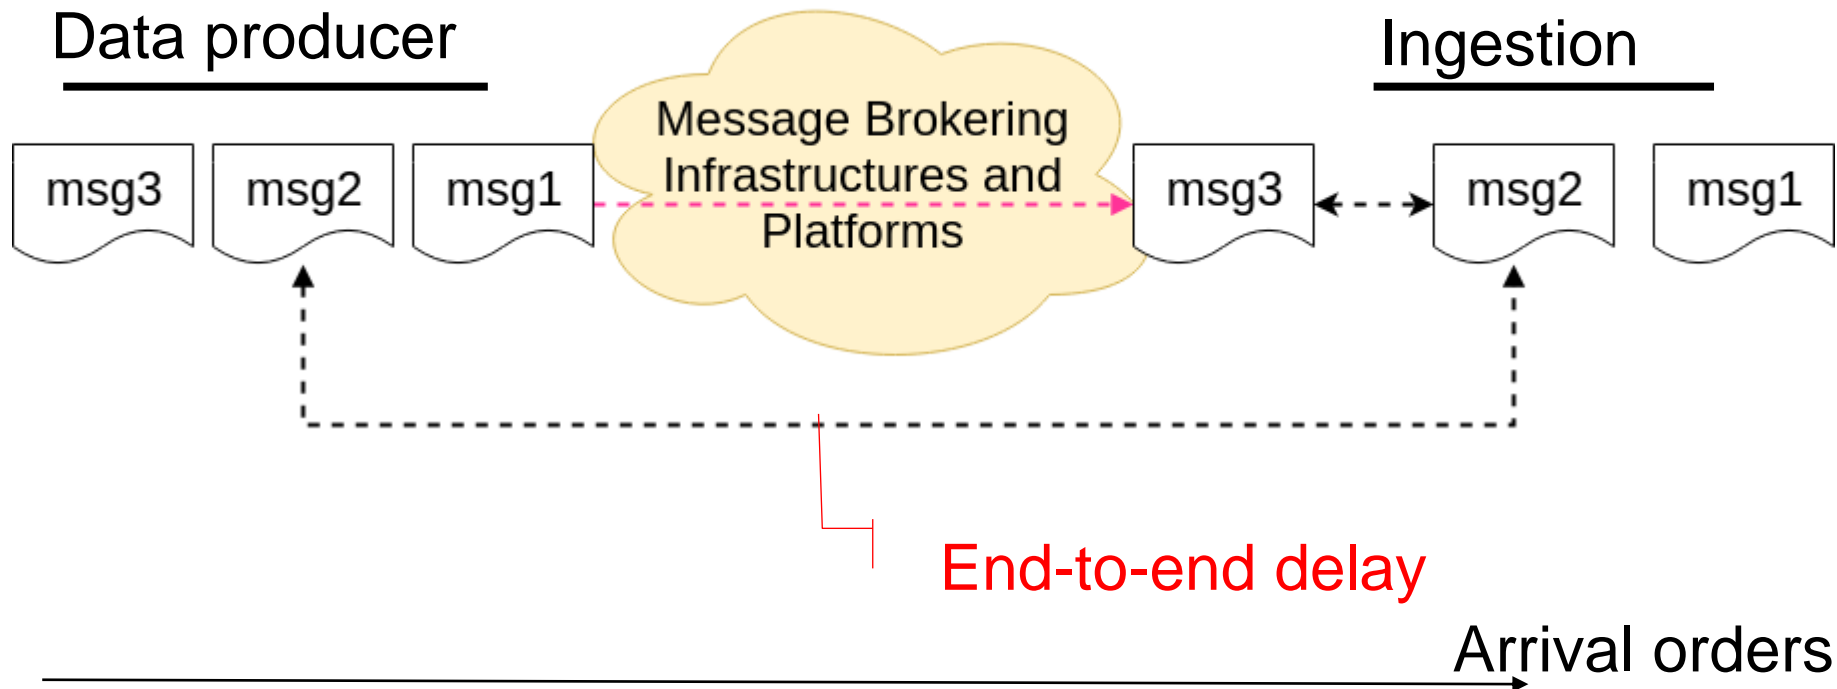
- **Moving streaming data**
- **Unbounded data, amount of data varies, fast ingestion**

# Example





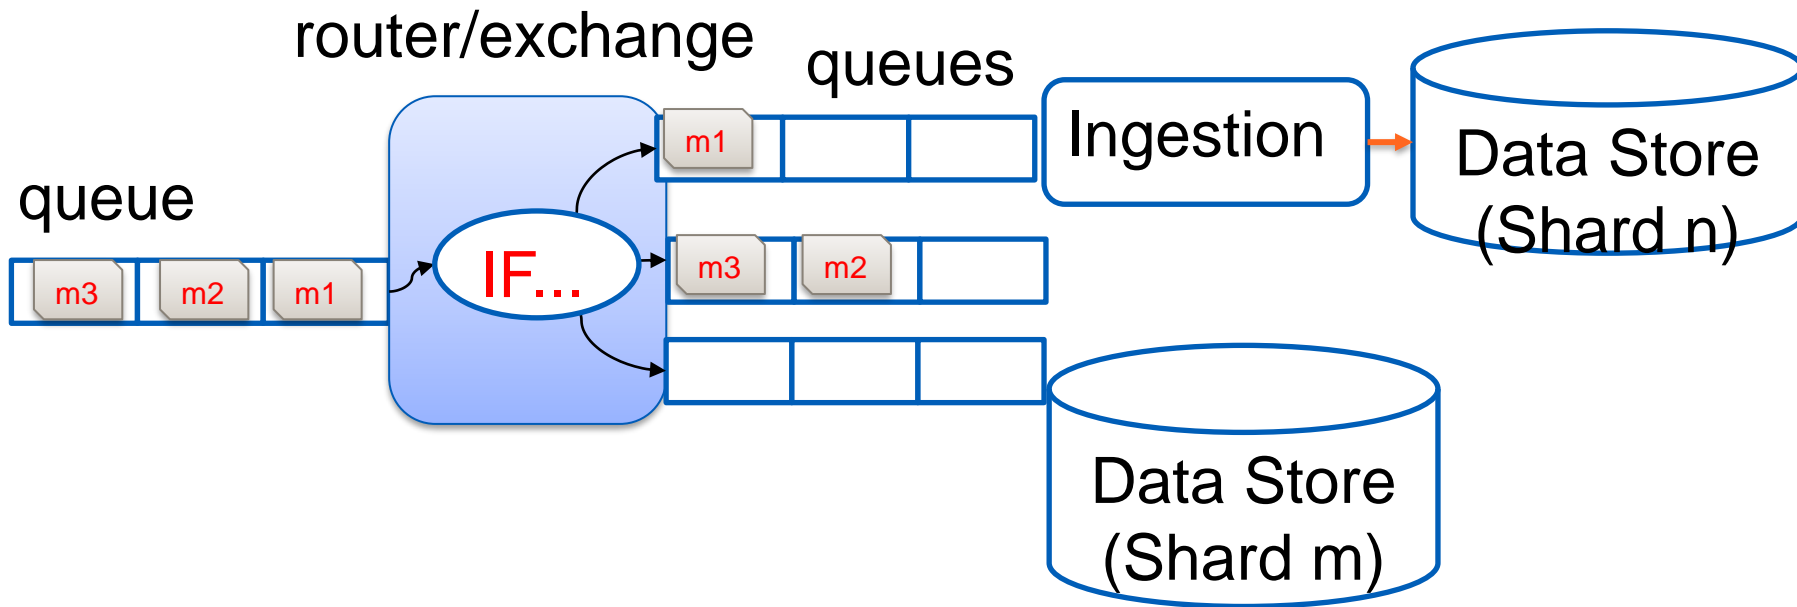
# Key issues in streaming data ingestion



# Some key issues

- **Late data, data out of order?**
- **Exactly once?**
- **Back pressure and retention**
  - for individual components or the whole pipelines
- **Scalability and elasticity**
  - changes in data streams can be unpredictable

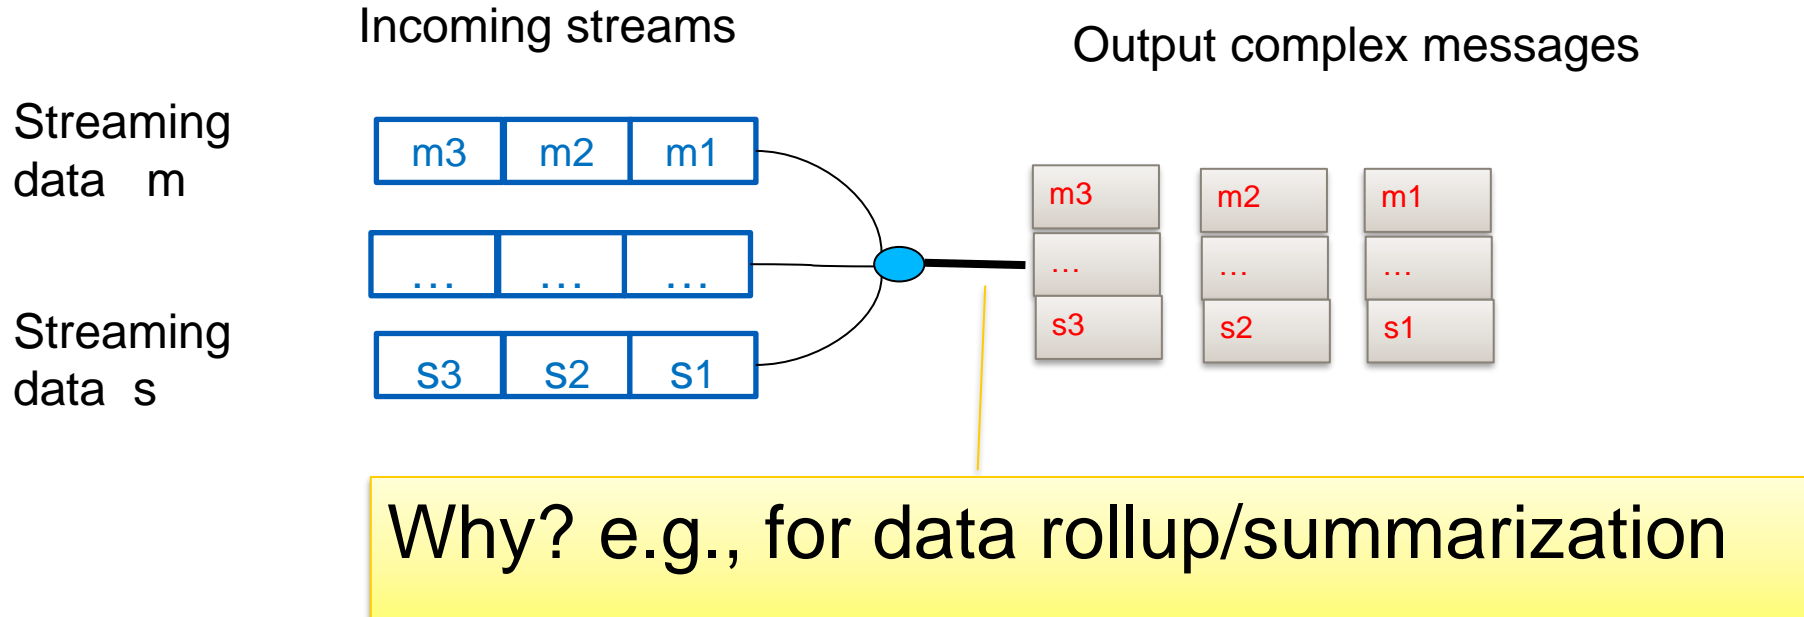
# Split (pub/sub) and partition with ingestion



# Some key issues

- **Multiple topics/streams of data**
  - amount of data per topic varies
  - should not have duplicate data in data store
- **How to distribute topic/data to ingestion clients?**
- **Where should we run the message broker?**
- **Where should the elasticity be applied?**

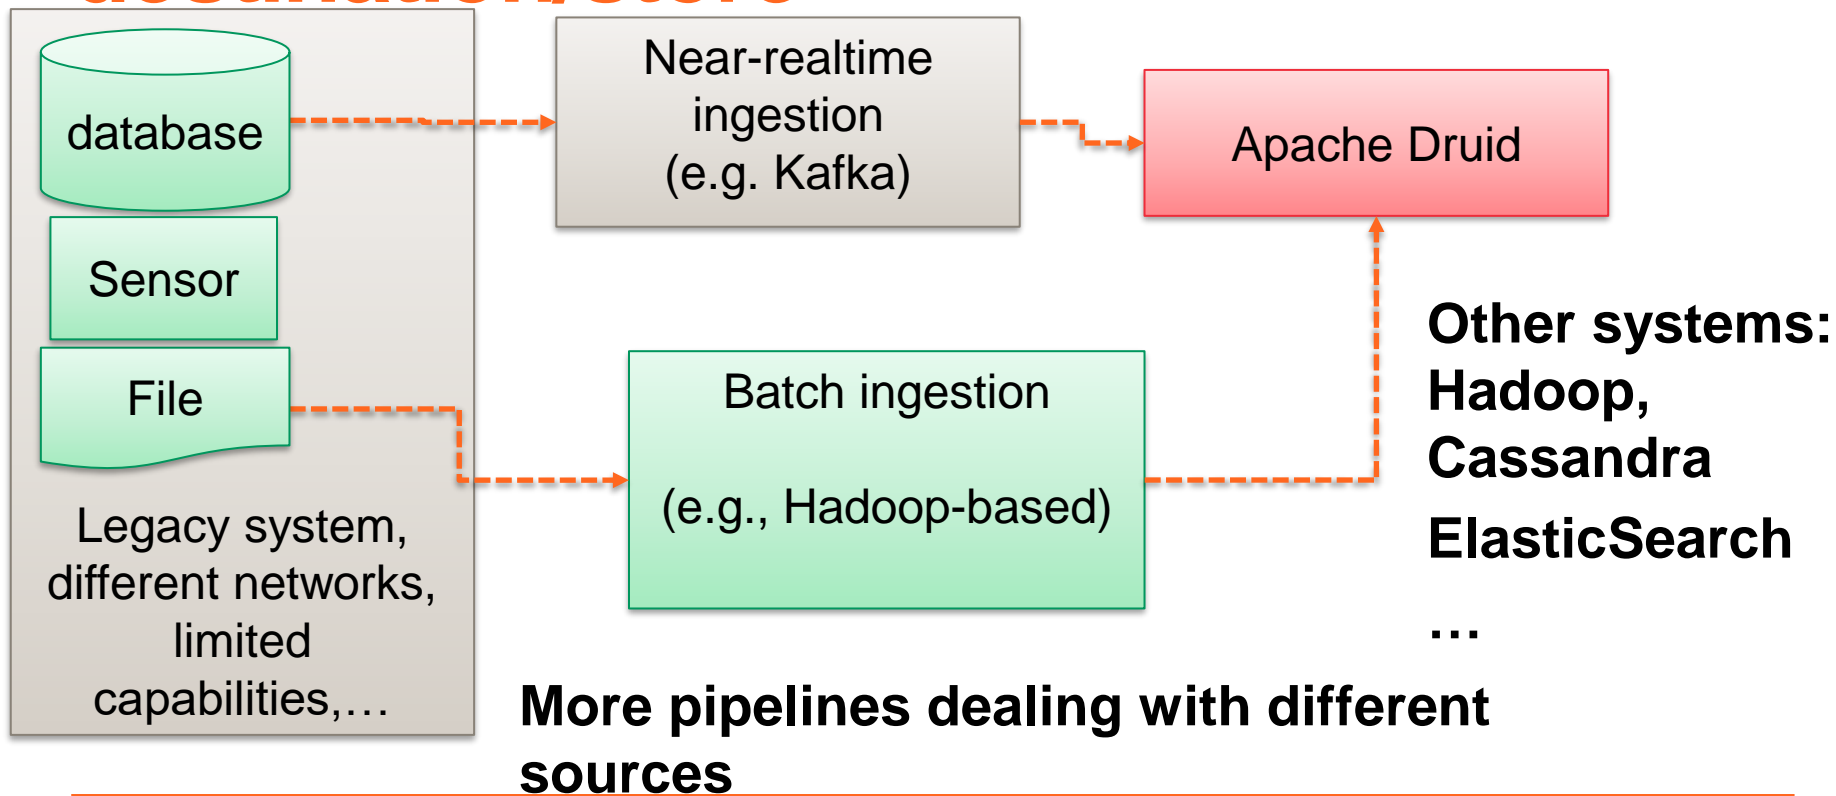
# Do we have to merge data before ingestion



# Complex ingestion pipelines in big data platforms


- Multiple types of pipelines for multiple types of customers
  - A customer might need different integrated pipelines
- Both batch and near-realtime ingestion are supported
- Complex architectural designs
    - Ingestion pipeline-to-pipeline needs “bridges”

# Multiple types of pipelines for the same destination/store



# Connecting different ingestion pipelines

A single tool might not be enough



**Real-world:**  
**both pipelines and their connection are complex**



# Data ingestion with (emerging) Data Lake

Data Lake provides single store for multiple types of data → reduce effort in building ingestion pipelines

**Example with  
Delta Lake**  
(<https://delta.io/>)

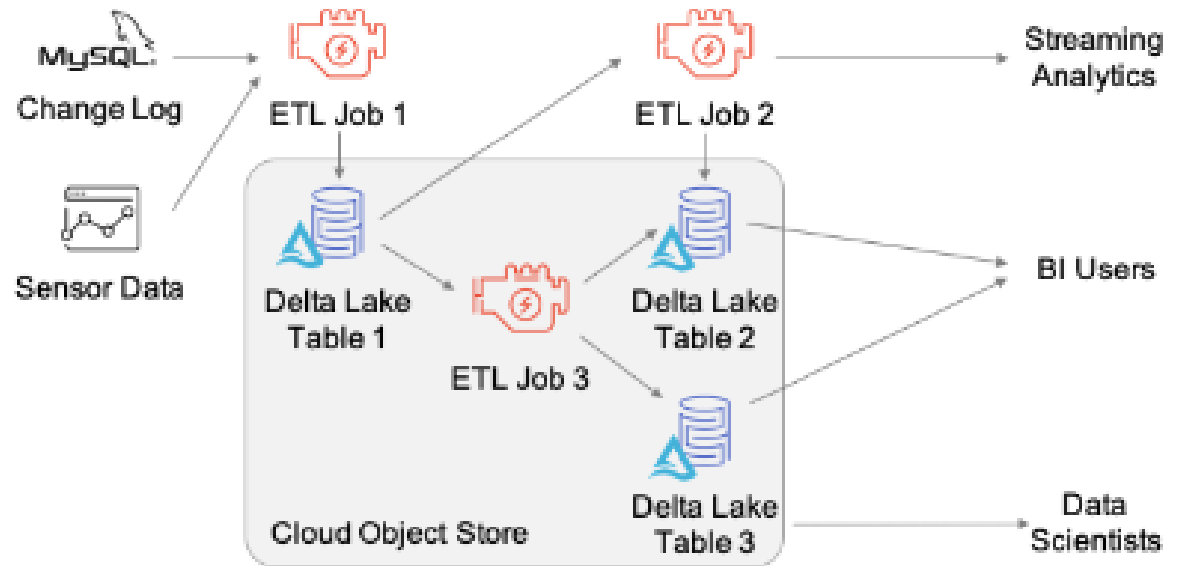


Figure source: “Delta Lake: High-Performance ACID Table Storage over Cloud Object Stores”, <https://databricks.com/wp-content/uploads/2020/08/p975-armbrust.pdf>

# Tooling and examples

# Tooling

- **Given different ingestion models, how do you deliver your ingestion tools/services?**
- **(Traditional) ways of REST API/specific client libraries**
  - Upload using put/get operations
- **Workflows**
  - Self-developed workflows vs automatically generated workflows
- **Pipelines are bundled into containers**
  - Self-developed vs generic pipelines based on user configurations

# Design tools for ingestion processes: Apache Kafka + various data sinks

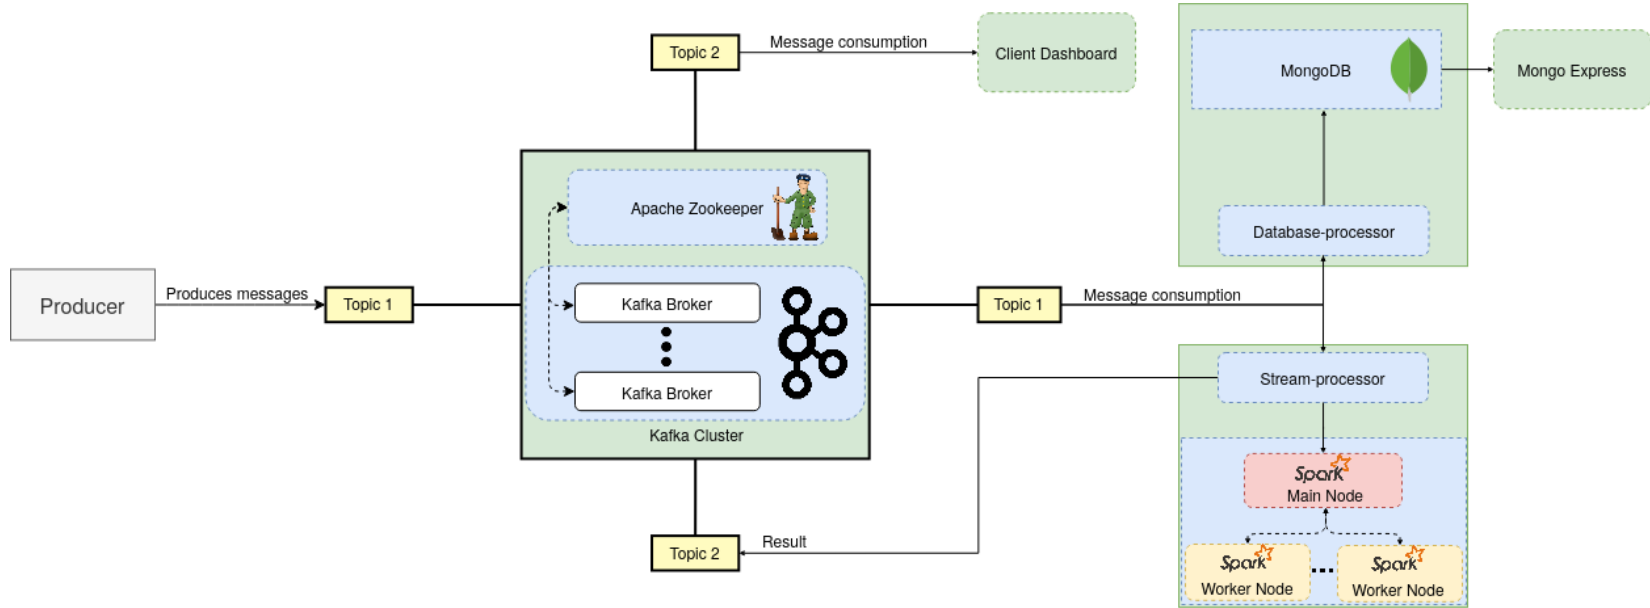


Figure source: <https://version.aalto.fi/gitlab/bigdataplatforms/cs-e4640/-/tree/master/tutorials/cloud-data-pipeline>

# Design tools for ingestion processes: Logstash

- **For managing logs and events**
  - Collect data from various connectors
  - And parse and store the results through various connectors
- **Programming**
  - Focus on making pipelines of pluggable components
  - Both programming and configuration deployment needed
- **Deployment**
  - Individual deployment or pipelines
- **Work very well with Elasticsearch**

# Design tools for ingestion processes: Logstash

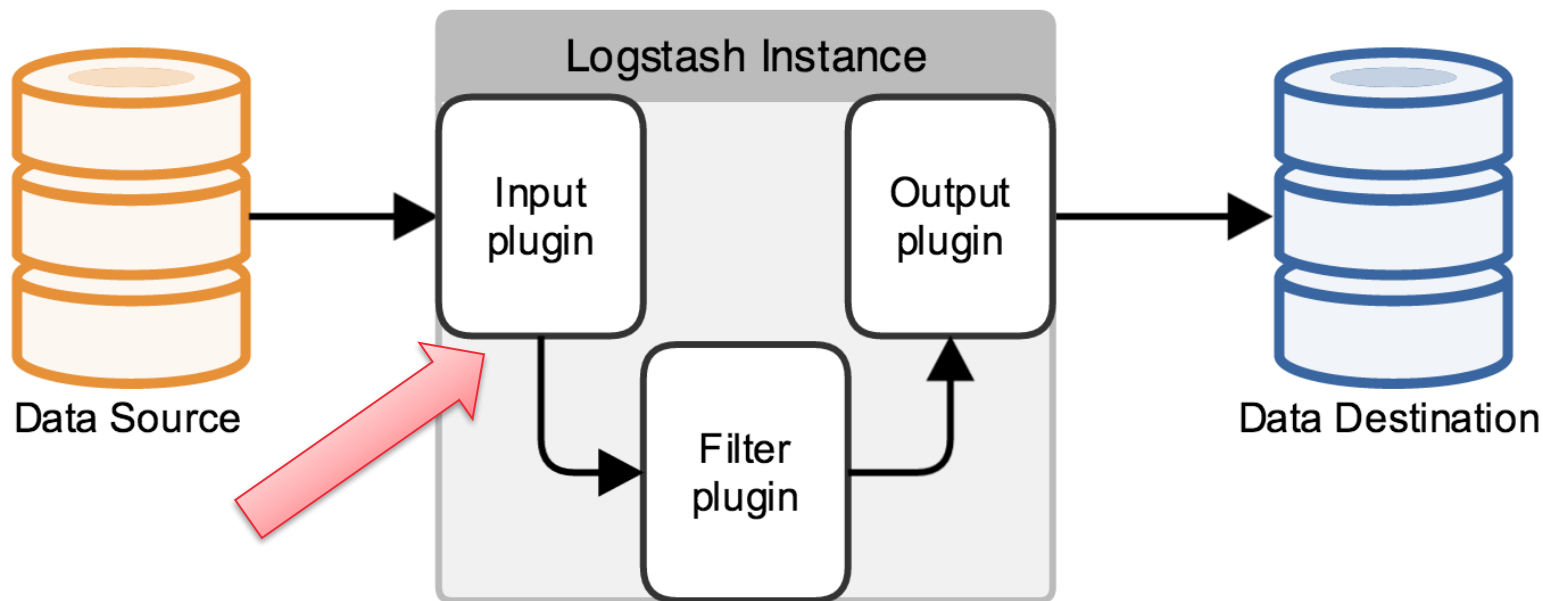
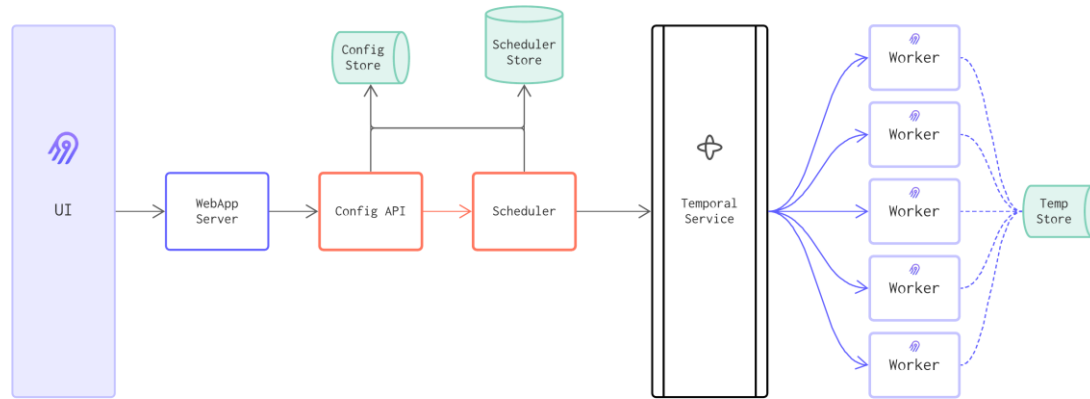


Figure source:  
<https://www.elastic.co/guide/en/logstash/current/advanced-pipeline.html>

## Pluggable approaches

# Design tools for ingestion processes: Airbyte

Allow the user defines input and output configuration then create and deploy containers including ingestion code



**Connectors for multiple data sources**

**Platform with scheduler, jobs, workers for data ingestion**

Figure source: <https://docs.airbyte.com/understanding-airbyte/high-level-view>

# Design tools for ingestion processes: Apache Druid

Allow the user to build the plan: select tasks, configuration, etc.  
and then generate ingestion pipelines

The screenshot displays the Apache Druid ingestion tool interface, which is divided into four main tabs: "Connect and parse raw data", "Transform data and configure schema", "Tune parameters", and "Verify and submit". The "Transform data and configure schema" tab is currently active, showing a list of tasks and their configurations. The tasks are listed in a table with columns for task ID, start time, end time, and status. The tasks are: VendorID, tpep\_pickup\_datetime, tpep\_dropoff\_datetime, passenger\_count, trip\_distance, RatecodeID, store\_and\_fwd\_flag, PULocationID, DOLocationID, payment\_type, and trip\_start\_datetime. The tasks are listed in a table with columns for task ID, start time, end time, and status. The tasks are: VendorID, tpep\_pickup\_datetime, tpep\_dropoff\_datetime, passenger\_count, trip\_distance, RatecodeID, store\_and\_fwd\_flag, PULocationID, DOLocationID, payment\_type, and trip\_start\_datetime. The tasks are listed in a table with columns for task ID, start time, end time, and status. The tasks are: VendorID, tpep\_pickup\_datetime, tpep\_dropoff\_datetime, passenger\_count, trip\_distance, RatecodeID, store\_and\_fwd\_flag, PULocationID, DOLocationID, payment\_type, and trip\_start\_datetime.

Druid ingests raw data and converts it into a custom, indexed format that is optimized for analytic queries.

To get started, please specify what data you want to ingest.

[Learn more](#)

Source type: local

Base directory: /opt/data/rawdata/bdp

File filter: \*.csv

This path must be available on the local filesystem of all Druid services.

Apply



# Design tools for ingestion processes: Apache Nifi

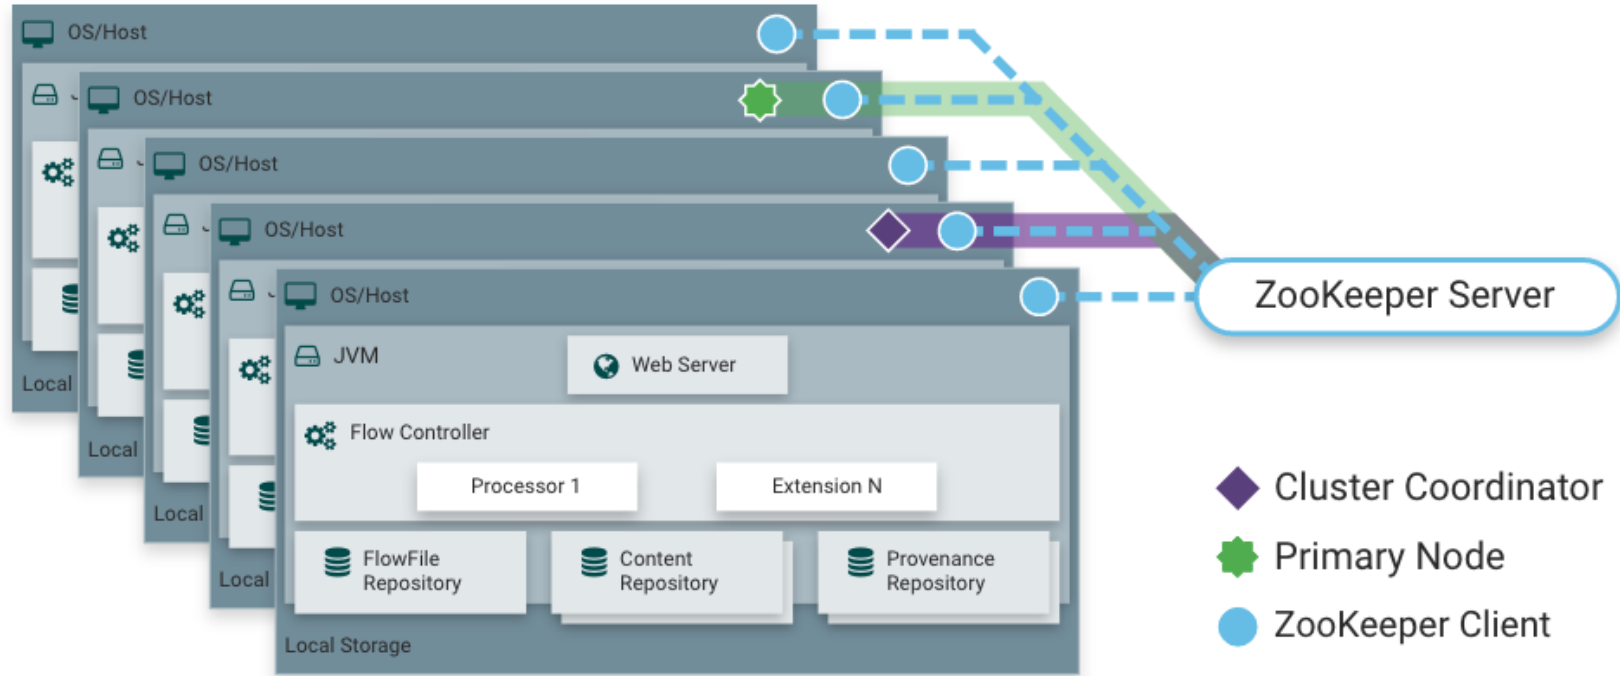


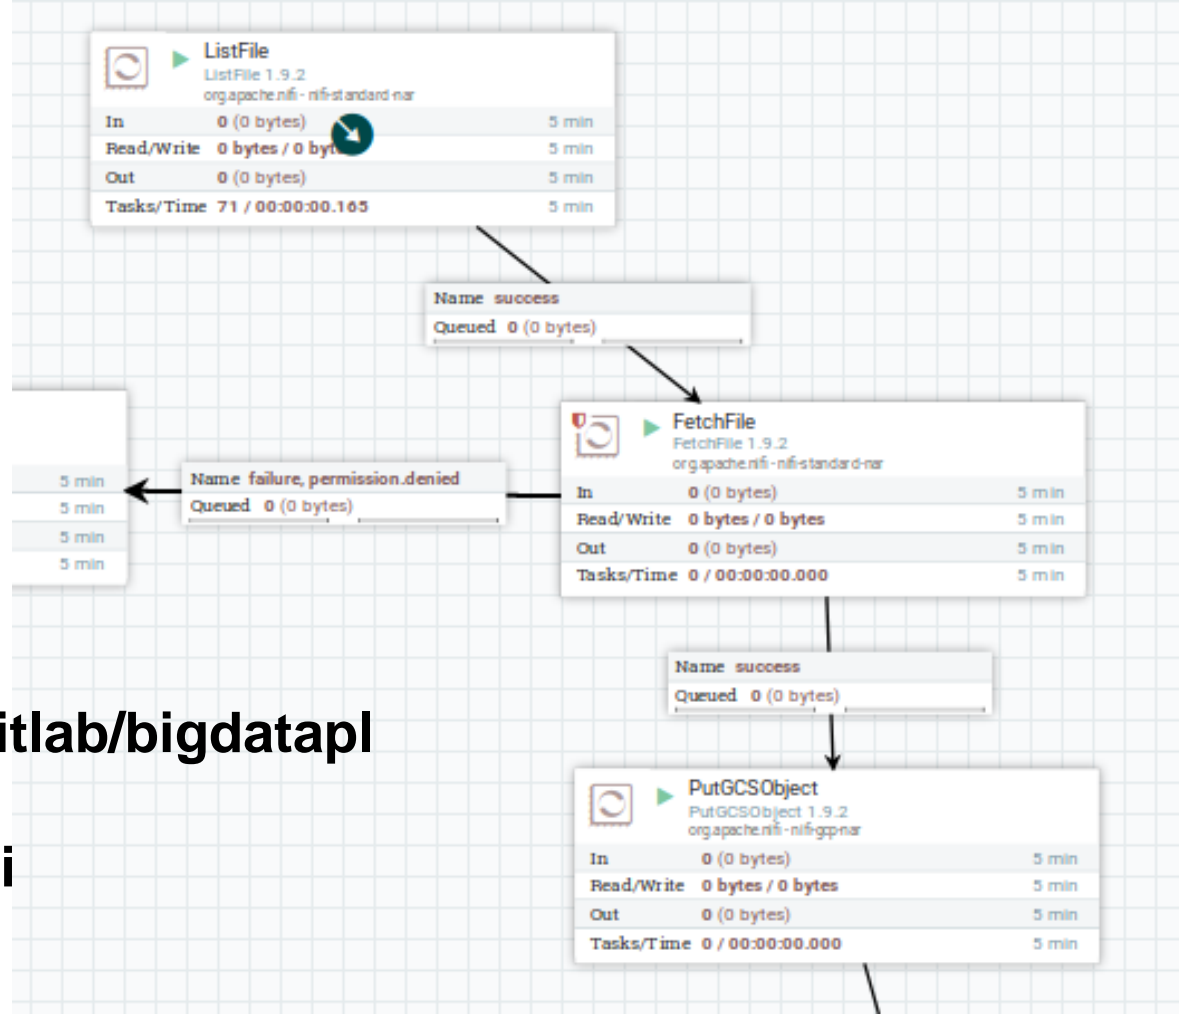
Figure source: <https://nifi.apache.org/docs.html>

# Design tools for ingestion processes: Apache Nifi - key concept

- Data is encapsulated into “FlowFile”
- **Processor** (Component) performs tasks
- **Processor** handle FlowFile and has different states
  - Each state indicates the results of processing that can be used for establishing relationships to other components
- **Processors** are connected by **Connection**
- **Connection** can have many **relationships** based on states of upstream Processors

# Design tools for ingestion processes: Apache Nifi

See the tutorial:  
<https://version.aalto.fi/gitlab/bigdataplatforms/cs-e4640/-/tree/master/tutorials/nifi>



# Thanks!

**Hong-Linh Truong**  
**Department of Computer Science**

**rdsea.github.io**