



Aalto University  
School of Science

# Hadoop and its Big Data Ecosystems

*Hong-Linh Truong*

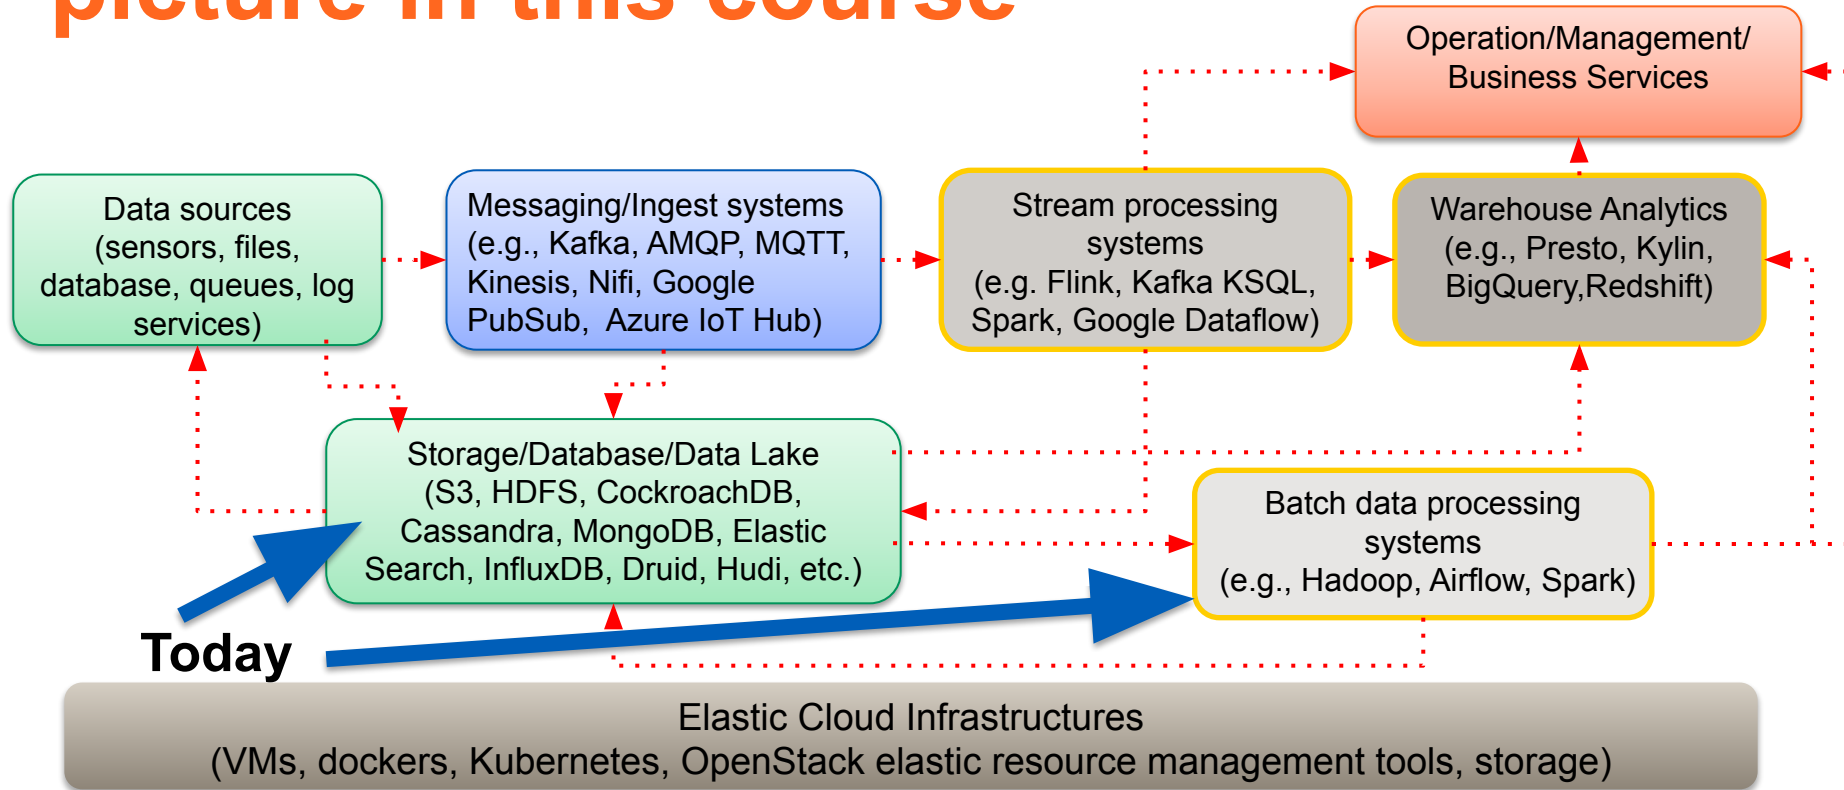
*Department of Computer Science*

*[linh.truong@aalto.fi](mailto:linh.truong@aalto.fi)*, *<https://rdsea.github.io>*

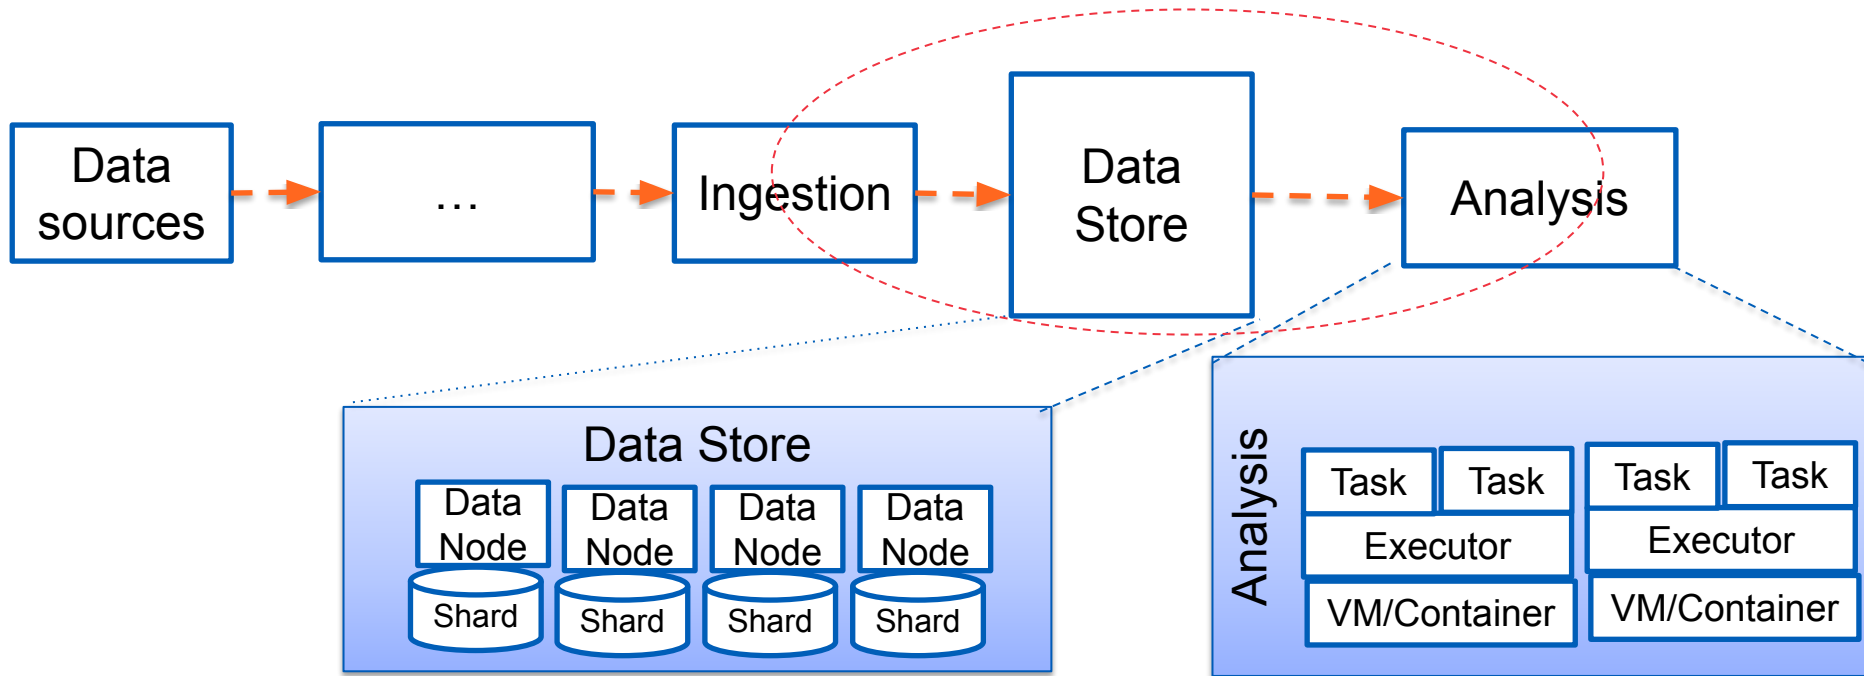
# Learning objectives

- Understand massive scale data management and processing with Hadoop
- Understand and able to use Hadoop components for big data platform designs
- Able to integrate Hadoop with other frameworks for data ingestion and analytics systems

# Big data at large-scale: the big picture in this course

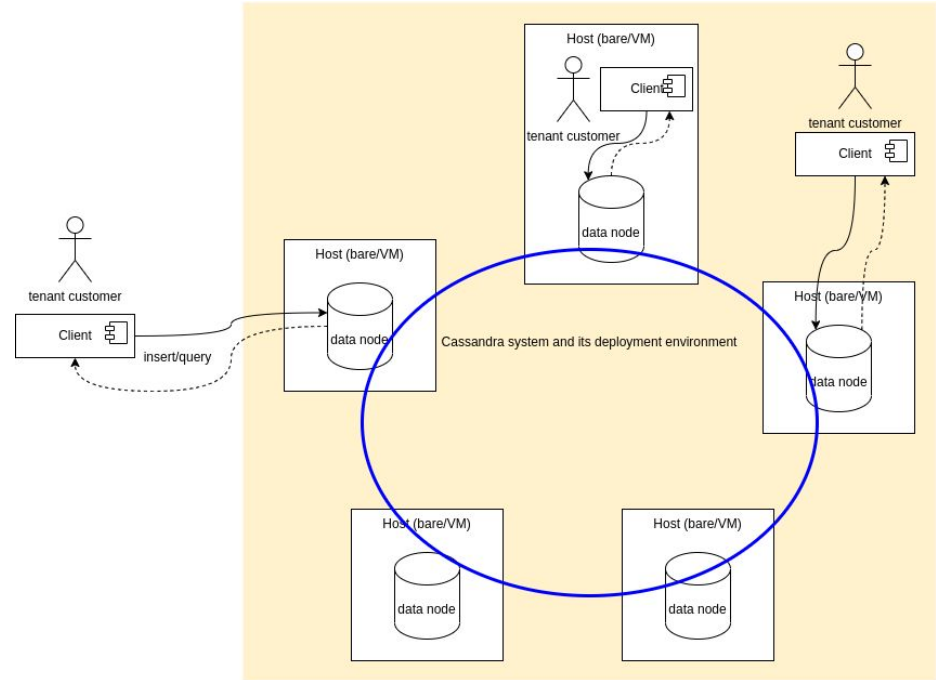


# Consider again our big data pipelines



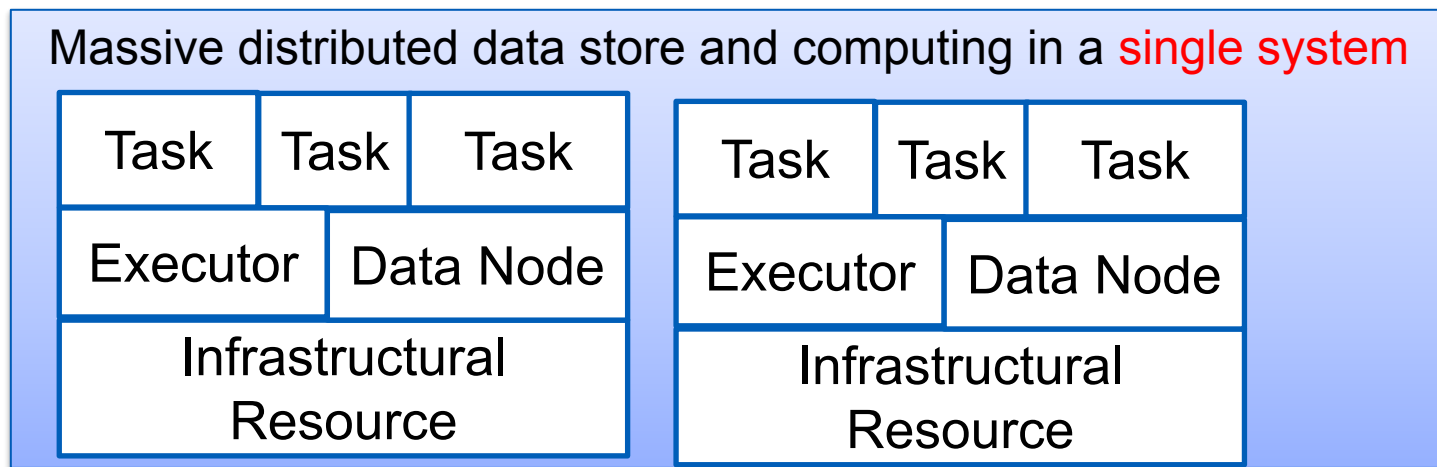
Except  
ingestion/queries,  
most analytics are  
done outside the  
data system

Data store: E.g., Cassandra



# Combining data storage and analysis

**Data locality, massive parallel and distributed data processing with large-scale data store**

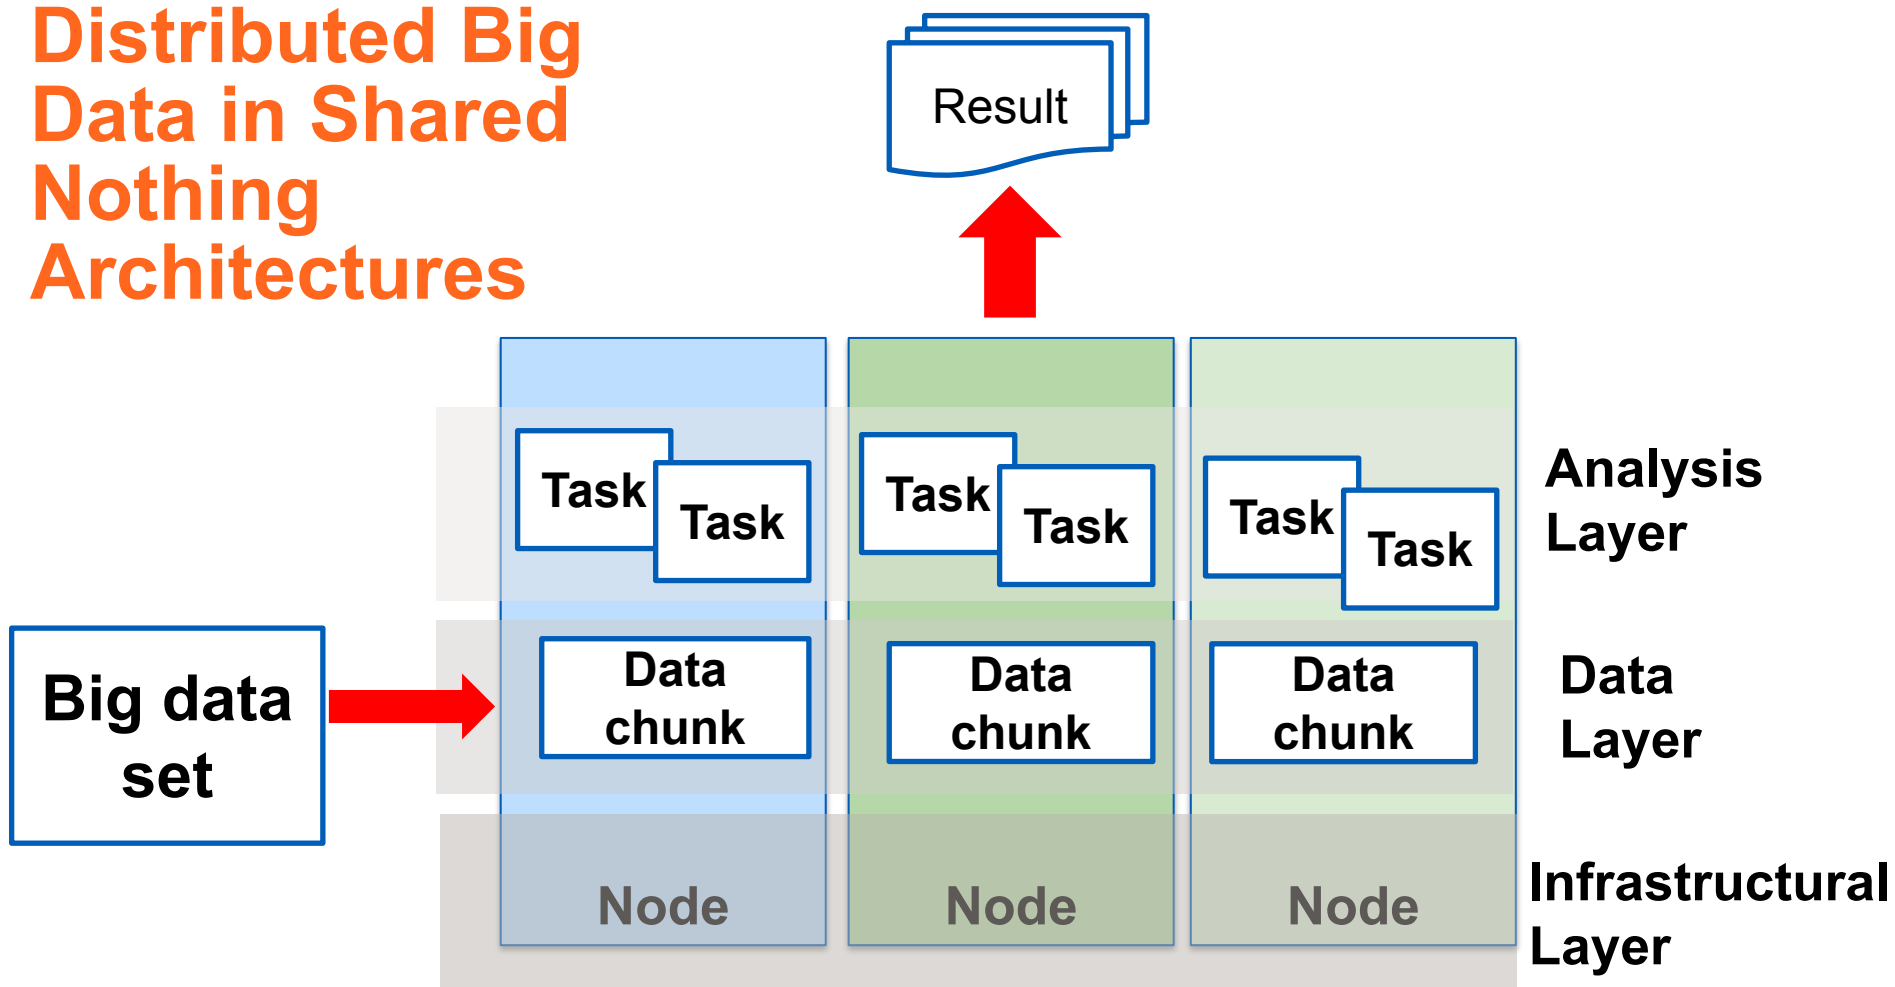


**Single system: **many machines** connected into a type of “cluster”**

# Shared nothing architectures

- **Machine as a node:**
  - completely independent from other nodes
  - no shared physical resources (CPUs, memory, hard disk, ...)
- **A distributed system of machines with shared nothing architectures**
  - allow to scale and extend physical resources easily
  - enable high fault tolerance designs
  - quite common and easy-to-be-setup for enterprises
    - *think about commodity clusters*

# Distributed Big Data in Shared Nothing Architectures





# Benefit

- Reduce cost and increase performance when moving data is much more expensive than moving computation
- Consolidate and integrate various types of data
- Save shared infrastructural resources
- Provide suitable solutions for different types of customers/clients and different infrastructures

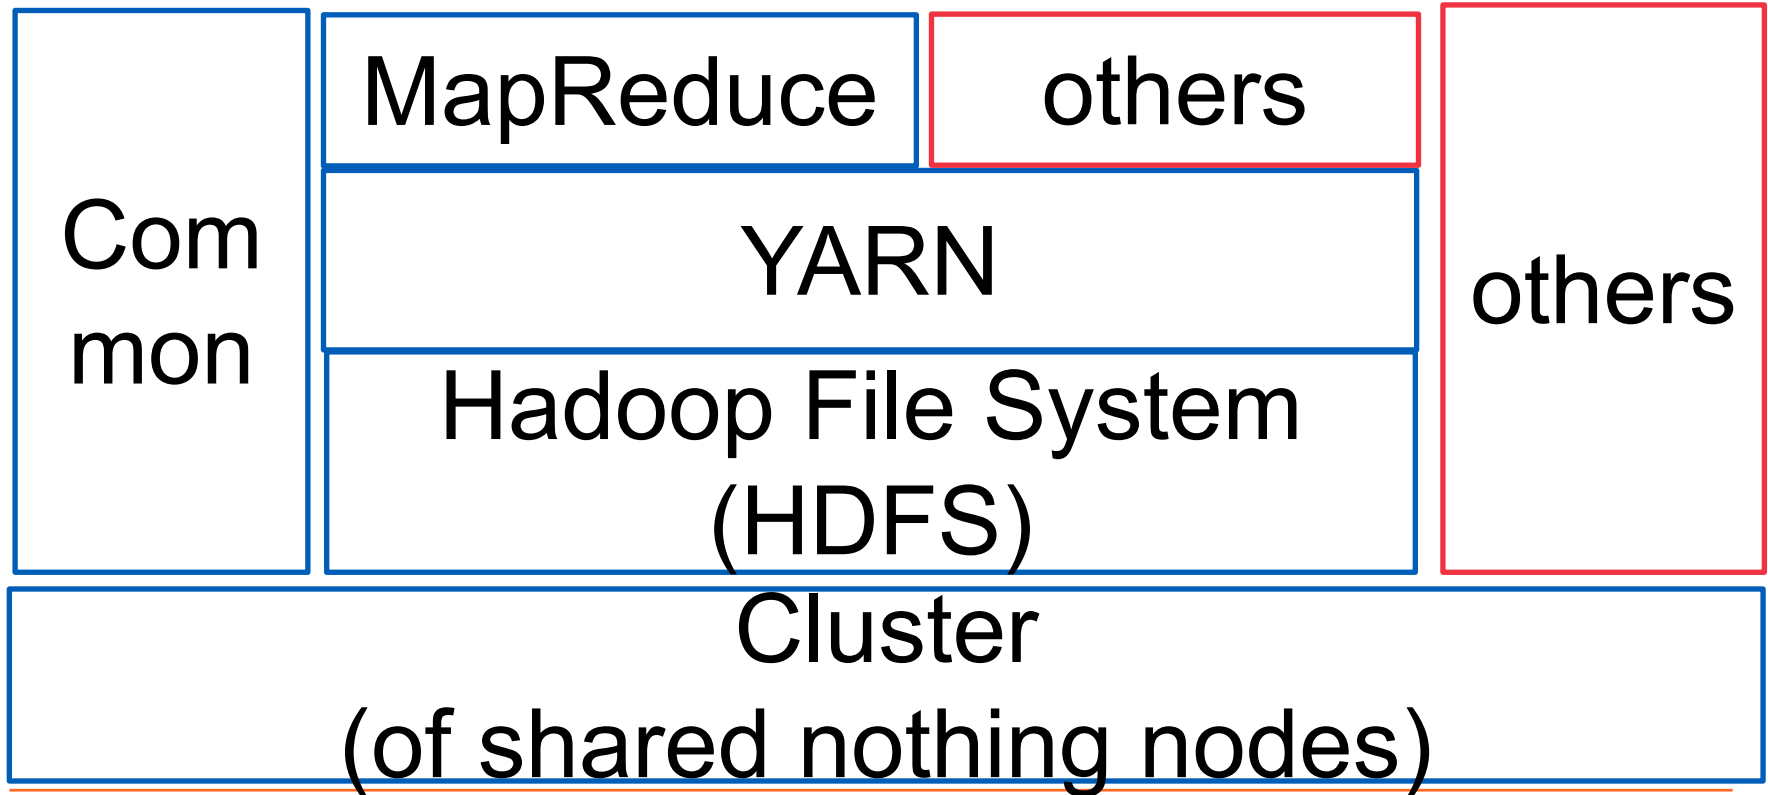
# Benefit - Variety

- **Multiple data formats in files**
  - direct access by applications or low-level storage layer
- **In big data we have a lot of files**
  - do not need to transform them into a single format or database
  - file storage as the data lake storage layer
- **Processing files directly is a benefit!**
- **The data lake concept:**
  - multiple types of big data analytics with high concurrent/parallel data writes/reads
  - dealing with different data access frequencies: **hot, warm and cold data**

# Hadoop

- <http://hadoop.apache.org/>, original from Yahoo
- The goal is to combine storage and processing in the same cluster system
- Designed for *massive scale of data and computing with shared nothing architecture*
  - commodity hardware, highly scalability, fault tolerance, easy to extend
- Suitable for both on-premise and clouds
- There are very rich software ecosystems centered around Hadoop

# Hadoop: layers



# Hadoop key components

- **HDFS as a distributed file system**
  - for managing data
- **YARN as a resource management system**
  - for executing and managing analytics tasks
- **MapReduce as one programming model**
  - for MapReduce applications
- **Coordination (ZooKeeper)**
  - for fault tolerance and metadata

# Hadoop File System (HDFS)

- **For handling very big data files**
  - GBs of data within a single file
- **Assumption of data handling**
  - write-once-read-many
  - not suitable for random-access update → analytics data
- **Deal with hardware failures, support data locality, reliability**
  - “high fault tolerance” with “low-cost/commodity hardware”

# Example

e.g., 112M rows

NYC OpenData

HomeDataAbout ▾Learn ▾AlertsContact UsBlog

q

Sign In

2018 Yellow Taxi Trip Data

Transportation

View DataVisualize ▾ExportAPI⋮

The yellow and green taxi trip records include fields capturing pick-up and drop-off dates/times, pick-up and drop-off locations, trip distances, itemized fares, rate types, payment types, and driver-reported passenger counts. The data used in the attached datasets were collected and provided to the NYC Taxi and Limousine Commission (TLC) by

[More](#)

Updated

February 8, 2020

Data Provided by

Taxi and Limousine Commission (TLC)

About this Dataset

Mute Dataset

Updated

February 8, 2020

Data Last Updated

April 5, 2019

Metadata Last Updated

February 8, 2020

Date Created

September 24, 2018

Views

33.1K

Downloads

3,703

Data Provided by

Taxi and Limousine Commission (TLC)

Dataset Owner

NYC OpenData

Update

Update Frequency

Historical Data

Automation

No

Date Made Public

10/19/2018

Dataset Information

Agency

Taxi and Limousine Commission (TLC)

Attachments

data\_dictionary\_trip\_records\_yellow.pdf

Topics

Category

Transportation

Tags

This dataset does not have any tags

Snapshot from: <https://data.cityofnewyork.us/Transportation/2018-Yellow-Taxi-Trip-Data/t29m-gskq>

# HDFS - data blocks

- **Files are stored in many nodes**
  - but the application accesses HDFS files just like “typical file systems”
- **A file includes many blocks**
- **File blocks are replicated and distributed across nodes**
- **Conventional way of access data**
  - naming resolving: **hdfs://**
  - common operations: list, put, get, ...



# HDFS Architecture

Crucial components

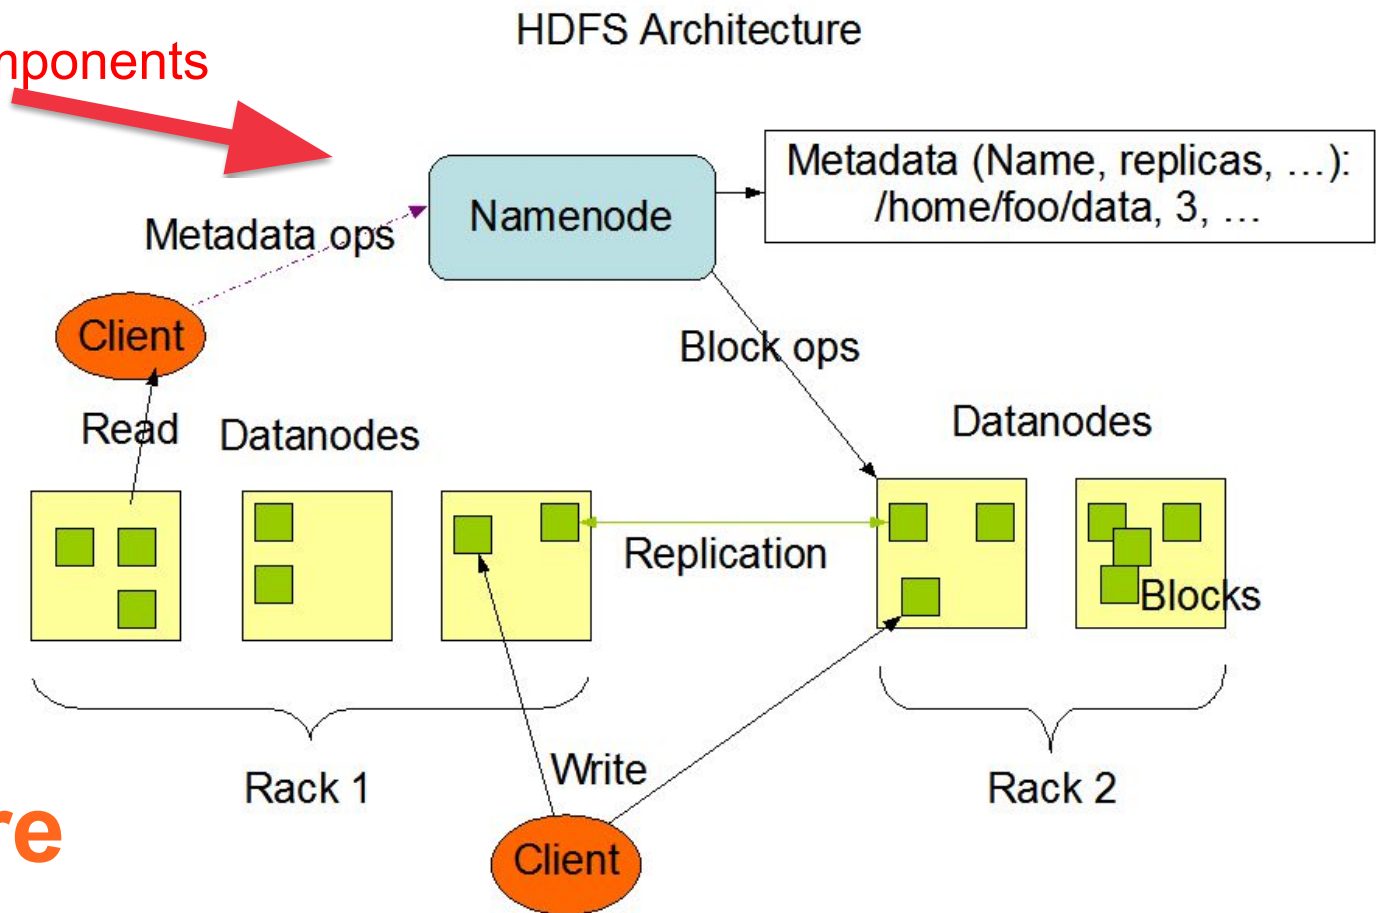


Figure source: <https://hadoop.apache.org/docs/stable/hadoop-project-dist/hadoop-hdfs/HdfsDesign.html>

# File blocks, metadata and data replication

- **Block size is 128MB (default)**
  - can be configurable but should not be small size
  - all blocks of the same file are the same, except the last one
- **Data is replicated across the cluster**
  - usually, replication factor is 3
- **NameNode** manages file system metadata

# HDFS fault tolerance

- **Data blocks**
  - file blocks are replicated and distributed across nodes
- **Rack awareness**
  - avoid communication problems between nodes in different racks
- **Monitoring**
  - **DataNode** reports to **NameNode**
- **Read and write**
  - using **NameNode** for metadata and for information of DataNodes
  - NameNode has replication (master-worker)

# Compatible file systems with HDFS

- **Many file systems are compatible with HDFS**
  - for integration and analysis purpose: m
- **Examples:**
  - Amazon S3
  - Azure Blob Storage
  - Azure Data Lake Storage
  - OpenStack Swift

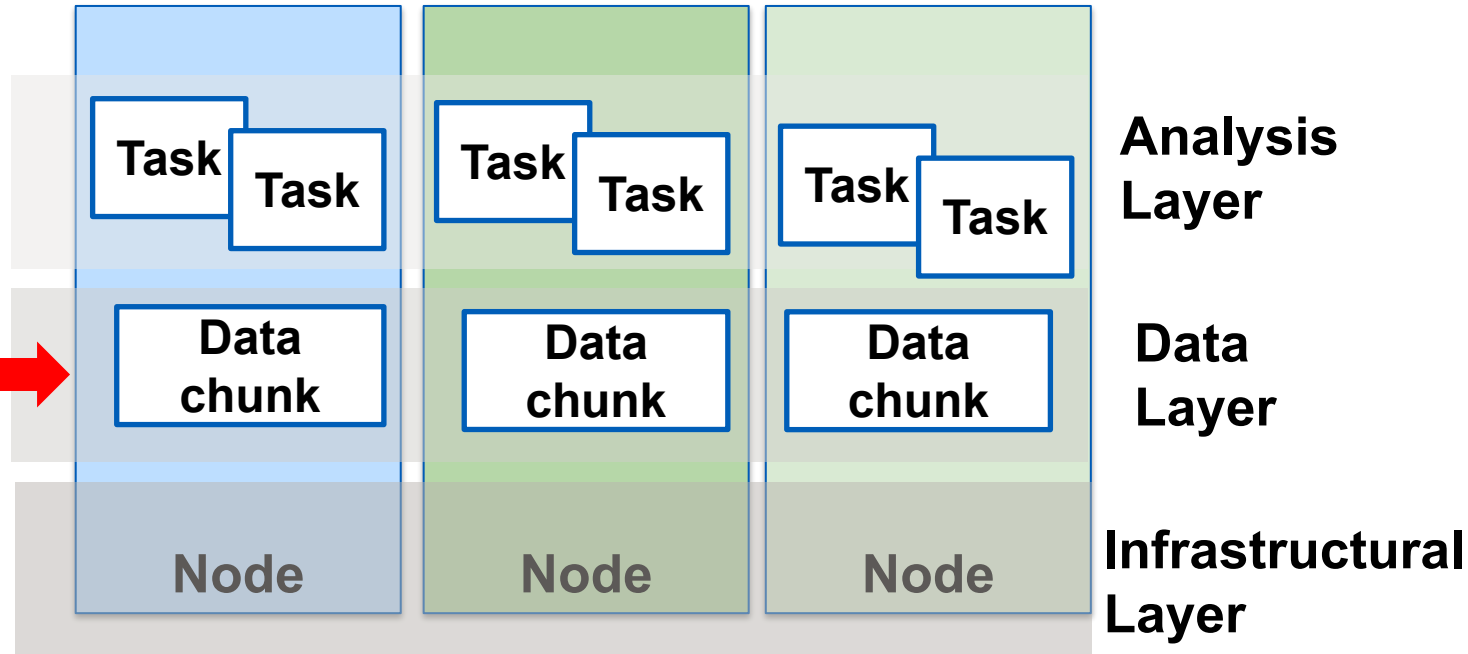
# YARN (Yet Another Resource Negotiator)

**If HDFS can be used to store different files, what would be the good way to enable “data processing” atop HDFS?**

**Key requirements:**

**scalability/elasticity, high utilization of infrastructural resources, serviceability for multiple users/application types**

Take a  
look again



Big data  
set



**How to leverage the same infrastructure and data management layers to perform data analysis?**

**How to enable different (distributed) programming models for data analytics?**



# YARN (Yet Another Resource Negotiator)

- **Manage resources for processing tasks**
  - each node in the cluster provides resources for executing tasks
- **Computing resource types:**
  - CPU, Memory and Disks
  - also support GPU and FPGA Node
- **Computing resources are abstracted into “Containers”**
  - don't be confused: it is not a (Docker) container!
- **Multi-tenancy support**

# YARN Components

- **Resource Manager**

- scheduler: how to schedule analytics tasks atop computing resources

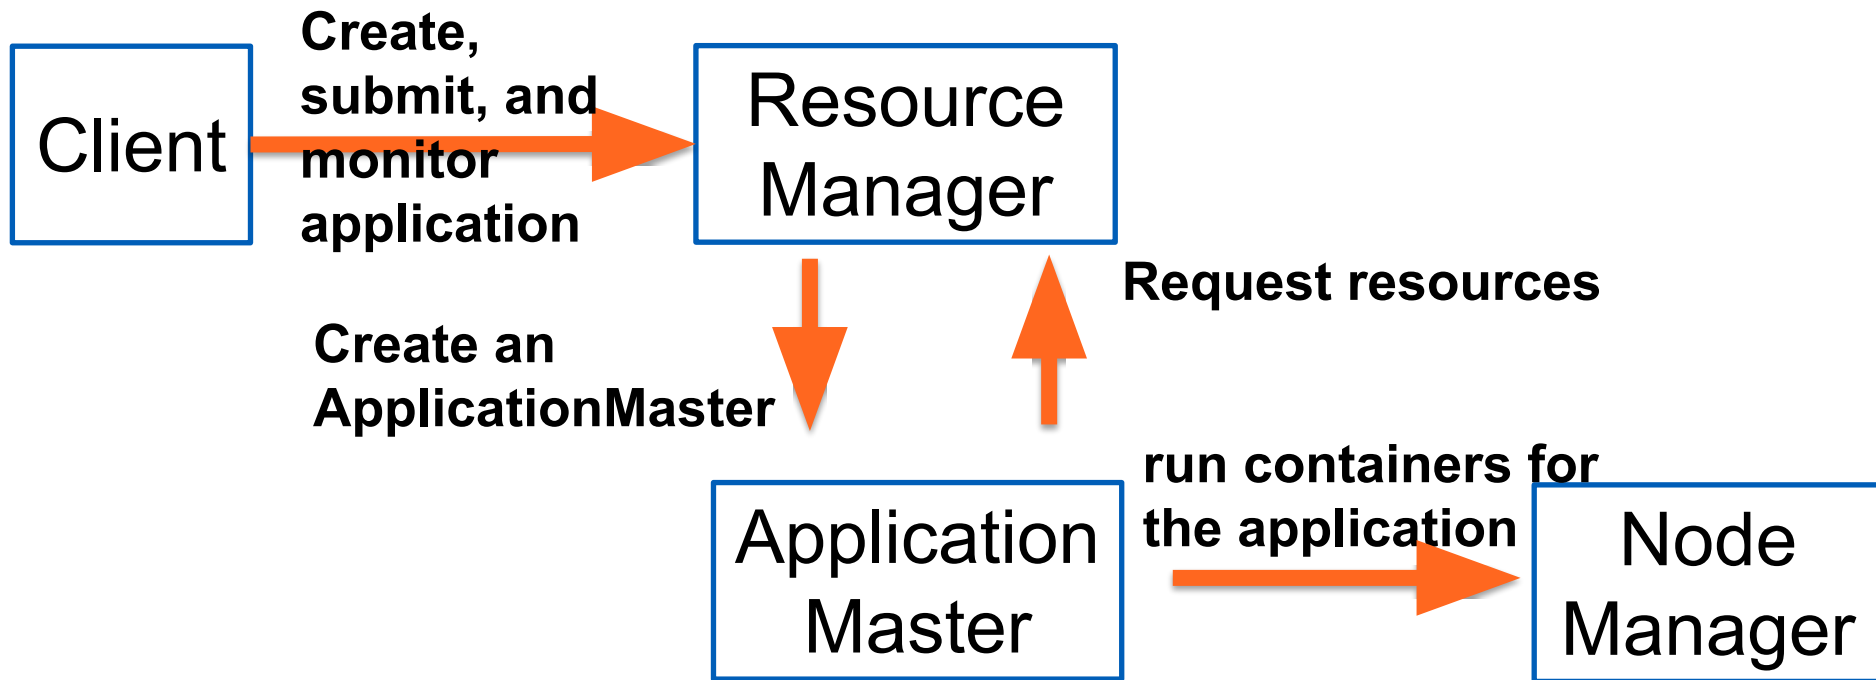
- **NodeManager**

- for managing resources of execution tasks in a node

- **ApplicationMaster**

- application-specific manager for each application
- work with Resource Manager and NodeManager provisioning resources and manage application execution
- handle application-specific tasks

# YARN basic model



# YARN Architecture

User applications running within “Containers”:  
This model **decouples** from specific types of applications

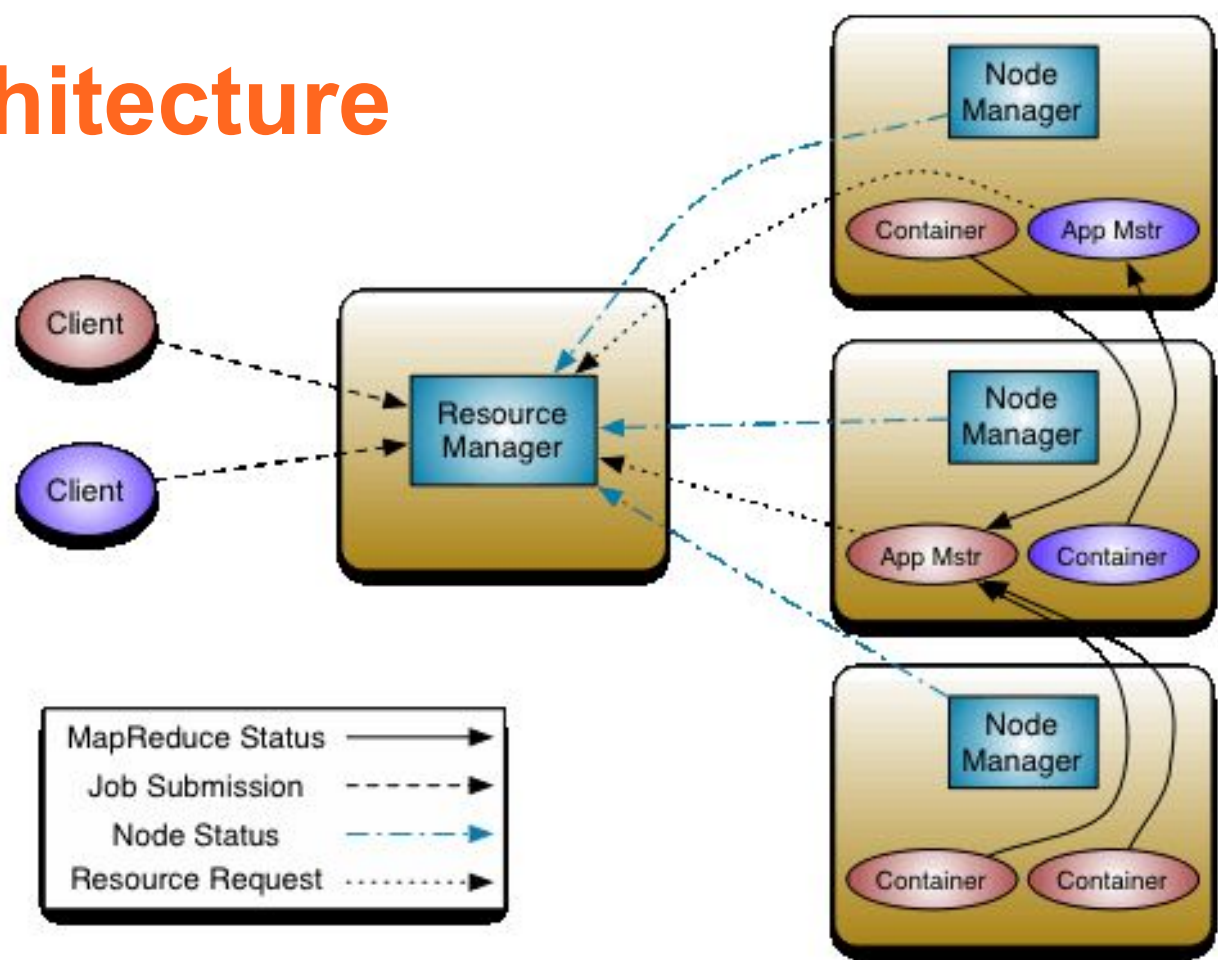


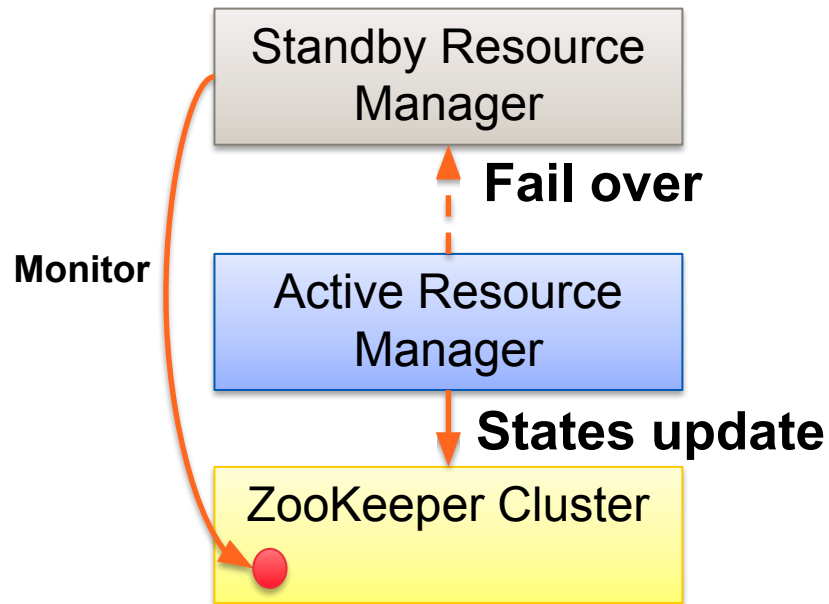
Figure source: <https://hadoop.apache.org/docs/stable/hadoop-yarn/hadoop-yarn-site/YARN.html>

# Scheduling of tasks

- **Scheduling**
  - FIFO
  - CapacityScheduler
    - *Use multiple queues, each with a limit of resources*
  - FairScheduler
    - *All apps will get an average share of resources over time*
- **You can research and add new types of scheduling algorithms**

# Fault tolerance

- **Resource Manager is a critical component**
  - active-passive Resource Manager
  - Zookeeper quorum failover
- **ApplicationMaster**
  - ApplicationMaster is application-specific
  - Resource Manager restarts AM
- **Node**
  - remove out of the cluster



# Programming models

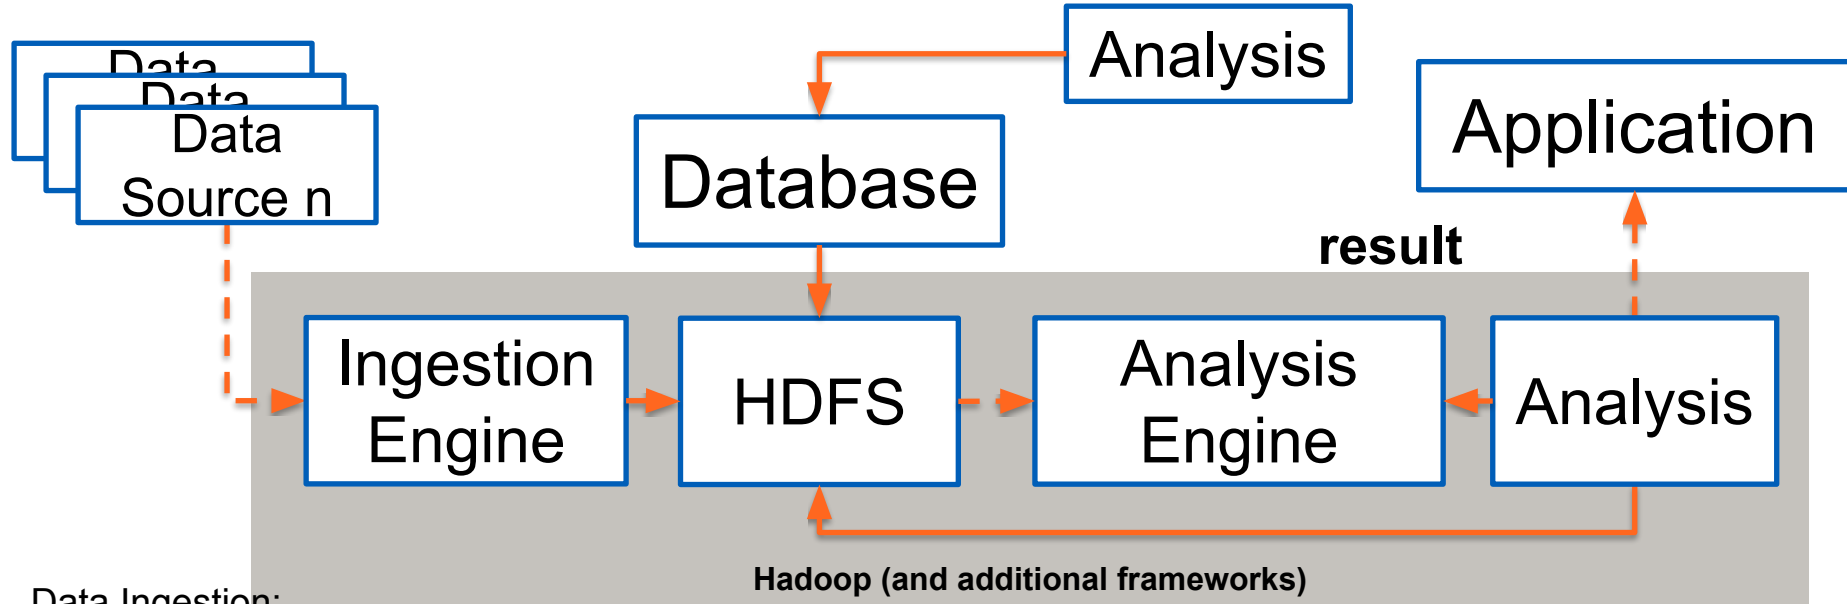
- YARN is an execution environment
- YARN allows different **programming models** for applications
  - MapReduce
  - Apache Spark
  - Workflows
    - *E.g., Apache Tez for DAG (direct acyclic graph) of data processing tasks*

# Integration models

- **Using Hadoop for developing large-scale data analysis**
  - Apache Spark, HBase, Hive, Apache Tez
- **Using Hadoop HDFS as components in a big data system**
  - Hadoop HDFS can become data storage → storage layer
  - emerging data lake models, combined batch and stream ingestions for incremental data processing
    - *e.g., Apache Hudi*



# ETL and Analytics with Hadoop/HDFS



- ## Data Ingestion:

- HDFS Client/Hadoop Streaming
- Spark Streaming
- Kafka Connect
- Apache Nifi

- HDFS as storage for databases

- Accumulo, Druid, etc.

- Computing/Data Processing Framework

- Apache Spark
- Hadoop MapReduce
- Apache Tez

# Is that Hadoop dead/unattractive?

- **Hadoop is for shared nothing architecture**
- **Shared nothing architectures are not “modern”?**
  - new, advanced developments w.r.t. memory, networks, ...
  - in-memory processing, move data to computing nodes
- **Many tools seem not powerful for data science/ML?**
  - Java vs Python development for data science and ML
  - services in ecosystem
  - near realtime analytics use cases
- **So why do we still study it?**
  - foundational designs & still many Hadoop services are important (data warehouse, deep storage)
  - well-supported by many cloud providers for big data workloads

e.g.,

**Azure HDInsight**

Provision cloud Hadoop, Spark and HBase clusters.

[Products](#) / [Analytics](#) / [Amazon EMR](#)

**Apache Hadoop on Amazon EMR**

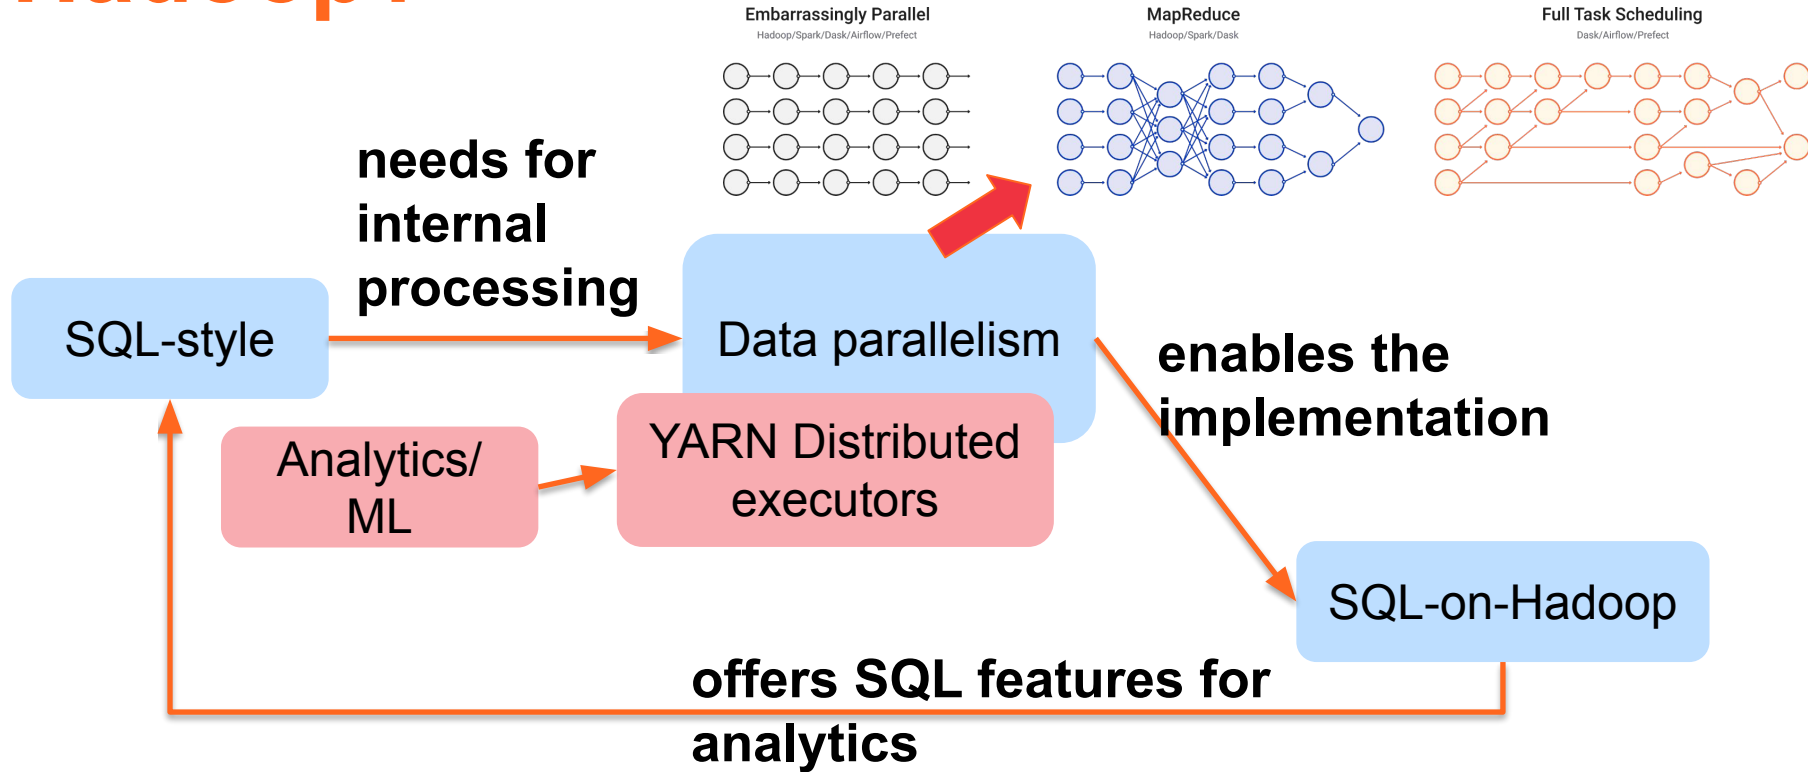
**Dataproc**

# Why do we still need to learn Hadoop?

- **Hadoop file systems as a storage for many big data services**
  - databases: Accumulo, Apache Druid
  - data warehouses/Data Lakes: Hive, Hudi
- **Support different models of access/analytics**
  - via SQL styles atop big data platforms: You can do extract/transform/load (ETL), reporting, and data analysis using SQL styles
  - large-scale parallel programming (e.g., Spark)
- **Still many needs, given shared nothing architectures and on-premise conditions**

# Why do we still need to learn Hadoop?

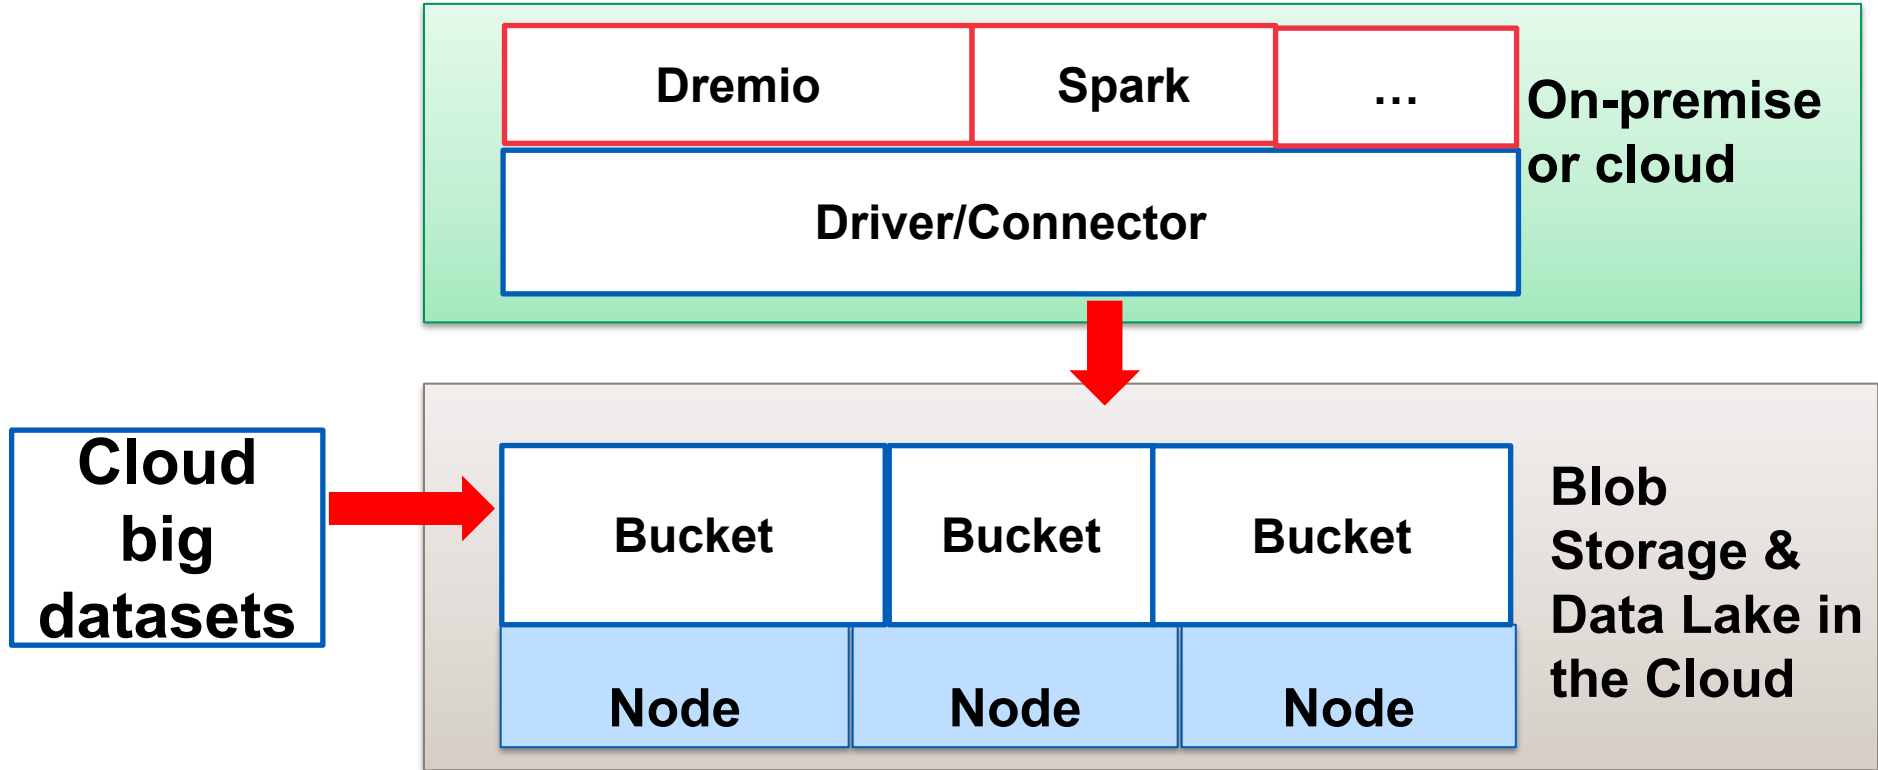
Figure source: <https://docs.dask.org/en/stable/graphs.html>



# Enabling analysis big data using SQL style

- Provide command line tools & JDBC and server for integration
- “SQL-on-Hadoop”
- Examples
  - Apache Hive, <https://hive.apache.org/>
    - *Data warehouse, access data in HDFS or Hbase*
  - Apache Druid (using HDFS as a deep storage)
  - Spark SQL for various types of files

# Integration with other cloud storage services



# Hadoop-native big database/data warehouse systems

# HBase

- **NoSQL database atop Hadoop**
  - use HDFS for storing data
  - use YARN for running jobs on data
- **Follow a master-based architecture**

Reading – Why HBase?

<https://engineering.fb.com/2010/11/15/core-data/the-underlying-technology-of-messages/>

<https://engineering.fb.com/2014/06/05/core-data/hydrabase-the-evolution-of-hbase-facebook/>

<https://engineering.fb.com/2018/06/26/core-data/migrating-messenger-storage-to-optimize-performance/>



# Recall: Big data: column-family data model

Many situations we aggregate and scan few columns of million rows of data  $\Rightarrow$  store big data in columns enable fast scan/retrieval/aggregation

**Column Family = (Column, Column, ...):** for similar type of data & access patterns

**Column Key = Family: qualifier**

**Data = (Key, Value)** where Key =(Row Key, Column Key, Timestamp)

# Example of a data model in HBase

**Column**  
**Column family (e.g., birdinfo)**  
**Row key**  
**Cell: versioning**

1	species	birdinfo:country	birdinfo:english_name	duration					
	Aberti	US	...	3					
2	species	birdinfo:country	birdinfo:english_name	duration	name	url	latitude	longitude	Text

**Row**

Sparsely stored: not all rows have the same number of columns

# Example of a data model

- Example with families: birdinfo, songinfo, location

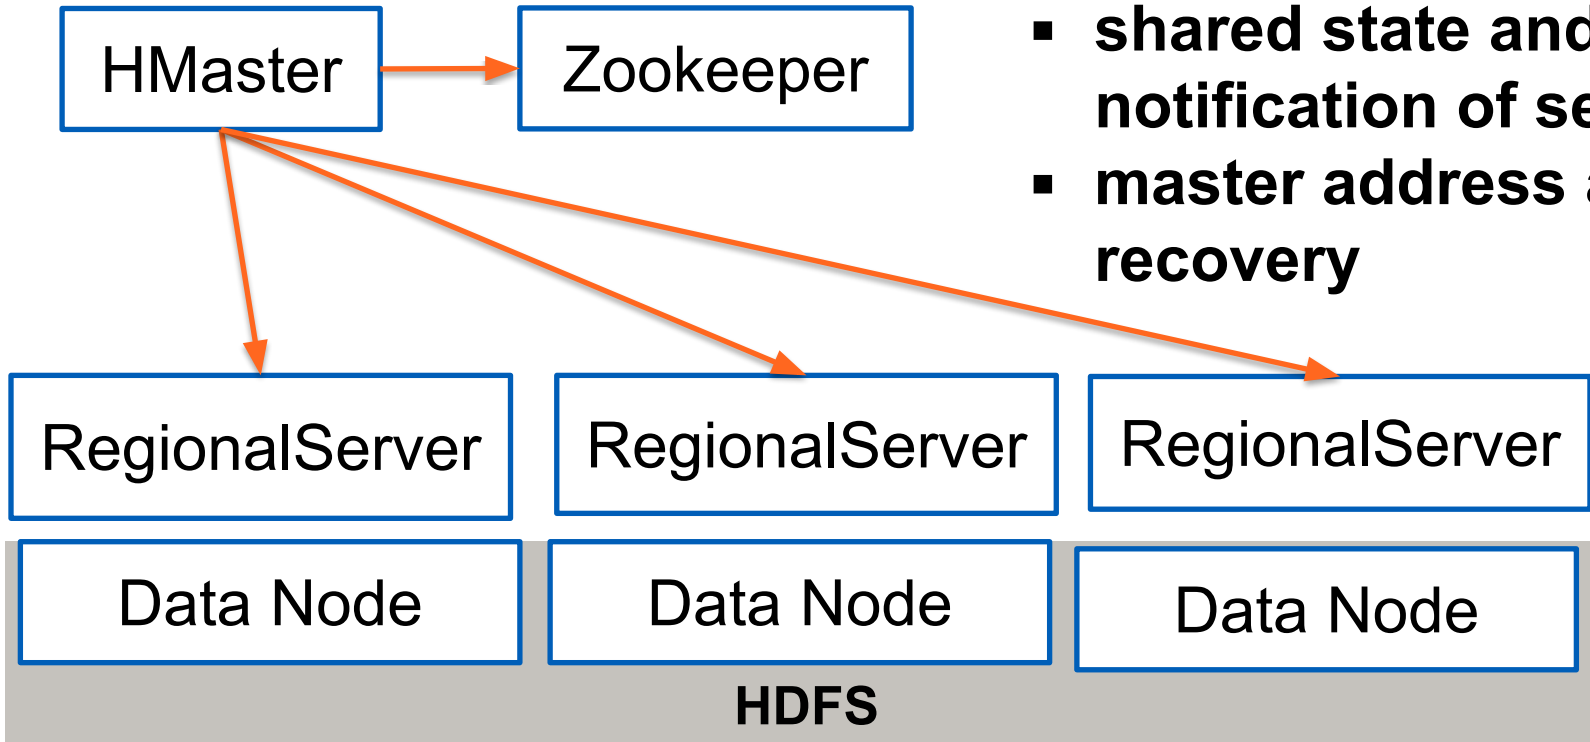
```
hbase:038:0> scan 'hbird0'
ROW                                COLUMN+CELL
17804                             column=birdinfo:english_cname, timestamp=2022-02-05T11:56:09.114, value=Aberts Towhee
17804                             column=songinfo:duration, timestamp=2022-02-05T11:56:09.154, value=3
17804                             column=songinfo:file_id, timestamp=2022-02-05T11:56:09.178, value=17804
17804                             column=songinfo:file_name, timestamp=2022-02-05T11:56:09.205, value=XC17804.mp3
71852                             column=birdinfo:country, timestamp=2022-02-05T11:56:09.415, value=Mexico
71852                             column=birdinfo:english_cname, timestamp=2022-02-05T11:56:09.445, value=Ash-throated Flycatcher
71852                             column=birdinfo:species, timestamp=2022-02-05T11:56:09.488, value=cinerascens
71852                             column=location:latitude, timestamp=2022-02-05T11:56:09.504, value=32.156
71852                             column=location:longitude, timestamp=2022-02-05T11:56:12.253, value=-115.79299999999999
71852                             column=songinfo:duration, timestamp=2022-02-05T11:56:09.430, value=28
71852                             column=songinfo:file_id, timestamp=2022-02-05T11:56:09.459, value=71852
2 row(s)
Took 0.0206 seconds
```

**Enable analytics based on column families (as well as data management)**

# HBase data model – sharding and storage

- **Table** includes multiple **Regions**
  - a **Region** keeps related row data of a **Table** (*partitioning*)
- **Auto-sharding**
  - **Regions** are spitted based on policies
- **Region** has data of multiple column families
  - Different **column families** will be stored in different files
  - **HFiles** are used to store real data (also include index data)

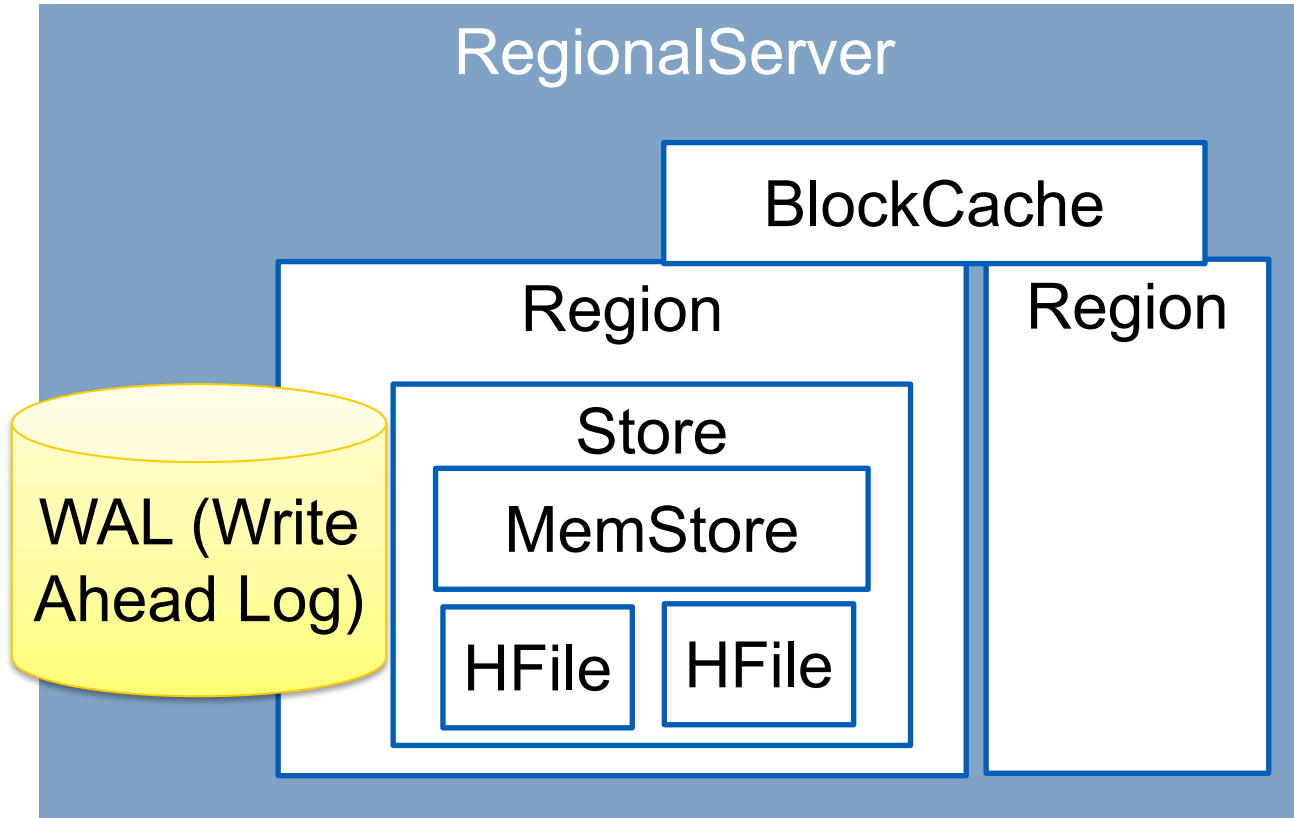
# HBase architecture



## Zookeeper:

- **shared state and failure notification of servers**
- **master address and recovery**

# HBase architecture



**MemStore:** **write cache** for data in memory before written into files

**BlockCache:** for **read cache**

**WAL is for durability**

# ACID

- **Atomic within a row**
- **Consistency**
  - can be programmed: e.g.,
    - read: STRONG (read performed by the primary region) and TIMELINE (primary region first, if not, then the secondary region)
- **Durability**
  - can be programmed
  - WAL (write ahead log)

# Example of using HBase

## Open Scalable Open Time Series Database (OpenTSDB)

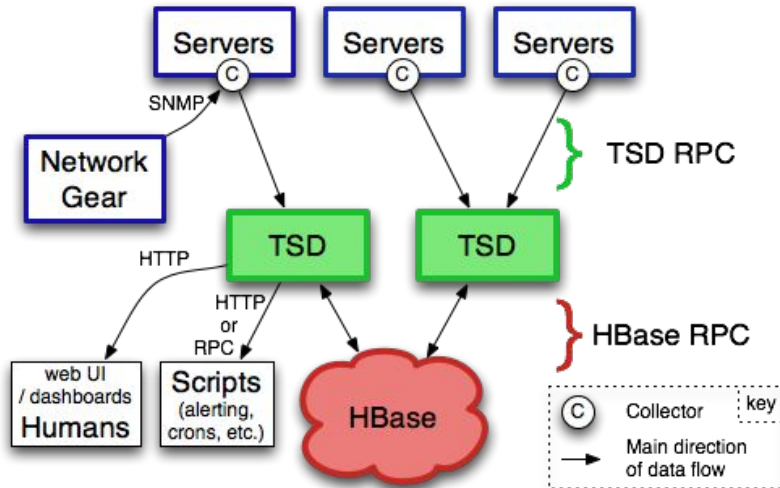


Figure source: <http://opentsdb.net/overview.html>

## Salesforce

- **Salesforce Shield:** for supporting data security/compliance
- **Field Audit Trail:** for track changes
- **Event Monitoring:** log/event
- **Argus:** time series data and alerts

### Source:

<https://engineering.salesforce.com/investing-in-big-data-apache-hbase-420edfba2d30/>



# Apache Hive

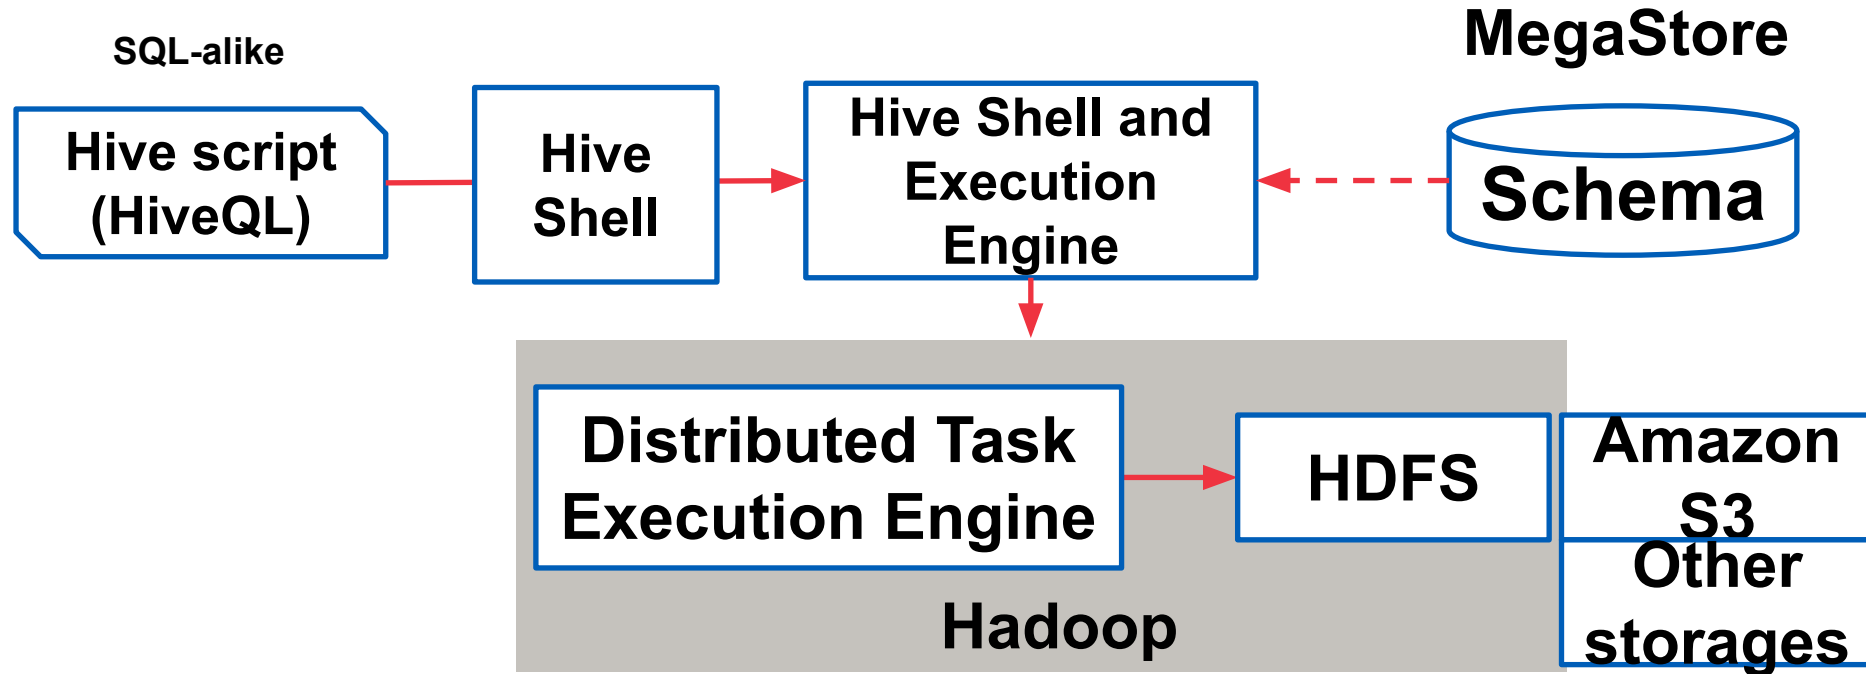
# Apache Hive

- **<https://hive.apache.org/>, on top of Hadoop**
  - data warehouse at a very large-scale
  - access data in large-scale storage like HDFS or S3
- **Support access to data via SQL styles**
  - extract/transform/load (ETL), reporting, and data analysis using SQL styles
- **Provide command line tools & JDBC and server for integration**

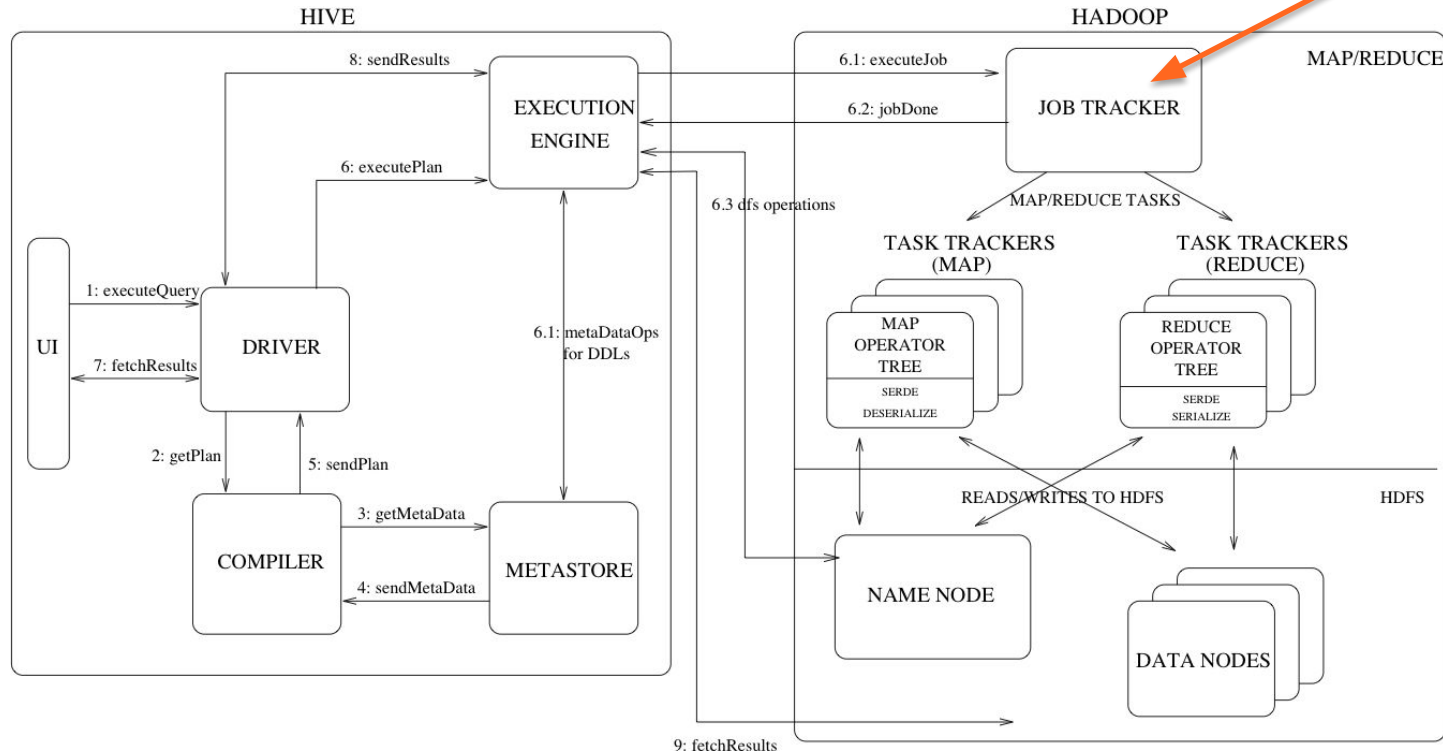
Reading – Hive in Facebook

<https://engineering.fb.com/2009/06/10/web/hive-a-petabyte-scale-data-warehouse-using-hadoop/>

# High-level data flow language & programs



# Hive building blocks



**Distributed tasks with MapReduce, Tez (Workflow) or Spark**

Figure source: <https://cwiki.apache.org/confluence/display/Hive/Design>

# Hive data organization

- **Schema-on-read approach**
  - schema inferred from data, structured and unstructured data
- **Databases**
- **Table**
  - **Managed** table versus **External** tables
  - External table: real data is **stored outside** Hive and **referenced** so delete only table metadata but not the data (read data only)
  - Table is mapped to a directory in HDFS
- **Low-level files: CSV, Apache Parquet, ORC, ...**
- **Managed vs External tables concept:**
  - also in BigQuery, Snowflake, Databricks Lakehouse

# Example

Large tables lead  
to performance  
issues

```
0: jdbc:hive2://localhost:10000> describe taxiinfo;
+-----+-----+-----+
| col_name | data_type | comment |
+-----+-----+-----+
| vendorid | int       |         |
| tpep_pickup_datetime | timestamp |         |
| tpep_dropoff_datetime | timestamp |         |
| passenger_count | int      |         |
| trip_distance | float    |         |
| ratecodeid | int      |         |
| store_and_fwd_flag | int      |         |
| pulocationid | string   |         |
| dolocationid | int      |         |
| payment_type | int      |         |
| fare_amount | float    |         |
| extra | float    |         |
| mta_tax | float    |         |
| tip_amount | float    |         |
| tolls_amount | float    |         |
| improvement_surcharge | float    |         |
| total_amount | float    |         |
+-----+-----+-----+
```

# Hive data organization for performance optimization

- **Partitioning: using value of a column as a partitioning key**
  - partition keys determine how data in Table will be divided
    - *E.g. date or countries*
  - each partition is stored as a subdirectory
    - *Avoid many sub directories!*
- **Buckets: using a hash function of a column for grouping records into the same bucket**
  - avoid large number of small partitions
  - each bucket is stored in a file
  - bucket columns should be optimized for join/filter operations

# Example of partitions

```
CREATE TABLE taxiinfo1 ( ... )  
PARTITIONED BY (year int, month int)  
...;
```

Indicate partition info

Define partition names

```
LOAD DATA LOCAL INPATH .... INTO TABLE taxiinfo1  
PARTITION (year=2019, month=11);
```

```
truong@aaltosea:/opt/hadoop$ bin/hdfs dfs -ls /user/hive/warehouse/taxiinfo1
Found 4 items
drwxr-xr-x - truong supergroup 0 2021-03-02 22:37 /user/hive/warehouse/taxiinfo1/year=2017
drwxr-xr-x - truong supergroup 0 2021-03-02 22:37 /user/hive/warehouse/taxiinfo1/year=2018
drwxr-xr-x - truong supergroup 0 2021-03-02 22:36 /user/hive/warehouse/taxiinfo1/year=2019
drwxr-xr-x - truong supergroup 0 2021-03-02 22:33 /user/hive/warehouse/taxiinfo1/year=__HIVE_DEFAULT_PA
truong@aaltosea:/opt/hadoop$ bin/hdfs dfs -ls /user/hive/warehouse/taxiinfo1/year=2019
Found 2 items
drwxr-xr-x - truong supergroup 0 2021-03-02 22:36 /user/hive/warehouse/taxiinfo1/year=2019/month=11
drwxr-xr-x - truong supergroup 0 2021-03-02 22:36 /user/hive/warehouse/taxiinfo1/year=2019/month=12
```



# Example of buckets

```
CREATE TABLE taxiinfo2 (VendorID int, ....)
  CLUSTERED BY (VendorID) INTO 2 BUCKETS
.....;
```

Identify bucket column



```
truong@aaltosea:/opt/hadoop$ bin/hdfs dfs -ls /user/hive/warehouse/taxiinfo2
Found 2 items
-rw-r--r--  1 truong supergroup    66784 2021-03-02 22:54 /user/hive/warehouse/taxiinfo2/000000_0
-rw-r--r--  1 truong supergroup    31339 2021-03-02 22:54 /user/hive/warehouse/taxiinfo2/000001_0
```

## Combine partitions with buckets



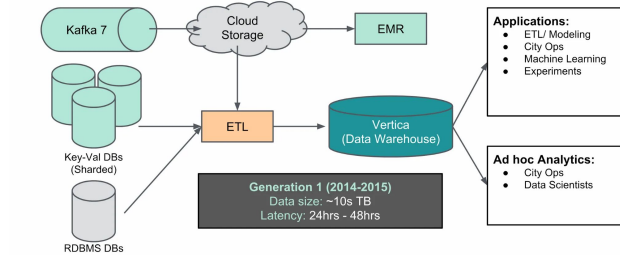
```
truong@aaltosea:/opt/hadoop$ bin/hdfs dfs -ls /user/hive/warehouse/taxiinfo3/year=2019/month=11
Found 2 items
-rw-r--r--  1 truong supergroup    66784 2021-03-02 23:04 /user/hive/warehouse/taxiinfo3/year=2019/month=11/000000_0
-rw-r--r--  1 truong supergroup    31339 2021-03-02 23:04 /user/hive/warehouse/taxiinfo3/year=2019/month=11/000001_0
```

# ACID

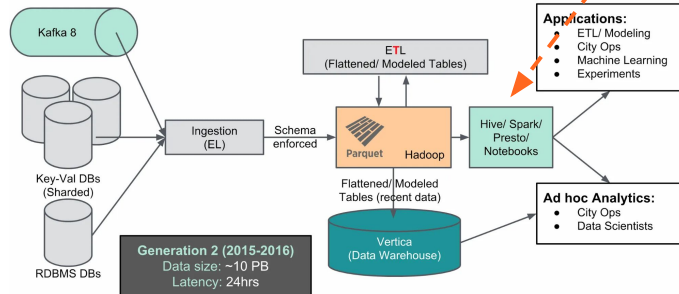
- **ACID support**
  - for managed tables
  - different for CRUD and insert-only tables
- **Locks are used for data isolation**
  - shared lock: for concurrent read of tables/partitions
  - exclusive lock: for modifying table/partition

# Uber example

Generation 1 (2014-2015) - The beginning of Big Data at Uber

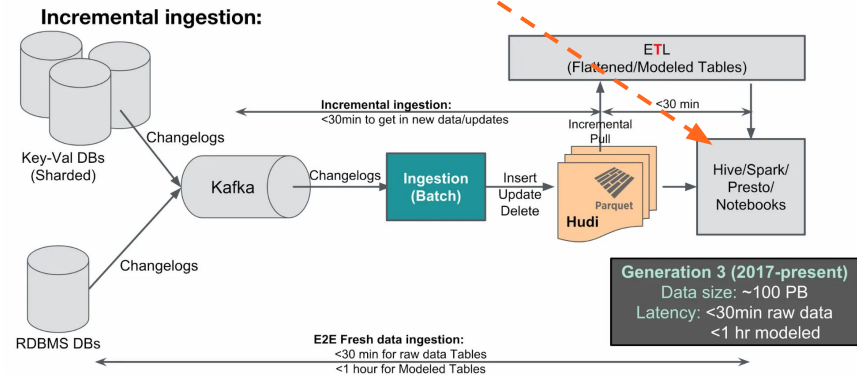


Generation 2 (2015-2016) - The arrival of Hadoop



Hive for analytics: “extremely large queries”, “20,000 Hive queries per day” (see the source)

Generation 3 (2017-present) - Let's rebuild for long term



Figures source:

<https://www.uber.com/en-FI/blog/uber-big-data-platform/>

# Hive in the age of ML

- **Hive has been around for a while → many goals have been diminished**
- **But Hive Megastore continues an important role in data lakes and distributed data querying**
  - power data management and discovery
  - support unified data management
  - enable “feature store”, discovery features used for machine learning

# Summary

- **Hadoop software ecosystem is very powerful**
    - many applications and use cases have been developed
    - Hadoop File System (HDFS) is a crucial subject
  - **Managed Hadoop ecosystem services by cloud providers**
    - try to look at Azure HDInsight, Google Dataproc, and Amazon EMR
  - **High-level distributed query engineering using Hadoop components (HDFS, Hive)**
  - **Understand the combination of data management with data processing techniques in the same system with Hadoop that simplify your big data tasks**
-

# Thanks!

**Hong-Linh Truong**  
**Department of Computer Science**

**rdsea.github.io**