



Aalto University
School of Science

Big Data Processing - The Spark Programming Model

Hong-Linh Truong

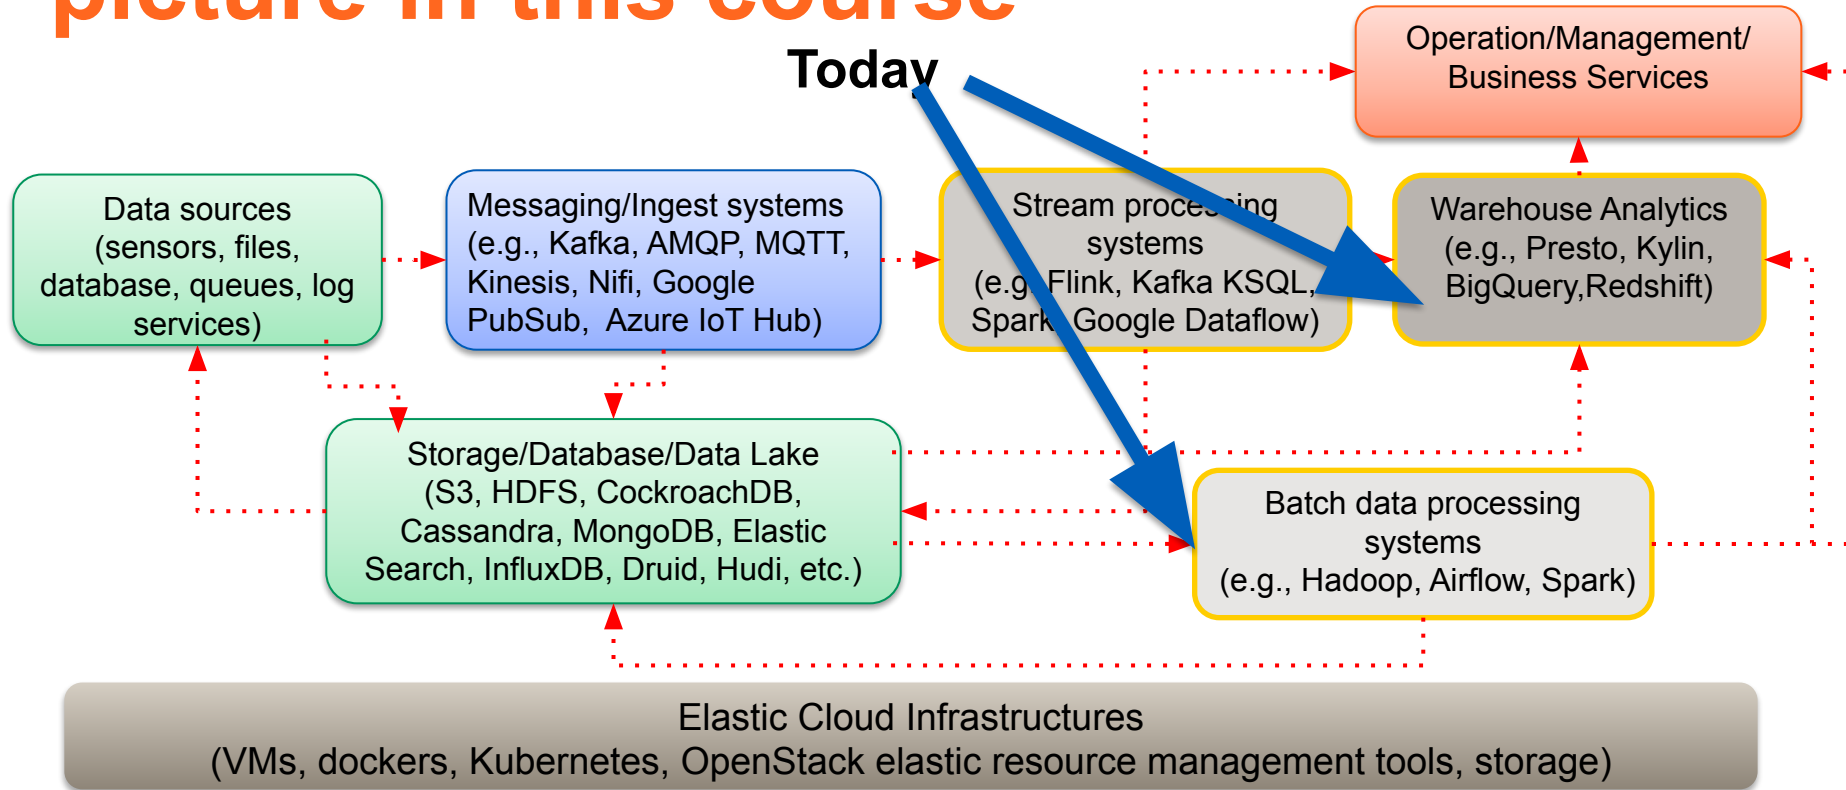
Department of Computer Science

linh.truong@aalto.fi, *<https://rdsea.github.io>*

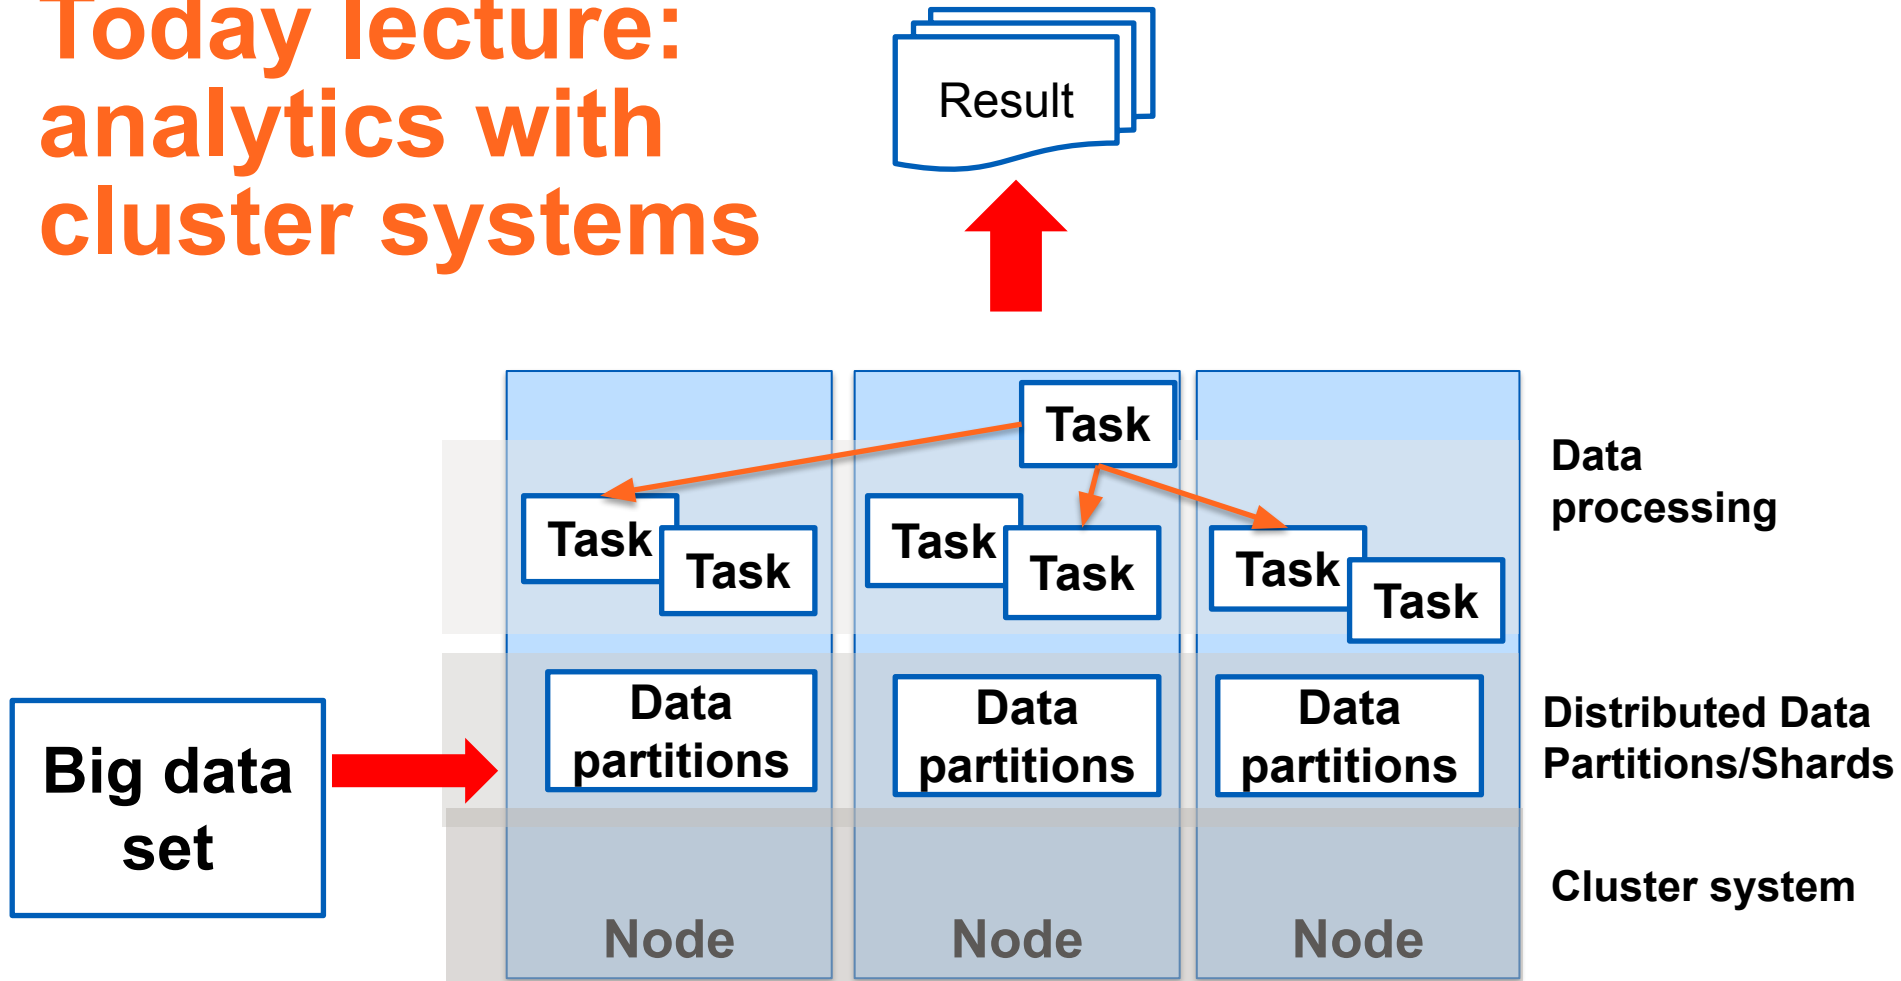
Learning objectives

- **Be familiar with big data processing models using multiple nodes/clusters**
- **Understand the Spark programming model for big data processing**
- **Able to perform practical programming features with Apache Spark**
- **Able to design and apply Spark data processing for data in lake storage**

Big data at large-scale: the big picture in this course



Today lecture: analytics with cluster systems

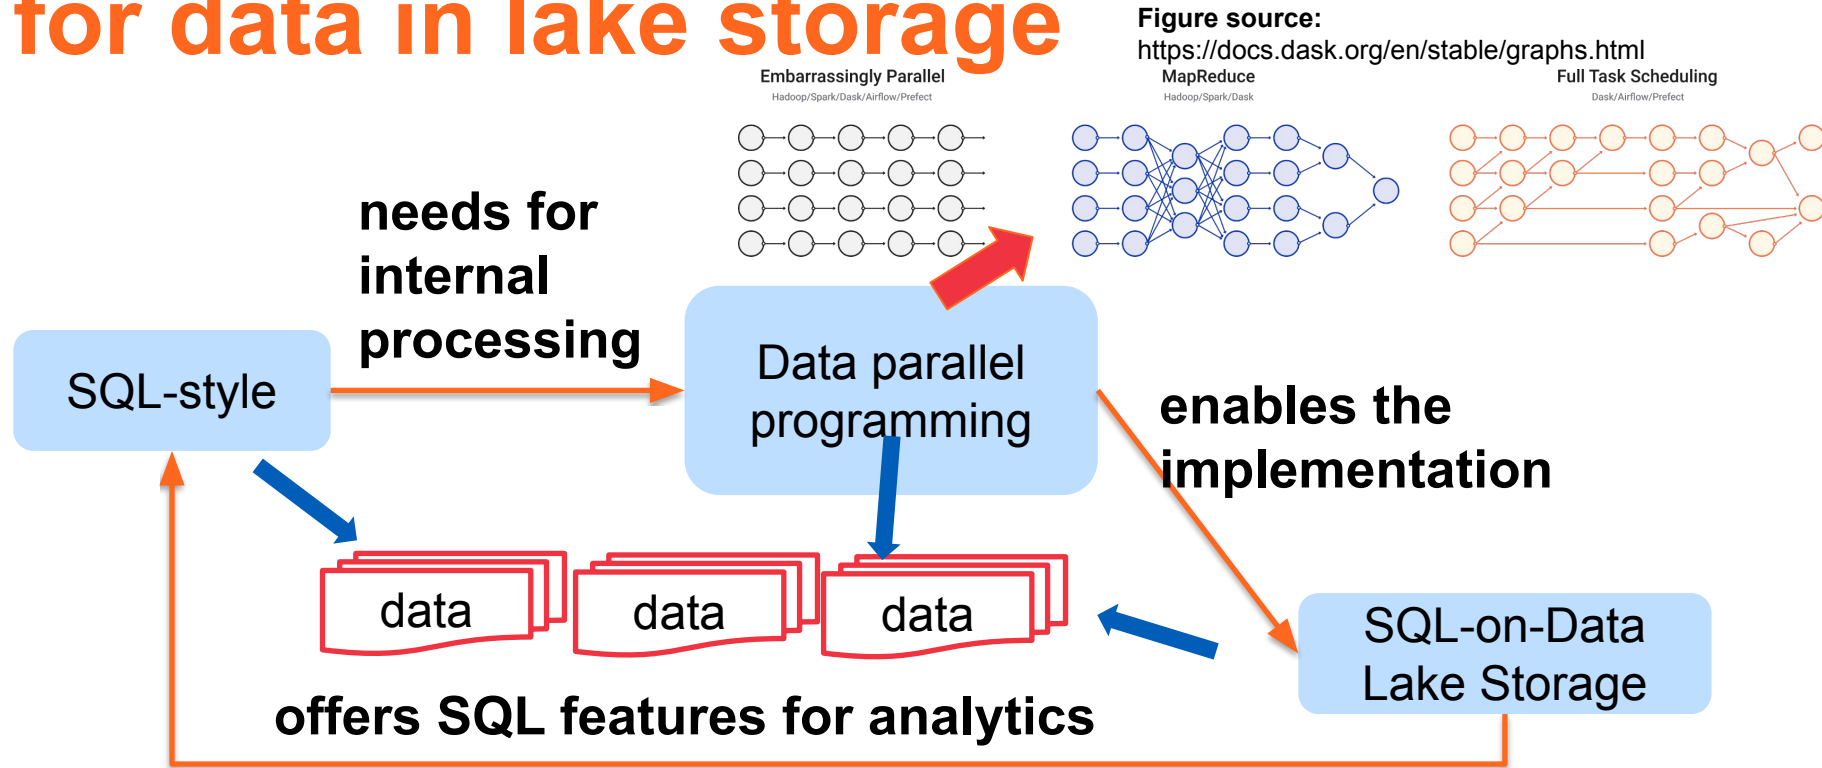


Our first focus: big data analytics for data at rest

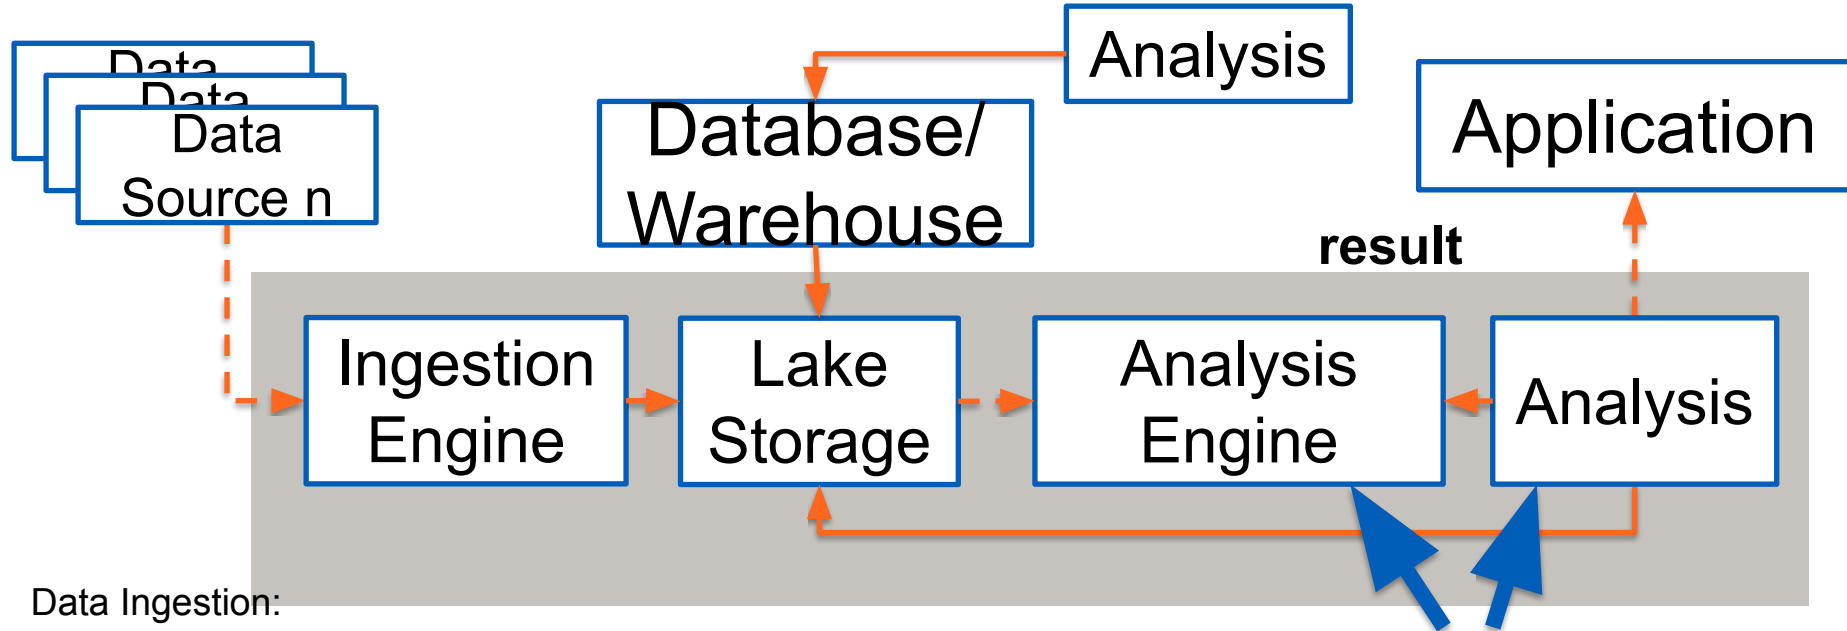
Recall: Data at rest

- **At rest**
 - distributed file systems/object storages
 - *in big data we have a lot of files with different data formats*
 - data in a set of databases
- **Multiple types of big data analytics with high concurrent/parallel data writes/reads**
- **Dealing with different data access/analytics frequencies:**
 - e.g., data organized into **hot, warm and cold data**

SQL-style/data parallel processing for data in lake storage



ETL and Analytics with Lake Storage



Data Ingestion:

- Spark Streaming
- Kafka Connect
- Apache Nifi

- HDFS, AWS S3, Google Storage, Azure Data Lake Storage, etc., as storage

- Computing/Data Processing Framework
 - Apache Spark
 - Hadoop MapReduce

DataFrame/Table view of data

Example taxi records: named columns

passenger_count	trip_distance	RatecodeID	store_and_fwd_flag	PULocationID	DOLocationID	payment_type	fare_amount	extra_mta_tax	tip_amount	tolls_amount	improvement_surcharge	total_amount
1	1.34	1	N	238	236	2	10.0	0.0	0.5	0.0	0.0	10.8
1	1.34	1	N	238	236	2	10.0	0.0	0.5	0.0	0.0	10.8
1	0.32	1	N	238	238	2	4.0	0.0	0.5	0.0	0.0	4.8
1	0.32	1	N	238	238	2	4.0	0.0	0.5	0.0	0.0	4.8
1	1.85	1	N	236	238	2	10.0	0.0	0.5	0.0	0.0	10.8
1	1.85	1	N	236	238	2	10.0	0.0	0.5	0.0	0.0	10.8
1	1.65	1	N	68	237	2	12.5	0.0	0.5	0.0	0.0	13.3
1	1.65	1	N	68	237	2	12.5	0.0	0.5	0.0	0.0	13.3
1	1.07	1	N	170	68	2	9.0	0.0	0.5	0.0	0.0	9.8
1	1.07	1	N	170	68	2	9.0	0.0	0.5	0.0	0.0	9.8
1	1.3	1	N	107	170	2	7.5	0.0	0.5	0.0	0.0	8.3
1	1.3	1	N	107	170	2	7.5	0.0	0.5	0.0	0.0	8.3
1	1.85	1	N	113	137	2	10.0	0.0	0.5	0.0	0.0	10.8
1	1.85	1	N	113	137	2	10.0	0.0	0.5	0.0	0.0	10.8
1	0.62	1	N	231	231	2	4.5	0.0	0.5	0.0	0.0	5.3
1	0.62	1	N	231	231	2	4.5	0.0	0.5	0.0	0.0	5.3
1	0.0	1	N	264	264	2	0.0	0.0	0.0	0.0	0.0	0.0
1	0.29	1	N	162	162	2	4.0	0.0	0.5	0.0	0.0	4.8
1	0.29	1	N	162	162	2	4.0	0.0	0.5	0.0	0.0	4.8
1	1.34	1	N	239	151	2	7.0	0.5	0.5	0.0	0.0	8.3

- Very common we analyze **big data files** based on this view
- Streaming data can be also represented as **unbounded tables**

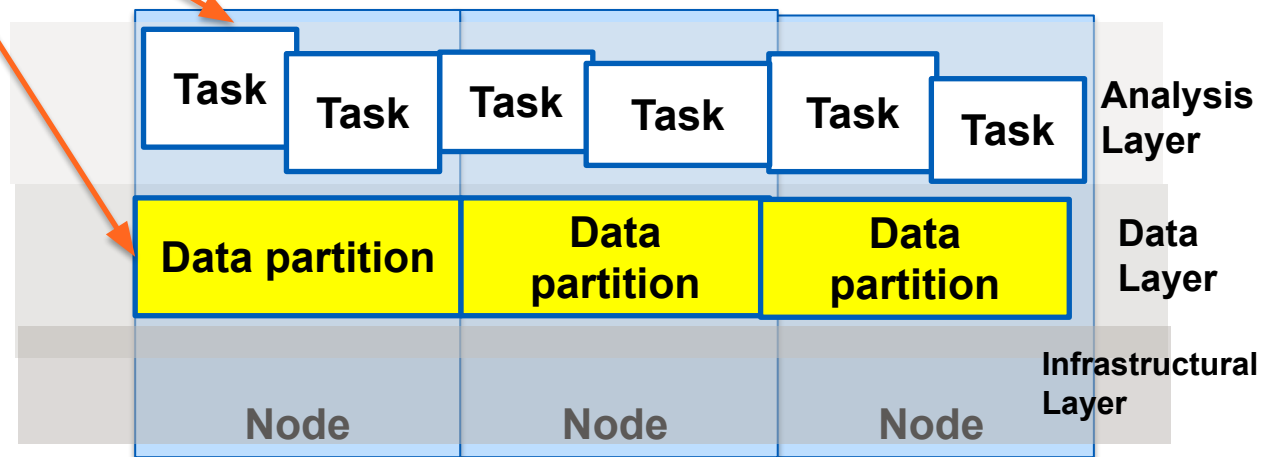
```

inputFile =args.input_file
## hadoop inputFile="hdfs://"
df =spark.read.csv(inputFile,header=True,inferSchema=True)
#df.show()
print("Number of trips", df.count())
#number of passenger count per vendor and total amount of money
passenger_exprs = {"passenger_count":"sum","total_amount":"sum"}
df2 = df.groupBy('VendorID').agg(passenger_exprs)
# Where do you want to write the output
df2.repartition(1).write.csv(args.output_dir,header=True)

```



What we need
when we
develop
analysis
programs for
big data



Big data processing techniques in our focus for data at rest

- **Programming models**

- MapReduce/Spark
- Workflows
- (Distributed) SQL-style processing

- **Studied frameworks**

- Apache Hadoop/Spark, Dask
- Apache Airflow

- **Not in our focus:**

- Bulk synchronous parallel (BSP)
 - HPC MPI (Message Passing Interface)
-

Apache Spark

<https://spark.apache.org/>

Apache Spark

- **Cluster-based high-level computing framework**
- **“unified engine” for different types of big data processing**
 - SQL/structured data processing
 - Machine learning
 - Graph processing
 - Streaming processing
- **It is a powerful computing framework and system \Rightarrow an important service that a big data platform should support**
 - public cloud: Google DataProc, Azure HDInsight, Amazon EMR
 - data lake systems: e.g., Hudi and Delta Lake

Apache Spark

Can be run a top

- Hadoop (using HDFS and YARN)
- Mesos cluster
 - <http://mesos.apache.org/>
- Kubernetes
- Standalone machines

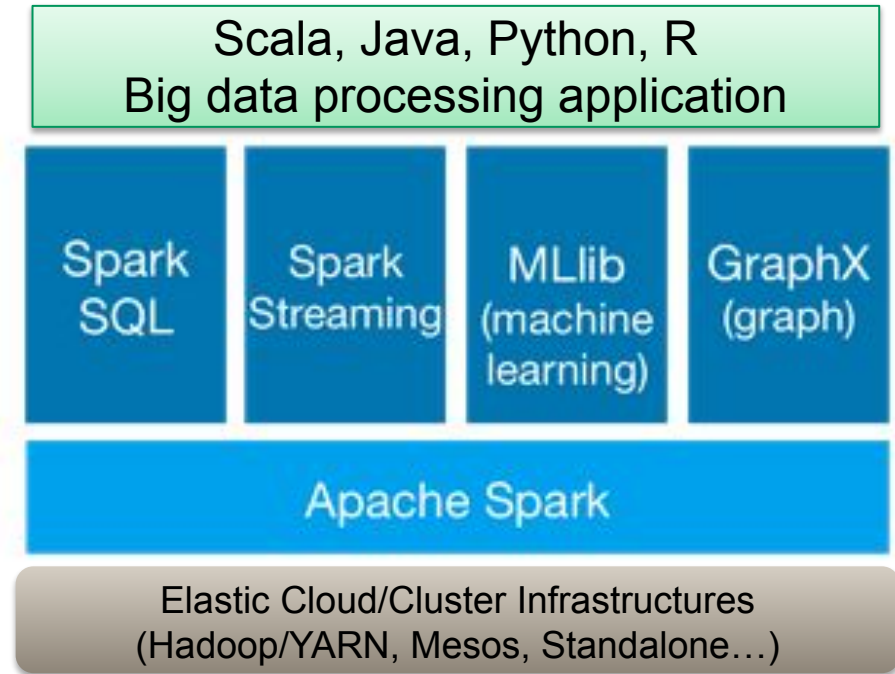
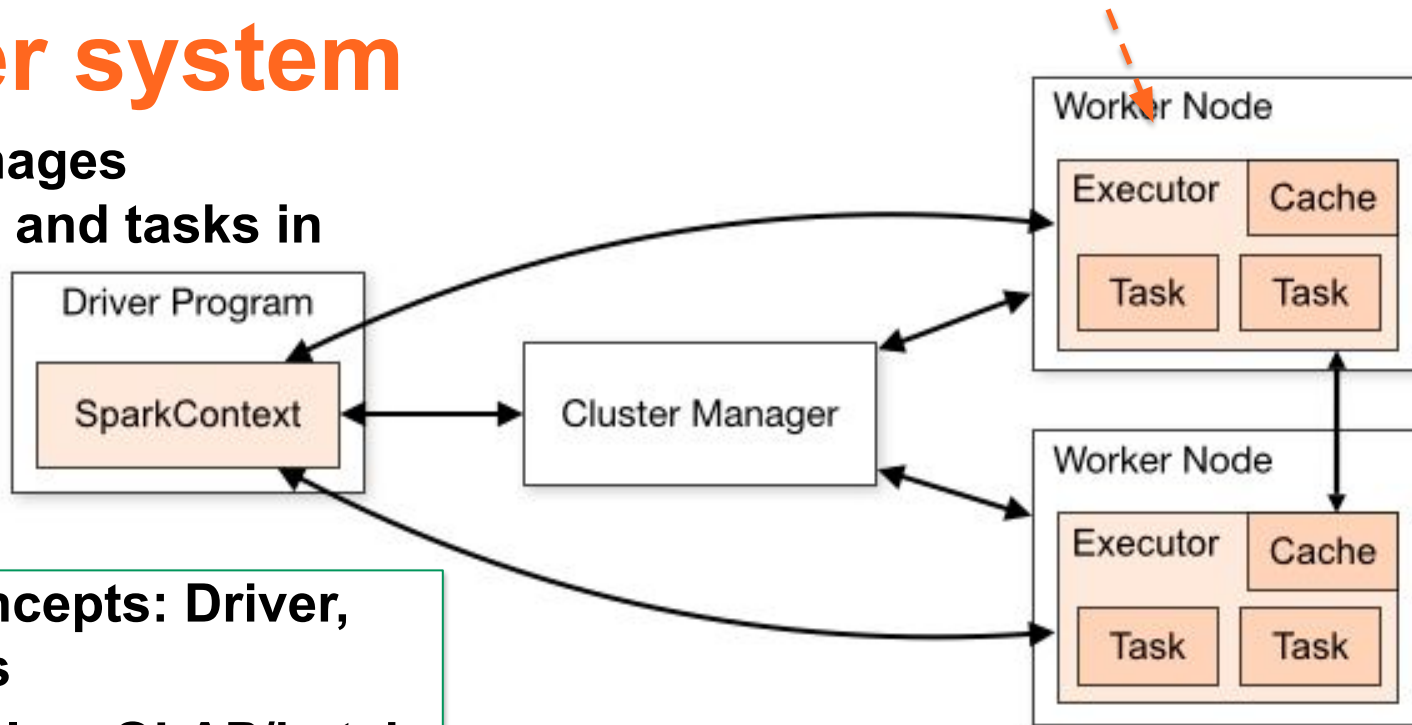


Figure source: <http://spark.apache.org/>

Execution model in a cluster system

Driver manages operations and tasks in nodes

Computing resources in a cluster node



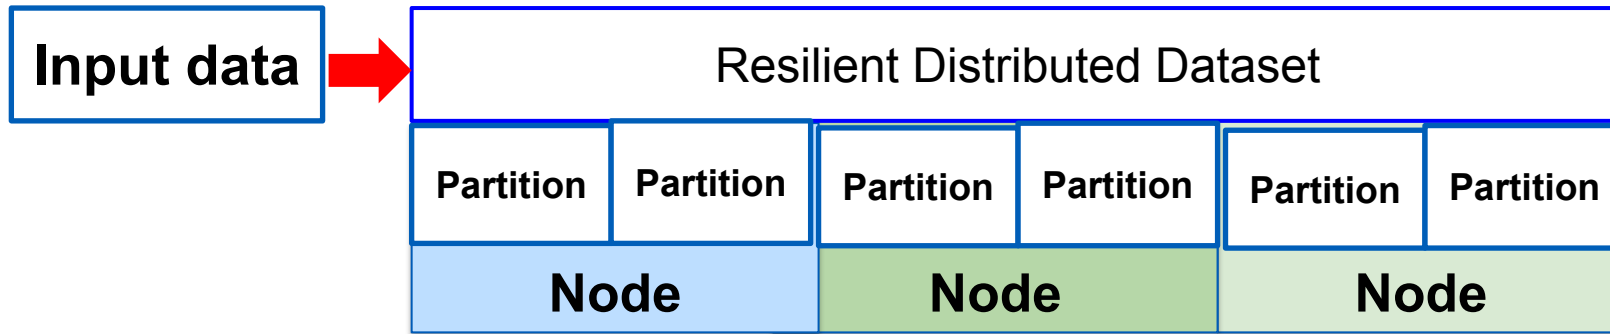
Common concepts: Driver, Nodes, Tasks

Workload styles: OLAP/batch jobs with a lot of data

Figure source:

<http://spark.apache.org/docs/latest/cluster-overview.html>

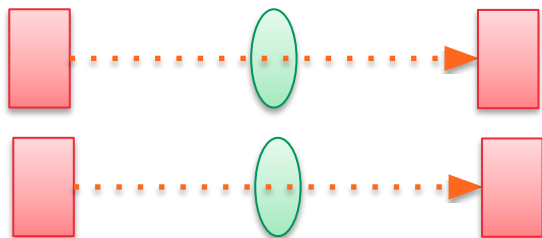
Key features



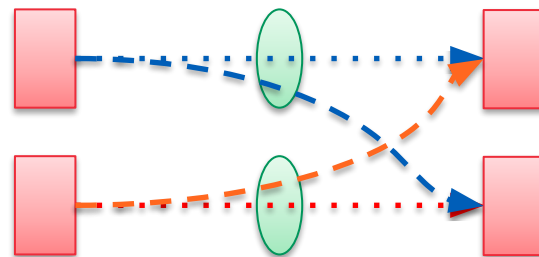
- Input data is **distributed** in different nodes for processing
 - Support partitions for data processing: a node keeps one or n partitions, a partition resides only in a node \Rightarrow for computing
- Key operations: **transformations** and **actions** on data
- Leverage parallel computing concepts to run **multiple tasks**
 - Operation \rightarrow task executed by executor
 - Parallel tasks, task pipeline, DAG of processing stages
- Persistent data in memory/disk for operations

Transformation operations

- **Transformation:**
 - Instructions about how to transform a data in a form to another form \Rightarrow it will not change the original data (immutability)
- **Only tell what to do: to build a DAG (direct acyclic graph): a lineage of what to do**
- **lazy approach \Rightarrow real transformation will be done at **action operators****




**Narrow transformation,
no data shuffle**



**Wide transformation, cross data
partitions, requires a shuffle**

Action operations

- **Compute the results for a set of transformations**
 - Examples: count or average
- **Actions: view, collect, write, calculation**



Lazy approach: an action triggers execution of transformation operations \Rightarrow enable various types of optimization

Spark program: programming elements

- **SparkSession**

- Act as a program driver to manage the execution of tasks
- SparkContext: manages connection to cluster, manage internal services

- **Data APIs**

- Low-level Resilient Distributed Dataset (RDD)
- High-level DataFrames/DataSets
- Load and hold distributed data
- Transformation and action functions

- **ML, graph and streaming functions and pipelines**

Spark application management: high-level view

- **Submission/Request**

- submit the Spark application for running
- resource is provided for running the Driver

- **Launch**

- the Driver requests resources for executors (through SparkContext)
- establish executors across worker nodes

- **Execution**

- the Driver starts to execute code and move data

- **Finish/Completion:**

- finish, release executors

Spark program logic: typical steps

- **Load data and distribute data**
 - data is **immutable** after created
 - data partition in Spark: a partition is allocated in a node
- **Perform **transformations and actions** operations**
 - *transformations*: build plans for transforming data models
 - *actions*: perform computation on data

Resilient distributed dataset (RDD)

- **Low-level data structure**

- collection of data elements partitioned across nodes in the cluster
- with data sharing, parallel operations, fault-tolerant features

- **Create RDD**

- created by loading data from files (text, sequence file) including local file systems, HDFS, Cassandra, HBase, Amazon S3, etc.

- **Persist RDD**

- in memory or to files

RDD transformations and actions

Transformations

- **map**
- **filter**
- **sample**
- **intersection**
- **groupByKey**

Actions

- **reduce()**
- **collect()**
- **count()**
- **saveAs...File()**

Example with RDD

VendorID,tpep_pickup_datetime,tpep_dropoff_datetime,passenger_count,trip_distance,RatecodeID,store_and_fwd_flag,PULocationID,DOLocationID,payment_type,fare_amount,extra,mta_tax,tip_amount,tolls_amount,improvement_surcharge,total_amount

2,11/04/2084 12:32:24 PM,11/04/2084 12:47:41 PM,1,1.34,1,N,238,236,2,10,0,0.5,0,0,0.3,10.8

2,11/04/2084 12:32:24 PM,11/04/2084 12:47:41 PM,1,1.34,1,N,238,236,2,10,0,0.5,0,0,0.3,10.8

2,11/04/2084 12:25:53 PM,11/04/2084 12:29:00 PM,1,0.32,1,N,238,238,2,4,0,0.5,0,0,0.3,4.8

as a text file

```
conf = SparkConf().setAppName("cse4640-rddshow").setMaster(args.master)
sc = SparkContext(conf=conf)
##modify the input data
rdd=sc.textFile(args.input_file)
## if there is a header we can filter it otherwise comment two lines
csvheader = rdd.first()
rdd = rdd.filter(lambda csventry: csventry != csvheader)
## using map to parse csv text entry
rdd=rdd.map(lambda csventry: csventry.split(","))
rdd.repartition(1)
rdd.saveAsTextFile(args.output_dir)
```


Shared variables

- **A function is executed a remote and various tasks running in parallel**
 - how do tasks share variables? common patterns in parallel computing: *broadcast and global variable/counter*
- **Variables used in parallel operations**
 - variables are copied among parallel tasks
 - shared among tasks or between tasks and the driver
- **Types of variables**
 - broadcast variables: cache a value in all nodes
 - accumulators: a global counter shared across processes

Examples

```
conf = SparkConf().setAppName("CS-E4640-Broadcast").setMaster("ygs:master")
sc = SparkContext(conf=conf)
bVar = sc.broadcast([5,10])
print("The value of the broadcast",bVar.value,sep=" ")
counter = sc.accumulator(0)
sc.parallelize([1, 2, 3, 4]).foreach(lambda x: counter.add(bVar.value[0]))
print("The value of the counter is ",counter.value,sep=" ")
```

Use cases:

- **Broadcast variables: lookup tables**
- **Accumulators: monitoring/checkpoint counters**

Spark SQL and DataFrames

- **High-level APIs**
 - design with common programming patterns in data analysis, multi-language support
- **SparkSQL: enable dealing with structured data**
 - SQL query execution, Hive, JDBC/ODBC
- **DataFrame**
 - distributed data organized into named columns, similar to a table in relational database
 - Pandas and Spark DataFrames have similar design concepts

DataFrame

```
inputFile =args.input_file
df =spark.read.csv(inputFile,header=True,inferSchema=True)
print("Number of partition",df.rdd.getNumPartitions())
df.show()
```

PROVINCECODE	DEVICEID	IFINDEX	FRAME	SLOT	PORT	ONUINDEX	ONUID	TIME	SPEEDIN	SPEEDOUT
YN 1	3023	528	1	2	7	39 10	07039	01/08/2019 00:04:07	148163	49018
YN 1	3023	528	1	2	7	38 10	07038	01/08/2019 00:04:07	1658	1362
YN 1	3023	528	1	2	7	9 10	07009	01/08/2019 00:04:07	6693	5185
YN 1	3023	528	1	2	7	8 10	07008	01/08/2019 00:04:07	640	544
YN 1	3023	528	1	2	7	11 10	07011	01/08/2019 00:04:07	118	114
YN 1	3023	528	1	2	7	10 10	07010	01/08/2019 00:04:07	28514	12495
YN 1	3023	528	1	2	7	13 10	07013	01/08/2019 00:04:07	868699	23400
YN 1	3023	528	1	2	7	15 10	07015	01/08/2019 00:04:07	1822	1120
YN 1	3023	528	1	2	7	17 10	07017	01/08/2019 00:04:07	998069	117345
YN 1	3023	528	1	2	7	16 10	07016	01/08/2019 00:04:07	22402	1804
YN 1	3023	528	1	2	7	19 10	07019	01/08/2019 00:04:07	640	791
YN 1	3023	760	1	1	10	49 10	10049	01/08/2019 00:04:07	662	494
YN 1	3023	760	1	1	10	48 10	10048	01/08/2019 00:04:07	2158	759
YN 1	3023	528	1	2	7	21 10	07021	01/08/2019 00:04:07	0	0
YN 1	3023	760	1	1	10	51 10	10051	01/08/2019 00:04:07	2600890	54153
YN 1	3023	528	1	2	7	20 10	07020	01/08/2019 00:04:07	330	184



Create DataFrame

DataFrames can be created from a Hive table, from Spark data sources, or another DataFrame

Load and save

- From Hive, JSON, CSV
- HDFS, cloud object storage (AWS S3, Google Cloud Storage, Azure Blob Storage), local files, etc.



Figure source:

<https://databricks.com/blog/2015/02/17/introducing-dataframe-s-in-spark-for-large-scale-data-science.html>

DataFrame Transformations & Actions

- **Several transformations can be done**
 - *Think transformation for relational database or matrix*
- **Select**
 - *df.select*
- **Filter**
 - *df.filter*
- **Groupby**
 - *df.groupBy*
- **Handle missing data**
 - *Drop duplicate rows, drop rows with NA/null data*
 - *Fill NA/null data*

Actions

- Return values calculated from DataFrame

Examples

- reduce, max, min, sum, variance and stdev

⇒ **Distributed and parallel processing but it is done by the framework**

Example of a Spark

```
#!/usr/bin/env python2
#encoding: UTF-8
# CS-E4640
import csv
import sys
from datetime import datetime
from pyspark.sql import SparkSession
import numpy as np
from pyspark.sql import functions as F
import argparse
parser = argparse.ArgumentParser()
parser.add_argument('--input_file', help='input data file')
parser.add_argument('--output_dir', help='output dir')
args = parser.parse_args()

##define a context
spark = SparkSession.builder.appName("cse4640-onu").getOrCreate()
#NOTE: using hdfs:///..... for HDFS file or file:///
inputFile = args.input_file
df = spark.read.csv(inputFile, header=True, inferSchema=True)
#df.show()
print("Number of records", df.count())
exprs = {"SPEEDIN": "avg"}
df2 = df.groupBy('ONUID').agg(exprs)
df2.repartition(1).write.csv(args.output_file, header=True)
```

Session/Driver



Read data



Apply operations



Spark application runtime view

- **Tasks:**

- a unit of work executed in an executor: e.g., performing transformations of a data partition

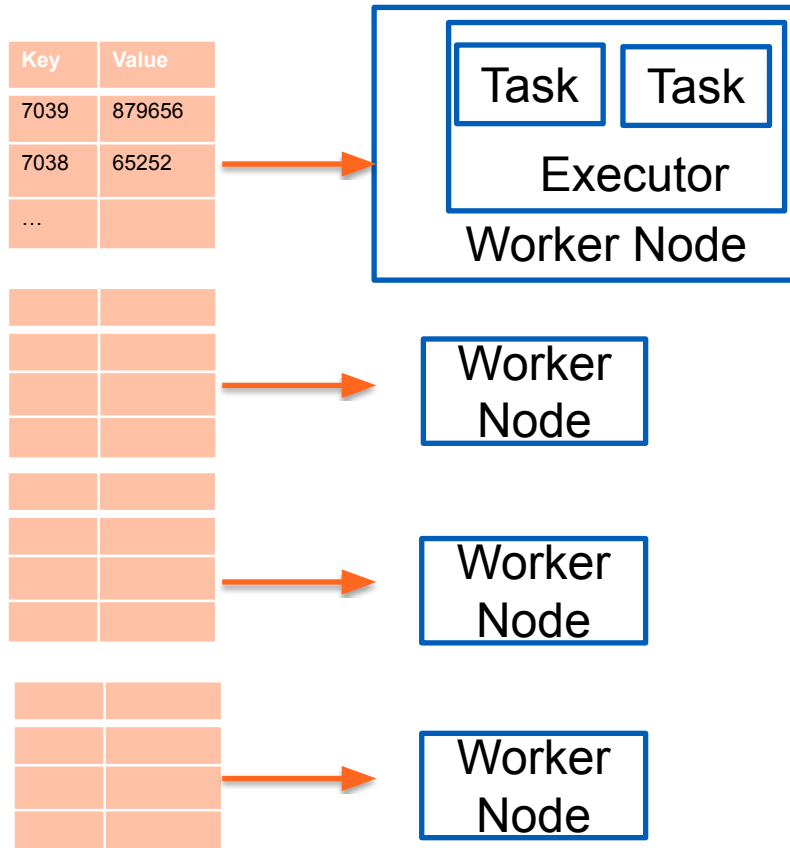
- **Stage: Shuffle Map Stage & Result Stage**

- a set of tasks executed in many nodes for performing the same operation
- move to a new stage: through a shuffle to produce output partitions or an action to produce results

- **Job**

- runtime view of an action operation (actual computation produces a result), includes many stages of tasks

Data Distribution



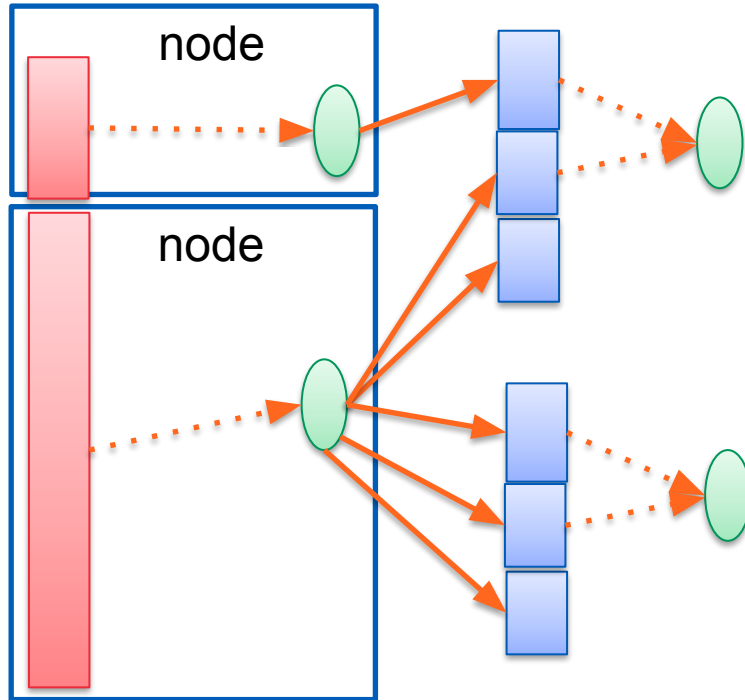
One task works on a partition at a time

⇒ Parallelism and performance are strongly dependent on number of partitions, tasks, CPU cores

Data Distribution: Load balance

Imbalance
processing

more data shuffle

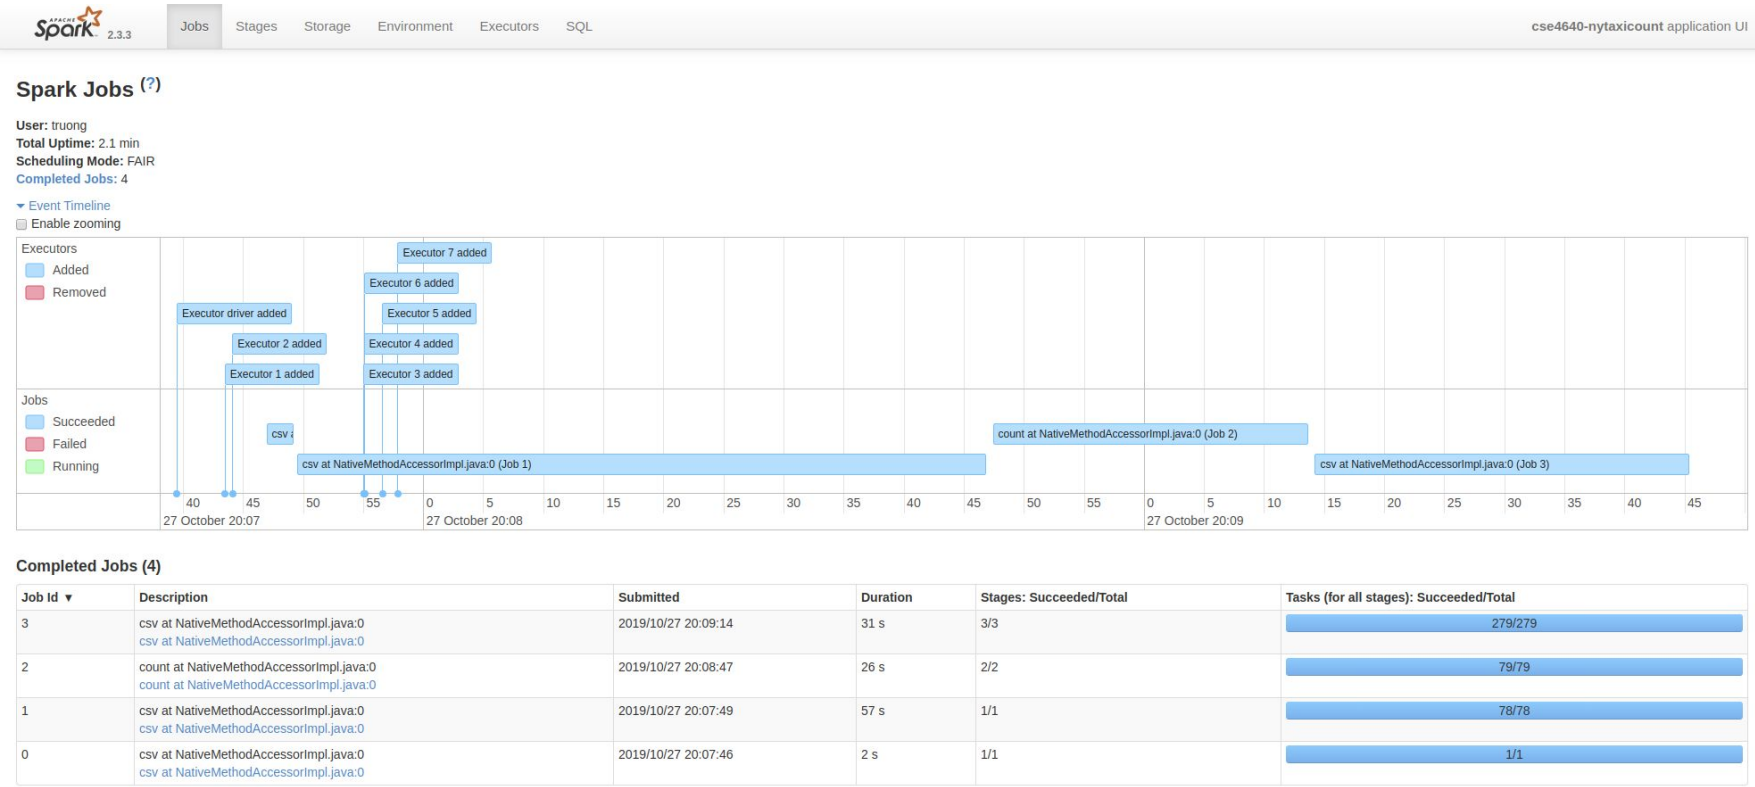


- It is important to have well-balanced data distribution across nodes
- **Detection:**
 - look at runtime execution time to see problems or check your data
- **Examples of solution:**
 - repartition
 - change group keys

Pipelining, Shuffle and DAG

- **Operations work in a pipeline without moving data across nodes**
 - e.g., map->filter, select->filter
- **Shuffle persistent**
 - shuffle needs move data across nodes
 - source tasks save shuffle files into local disks for data shuffle, then the target tasks will read data from source nodes
 - *Save time, recovery, fault tolerance*

Monitoring Spark: Executors and tasks



Executors and tasks

Shuffle Write: 216.0 B / 3

▶ DAG Visualization

▶ Show Additional Metrics

▼ Event Timeline

☐ Enable zooming

■ Scheduler Delay ■ Executor Computing Time ■ Getting Result Time
■ Task Deserialization Time ■ Shuffle Write Time
■ Shuffle Read Time ■ Result Serialization Time



Other important support of Spark

- **MLlib - Machine learning**

- Distributed and parallel machine learning algorithms with big data and clusters

- **Streaming: data processing in near real-time**

- *Related to our topic: stream data processing*

- **Graph Processing: Spark GraphX**

- Parallel computation for graphs

- **Many third-party frameworks, e.g.,**

- SparkOCR (<https://www.johnsnowlabs.com/spark-ocr/>), SparkNLP (<https://nlp.johnsnowlabs.com/>)
- PyDeequ (<https://pydeequ.readthedocs.io/en/latest/README.html#>) Data quality
 - check our example:
<https://github.com/rdsea/bigdataplatforms/tree/master/tutorials/dataquality>

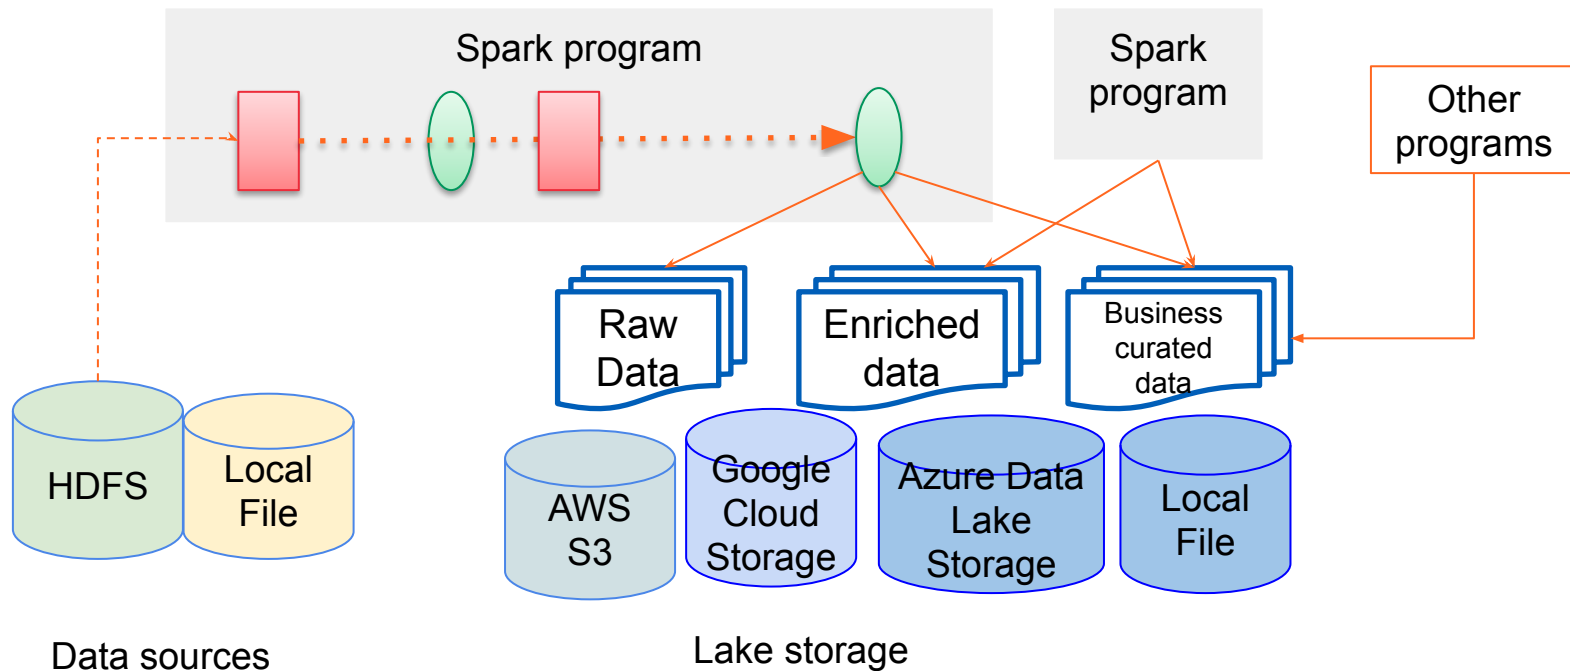
Spark as a key programming model/analytic engine for Data Lake

- **Modern lake data: cloud or on-premise**
 - *multiple types of data* from different sources (databases, files, sensors, etc)
 - *different forms in storage*: raw data, enriched/processed/cleansing, application-/business curated data, sandbox data (for testing, collaboration)
 - *common, standard, cost optimal storage*: object storage (S3, Azure), (distributed) file storage (Hadoop FS), ...
- **Data Lake Core**
 - Data tables, metadata and catalogs
 - Open standards: Parquet, ORC, Iceberg tables, Delta Lake formats
 - Many processing and governance tasks

Spark as a key programming model/analytic engine for Data Lake

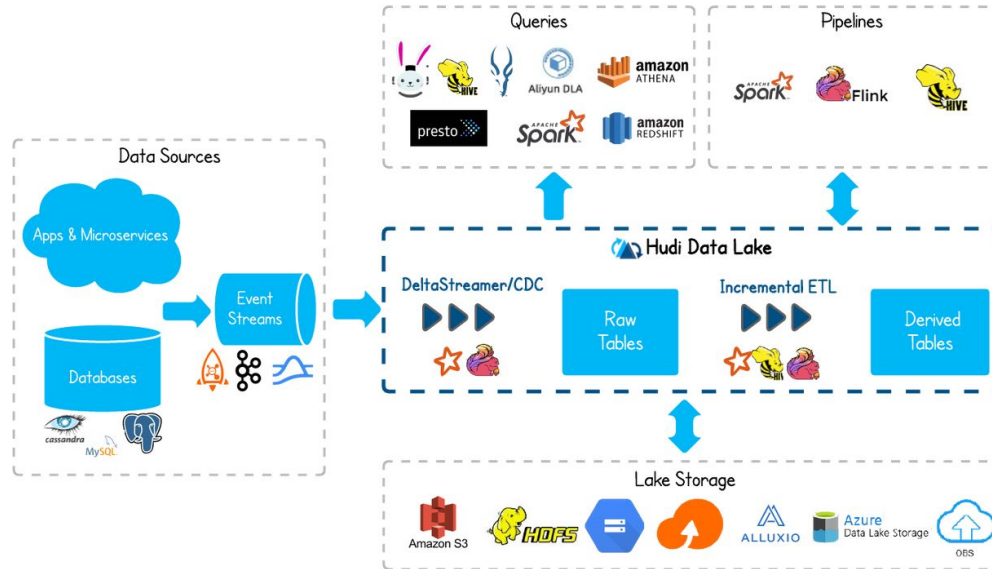
- **Many tasks required:**
 - Ingestion (insert/update)
 - Transformation
 - Query
 - Quality controls
- **Spark as an important engine**
 - for batch processing and stream processing (next lecture)
 - deal with different data formats
 - work with different lake storage
 - still in the same framework
- Core engine for Data lake platforms: Apache Hudi, Delta Lake

Spark programs for ingestion/analytics of lake data



Example: Hudi Data Lake

Figure source: <https://hudi.apache.org/>



different analytics

lake core

different storage

Problems: hard to govern, technology-centric, lack of considerations of data products, cost of effort, etc.

Example

Spark program with Spark Delta for processing data and store the processed data into a cloud data lake storage

```
## hadoop inputFile="hdfs://"
spark_df =spark.read.csv(inputFile,header=True,inferSchema=True)
print(spark_df.head(10))
#do many things, before producing data for datalake
spark_df.write.format("delta").mode("append").save(lake_table_path)
```

A program to read and write data from/to the same lake (delta-rs package, not Spark)

```
if args.read_only != "yes":
    # read data from csv file, no error checking
    df = pd.read_csv(args.input_file)
    write_deltalake(args.lake_table_path, df)
# Read from the lake and print out the first 100 entries
# Load data from the delta table
lake_table_data= DeltaTable(args.lake_table_path)
df_result = lake_table_data.to_pandas()
print(df_result.head(100))
```

E.g., Data lake storage based on Google Cloud Storage

Summary

- **Facts:**

- Spark is an important framework
- A user/developer needs to learn to develop Spark applications
- A platform operator/provider offer services for managing resources, processing and monitoring

- **Thoughts:**

- Think about the success of Apache Spark: rich ecosystems!
- Think if you combine data, different distributed programming supports for your big data platform

Thanks!

Hong-Linh Truong
Department of Computer Science

rdsea.github.io